



Building a data warehouse with Amazon Redshift

... and a quick look at Amazon Machine Learning

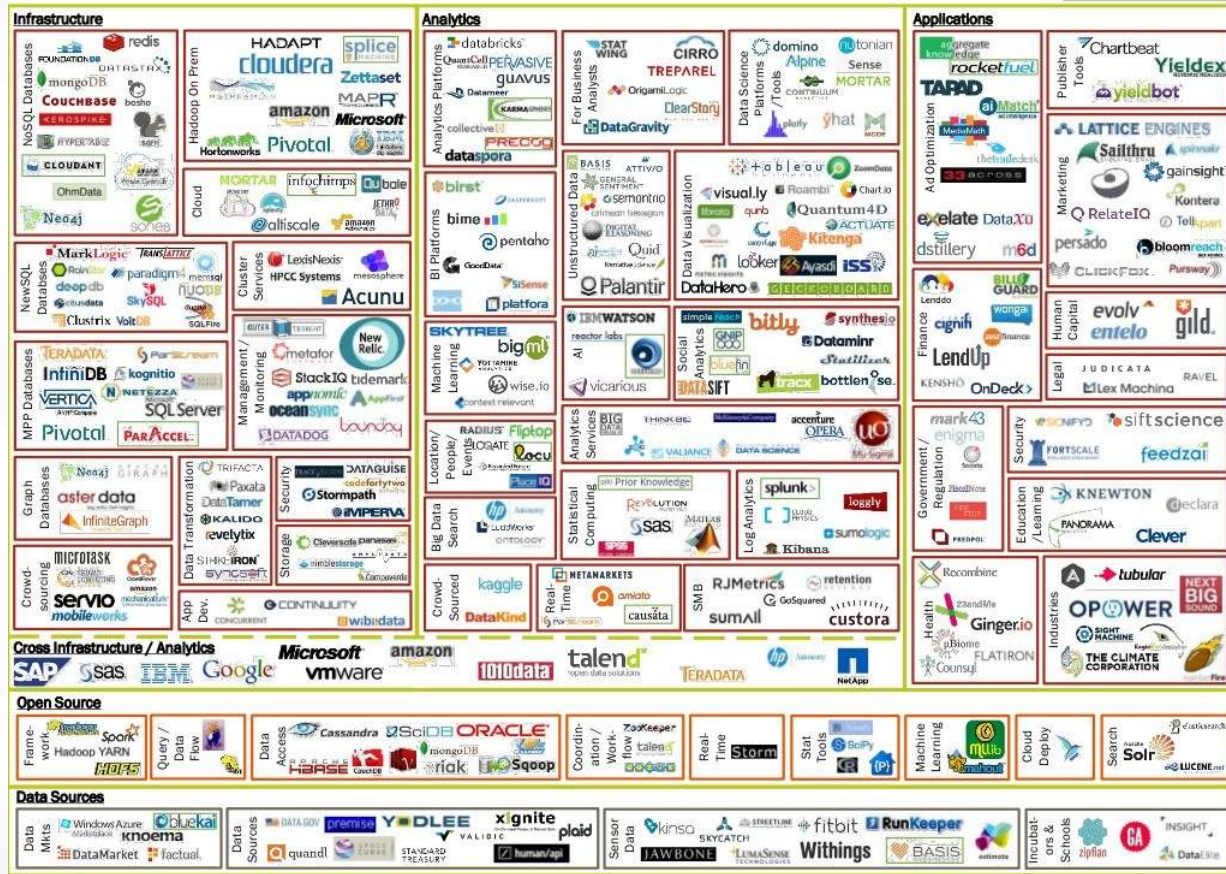
Julien Simon, Principal Technical Evangelist


@julsimon

Navigating the seven seas of Big Data

BIG DATA LANDSCAPE, VERSION 3.0

Exited: Acquisition or IPO






You

Your boss

Your infrastructure





Amazon Redshift

a relational, petabyte scale, fully managed data warehousing service

- SQL is all you need to know
- ODBC and JDBC drivers available
- Parallel processing on multiple nodes
- No system administration
- Free tier: 750 hours / month for 2 months
- Available on-demand from \$0.25 / hour / node
- As low as \$1,000 / Terabyte / year



“ Come for the cost, stay for the performance ”

A close-up photograph of a red, oval-shaped button with the words "ENGINE START" in white, sans-serif capital letters. The button is set within a silver-colored metallic ring. This ring is mounted on a black plastic base with concentric circular ridges. The entire assembly is on a dark grey or black dashboard. To the right, a portion of a black rotary knob is visible.

ENGINE
START

Amazon Redshift architecture

Parallel processing


Columnar data storage

Data compression

Query optimization

Compiled code

Workload management



Row storage vs columnar storage

SSN	Name	Age	Addr	City	St
101259797	SMITH	88	899 FIRST ST	JUNO	AL
892375862	CHIN	37	16137 MAIN ST	POMONA	CA
318370701	HANDU	12	42 JUNE ST	CHICAGO	IL

101259797 SMITH 88 899 FIRST ST JUNO AL	892375862 CHIN 37 16137 MAIN ST POMONA CA	318370701 HANDU 12 42 JUNE ST CHICAGO IL
---	---	--

Block 1	Block 2	Block 3
---------	---------	---------

SSN	Name	Age	Addr	City	St
101259797	SMITH	88	899 FIRST ST	JUNO	AL
892375862	CHIN	37	16137 MAIN ST	POMONA	CA
318370701	HANDU	12	42 JUNE ST	CHICAGO	IL

101259797 892375862 318370701	468248180 378568310 231346875 317346551 770336528 277332171 455124598 735885647 387586301
-----------------------------------	---

Block 1

Instance types

Dense Storage Node Types

Node Size	vCPU	ECU	RAM (GiB)	Slices Per Node	Storage Per Node	Node Range	Total Capacity
ds1.xlarge	2	4.4	15	2	2 TB HDD	1-32	64 TB
ds1.8xlarge	16	35	120	16	16 TB HDD	2-128	2 PB
ds2.xlarge	4	13	31	2	2 TB HDD	1-32	64 TB
ds2.8xlarge	36	119	244	16	16 TB HDD	2-128	2 PB

Dense Compute Node Types

Node Size	vCPU	ECU	RAM (GiB)	Slices Per Node	Storage Per Node	Node Range	Total Capacity
dc1.large	2	7	15	2	160 GB SSD	1-32	5.12 TB
dc1.8xlarge	32	104	244	32	2.56 TB SSD	2-128	326 TB

Case study: Photobox



<http://www.lemagit.fr/etude/Photobox-consolide-et-analyse-ses-donnees-avec-AWS-RedShift>

Maxime Mezin, Data & Photo Science Director:

“L’entrepôt de données ne comportait que les données du site e-commerce liées aux ventes. Alors que nous avons la volonté d’intégrer des données du service clients et des données d’analyse (...) Nous avons atteint la limite du stockage de la base Oracle, et cela ne marchait pas très bien en termes de performances”

“Avec Redshift, la rapidité d’exécution des traitements a été multipliée par 10. Sans parler de la vitesse de chargement des données”

“On paie en fonction de la quantité de données que l’on va stocker. Chez Google, cela était plus compliqué”

- 2 Redshift clusters : 1 for historical data, 1 for real-time processing (SSD)
- TCO divided by 7 (90K€→13K€)

Case study: Financial Times

<https://aws.amazon.com/solutions/case-studies/financial-times/>



- BI analysis of reader traffic, in order to decide which stories to cover
- Conventional data warehouse running on Microsoft technologies
- Scalability issues, impossible to perform real-time analytics → Amazon Redshift PoC
- Amazon Redshift performed so quickly that some analysts thought it was malfunctioning 😊

John O'Donovan, CTO: *"Amazon Redshift is the single source of truth for our user data."*

"Some of the queries we're running are 98 percent faster, and most things are running 90 percent faster (...) and the ability to try Redshift out before having to invest a significant amount of capital was a huge bonus."

"Being able to explore near-real-time data improves our decision making massively. We can make decisions based on what's happening now rather than what happened three or four days ago."

- Total Cost of Ownership divided by 4




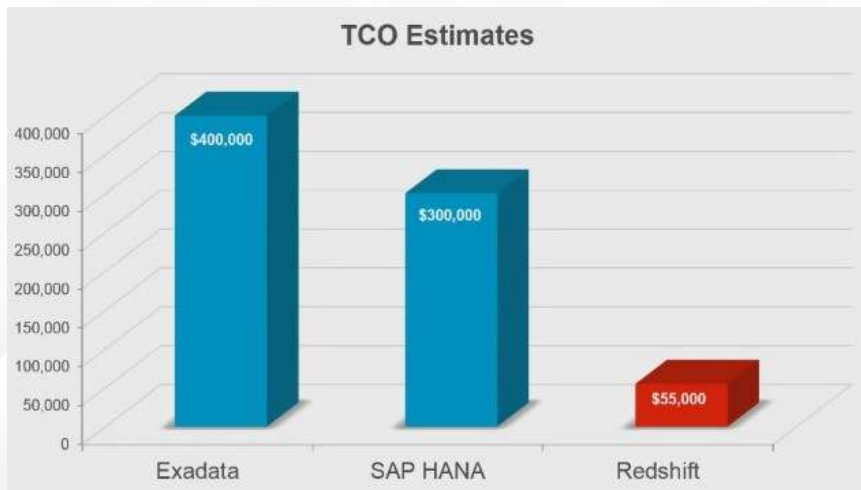
Case study: Boingo

<https://www.youtube.com/watch?v=58URZbp1voY>



- Largest operator of airport wireless hotspots in the world: 1M+ hotspots, 100+ countries
- About 15 TB of data, growing at 2-3 TB per year
- Platform running on SAP (ETL) & Oracle 11g: low performance, heavy admin, high cost
- Evaluated Oracle Exadata, SAP HANA and Amazon Redshift
- Selected Amazon Redshift and migrated in 2 months

6-7x less expensive than alternatives



Queries
180x faster

Data load
480x faster



Amazon Redshift performance

No indexes, no partitioning, no wizardry.

Distribution key

- How data is spread across nodes
- EVEN (default), ALL, KEY

Sort key

- How data is sorted inside of disk blocks
- Compound and interleaved keys are possible

Both are crucial to query performance!



Universal Pictures

DEMO #1

*Demo gods, I'm your humble servant,
please be good to me*

8-node cluster


1 billion lines of CSV data (45GB)

→ 1 billion rows in a single table



Amazon Machine Learning

*A managed service for building ML models
and generating predictions*



Integration with Amazon S3, Redshift and RDS
Data transformation, visualization and exploration
Model evaluation and interpretation tools
API for batch and real-time predictions

\$0.42 / hour for analysis and model building (eu-west-1)
\$0.10 per 1000 batch predictions
\$0.0001 per real-time prediction

Case study: BuildFax

<https://aws.amazon.com/solutions/case-studies/buildfax/>

BuildFax:
On-Demand Property Condition.

More than
8,000
Cities &
Counties


90+
MILLION
Properties

185+
MILLION
Building
Permits

10+
BILLION
Data Points

“Amazon Machine Learning democratizes the process of building predictive models. It's easy and fast to use, and has machine-learning best practices encapsulated in the product, which lets us deliver results significantly faster than in the past”

Joe Emison, Founder &
Chief Technology Officer



DEMO #2

*Demo gods, I know I'm pushing it,
but please don't let me down now*

Load data from Amazon Redshift

Train and evaluate a regression model with Amazon ML

Create a real-time prediction API

Perform real-time prediction from Java app



AWS User Groups




Lille
Paris
Rennes
Nantes
Lyon



facebook.com/groups/AWSFrance/



[@aws_actus](https://twitter.com/aws_actus) [@julsimon](https://twitter.com/julsimon)



Interested in starting a group? Talk to us!

Save the date: AWS Paris Summit, 31/05/2016



Events

AWSome Day



March 16, June 28, September 27, December 6

IoT
world

March 23-24



April 6-7 (Lyon)

DEVOXX France

April 20-22



April 25



Thank you !