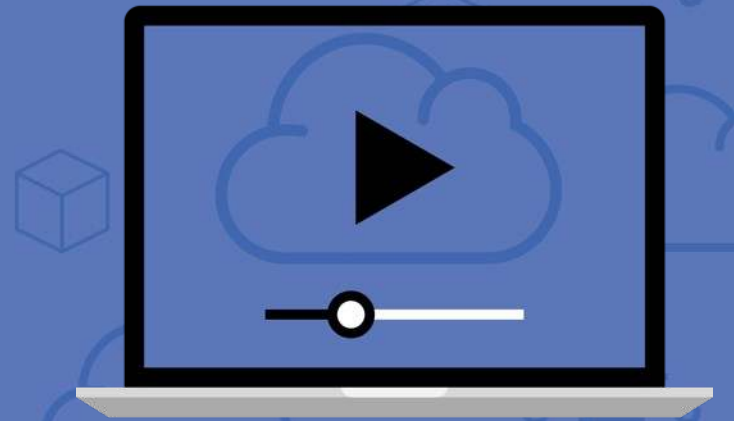




An introduction to Deep Learning

**Julien Simon, AI Evangelist,
EMEA
@julsimon**



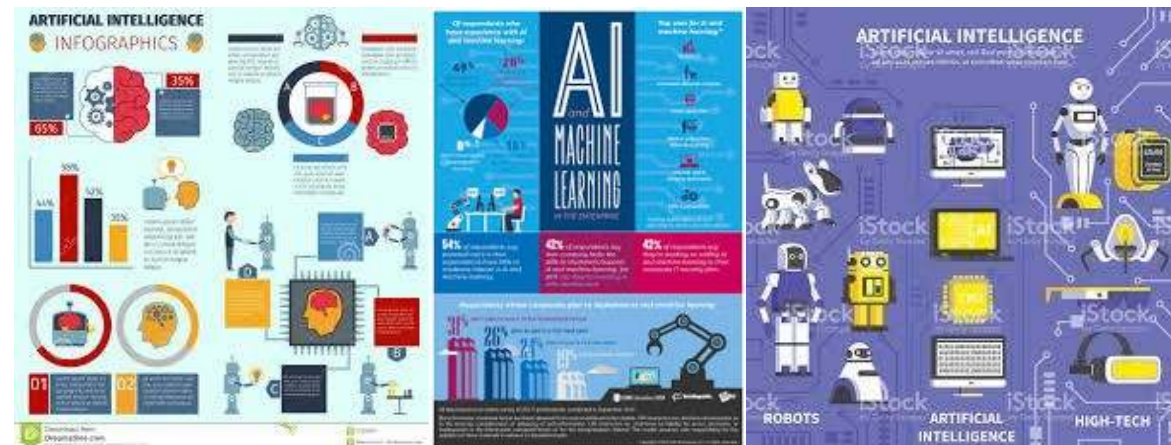
What to expect

- Artificial Intelligence, Machine Learning, Deep Learning
- The 5 myths of AI
- Deep Learning in action
- Basics of Deep Learning
- NVIDIA Volta V100 and AWS P3
- Q&A

- **Artificial Intelligence**: design software applications which exhibit human-like behavior, e.g. speech, natural language processing, reasoning or intuition
- **Machine Learning**: teach machines to learn without being explicitly programmed
- **Deep Learning**: using neural networks, teach machines to learn from data where features cannot be explicitly expressed

The 5 Myths of AI

Myth #1 - AI is the flavour of the month



Fact #1 - AI is 60 years old

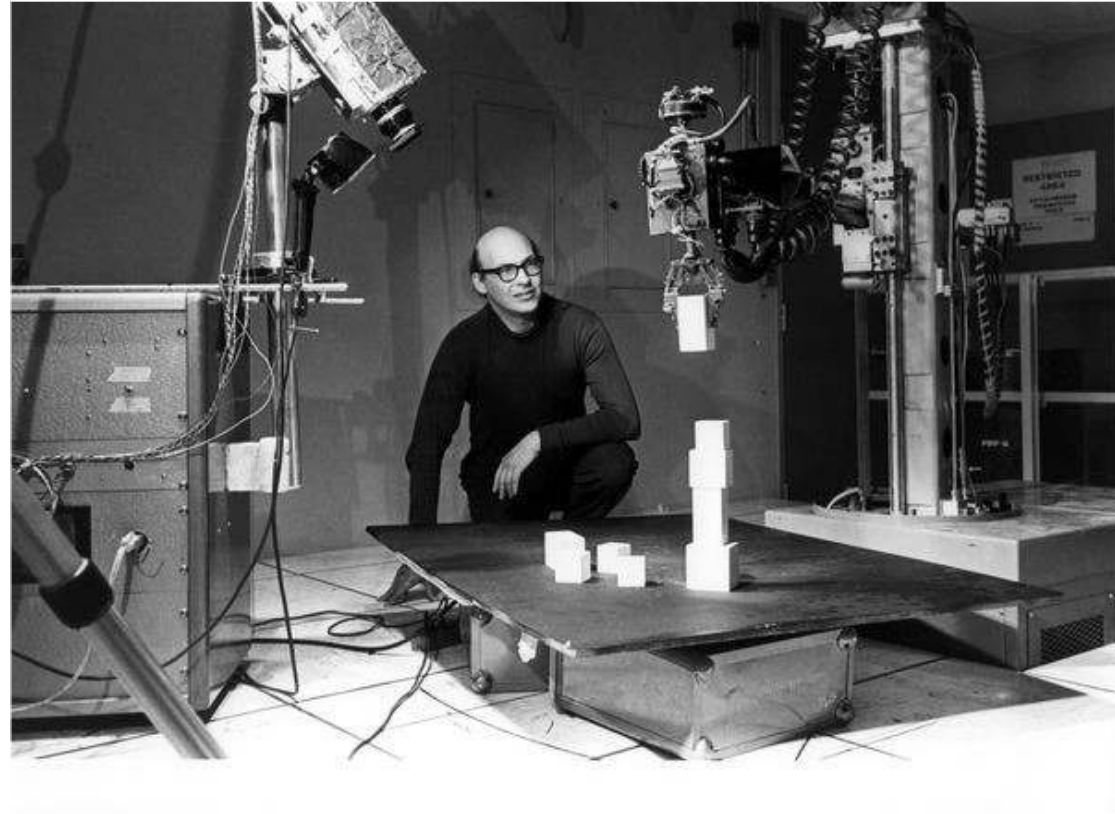


John McCarthy (1927-2011)

1956 - Coined the term “Artificial Intelligence”

1958 - Invented LISP

1971 - Received the Turing Award



Marvin Minsky (1927-2016)

1959 - Co-founded the MIT AI Lab

1968 - Advised Kubrick on “2001: A Space Odyssey”

1969 - Received the Turing Award

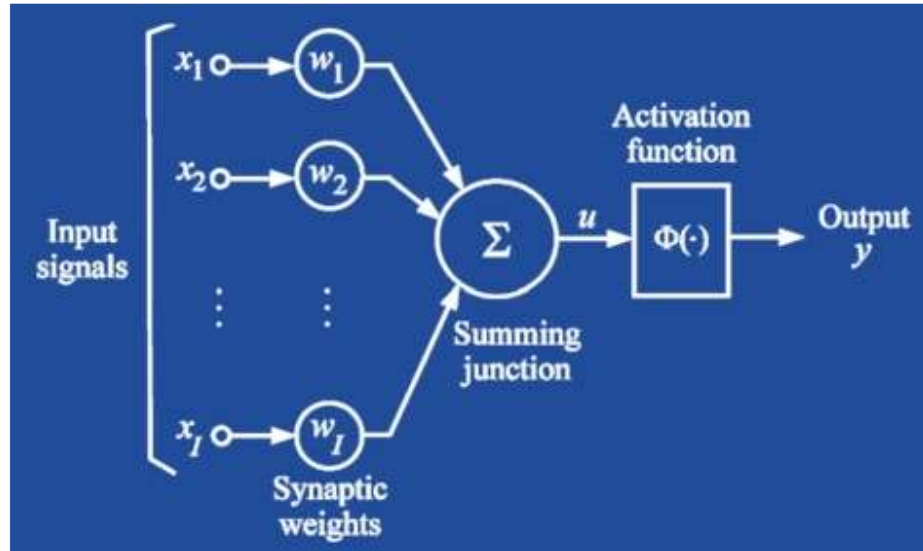
Myth #2 - AI is dark magic

aka « You're not smart enough »



Fact #2 - AI is math, code and chips

A bit of Science, a lot of Engineering



```
data = mx.symbol.Variable('data')
conv1 = mx.sym.Convolution(data=data, kernel=(5,5), num_filter=20)
relu1 = mx.sym.Activation(data=conv1, act_type="relu")
pool1 = mx.sym.Pooling(data=relu1, pool_type="max", kernel=(2,2), stride=(2,2))
conv2 = mx.sym.Convolution(data=pool1, kernel=(5,5), num_filter=50)
relu2 = mx.sym.Activation(data=conv2, act_type="relu")
pool2 = mx.sym.Pooling(data=relu2, pool_type="max", kernel=(2,2), stride=(2,2))
flatten = mx.sym.Flatten(data=pool2)
fc1 = mx.symbol.FullyConnected(data=flatten, num_hidden=500)
relu3 = mx.sym.Activation(data=fc1, act_type="relu")
fc2 = mx.symbol.FullyConnected(data=relu3, num_hidden=10)
lenet = mx.sym.SoftmaxOutput(data=fc2, name='softmax')
```



Myth #3 – The “cognitive” unicorn

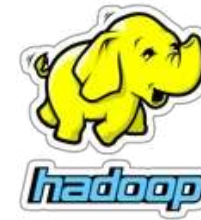


Myth #3 – The “cognitive” unicorn



Fact #3: AI is a wide range of techniques and tools

- Machine Learning
- Natural Language Processing
- Speech
- Vision
- Expert Systems
- And more



python™



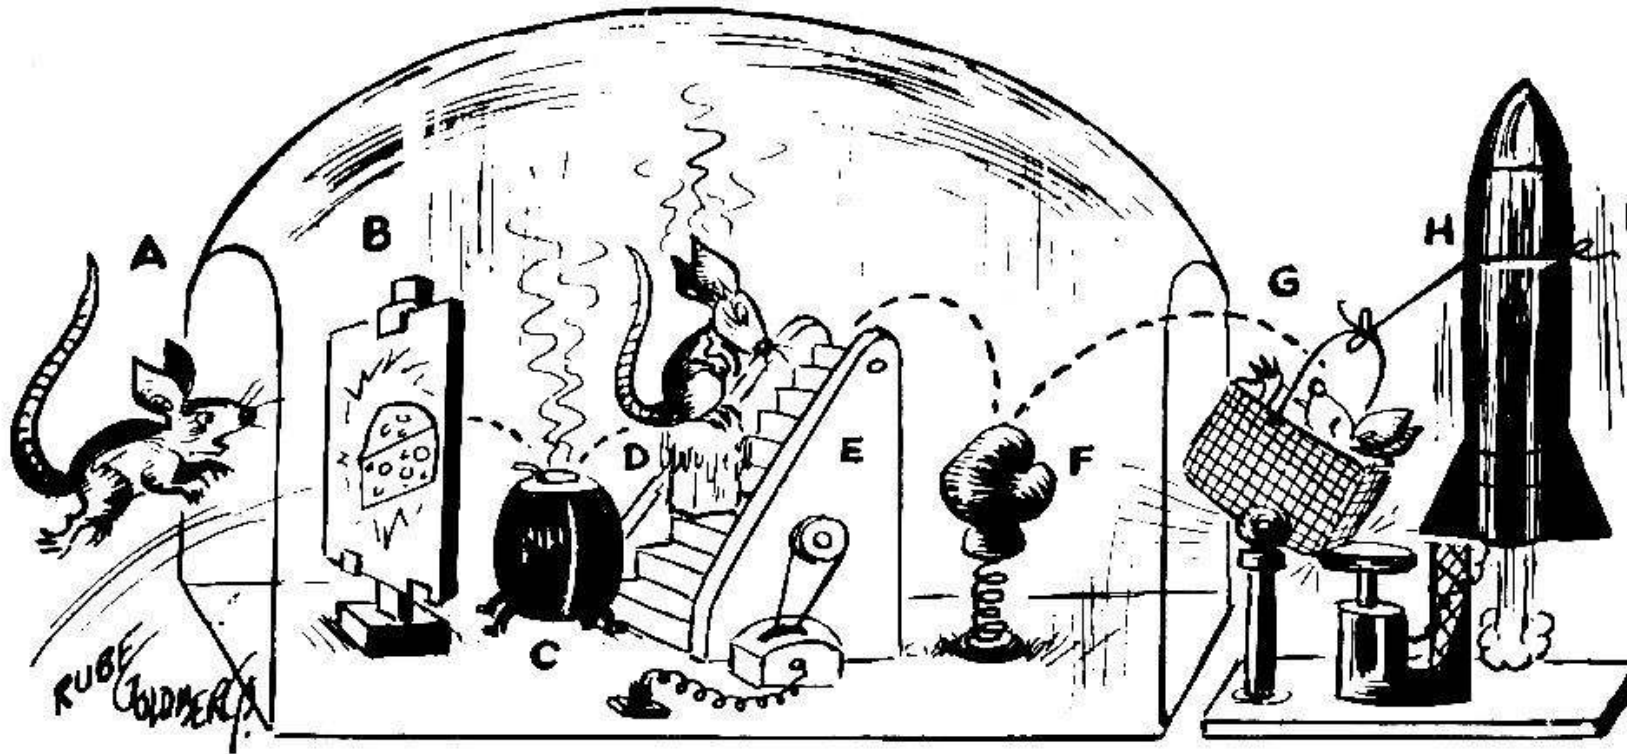
C++



Caffe

Myth #4 - AI is for esoteric use cases

How to Get Rid of a Mouse



Drawn for Newsweek by Rube Goldberg

The best mousetrap by Rube Goldberg: Mouse (A) dives for painting of cheese (B), goes through canvas and lands on hot stove (C). He jumps on cake of ice (D)

to cool off. Moving escalator (E) drops him on boxing glove (F) which knocks him into basket (G) setting off miniature rocket (H) which takes him to the moon.

Fact #4: AI shines on intuitive problems



Amazon AI is based on Deep Learning

Vision Services

Amazon Rekognition Image

Deep learning-based image analysis

[Learn more »](#)

Amazon Rekognition Video

Deep learning-based video analysis

[Learn more »](#)



Conversational chatbots

Amazon Lex

Build chatbots to engage customers

[Learn more »](#)

Language Services

Amazon Comprehend

Discover insights and relationships in text

[Learn more »](#)



Amazon Translate

Fluent translation of text

[Learn more »](#)



Amazon Transcribe

Automatic speech recognition

[Learn more »](#)



Amazon Polly

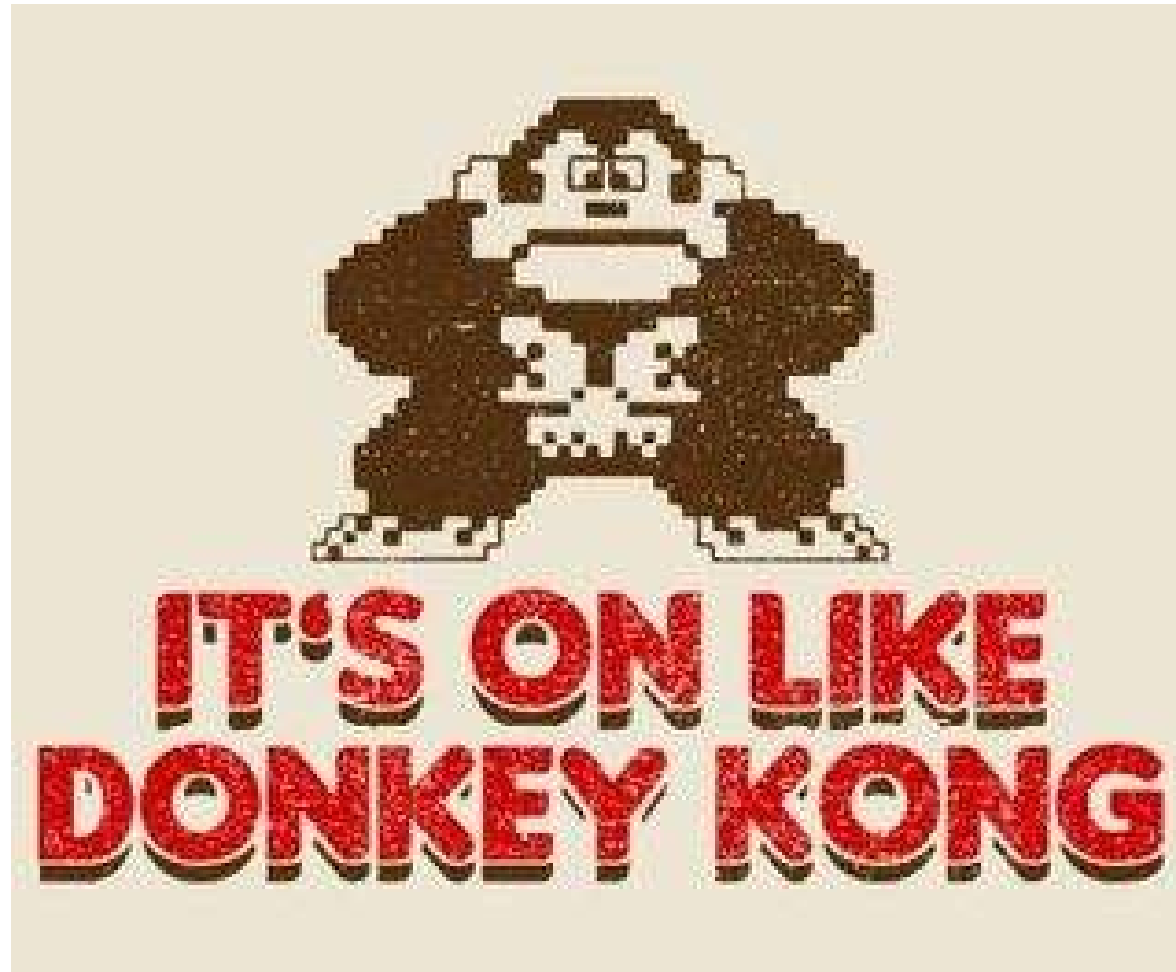
Natural sounding text to speech

[Learn more »](#)

Myth #5 - AI is not production-ready



Fact #5: AI means business



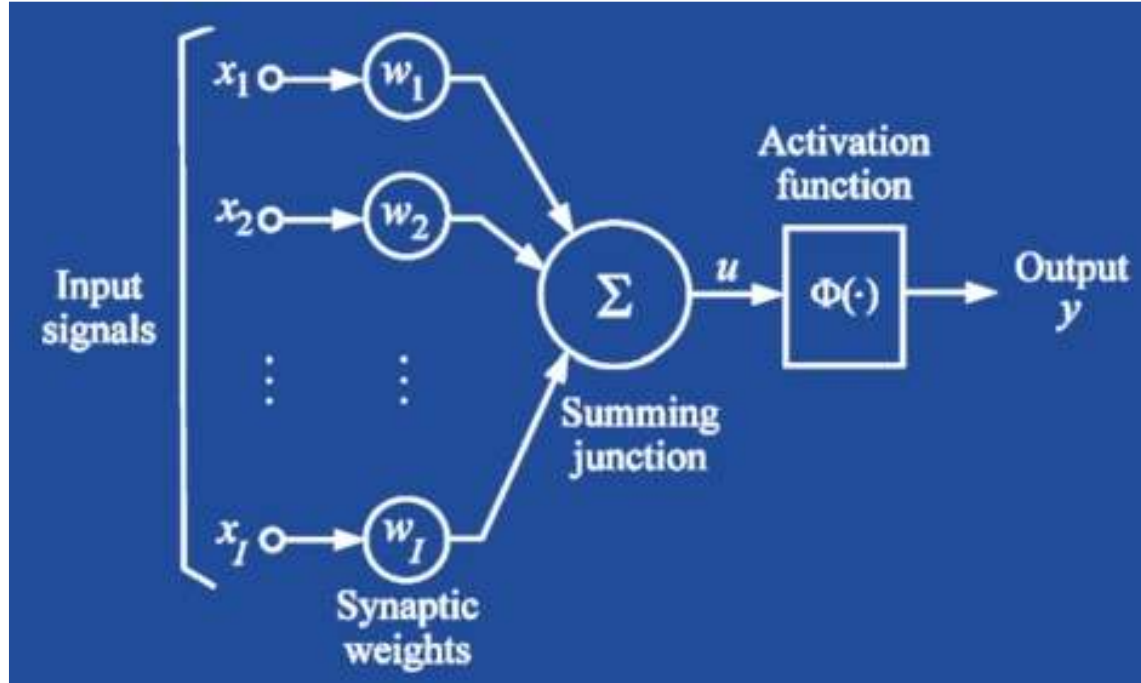
Deep Learning in action

Selected customers running AI on AWS



Basics of Deep Learning

The neuron



$$\sum_{i=1}^I x_i * w_i = u$$

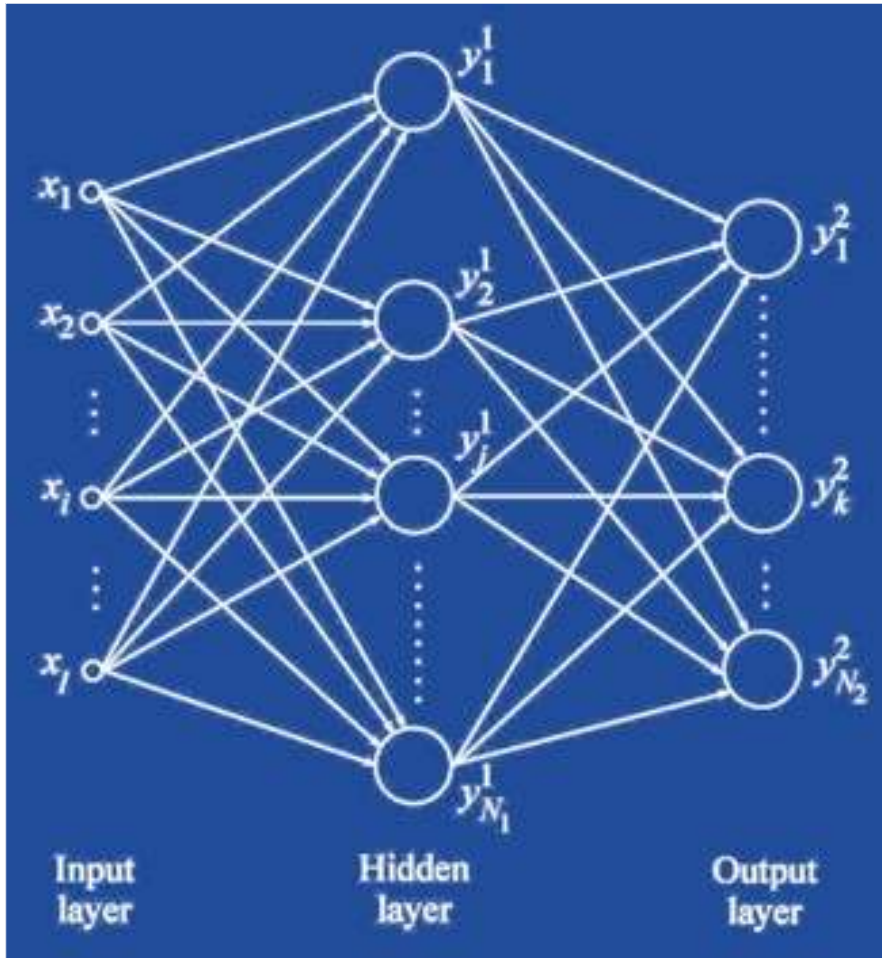
”Multiply and Accumulate”

Activation functions

Name	Plot	Equation
Identity		$f(x) = x$
Binary step		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Logistic (a.k.a. Soft step)		$f(x) = \frac{1}{1 + e^{-x}}$
TanH		$f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$
ArcTan		$f(x) = \tan^{-1}(x)$
Softsign [7][8]		$f(x) = \frac{x}{1 + x }$
Rectified linear unit (ReLU) [9]		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$

Source: Wikipedia

Neural networks



$$X = \begin{bmatrix} X_{11}, X_{12}, \dots, X_{1l} \\ X_{21}, X_{22}, \dots, X_{2l} \\ \dots \dots \dots \\ X_{m1}, X_{m2}, \dots, X_{ml} \end{bmatrix}$$

l features

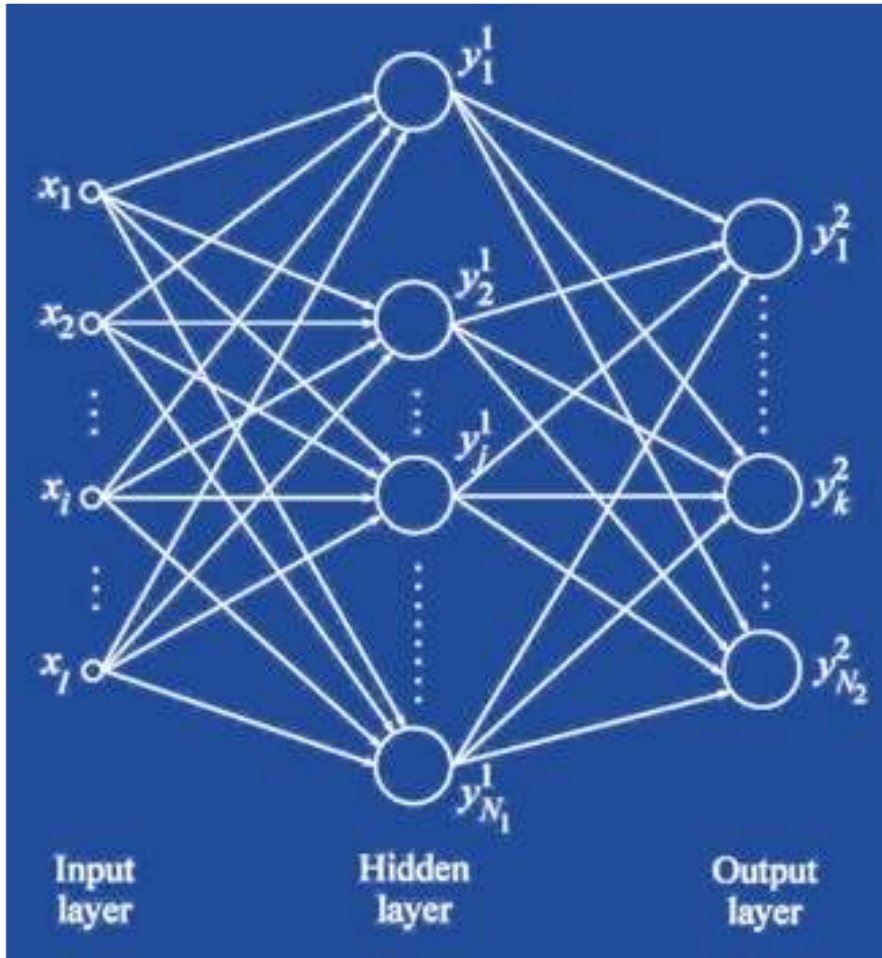
m samples

$$y = \begin{bmatrix} y_1 \\ y_2 \\ \dots \\ y_m \end{bmatrix} = \begin{bmatrix} 0,0,1,0,0,\dots,0 \\ 1,0,0,0,0,\dots,0 \\ \dots \\ 0,0,0,0,1,\dots,0 \end{bmatrix}$$

m labels,
 N_2 outputs

One-hot encoding

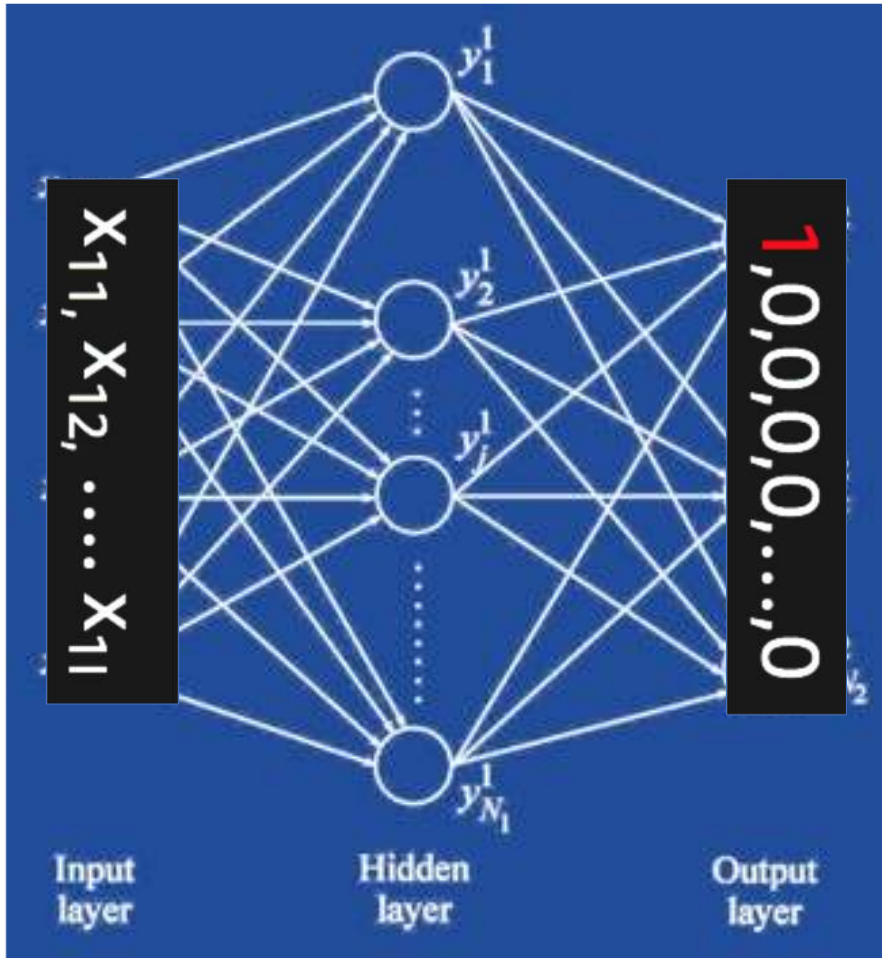
Neural networks



$$\begin{aligned}
 \mathbf{X} &= \begin{bmatrix} x_{11}, x_{12}, \dots, x_{1I} \\ x_{21}, x_{22}, \dots, x_{2I} \\ \vdots \\ x_{m1}, x_{m2}, \dots, x_{mI} \end{bmatrix} \quad \begin{matrix} I \text{ features} \\ m \text{ samples} \end{matrix} \\
 \mathbf{y} &= \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{bmatrix} = \begin{bmatrix} 0, 0, 1, \dots, 0 \\ 1, 0, 0, \dots, 0 \\ \vdots \\ 0, 0, 0, \dots, 0 \end{bmatrix} \quad \begin{matrix} m \text{ labels,} \\ N_2 \text{ categories} \end{matrix} \\
 &\quad \text{One-hot encoding}
 \end{aligned}$$

$$\text{Accuracy} = \frac{\text{Number of correct predictions}}{\text{Total number of predictions}}$$

Neural networks



Initially, the network will **not** predict correctly

$$f(X_1) = Y'_1$$

A **loss function** measures the difference between the **real label** Y_1 and the **predicted label** Y'_1

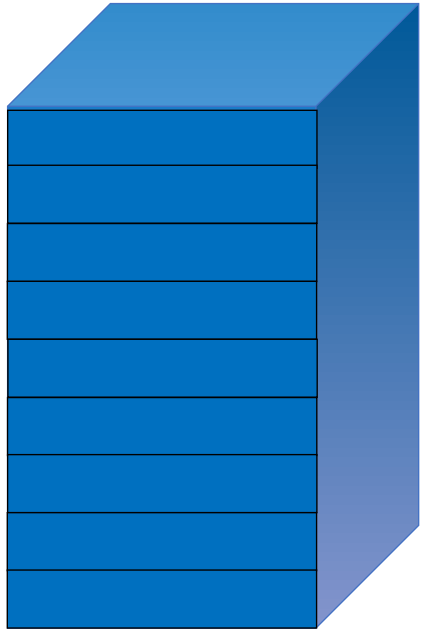
$$\text{error} = \text{loss}(Y_1, Y'_1)$$

For a **batch** of samples:

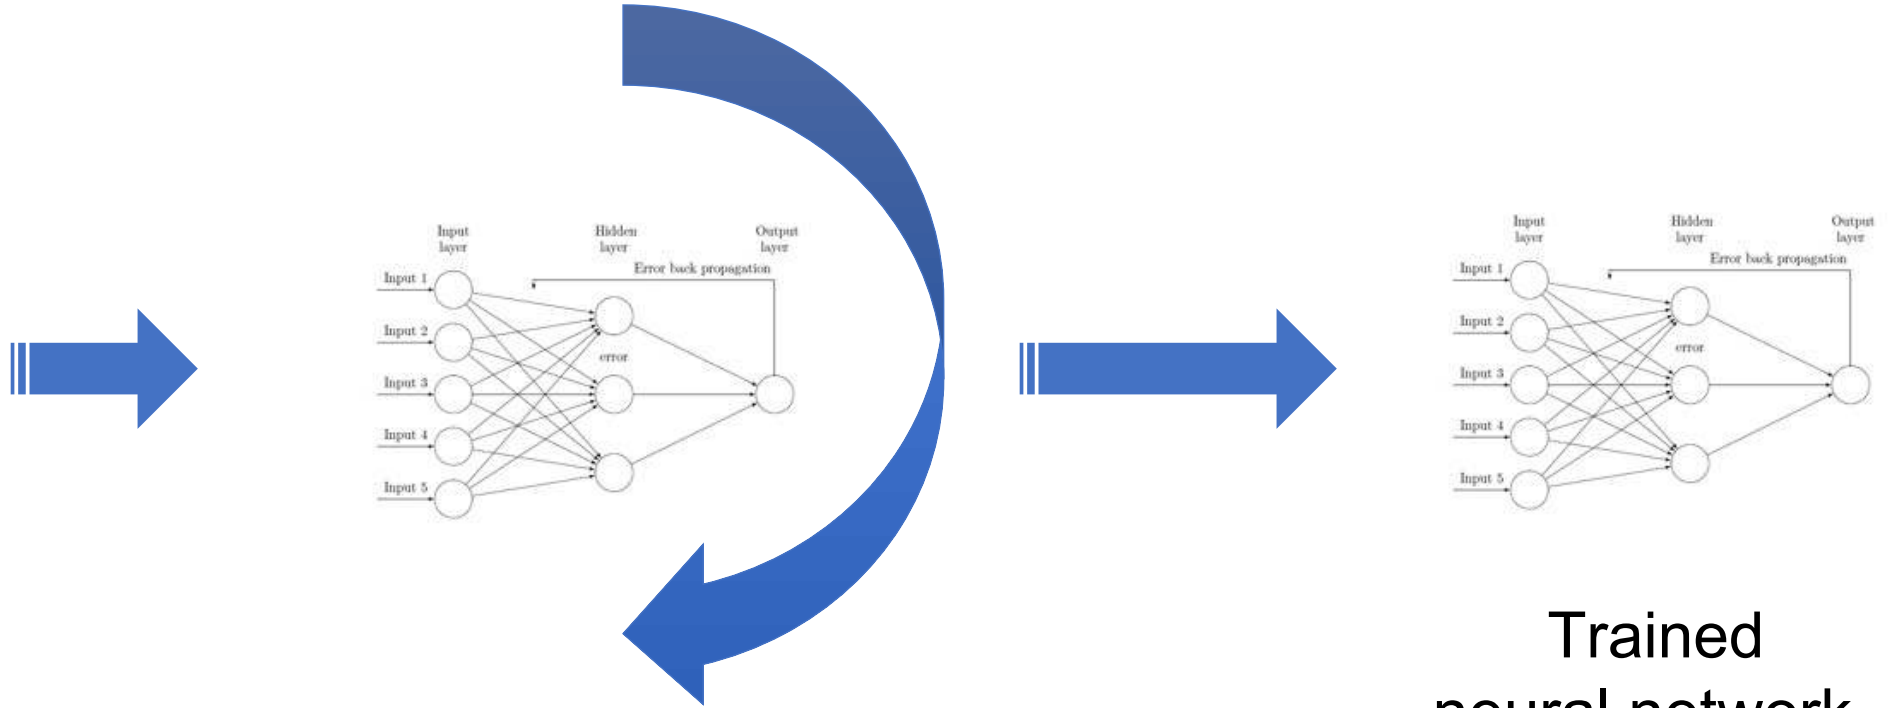
$$\sum_{i=1}^{\text{batch size}} \text{loss}(Y_i, Y'_i) = \text{batch error}$$

The purpose of the training process is to **minimize error** by gradually **adjusting weights**

Training



Training data set



Backpropagation

Batch size
Learning rate
Number of epochs

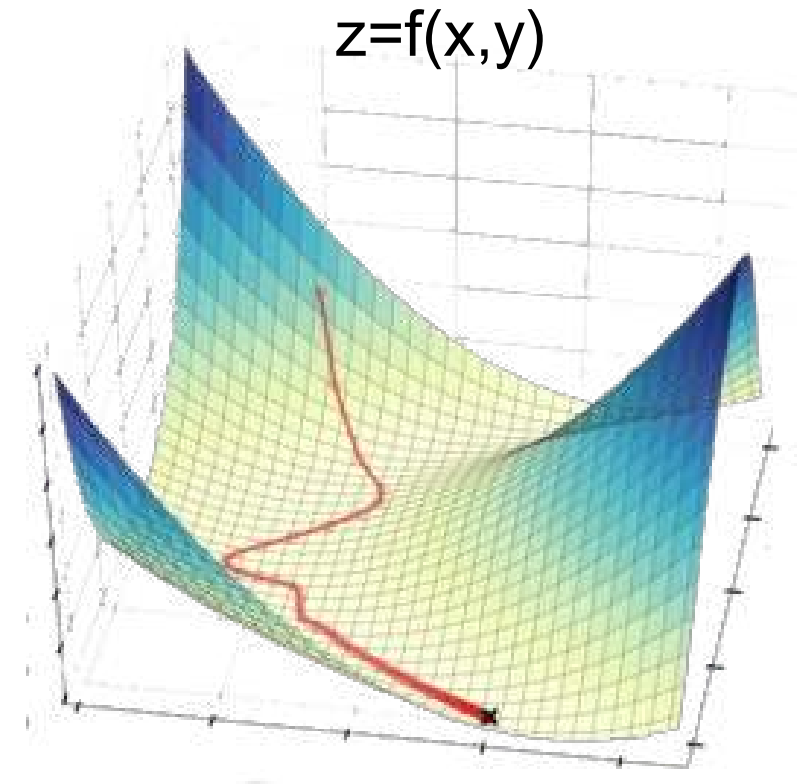
} Hyper parameters

Trained
neural network

Stochastic Gradient Descent (SGD)

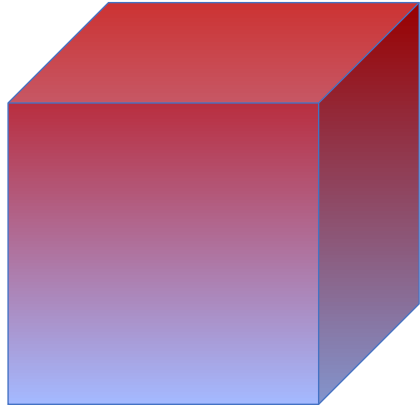
Imagine you stand on top of a mountain with skis strapped to your feet. You want to get down to the valley as quickly as possible, but there is fog and you can only see your immediate surroundings. How can you get down the mountain as quickly as possible? You look around and identify the steepest path down, go down that path for a bit, again look around and find the new steepest path, go down that path, and repeat—this is exactly what gradient descent does.

Tim Dettmers
University of Lugano
2015

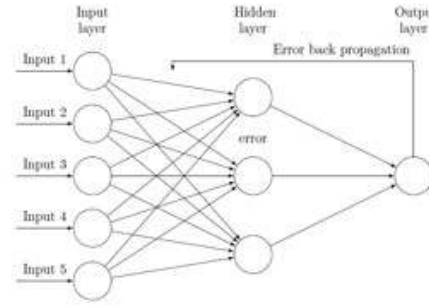


The « step size » is called the learning rate

Validation



Validation data set



Trained
neural network

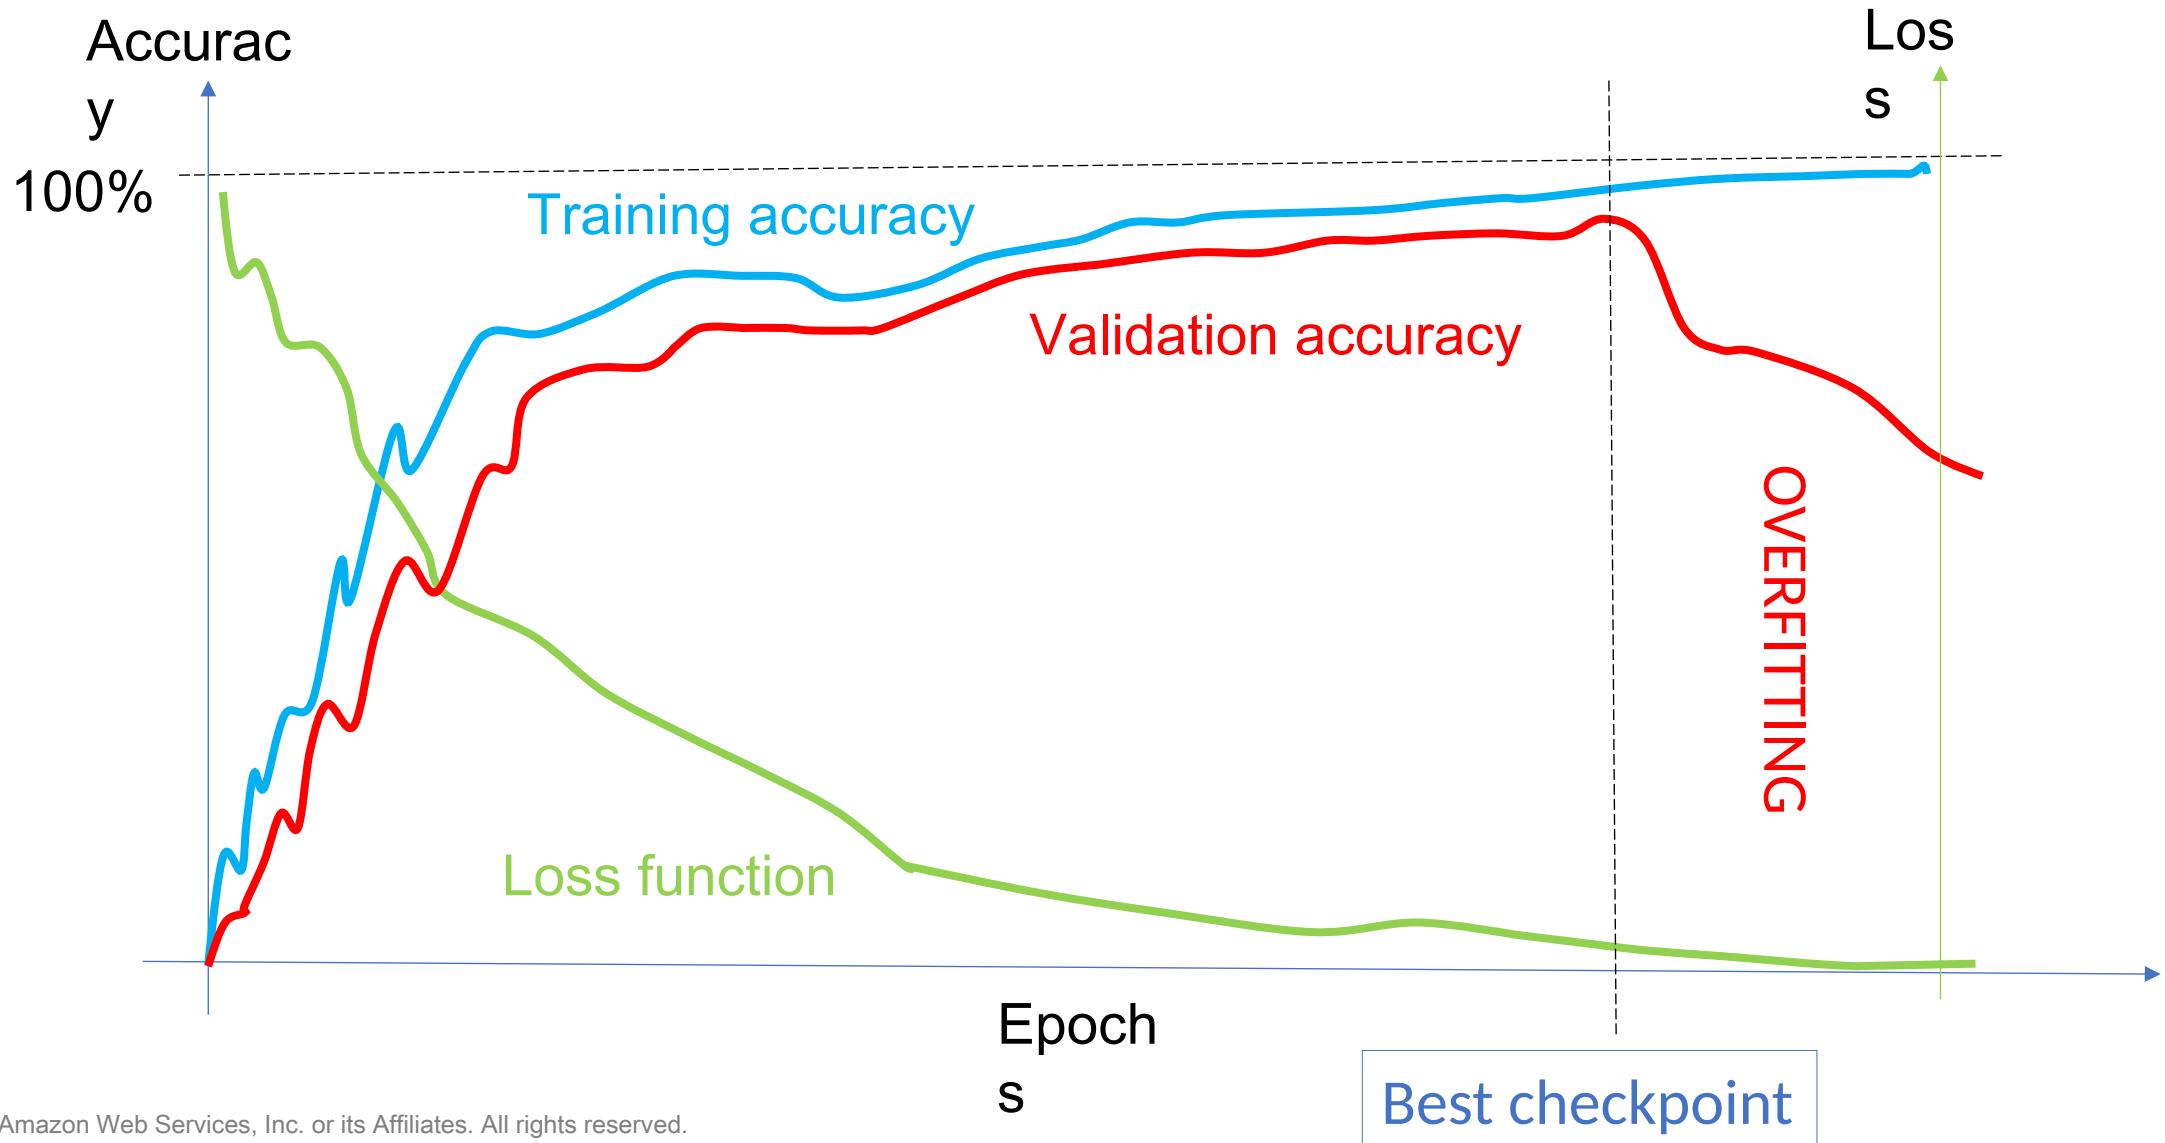


Validation
accuracy

Prediction at
the end of
each epoch

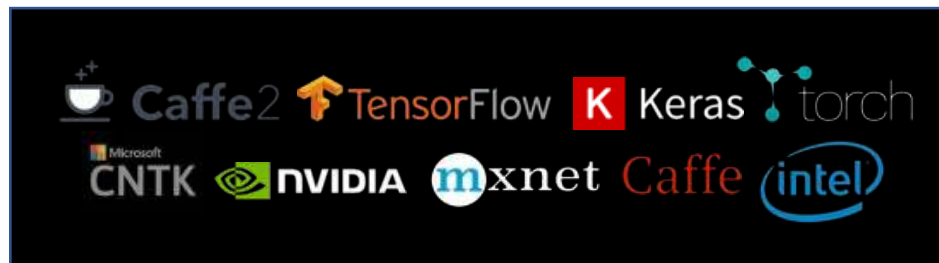
Save the model at the end of each epoch

Early stopping



Deep Learning in practice

- One-click launch
- Single node or distributed
- CPU, GPU, FPGA
- NVIDIA & Intel libraries
- Anaconda Data Science Platform
- Python w/ AI/ML/DL libraries



NVIDIA Volta and AWS P3

Resources

<https://aws.amazon.com/machine-learning>

<https://aws.amazon.com/blogs/ai>

<https://www.nvidia.com/en-us/deep-learning-ai/>

<https://www.nvidia.fr/dli>

<https://www.nvidia.fr/data-center/volta-gpu-architecture/>

<https://aws.amazon.com/ec2/instance-types/p3/>

<https://medium.com/@julsimon>



Thank you!

**Julien Simon, AI Evangelist,
EMEA
@julsimon**

