



# **Amazon Elastic Map Reduce: the concepts**

**Julien Simon, AI Evangelist, EMEA**  
**@julsimon**

# What to expect

- Amazon EMR
- Amazon S3 as HDFS
- Core node and Task nodes
- Elastic clusters
- Beyond Map Reduce: Spark
- Q&A

# What Hadoop is good at

Semi-structured/unstructured data


Disparate Data Sets

ETL at scale

Batch Analytics

Log Processing & Aggregation

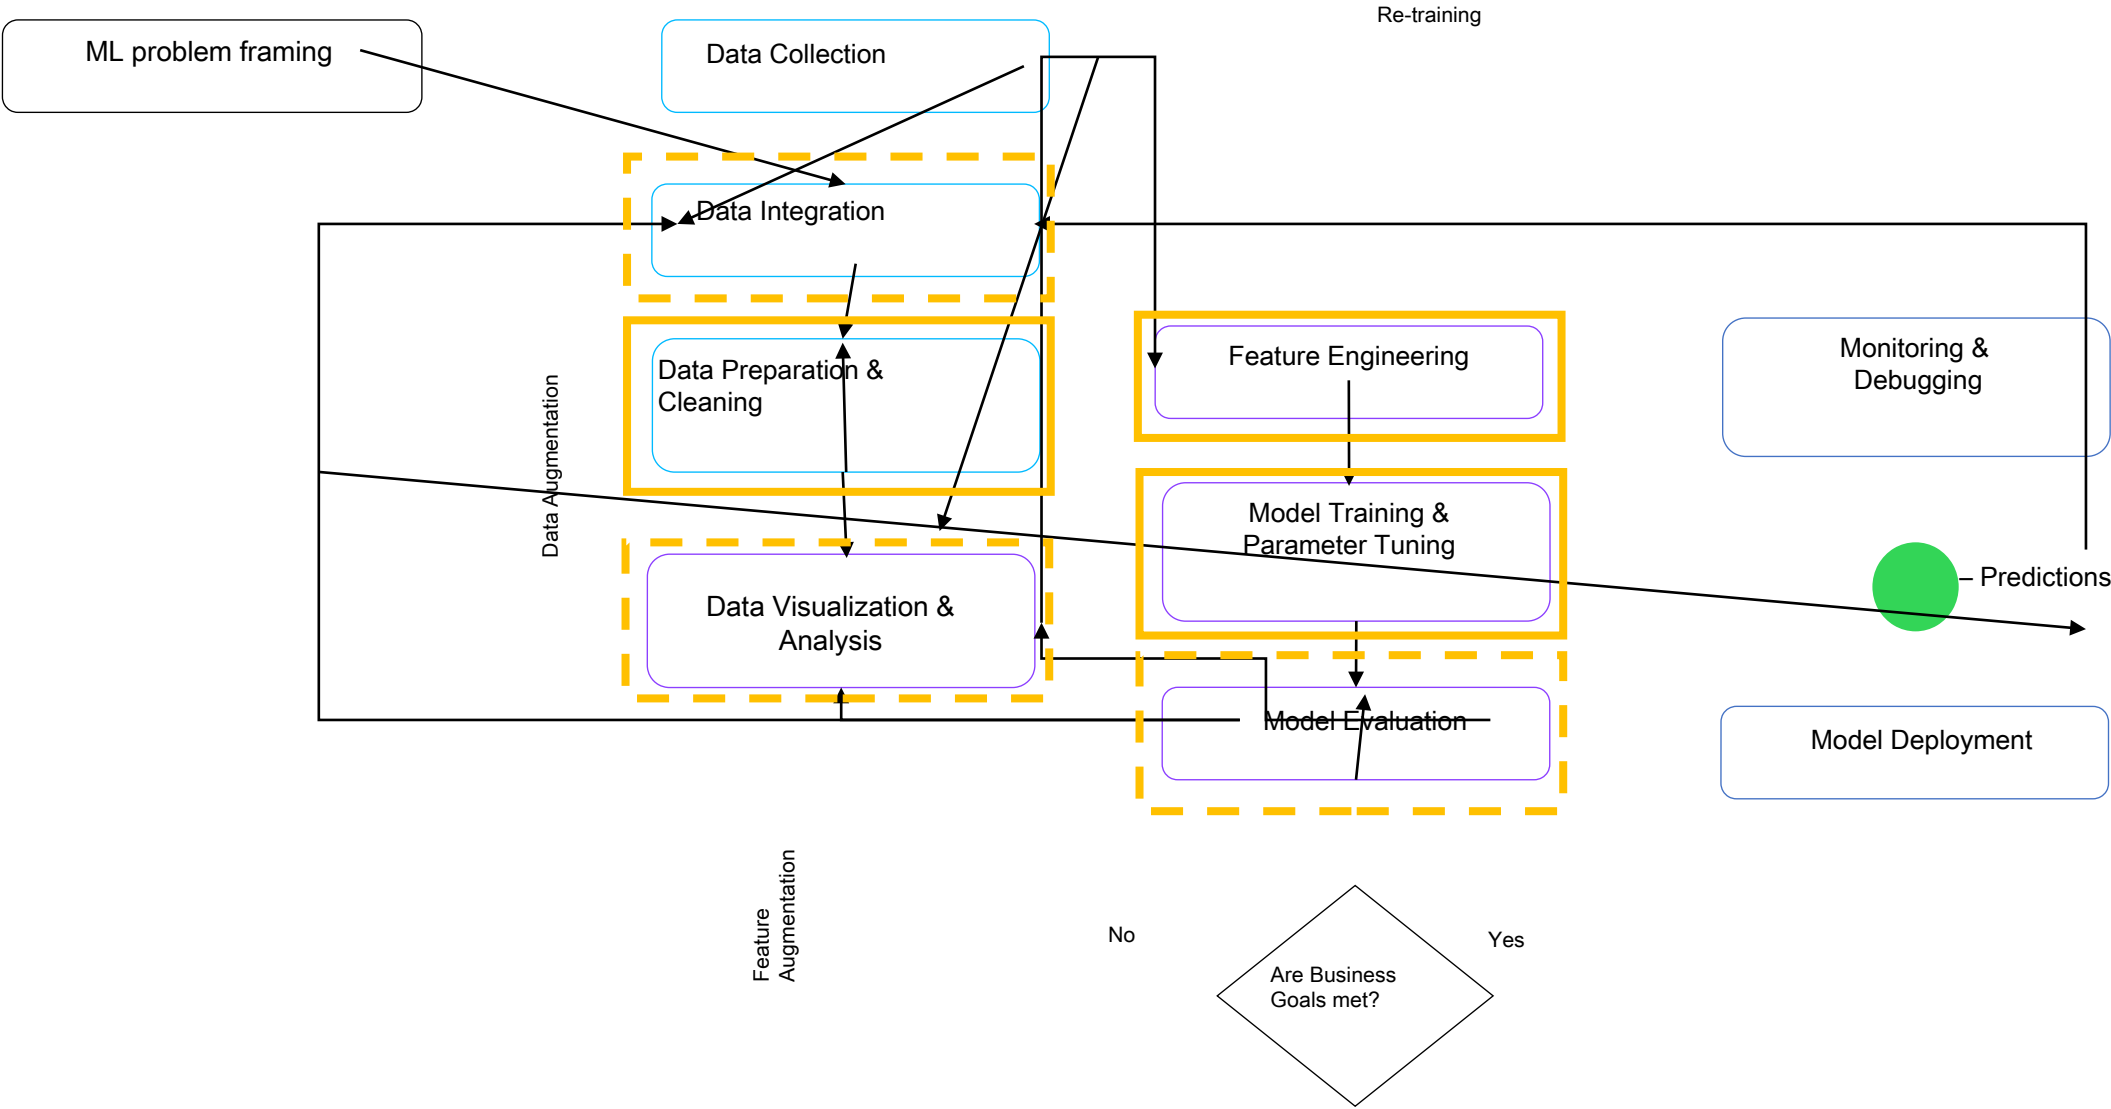
# Amazon EMR – Hadoop, Spark, Presto in the Cloud



- Managed platform
- Launch a cluster in minutes
- Leverage the elasticity of the cloud
- Baked in security features
- Pay by the hour and save with Spot
- Flexibility to customize




# Scope of Amazon EMR



# Amazon S3 as HDFS


# Amazon S3 as HDFS

- Use Amazon S3 as your **permanent** data store
- HDFS for **temporary** storage data between jobs
- No additional step to copy data to HDFS




# Understanding Decoupled Storage & Compute

## Old Clustering / Localized Model



**HDFS = 3X Replication**  
**500 TB Dataset = 1.5 PB cluster with replication**

## Amazon EMR Decoupled Model




**S3 as Streaming HDFS via EMRFS**




# Understanding Decoupled Storage & Compute

## Hadoop on EC2



- 3X HDFS Replication of Data across Nodes
- Protected against data loss from node failure
- Single AZ replication
- Paying for storage units via EC2 ephemeral drives or EBS volumes
- Data stored locally so cluster = 24x7 by default


## EMR




- Native S3 multi-AZ replication
- Protected against data loss from node failure, cluster failure or AZ failure
- Multiple physical facility replication
- Paying for storage units via S3 (cheap!!!)
- Spin clusters up and down or turn off!
- Stream multiple EMR clusters from same dataset in S3

# Amazon S3 as your persistent data store


- Designed for 99.999999999% durability
- Separate compute and storage
- Resize and shut down clusters with no data loss
- Point multiple clusters at the same data in S3
- Easily evolve your analytics infrastructure as technology evolves



# EMR with Amazon S3




# Disaster Recovery built in



# Benefits: Amazon S3 as HDFS

- No need to scale HDFS
  - Capacity?
  - Replication?
- Amazon S3 scales with your data
  - Both in **IOPS** and **storage**
  - Your data can grow **independent** of CPU and RAM



# Benefits: Amazon S3 as HDFS


- Take advantage of **S3 features**
  - Server-side encryption
  - Lifecycle policies
  - Versioning to protect against corruption
- Build **elastic** clusters
  - Add nodes to read from Amazon S3
  - Remove nodes with data safe on Amazon S3

Core nodes and task nodes

# Amazon EMR Core nodes

Run TaskTrackers  
(Compute)

Run DataNode  
(HDFS)






# Amazon EMR Core nodes

Can add core nodes


More HDFS space

More CPU/memory



# Amazon EMR Core nodes

Can't remove core nodes because of HDFS




# Amazon EMR Task nodes

Run TaskTrackers


No HDFS

Reads from core nodes



# Amazon EMR Task nodes

Can add task nodes




# Amazon EMR Task nodes

More CPU power


More memory

Jobs done faster



# Amazon EMR Task nodes


You can remove task nodes when processing is completed



# Elastic clusters

# Elastic Clusters


1. Start cluster with certain number of nodes






# Elastic Clusters

## 2. Set Auto-Scaling policies in CloudWatch



# Elastic Clusters

3. Leverage EMR Auto-Scaling to add nodes or gracefully remove nodes based on CloudWatch policies



# Scale up with Spot instances



10 node cluster running for 14 hours

$$\text{Cost} = \$1 * 10 * 14 = \$140$$

# Resize nodes with Spot instances



Add 10 more nodes on Spot

# Resize nodes with Spot instances



**20 node cluster running for 7 hours**

$$\begin{aligned}\text{Cost} &= \$1 * 10 * 7 && \$70 \\ &+ \$0.5 * 10 * 7 && \$35\end{aligned}$$

**Total \$105**

# Resize nodes with Spot instances



**50% less run-time: 14 → 7  
hours**


**25% less cost: \$140 → \$105**

# Instance fleets for advanced Spot provisioning


**Master Node**



**Core Instance Fleet**




**Task Instance Fleet**



- Provision from a list of instance types with Spot and On-Demand
- Launch in the most optimal Availability Zone based on capacity/price
- Spot Block support (Defined Duration)

# EMR Enables Transient Clusters

- Cluster lives for the **duration of the job**
- Shut down the cluster when the job is **done**
- Input and output data **persists** on Amazon S3




Data on Amazon S3



# Options to submit jobs

Submit a  
application



Amazon EMR  
Step API

Use AWS Lambda to  
submit applications to  
EMR Step API or directly  
on your cluster



AWS Lambda



Create a pipeline  
to schedule job  
submission or create  
complex workflows



AWS Data Pipeline



Airflow, Luigi, or other  
schedulers on EC2

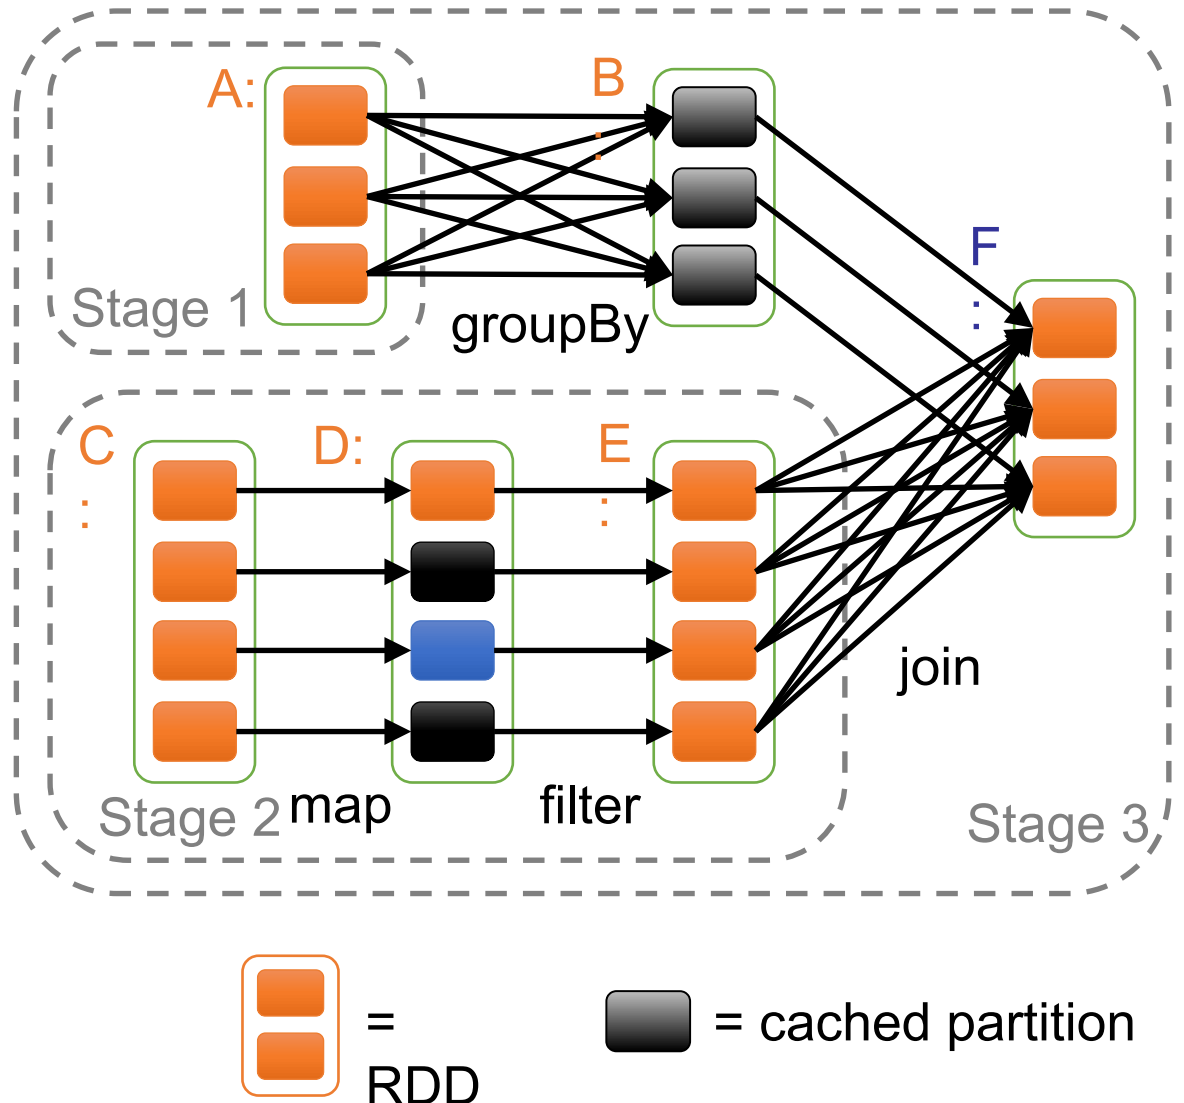


**Amazon  
EMR**  
Use Ecosystem on your  
cluster to build  
DAGs of jobs

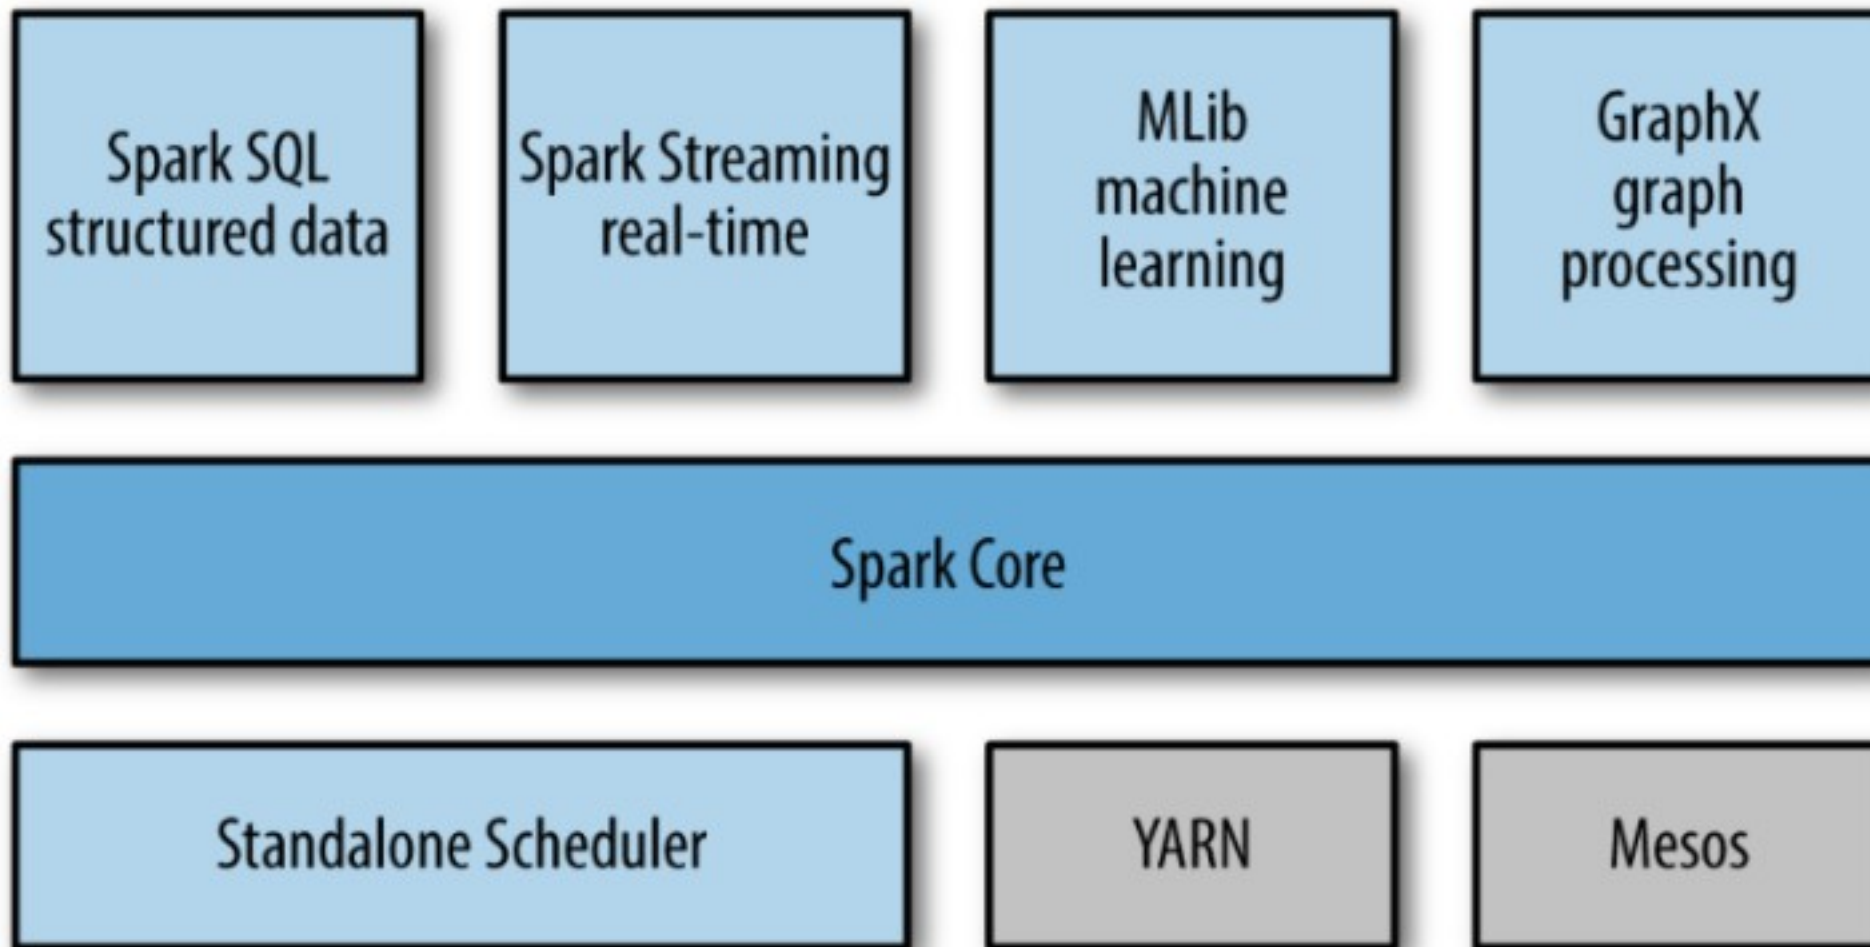
# Beyond Map Reduce: Spark

# Spark moves at interactive speed

- Massively **parallel**
- Uses **DAGs** instead of map-reduce for execution
- Minimizes I/O by storing data in **memory**
- Partitioning-aware to avoid network-intensive shuffle



# Spark components to match your use case



# Spark speaks your language



python



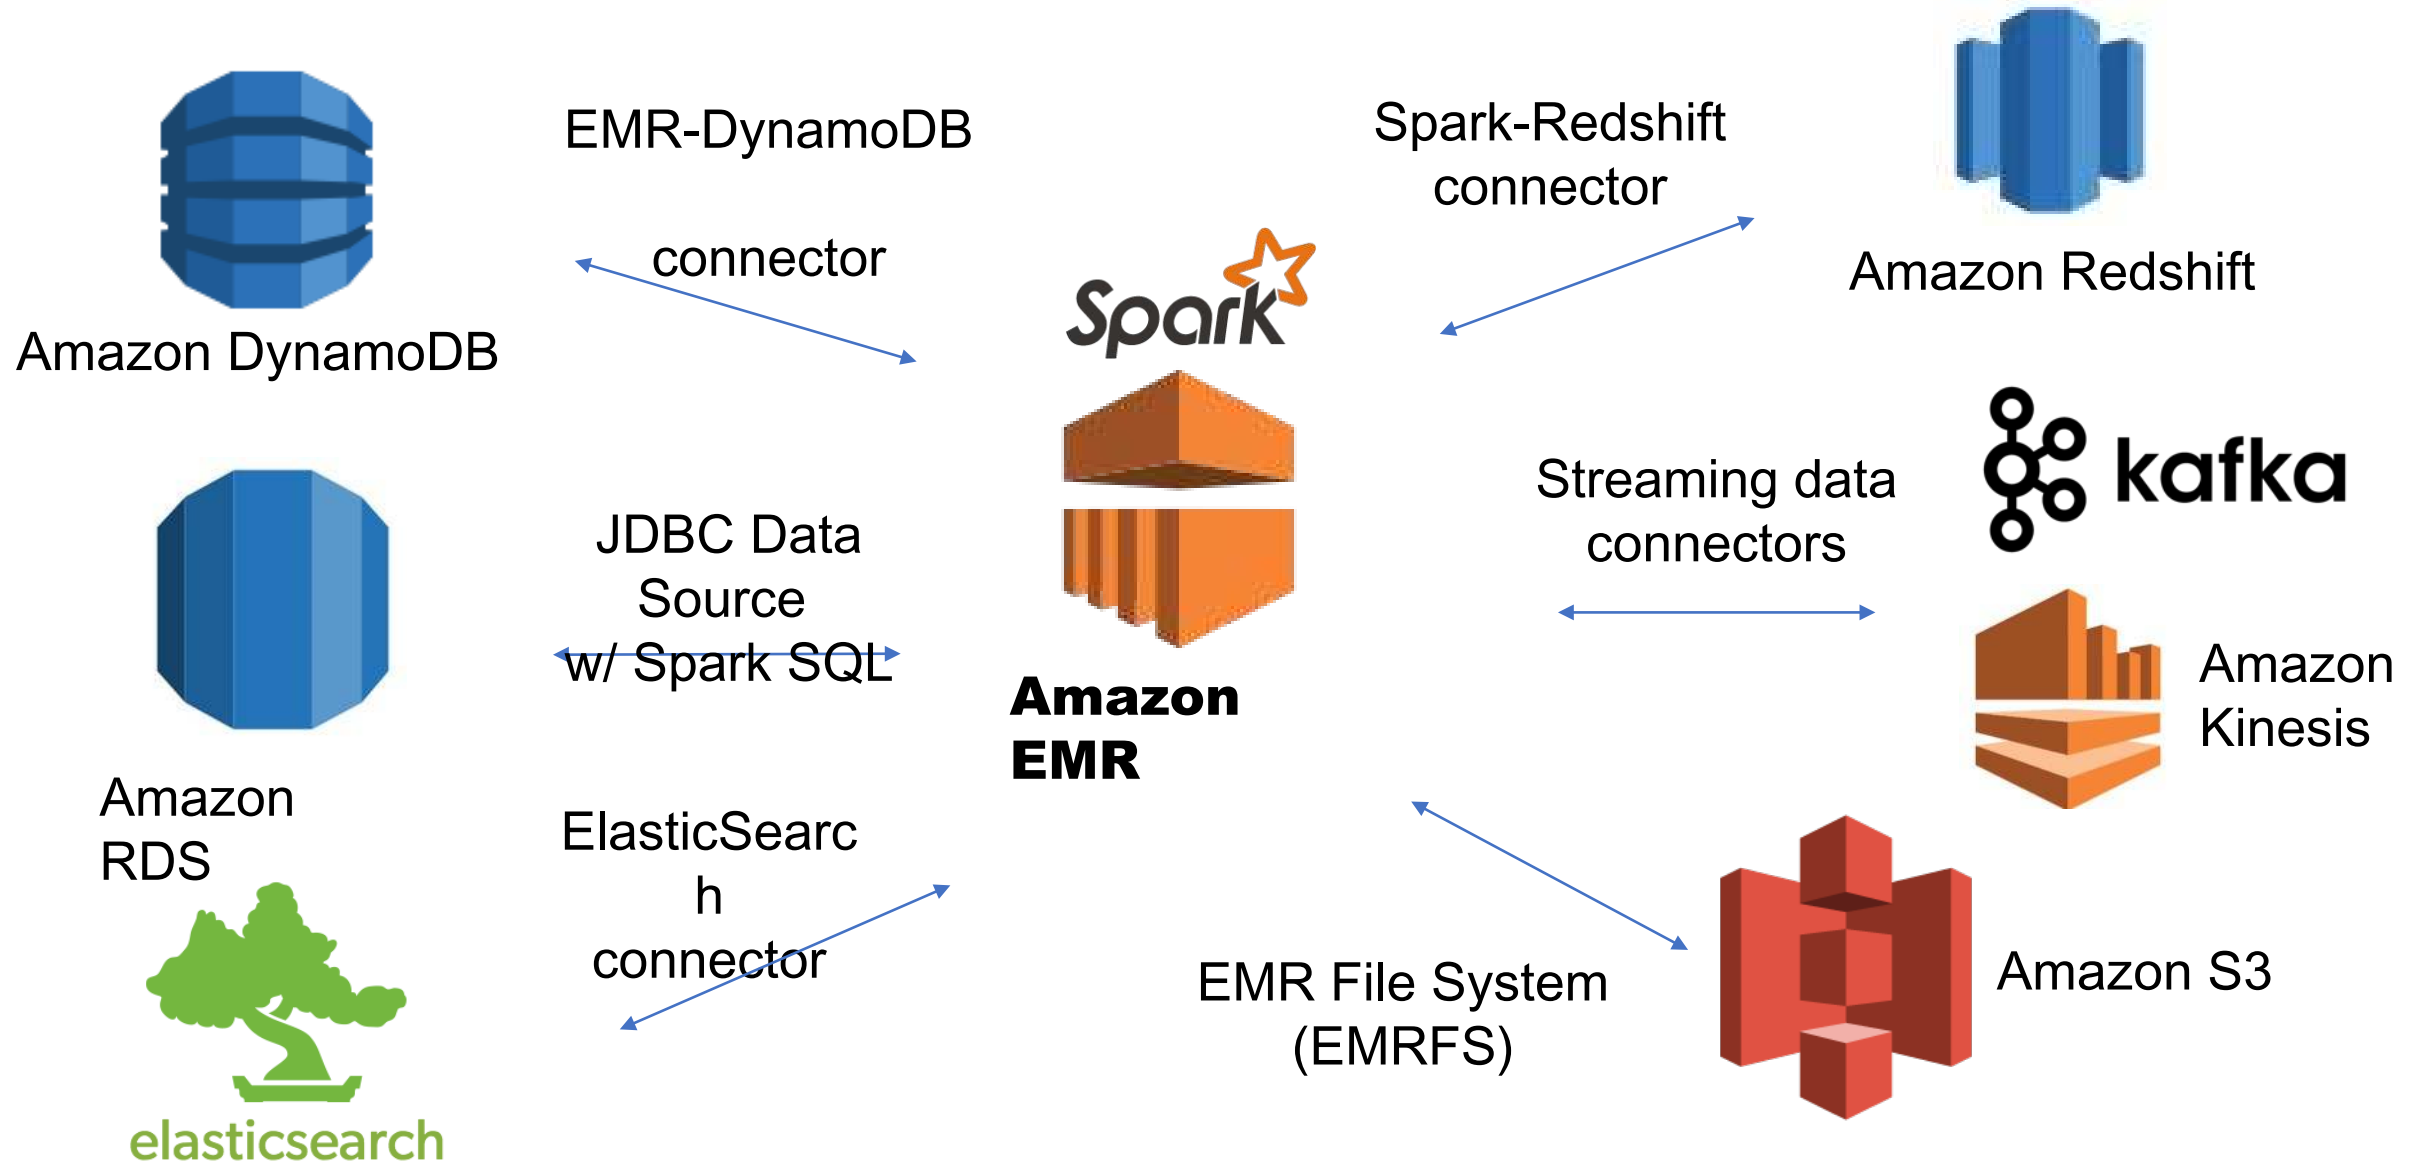
Java™



SQL



# Many storage layers to choose from



# Use DataFrames for machine learning

```
// Prepare training documents from a list of (id, text, label) tuples
val training = sqlContext.createDataFrame(Seq(
  (0L, "a b c d e spark", 1.0),
  (1L, "b d", 0.0),
  (2L, "spark f g h", 1.0),
  (3L, "hadoop mapreduce", 0.0)
)).toDF("id", "text", "label")

// Configure an ML pipeline, which consists of three stages: tokenizer, hashingTF, and lr.
val tokenizer = new Tokenizer()
  .setInputCol("text")
  .setOutputCol("words")
val hashingTF = new HashingTF()
  .setNumFeatures(1000)
  .setInputCol(tokenizer.getOutputCol)
  .setOutputCol("features")
val lr = new LogisticRegression()
  .setMaxIter(10)
  .setRegParam(0.01)
val pipeline = new Pipeline()
  .setStages(Array(tokenizer, hashingTF, lr))

// Fit the pipeline to the training data
val model = pipeline.fit(training)
```


- Spark ML (replacing MLlib) uses DataFrames as input/output for models
- Create ML pipelines with a variety of distributed algorithms

# Create DataFrames on streaming data





# Amazon EMR



## **Easy to Use**

Launch a cluster in minutes



## **Low Cost**

Pay an hourly rate



## **Elastic**

Easily add or remove capacity



## **Reliable**

Spend less time monitoring



## **Secure**

Manage firewalls



## **Flexible**

Customize the cluster

# Resources

<https://aws.amazon.com/emr/>

<https://aws.amazon.com/blogs/bigdata>

<https://aws.amazon.com/whitepapers/>

<https://medium.com/@julsimon>



**Thank you!**

**Julien Simon, AI Evangelist,  
EMEA  
@julsimon**