




Processing images with Deep Learning

**Julien Simon, AI Evangelist,
EMEA
@julsimon**



What to expect

- Amazon Rekognition or Apache MXNet?
- Github projects for image processing with Apache MXNet
- A deeper look at the Convolution operation
- Demos
- Q&A

Apache MXNet: Open Source library for Deep Learning




Programmable

Simple syntax,
multiple
languages




Most Open

Accepted into the
Apache Incubator




Portable

Highly efficient
models for
mobile
and IoT



High Performance

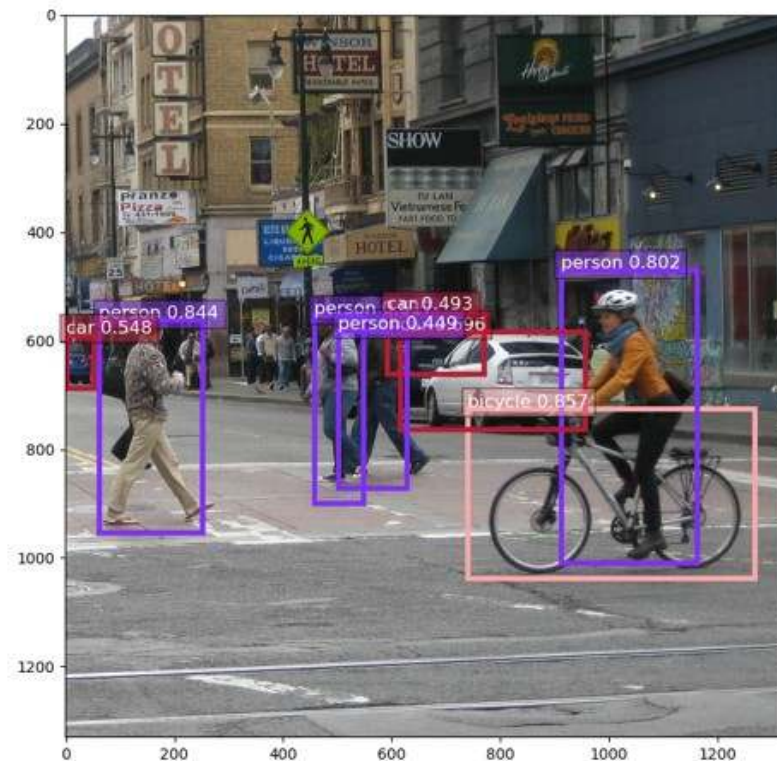


Near linear scaling
across hundreds of
GPUs



Best On AWS

Optimized for
Deep Learning on AWS


Object Detection



<https://github.com/precedenceguo/mx-rcnn>

<https://github.com/zhreshold/mxnet-yolo>

Object Segmentation




<https://github.com/TuSimple/mx-maskrcnn>

Text Detection and Recognition



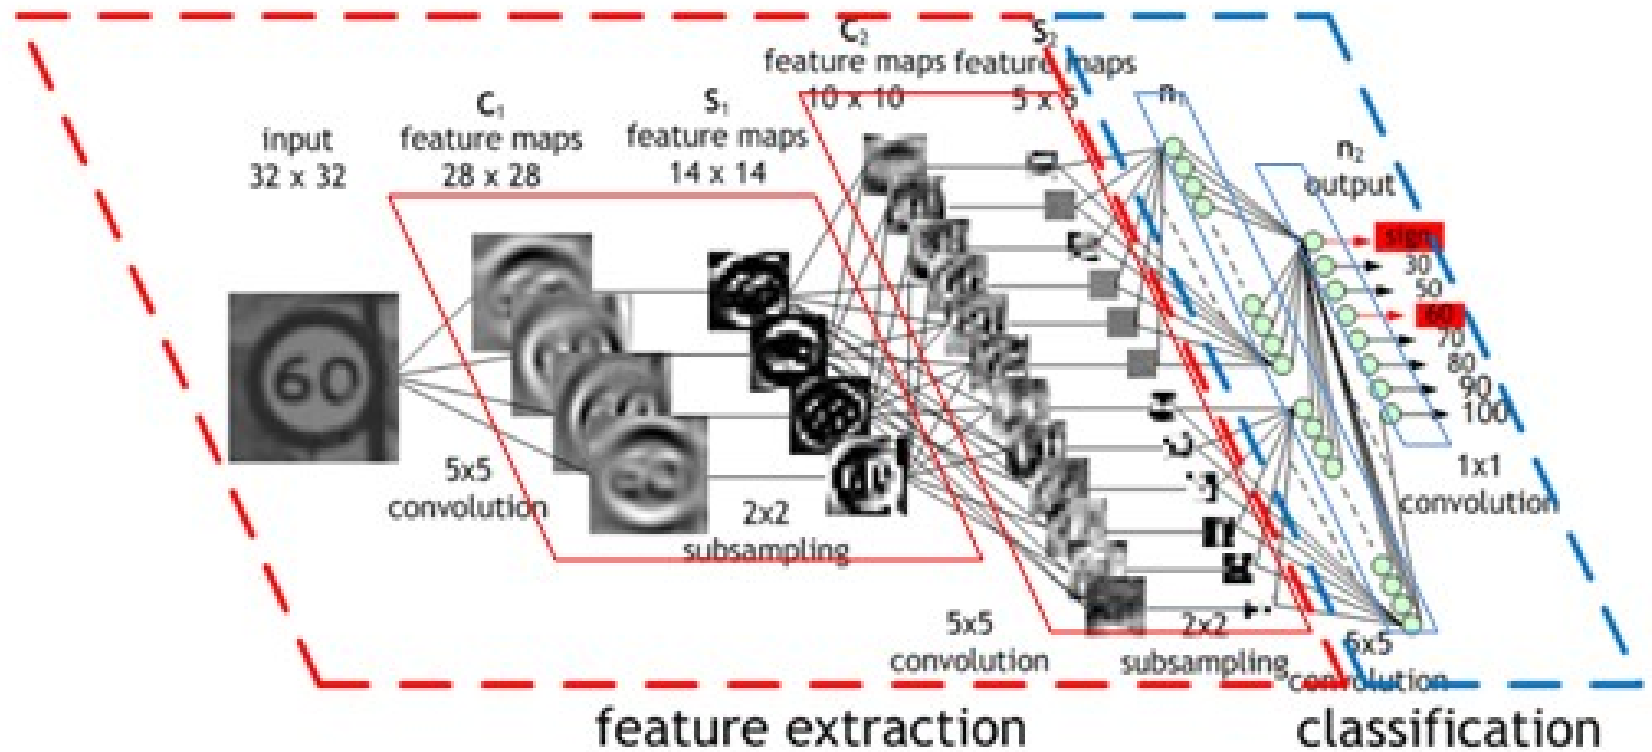

<https://github.com/Bartzi/stn-ocr>

Real-Time Pose Estimation

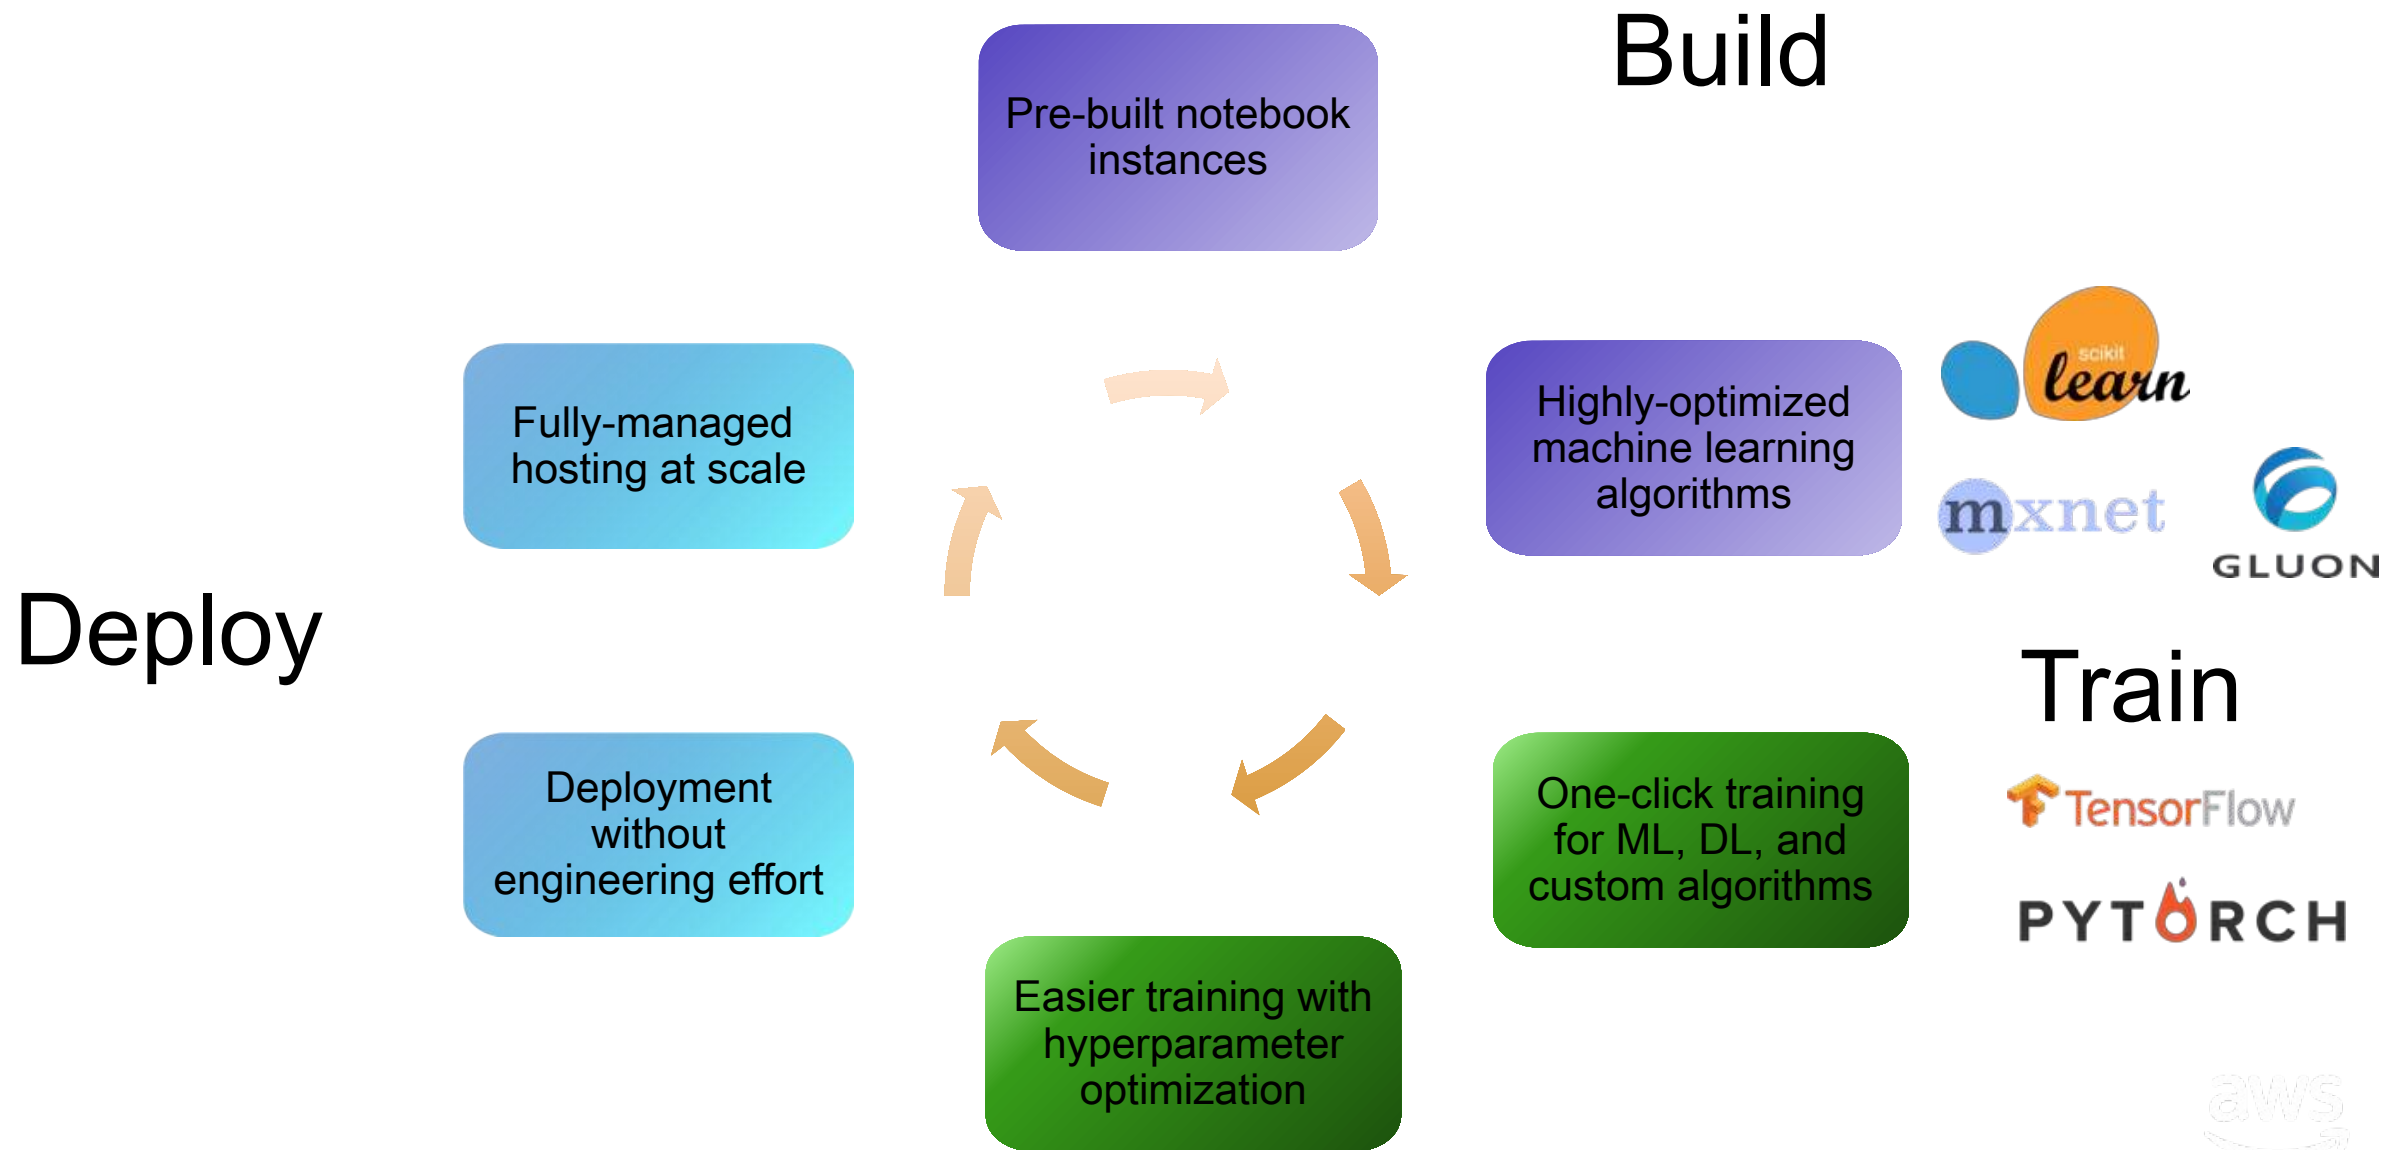


<https://github.com/dragonfly90/mxnet> Realtime Multi-Person Pose Estimation

Convolutional Neural Networks



Amazon SageMaker



Demos

<https://github.com/juliensimon/dlnotebooks>

<https://github.com/guyernest/TensorFlowTutorials>

- 1) Classifying MNIST with a CNN model (Keras)
- 2) Classifying images with pre-trained CNN models (MXNet)
- 3) Fine-tuning a pre-trained CNN model (Keras)
- 4) Generating new MNIST samples with a GAN (MXNet)

Demo #2 – Using a pre-trained model

*** VGG16

```
[(0.46811387, 'n04296562 stage'), (0.24333163, 'n03272010 electric guitar'), (0.045918692, 'n02231487 walking stick, walkingstick, stick insect'), (0.03316205, 'n04286575 spotlight, spot'), (0.021694135, 'n03691459 loudspeaker, speaker, speaker unit, loudspeaker system, speaker system')]
```

*** ResNet-152

```
[(0.8726753, 'n04296562 stage'), (0.046159592, 'n03272010 electric guitar'), (0.041658506, 'n03759954 microphone, mike'), (0.018624334, 'n04286575 spotlight, spot'), (0.0058045341, 'n02676566 acoustic guitar')]
```

*** Inception v3

```
[(0.44991142, 'n04296562 stage'), (0.43065304, 'n03272010 electric guitar'), (0.067580454, 'n04456115 torch'), (0.012423956, 'n02676566 acoustic guitar'), (0.0093934005, 'n03250847 drumstick')]
```



Demo #3 – Image classification: fine-tuning a model

- **CIFAR-10 data set**
 - 60,000 images in 10 classes
 - 32x32 color images
- **Initial training**
 - Resnet-50 CNN
 - 200 epochs
 - 82.12% validation
- **Cars vs. horses**
 - 88.8% validation accuracy

airplane



automobile



bird



cat



deer



dog



frog



horse



ship



truck



Demo #3 – Image classification: fine-tuning a model

- Freezing all layers but the last one
- Fine-tuning on « cars vs. horses » for 10 epochs
- 2 minutes on 1 GPU
- 98.8% validation accuracy

Epoch 10/10

10000/10000 [=====] - 12s

loss: 1.6989 - acc: 0.9994 - val_loss: 1.7490 - val_acc: 0.9880

2000/2000 [=====] - 2s

[1.7490020694732666, 0.9879999999999999]

Resources

<https://aws.amazon.com/machine-learning>

<https://aws.amazon.com/blogs/ai>

<https://mxnet.incubator.apache.org>

<https://github.com/apache/incubator-mxnet>

<https://github.com/gluon-api>

<https://devblogs.nvidia.com/parallelforall/deep-learning-nutshell-core-concepts/>

http://deeplearning.net/software/theano/tutorial/conv_arithmetic.html

<https://medium.com/@julsimon>



Thank you!

**Julien Simon, AI Evangelist,
EMEA
@julsimon**

