



Deep Dive: Amazon RDS

Julien Simon, Principal Technical Evangelist, AWS


julsimon@amazon.fr

@julsimon



What to expect

- Amazon RDS overview (super quick)
- Security
- Metrics and monitoring
- High availability
- Scaling on RDS
- Backups and snapshots
- Migrating to RDS



Amazon Relational Database Service (Amazon RDS)



Launch



No infrastructure
management




Cost-effective




Application
compatibility

AWS Free Tier



Instant provisioning



Scale up/down

Amazon RDS engines

Commercial

ORACLE®



Open source



PostgreSQL



MariaDB

Amazon Aurora

Amazon
Aurora

Selected Amazon RDS customers



vodafone

intuit®



SEGA®



Kempinski
HOTELIERS SINCE 1897



FCBARCELONA
més que un club

OUTBACK
STEAKHOUSE®



Newsweek
& THE DAILY BEAST



P
The
Washington Post
Company

Trinity Mirror plc

coursera



Selected Amazon Aurora customers




Trade-offs with a managed service

Fully managed host and OS

- No access to the database host operating system
- Limited ability to modify configuration that is managed on the host operating system
- No functions that rely on configuration from the host OS

Fully managed storage

- Max storage limits
 - Microsoft SQL Server—4 TB
 - MySQL, MariaDB, PostgreSQL, Oracle—6 TB
 - Aurora—64 TB
- Growing your database is a process



Amazon RDS: the fine print ☺



PostgreSQL

- Using the rds_superuser Role
- Supported PostgreSQL Database Versions
- Supported PostgreSQL Features and Extensions
- Limits for PostgreSQL DB Instances
- Upgrading a PostgreSQL DB Instance
- Using SSL with a PostgreSQL DB Instance
- Creating Roles
- Managing PostgreSQL Database Access
- Working with PostgreSQL Parameters
- Working with PostgreSQL Autovacuum on Amazon RDS
- Audit Logging for a PostgreSQL DB Instance
- Setting up PostGIS
- Using pgBadger for Log Analysis with PostgreSQL
- Viewing the Contents of pg_config

http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/CHAP_PostgreSQL.html#PostgreSQL.Concepts.General.FeatureSupport
<http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Appendix.PostgreSQL.CommonDBATasks.html>



- Killing a Session or Query
- Skipping the Current Replication Error
- Working with InnoDB Tablespaces to Improve Crash Recovery Times
- Managing the Global Status History

<http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Appendix.MySQL.CommonDBATasks.html>

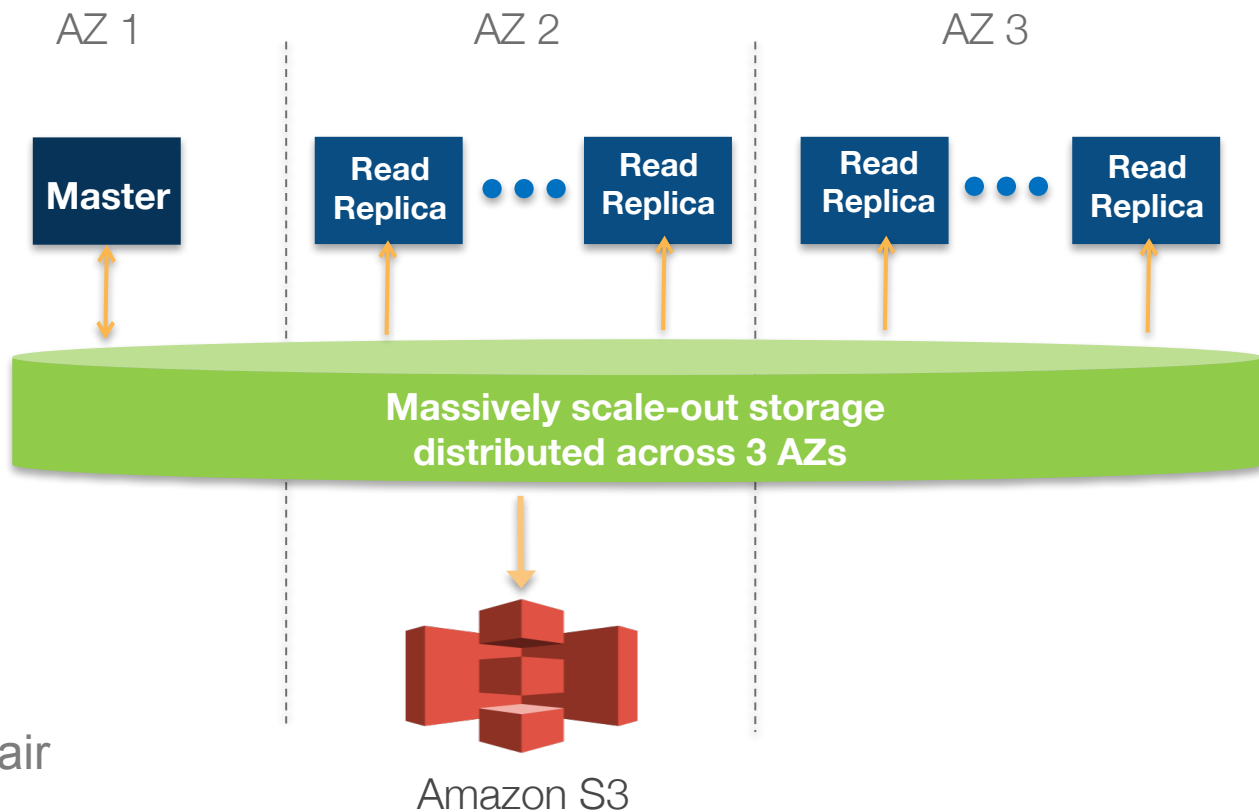


Appendix: Parameters for MariaDB

<http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Appendix.MariaDB.Parameters.html>

Aurora at a glance

- Compatible with **MySQL 5.6**
- Storage **automatically** grows up to 64 TB
- 6 copies** across 3 AZs
- Continuous backup** to Amazon S3
- Continuous monitoring** of nodes and disks for repair

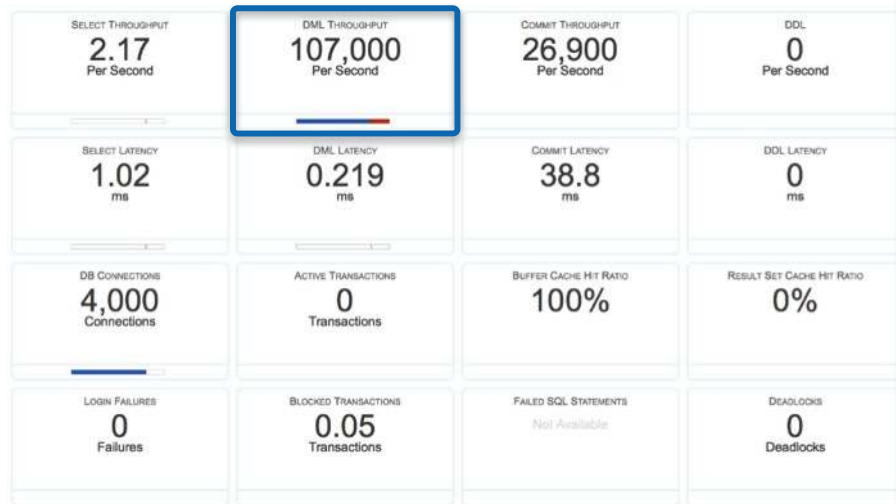


5X faster than MySQL

MySQL SysBench results

R3.8XL: 32 cores / 244 GB RAM

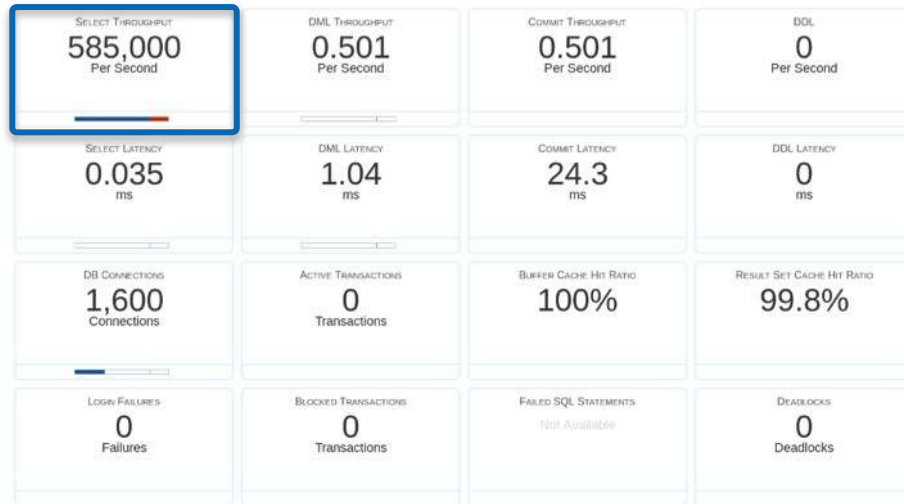
WRITE PERFORMANCE



4 client machines with 1,000 connections each


**Five times higher throughput than stock MySQL,
based on industry standard benchmarks.**

READ PERFORMANCE





Single client machine with 1,600 connections

Security



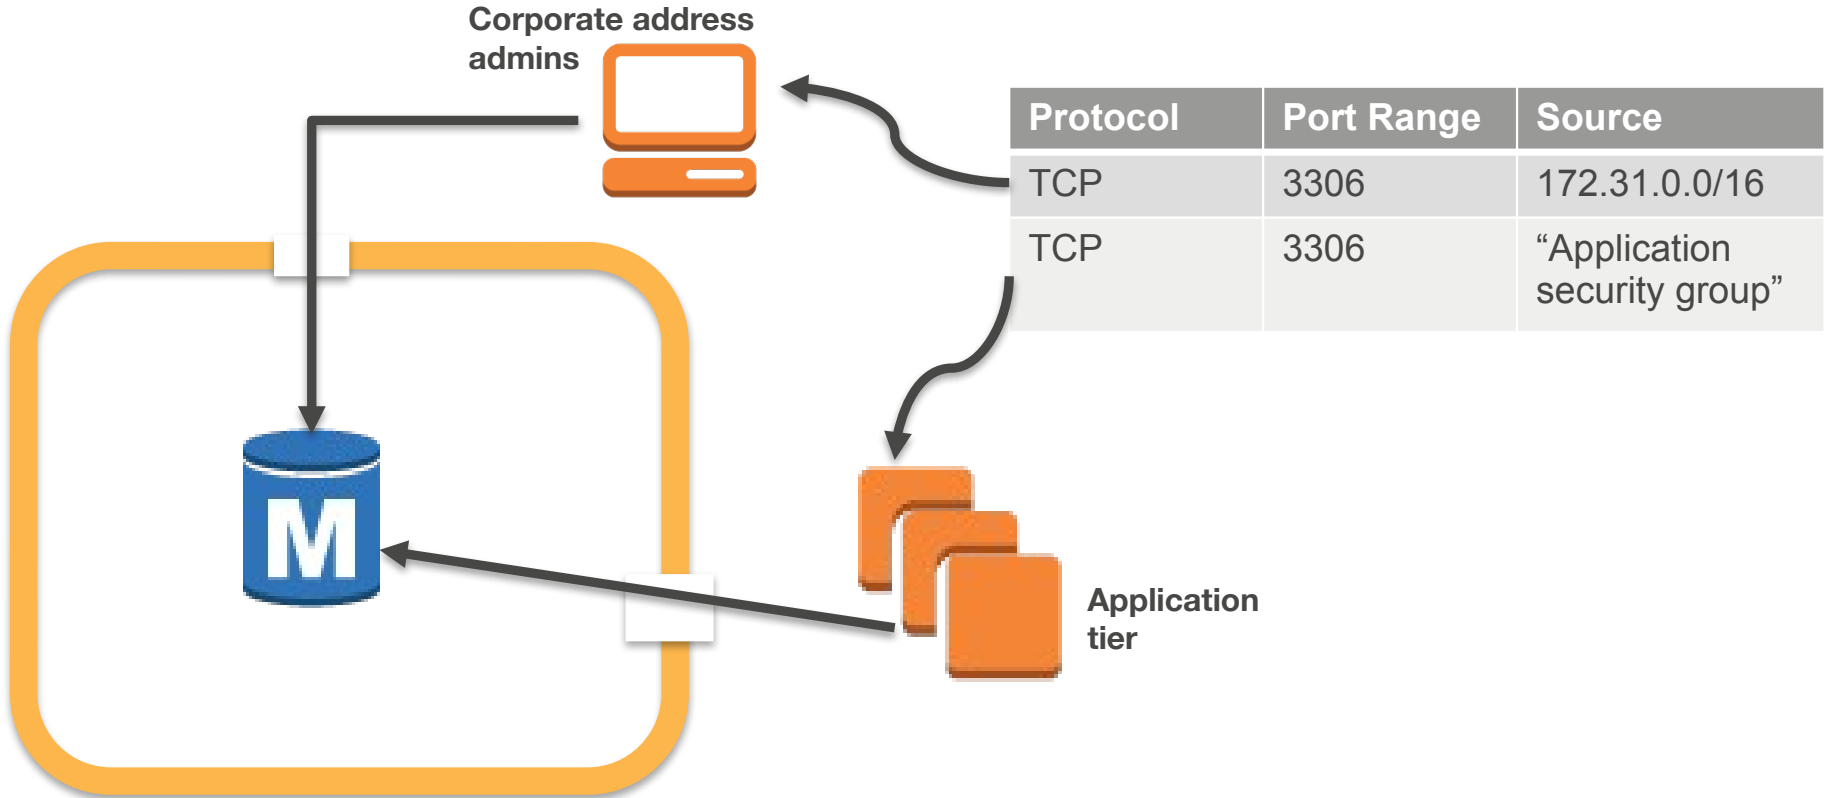
Amazon Virtual Private Cloud (VPC)

Securely control network configuration




Security groups

Database IP firewall protection



Compliance



Singapore MTCS



27001/9001
27017/27018

Compliance


MySQL, Oracle, Postgres

- SOC 1, 2, and 3
- ISO 27001/9001
- ISO 27017/27018
- PCI DSS
- FedRAMP
- HIPAA BAA
- UK government programs
- MTCS (Singapore)
- C5 (Germany)

SQL Server

- SOC 1, 2, and 3
- ISO 27001/9001
- ISO 27017/27018
- PCI DSS
- UK government programs
- MTCS (Singapore)
- C5 (Germany)

In-flight data encryption




SSL available for all six engines

New (Feb'17)
forced SSL for SQL Server

At-rest data encryption

- Based on Amazon KMS
- DB instance storage
- Logs, backups, snapshots
- Read Replicas



- Available for all six engines
- No additional cost
- Support compliance requirements
- TDE also available for Oracle / SQL Server

<http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Overview.Encryption.html>

<https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Appendix.Oracle.Options.AdvSecurity.html>

<https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Appendix.SQLServer.Options.TDE.html>

Amazon RDS encryption hints

- You can only encrypt on new database creation
- Encryption cannot be removed
- Master and Read Replica must be encrypted
 - (Jan'17) you can now replicate encrypted DB across regions
- Unencrypted snapshots can't be restored to encrypted DB
 - Aurora will allow this
 - You can create encrypted copies of your unencrypted snapshots

Enabling encryption with the AWS CLI

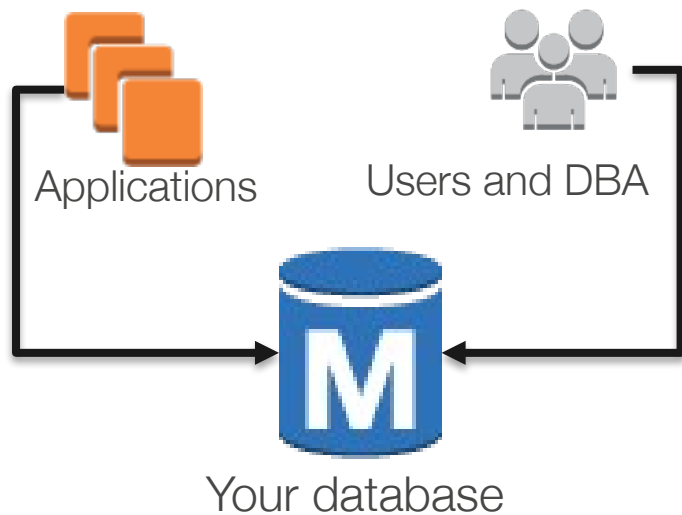
```
aws rds create-db-instance --region us-west-2 --db-instance-identifier sg-cli-test \  
--allocated-storage 20 --storage-encrypted \  
--db-instance-class db.m4.large --engine mysql \  
--master-username myawsuser --master-user-password myawsuser
```

```
aws rds create-db-instance --region us-west-2 --db-instance-identifier sg-cli-test1 \  
--allocated-storage 20 --storage-encrypted --kms-key-id xxxxxxxxxxxxxxxxxxxx \  
--db-instance-class db.m4.large --engine mysql \  
--master-username myawsuser --master-user-password myawsuser
```


IAM governed access

You can use AWS Identity and Access Management (IAM) to control who can perform actions on RDS


Controlled with database GRANTS



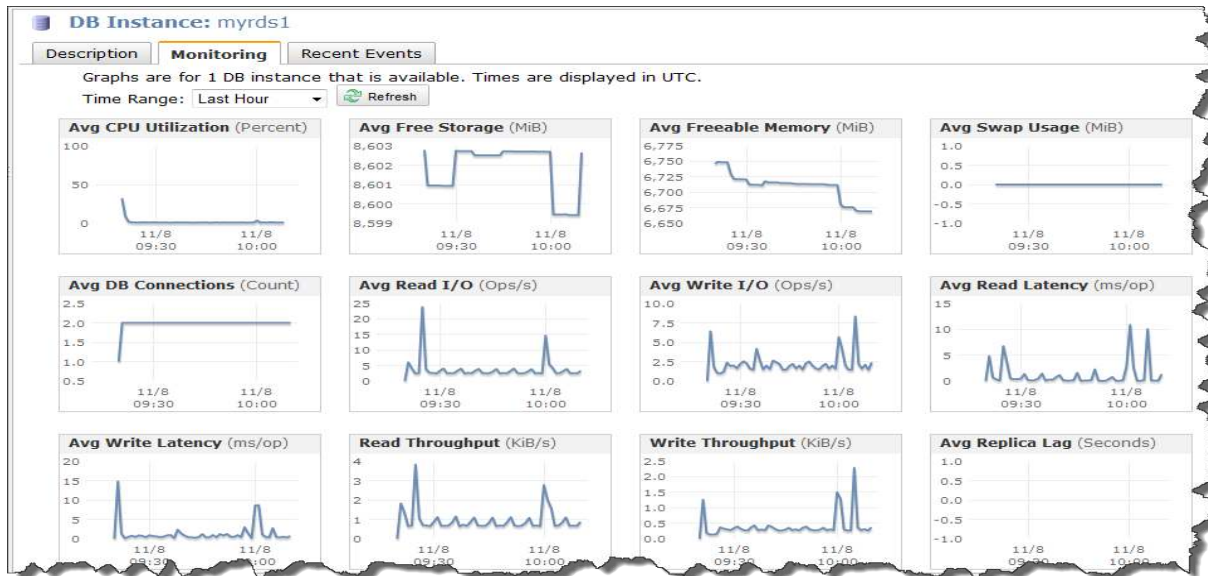
Controlled with IAM



Metrics and monitoring



Standard monitoring



Amazon CloudWatch metrics for Amazon RDS

- CPU utilization
- Storage
- Memory
- Swap usage
- DB connections
- I/O (read and write)
- Latency (read and write)
- Throughput (read and write)
- Replica lag
- Many more

Amazon CloudWatch Alarms

- Similar to on-premises custom monitoring tools

(Nov'16) price drop, longer retention & percentile monitoring

<https://aws.amazon.com/about-aws/whats-new/2016/11/announcing-cloudwatch-metrics-price-reduction-and-new-volume-based-pricing-tiers/>

<https://aws.amazon.com/blogs/aws/amazon-cloudwatch-update-percentile-statistics-and-new-dashboard-widgets/>

<https://aws.amazon.com/about-aws/whats-new/2016/11/cloudwatch-extends-metrics-retention-and-new-user-interface/>

Enhanced Monitoring

Access to over 50 new CPU, memory, file system, and disk I/O metrics as low as 1 second intervals (sent to CloudWatch Logs)



Monitoring

Enable Enhanced Monitoring Yes ▾

Monitoring Role Default ▾

Granularity 1 ▾ second(s)

☒ I authorize RDS to create the IAM role rds-monitoring-role.

Event notifications

- Uses **Amazon Simple Notification Service** (Amazon SNS) to notify users when an event occurs
- **17** different event categories (availability, backup, configuration change, and so on)

Create Event Subscription

Name ⓘ

Send notifications to ⓘ [create topic](#)

Source Type ⓘ

Enabled ☒ Yes ☐ No

Event Categories

☐ Select All
☒ Select specific


availability
backup
configuration change
creation
deletion
failover
failure
low storage
maintenance
notification
recovery
restoration

DB Instances

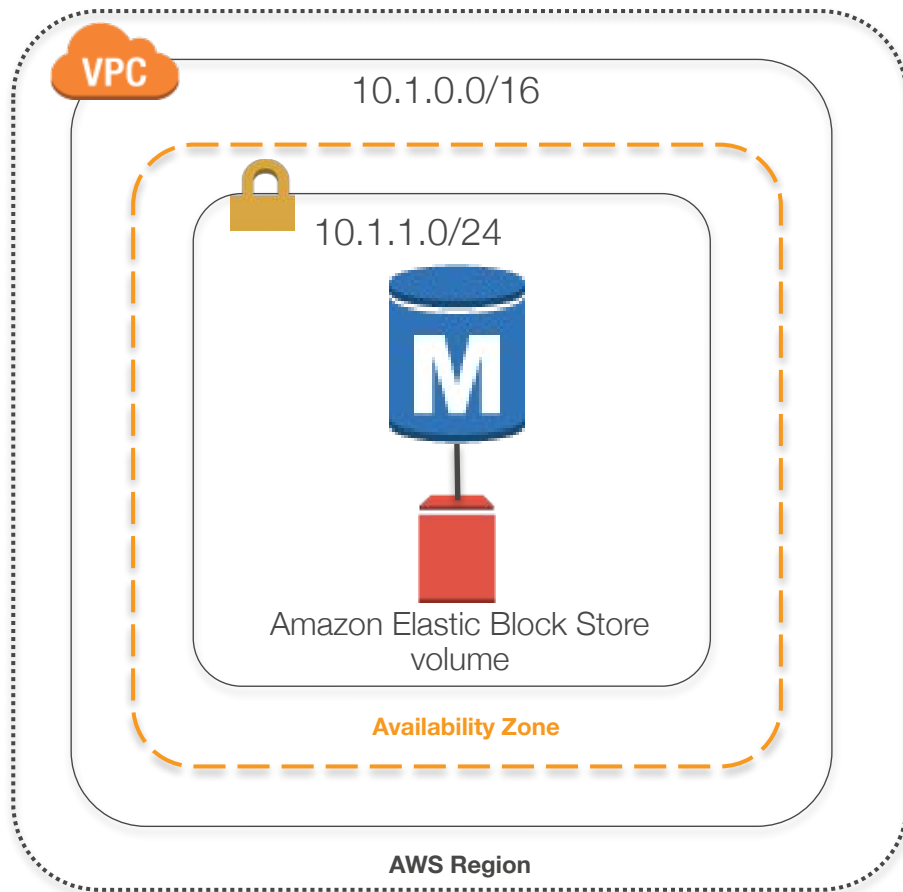
☐ Select All
☒ Select specific

djr-mysqlexampledb
djr-mysqlexampledb-restore
djr-mysqlexampledb-rr
djr-mysqlexampledb4
djr-rr-v2
djr-rr-v3
djr-sqltest
myvpcdbinstance
sg-dp-target
sg-oracle11204
sg-postgresql1
sg-rest-snap
sg-sqlsvr-ec2

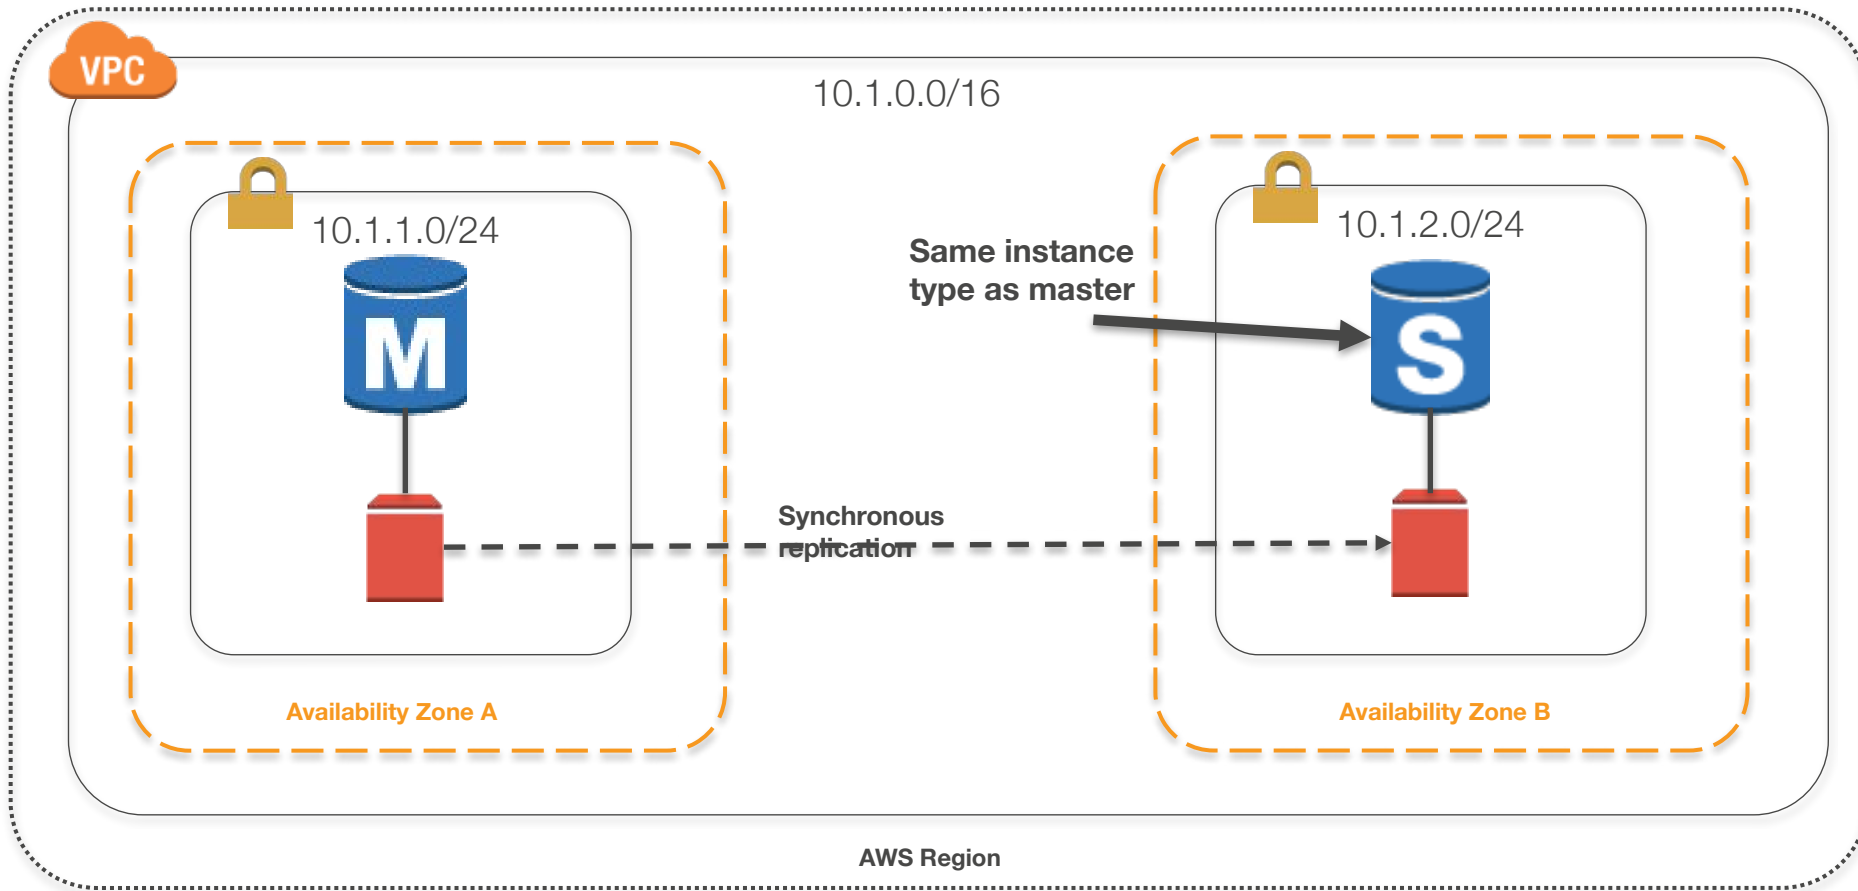
High availability




Minimal deployment—single AZ



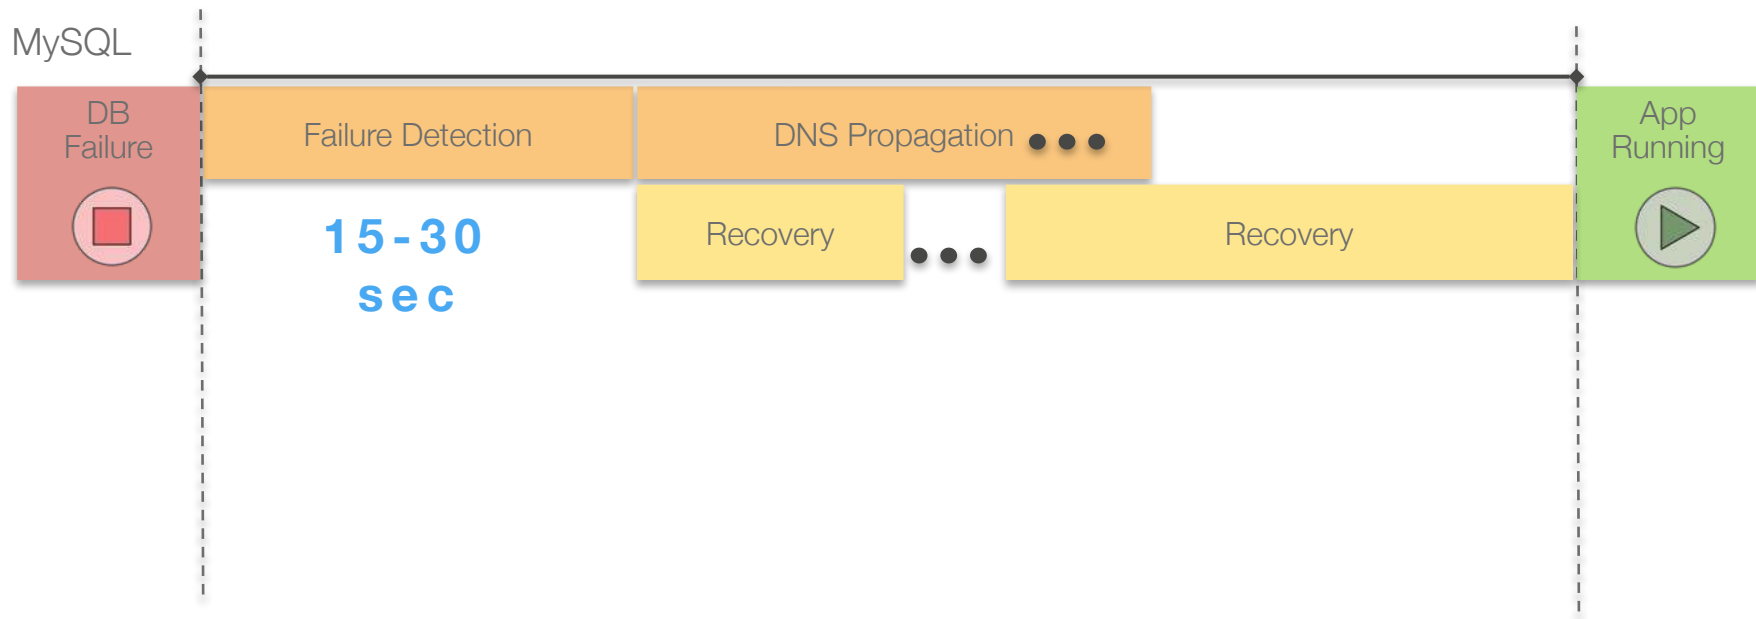
High availability – Multi-AZ



High availability—Multi-AZ to DNS

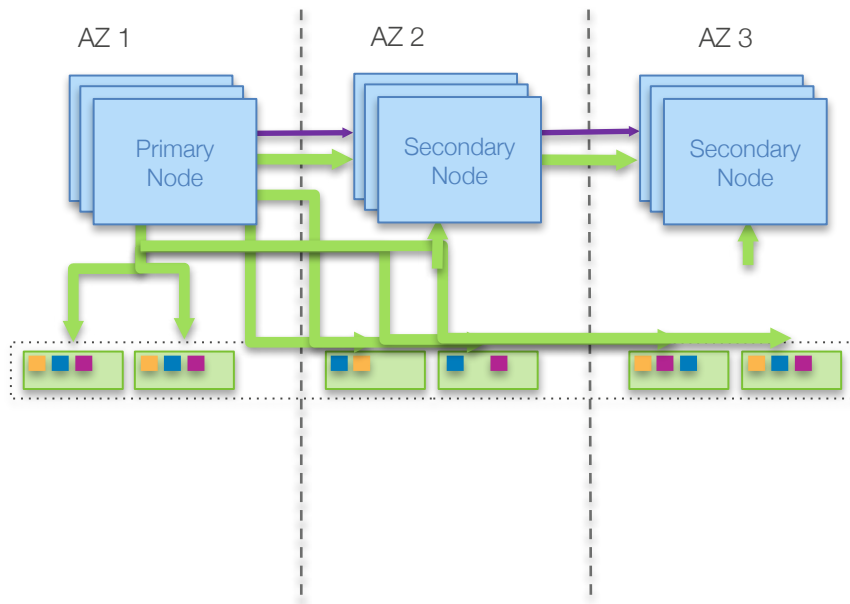


Failover – MySQL

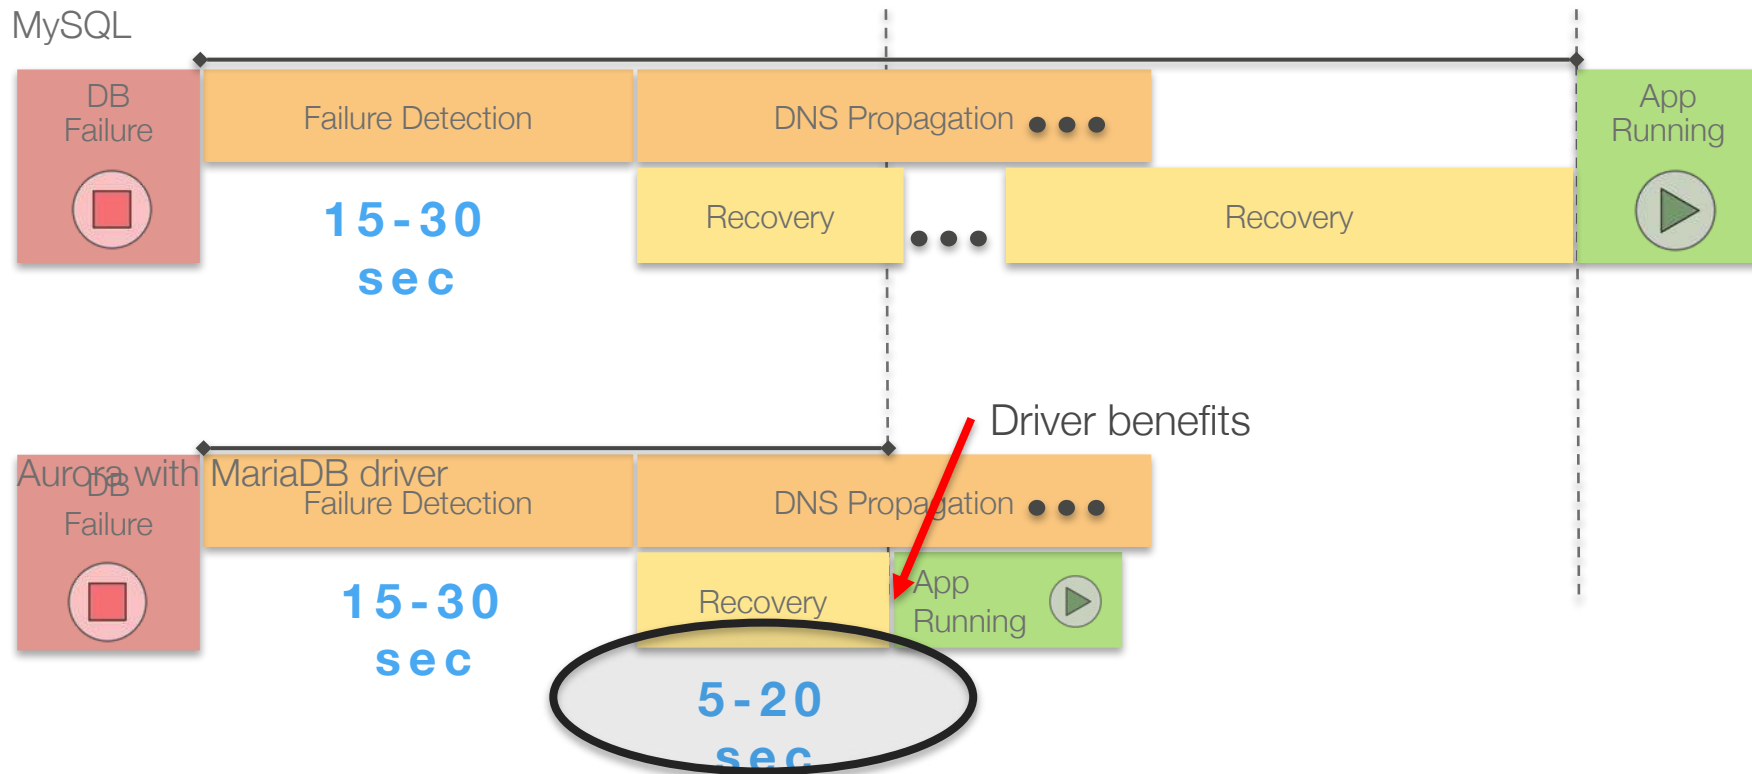


High availability in Aurora

- Aurora cluster contains primary node and **up to 15** secondary nodes (read-only)
- Failing nodes are **automatically** detected and replaced
- Failing database processes are **automatically** detected and recycled
- Secondary nodes **automatically** promoted on persistent outage, no single point of failure
- Customer application can **scale** out read traffic across secondary nodes



Failover – MySQL vs Aurora



Tips to improve recovery time with MySQL

- DO NOT use the IP address to connect to RDS!
- Set a low TTL on your own CNAME (beware if you use Java)
- Avoid large number of tables :
 - No more than 1000 tables using Standard Storage
 - No more than 10,000 tables using Provisioned IOPS
- Avoid very large tables in your database
- Avoid large transactions
- Make sure you have enough IOPS for recovery
- Use RDS Events to be notified

Simulating Amazon Aurora failures


ALTER SYSTEM CRASH [INSTANCE | DISPATCHER | NODE];

ALTER SYSTEM SIMULATE *percentage_of_failure* PERCENT

- **READ REPLICA FAILURE** [TO ALL | TO "replica name"]
- **DISK FAILURE** [IN DISK *index* | NODE *index*]
- **DISK CONGESTION** BETWEEN *minimum* AND *maximum* MILLISECONDS [IN DISK *index* | NODE *index*]

FOR INTERVAL *quantity* [YEAR | QUARTER | MONTH | WEEK | DAY | HOUR | MINUTE | SECOND];

Scaling on RDS



Read Replicas

Bring data close to your customer's applications in different regions

Relieve pressure on your master node for supporting reads and writes

Promote a Read Replica to a master for faster recovery in the event of disaster



Read Replicas

Read replicas are available for

- MySQL
- MariaDB
- PostgreSQL
- Aurora

They can be in the **same** region
or in a **different** region



Read Replicas—Oracle and SQL Server

Options



- Oracle GoldenGate
- Third-party replication products
- Snapshots

ORACLE®



Scaling up—or down

- Handle higher load or lower usage
- Control costs



Scaling up—or down

AWS Management Console

Instance Actions ▾

- See Details
- Create Read Replica
- Promote Read Replica
- Take Snapshot
- Restore to Point in Time
- Migrate Latest Snapshot
- Modify**
- Reboot
- Delete

Modify DB Instance: sg-cli-test

Instance Specifications

DB Engine Version MySQL 5.6.27 (default) ▾

DB Instance Class db.m4.large — 2 vCPU, 8 GiB RAM ▾

Multi-AZ Deployment No ▾

Storage Type General Purpose (SSD) ▾

Allocated Storage* 600 GB



Apply Immediately



Scaling—single AZ

With single AZ deployment, the master takes an **outage**

Alarms and Recent Events

TIME (UTC-7)	EVENT
Mar 26 7:01 AM	DB instance restarted
Mar 26 7:00 AM	Finished applying modification to DB instance class
Mar 26 6:53 AM	Applying modification to database instance class

Scaling—Multi-AZ

With Multi-AZ, the **standby** gets upgraded first

Alarms and Recent Events

TIME (UTC-7)	EVENT
Mar 26 6:34 AM	Finished applying modification to DB instance class
Mar 26 6:28 AM	Multi-AZ instance failover completed
Mar 26 6:28 AM	DB instance restarted
Mar 26 6:28 AM	Multi-AZ instance failover started
Mar 26 6:20 AM	Applying modification to database instance class

Scaling on a schedule – CLI or AWS Lambda

```
aws rds modify-db-instance
--db-instance-identifier sg-cli-test
--db-instance-class db.m4.large
--apply-immediately
```



```
#Scale down at 8:00 PM on Friday
0 20 * * 5
/home/ec2-user/scripts/scale_down_rds.s
h

#Scale up at 4:00 AM on Monday
0 4 * * 1
/home/ec2-user/scripts/scale_up_rds.sh
```

```
import boto3

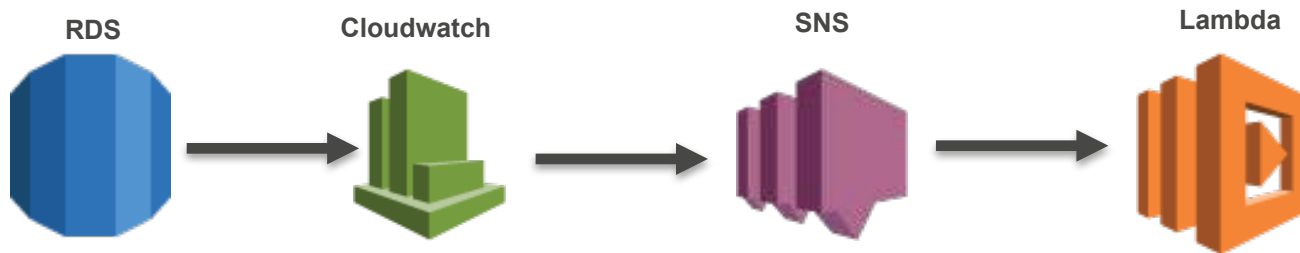
client=boto3.client('rds')

def lambda_handler(event, context):
    response=client.modify_db_instance(DBInstanceIdentifier='sg-cli-test',
                                       DBInstanceClass='db.m4.xlarge',
                                       ApplyImmediately=True)

    print response
```



Scaling on demand – Cloudwatch & AWS Lambda




```
import boto3
import json

client=boto3.client('rds')

def lambda_handler(event, context):
    message = event['Records'][0]['Sns']['Message']
    parsed_message=json.loads(message)
    db_instance=parsed_message['Trigger']['Dimensions'][0]['value']
    print 'DB Instance: ' + db_instance
    response=client.modify_db_instance(DBInstanceIdentifier=db_instance,
                                       DBInstanceClass='db.m4.large',
                                       ApplyImmediately=True)

    print response
```

Backups and snapshots



Backups

MySQL, PostgreSQL, MariaDB, Oracle, SQL Server

- Scheduled **daily** backup of **entire** instance
- Archive **database change** logs
- **35 day** retention for backups
- **Multiple copies** in each AZ where you have instances

Aurora

- **Automatic, continuous, incremental** backups
- **Point-in-time** restore
- **No impact** on database performance
- **35 day** retention



Restoring

- Restoring creates an entirely **new** database instance
- You define the instance configuration just like a new instance

Restore DB Instance

You are creating a new DB Instance from a source DB Instance at a specified time. This new DB Instance will have the default DB Security Group and DB Parameter Groups.

This feature is currently supported for InnoDB storage engine only. If you are using MyISAM, refer to details [here](#).

Use Latest Restorable Time

☒ March 8, 2016 at 12:10:00 PM UTC-8

Use Custom Restore Time

☐

MMMM d, y

00

:

00


:

00

UTC-8

Snapshots


- Full copies of your Amazon RDS database that are different from your scheduled backups
- Backed by Amazon S3
- May be used to create a new RDS instance
- Are encrypted if using encryption




Snapshots

Use cases

- Resolve production issues
- Build non-production environments
- Point-in-time restore
- Final copy before terminating a database
- Disaster recovery
- Cross-region copy
- Copy between accounts



Migrating onto RDS






AWS Database Migration Service



PostgreSQL

Amazon Aurora



mongoDB

ORACLE



MySQL



MariaDB



ASE



Microsoft
SQL Server

- ✓ Move data to the **same** or different **database** engine
- ✓ Move data to Redshift, DynamoDB or S3
- ✓ Keep your apps **running** during the migration
- ✓ Start your first migration in **10** minutes or less
- ✓ Replicate **within**, **to**, or **from** Amazon EC2 or RDS


<https://aws.amazon.com/dms/>

http://docs.aws.amazon.com/dms/latest/userguide/CHAP_Introduction.Sources.html

http://docs.aws.amazon.com/dms/latest/userguide/CHAP_Introduction.Targets.html

<https://aws.amazon.com/blogs/database/database-migration-what-do-you-need-to-know-before-you-start/>

Keep your apps running during the migration



Start a replication instance

Connect to source and target database


Select tables, schemas, or databases



Application Users

Let the AWS Database Migration Service create tables, load data, and keep them in sync

Switch applications over to the target at your convenience



AWS Schema Conversion Tool

- Move your **tables, views, stored procedures,** and **data manipulation language (DML)** to RDS or Amazon Redshift
- **Highlight** where manual edits are needed

Source Database	Target Database on Amazon RDS
Oracle	Amazon Aurora, MySQL, PostgreSQL, MariaDB
Oracle Data Warehouse	Amazon Redshift
Microsoft SQL Server	Amazon Aurora, MySQL, PostgreSQL, MariaDB
Teradata	Amazon Redshift
Netezza	Amazon Redshift
Greenplum	Amazon Redshift
MySQL and MariaDB	PostgreSQL
PostgreSQL	Amazon Aurora, MySQL, MariaDB
Amazon Aurora	PostgreSQL



Thank you

Julien Simon, Principal Technical Evangelist, AWS

julsimon@amazon.fr

@julsimon

