



Deep Learning for Developers

16/12/2017

Julien Simon, AI Evangelist, EMEA
@julsimon

What to expect

- AI ?
- An introduction to Deep Learning
- Common neural network architectures and use cases
- An introduction to Apache MXNet
- Demos using Jupyter notebooks on Amazon SageMaker
- Resources

- **Artificial Intelligence**: design software applications which exhibit human-like behavior, e.g. speech, natural language processing, reasoning or intuition
- **Machine Learning**: teach machines to learn without being explicitly programmed
- **Deep Learning**: using neural networks, teach machines to learn from complex data where features cannot be explicitly expressed


Myth: AI is dark magic

aka « You're not smart enough »



Fact: AI is math, code and chips

A bit of Science, a lot of Engineering



```
data = mx.symbol.Variable('data')
conv1 = mx.sym.Convolution(data=data, kernel=(5,5), num_filter=20)
relu1 = mx.sym.Activation(data=conv1, act_type="relu")
pool1 = mx.sym.Pooling(data=relu1, pool_type="max", kernel=(2,2), stride=(2,2))
conv2 = mx.sym.Convolution(data=pool1, kernel=(5,5), num_filter=50)
relu2 = mx.sym.Activation(data=conv2, act_type="relu")
pool2 = mx.sym.Pooling(data=relu2, pool_type="max", kernel=(2,2), stride=(2,2))
flatten = mx.sym.Flatten(data=pool2)
fc1 = mx.symbol.FullyConnected(data=flatten, num_hidden=500)
relu3 = mx.sym.Activation(data=fc1, act_type="relu")
fc2 = mx.symbol.FullyConnected(data=relu3, num_hidden=10)
lenet = mx.sym.SoftmaxOutput(data=fc2, name='softmax')
```



Amazon AI is based on Deep Learning

Vision Services

Amazon Rekognition Image

Deep learning-based image analysis

[Learn more »](#)

Amazon Rekognition Video

Deep learning-based video analysis

[Learn more »](#)

NEW

Conversational chatbots

Amazon Lex

Build chatbots to engage customers

[Learn more »](#)



Language Services

Amazon Comprehend

Discover insights and relationships in text

[Learn more »](#)

NEW

Amazon Translate

Fluent translation of text

[Learn more »](#)

NEW

Amazon Transcribe

Automatic speech recognition

[Learn more »](#)

NEW

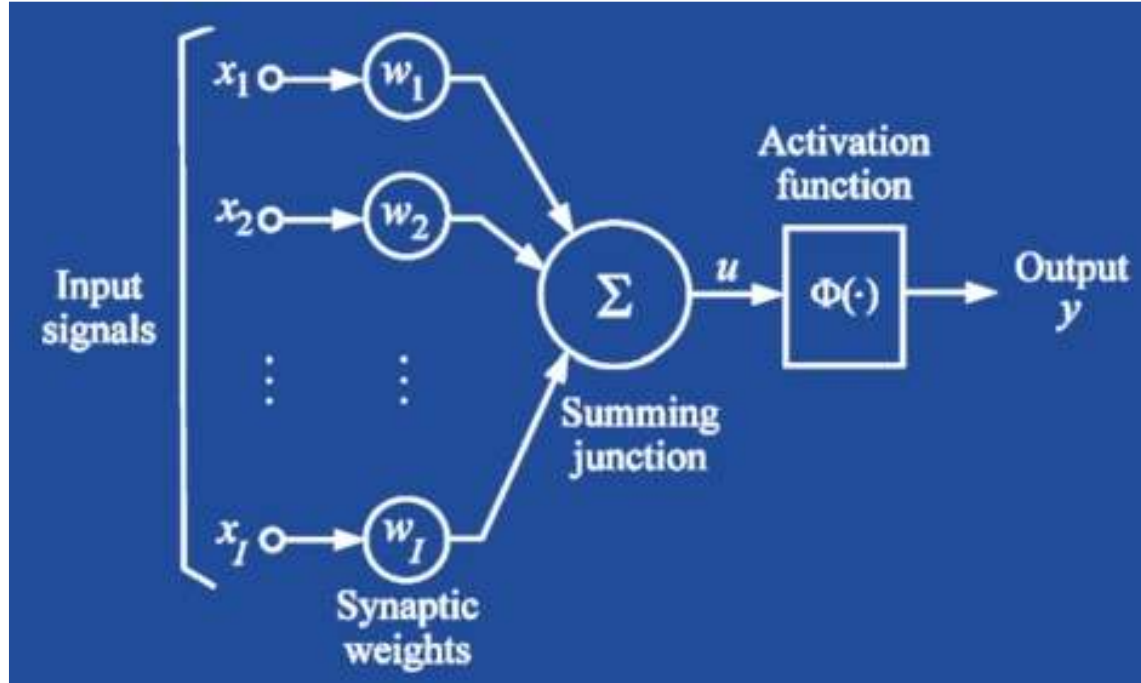
Amazon Polly

Natural sounding text to speech

[Learn more »](#)

An introduction to Deep Learning

The neuron



$$\sum_{i=1}^l x_i * w_i = u$$

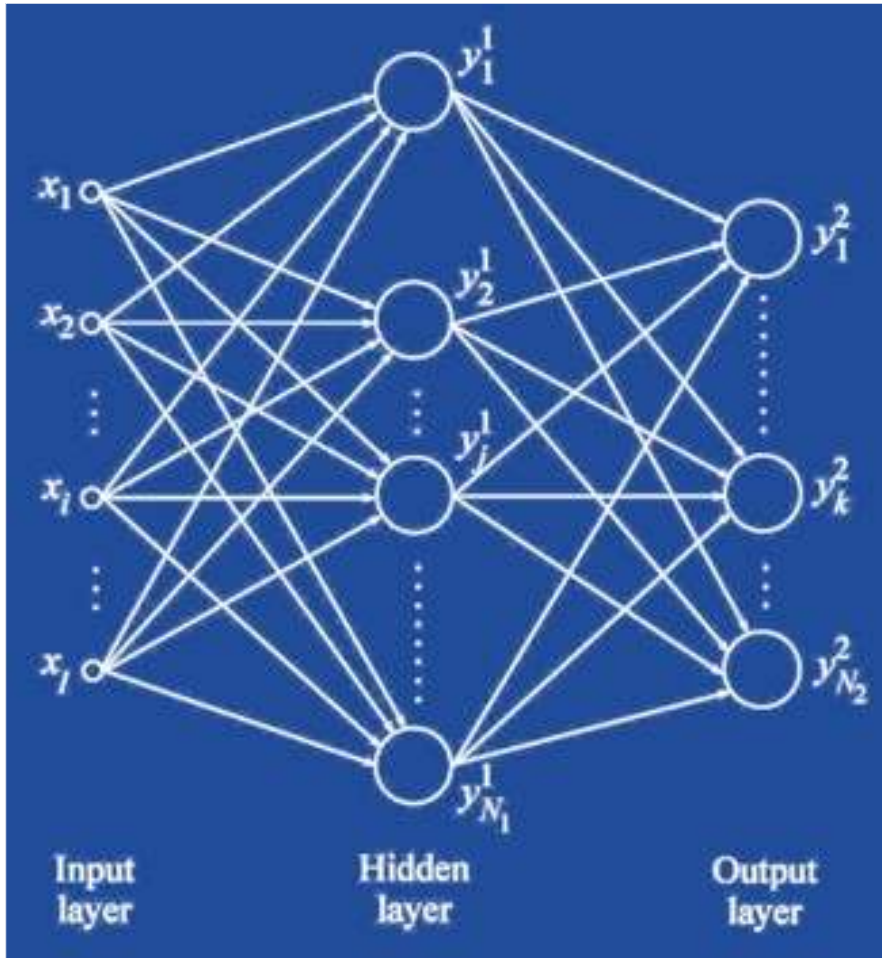
”Multiply and Accumulate”

Activation functions

Name	Plot	Equation
Identity		$f(x) = x$
Binary step		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Logistic (a.k.a. Soft step)		$f(x) = \frac{1}{1 + e^{-x}}$
TanH		$f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$
ArcTan		$f(x) = \tan^{-1}(x)$
Softsign ^{[7][8]}		$f(x) = \frac{x}{1 + x }$
Rectified linear unit (ReLU) ^[9]		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$

Source: Wikipedia

Neural networks



$$X = \begin{bmatrix} X_{11}, X_{12}, \dots, X_{1l} \\ X_{21}, X_{22}, \dots, X_{2l} \\ \dots \dots \dots \\ X_{m1}, X_{m2}, \dots, X_{ml} \end{bmatrix}$$

l features

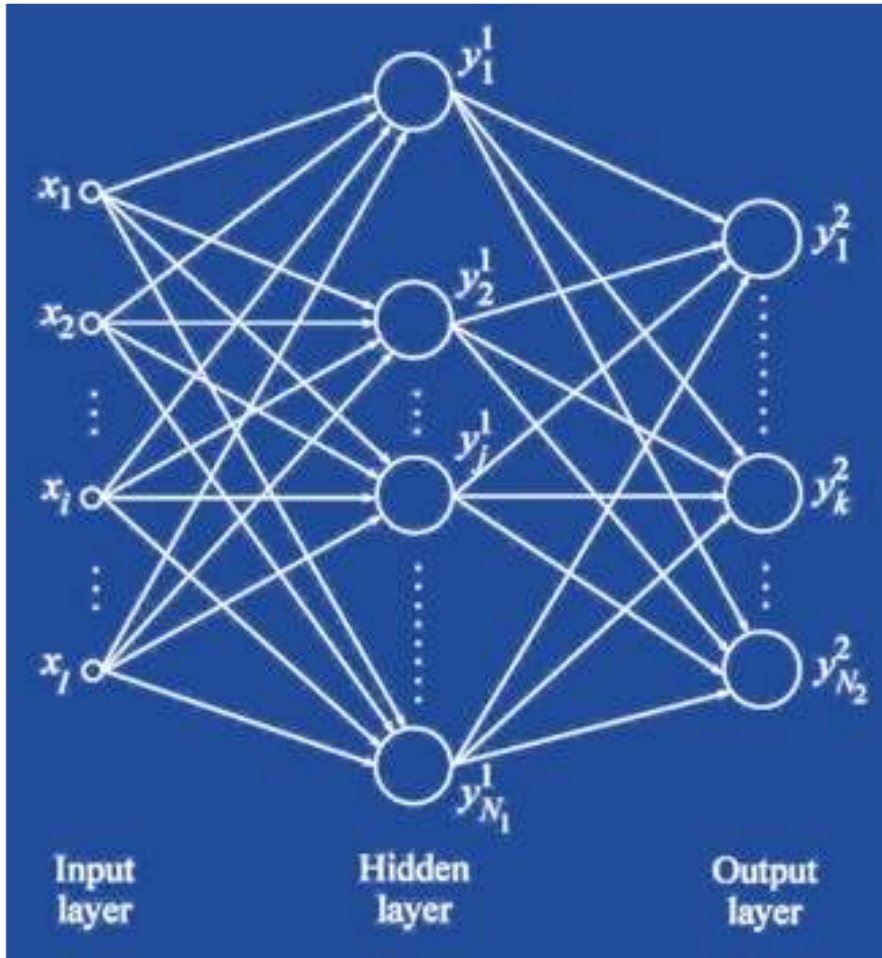
m samples

$$y = \begin{bmatrix} y_1 \\ y_2 \\ \dots \\ y_m \end{bmatrix} = \begin{bmatrix} 0,0,1,0,0,\dots,0 \\ 1,0,0,0,0,\dots,0 \\ \dots \\ 0,0,0,0,1,\dots,0 \end{bmatrix}$$

m labels,
 N_2 outputs

One-hot encoding

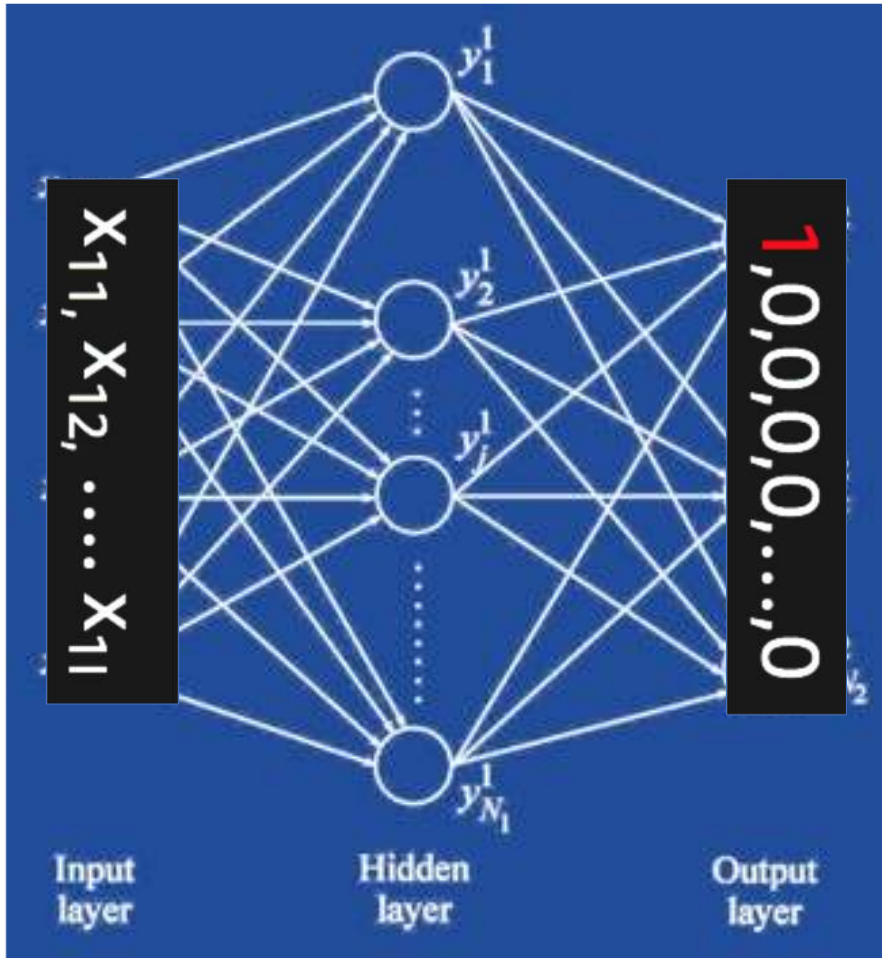
Neural networks



$$\begin{aligned}
 \mathbf{X} &= \begin{bmatrix} x_{11}, x_{12}, \dots, x_{1I} \\ x_{21}, x_{22}, \dots, x_{2I} \\ \vdots \\ x_{m1}, x_{m2}, \dots, x_{mI} \end{bmatrix} \quad \begin{matrix} I \text{ features} \\ m \text{ samples} \end{matrix} \\
 \mathbf{y} &= \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{bmatrix} = \begin{bmatrix} 0, 0, 1, \dots, 0 \\ 1, 0, 0, \dots, 0 \\ \vdots \\ 0, 0, 0, \dots, 0 \end{bmatrix} \quad \begin{matrix} m \text{ labels,} \\ N_2 \text{ categories} \end{matrix} \\
 &\quad \text{One-hot encoding}
 \end{aligned}$$

$$\text{Accuracy} = \frac{\text{Number of correct predictions}}{\text{Total number of predictions}}$$

Neural networks



Initially, the network will **not** predict correctly

$$f(X_1) = Y'_1$$

A **loss function** measures the difference between the **real label** Y_1 and the **predicted label** Y'_1

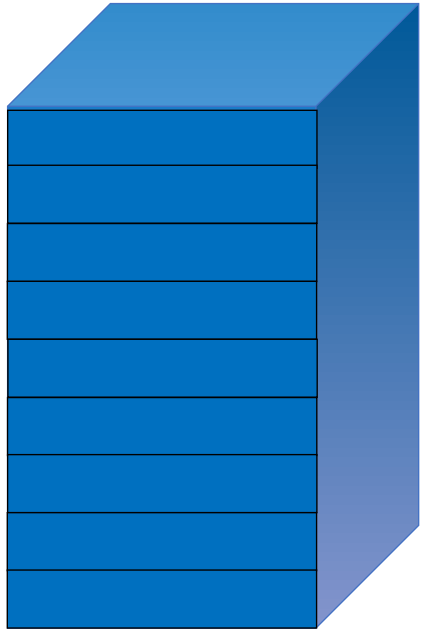
$$\text{error} = \text{loss}(Y_1, Y'_1)$$

For a **batch** of samples:

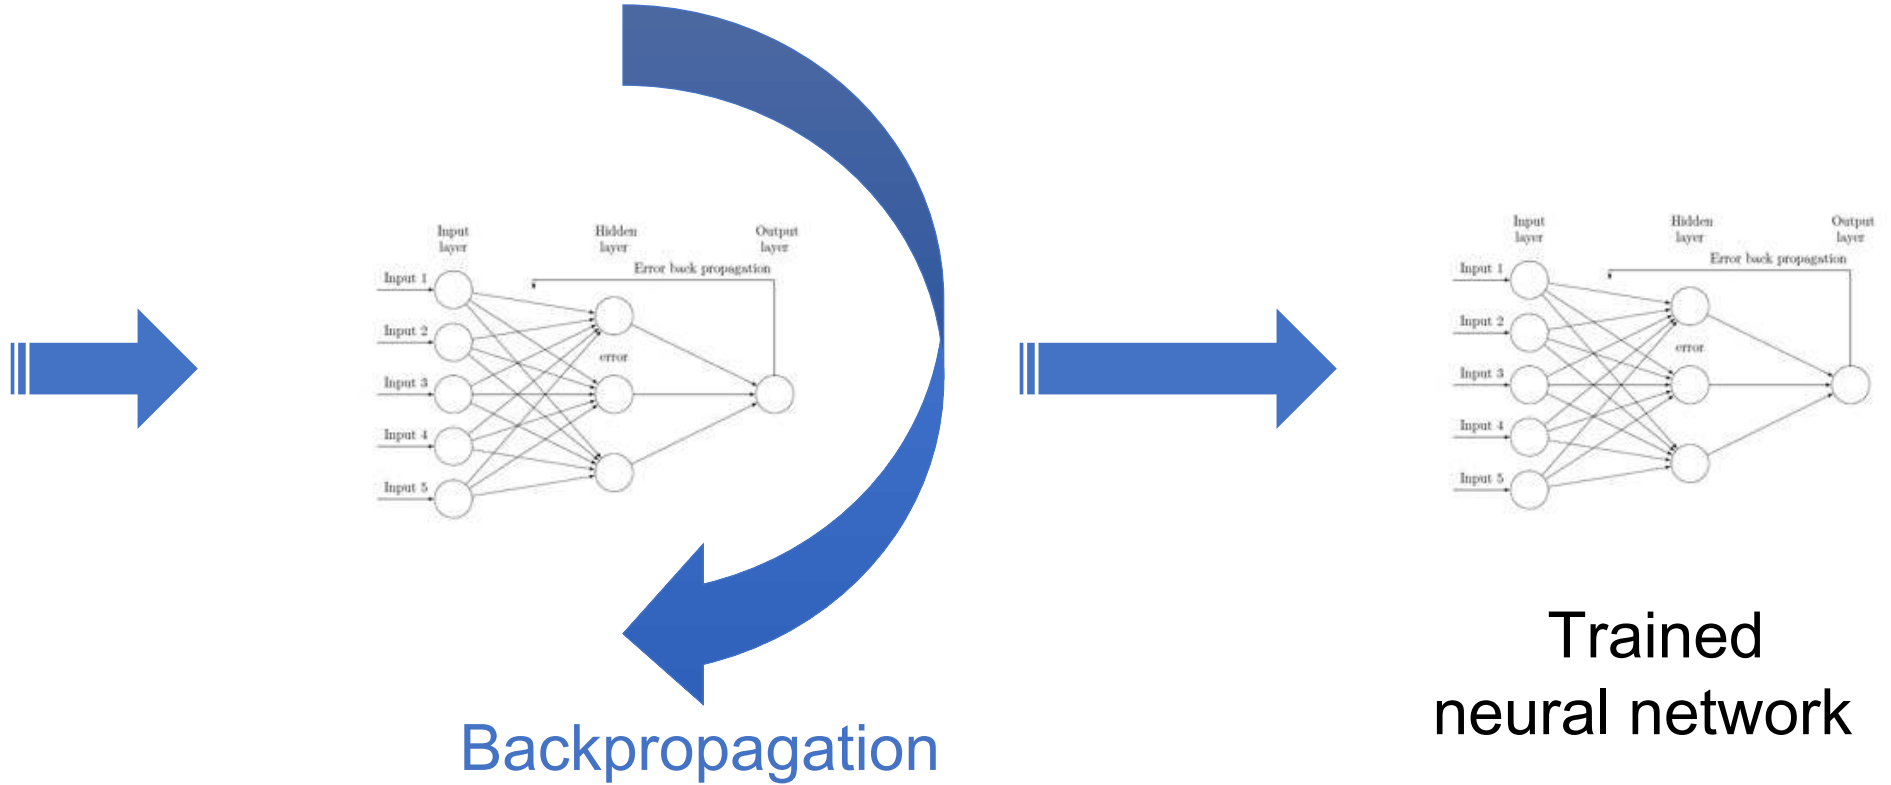
$$\sum_{i=1}^{\text{batch size}} \text{loss}(Y_i, Y'_i) = \text{batch error}$$

The purpose of the training process is to **minimize error** by gradually **adjusting weights**

Training



Training data set




Batch size
Learning rate
Number of epochs

} Hyper parameters

Stochastic Gradient Descent (SGD)

*Imagine you stand on top of a mountain with skis strapped to your feet. You want to get down to the valley as quickly as possible, but there is fog and you can only see your immediate surroundings. How can you get down the mountain as quickly as possible? You look around and *identify the steepest path down*, *go down that path for a bit*, *again look around* and *find the new steepest path*, *go down that path*, and *repeat*—this is exactly what gradient descent does.*


Tim Dettmers
University of Lugano
2015




The « step size » is called the **learning rate**

Alternatives to SGD:
rmsprod, adagrad, adadelata,
adam, etc.

Validation



Validation data set




Trained
neural network



Validation
accuracy

Prediction at
the end of
each epoch

Early stopping

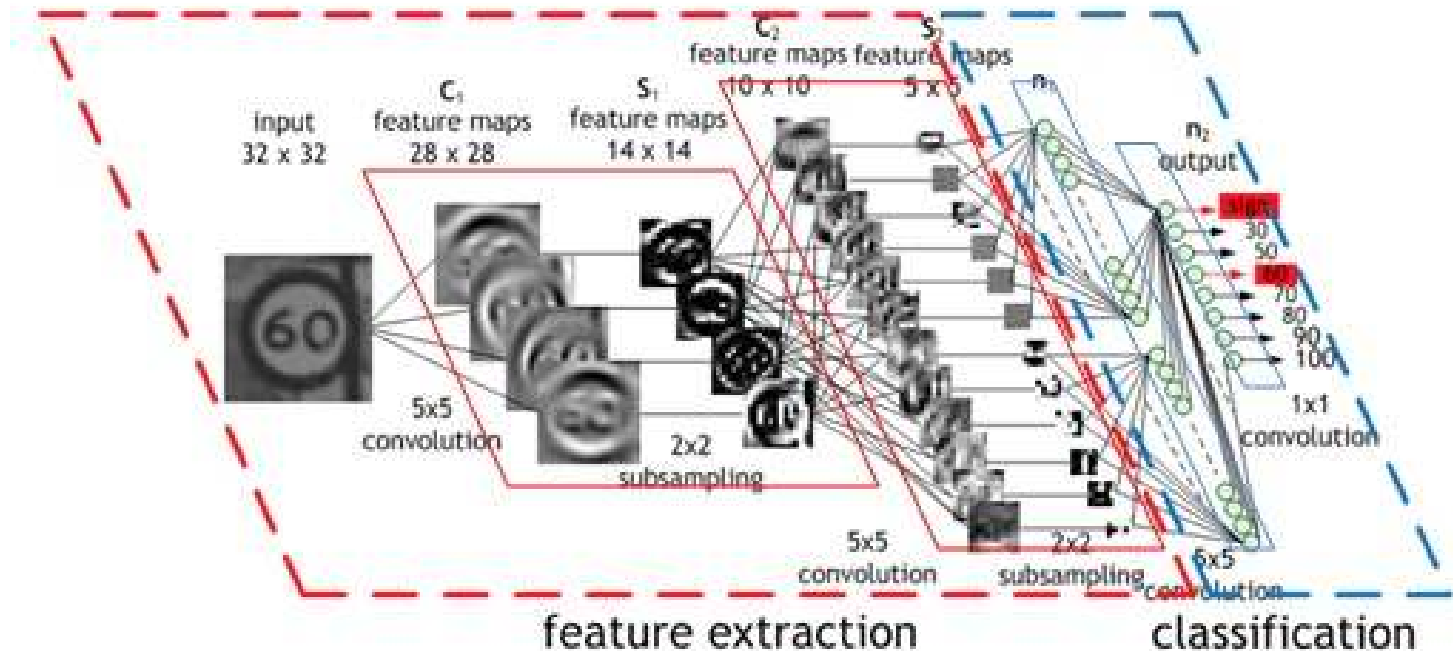



Common network architectures and use cases

Convolutional Neural Networks (CNN)

Le Cun, 1998: handwritten digit recognition, 32x32 pixels

Convolution and pooling reduce dimensionality



CNN: Object Classification

- Expedia have over 10 million images from 300,000 hotels
- Using great images boosts conversion
- Using Keras and EC2 GPU instances, they fine-tuned a pre-trained Convolutional Neural Network using 100,000 images
- Hotel descriptions now automatically feature the best available images


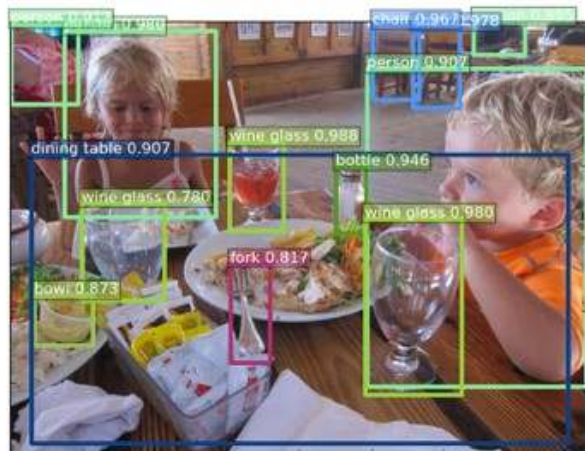

Some images are really good



Others not so much



CNN: Object Detection




<https://github.com/precedenceguo/mx-rcnn>

<https://github.com/zhreshold/mxnet-yolo>

CNN: Object Detection

- 17,000 images from Instagram
- 7 brands
- Deep Learning model pre-trained on ImageNet
- Fine-tuning with TensorFlow and EC2 GPU instances
- Additional work on color extraction

CONDÉ NAST




	Chanel	Coach	Gucci	Marc Jacobs	Kate Spade	No Handbag	Prada	Vuitton
Chanel	0.83	0.00	0.01	0.02	0.00	0.00	0.00	0.01
Coach	0.01	0.85	0.00	0.05	0.05	0.01	0.04	0.03
Gucci	0.01	0.00	0.85	0.02	0.00	0.01	0.01	0.02
Marc Jacobs	0.00	0.03	0.01	0.78	0.00	0.01	0.03	0.00
Kate Spade	0.00	0.01	0.01	0.01	0.87	0.00	0.00	0.00
No Handbag	0.09	0.06	0.08	0.09	0.04	0.97	0.04	0.09
Prada	0.03	0.03	0.02	0.03	0.01	0.00	0.85	0.01
Vuitton	0.01	0.00	0.00	0.02	0.00	0.01	0.01	0.81



<https://technology.condenast.com/story/handbag-brand-and-color-detection>

CNN: Object Segmentation



<https://github.com/TuSimple/mx-maskrcnn>

图森 **tu** Simple



Last June, tuSimple drove an autonomous truck
for 200 miles from Yuma, AZ to San Diego,
CA


CNN: Face Detection



```
attribution is:  
5_o_Clock_Shadow : No  
Arched_Eyebrows : No  
Attractive : No  
Bags_Under_Eyes : No  
Bald : No  
Bangs : No  
Big_Lips : No  
Big_Nose : No  
Black_Hair : No  
Blond_Hair : No  
Blurry : Yes  
Brown_Hair : No  
Bushy_Eyebrows : No  
Chubby : No  
Double_Chin : No  
Eyeglasses : No  
Goatee : No  
Gray_Hair : No  
Heavy_Makeup : No  
High_Cheekbones : No  
Male : Yes  
Mouth_Slightly_Open : No  
Mustache : No  
Narrow_Eyes : Yes  
No_Beard : Yes  
Oval_Face : No  
Pale_Skin : No  
Pointy_Nose : No  
Receding_Hairline : No  
Rosy_Cheeks : No  
Sideburns : No  
Smiling : No  
Straight_Hair : No  
Wavy_Hair : No  
Wearing_Earrings : No  
Wearing_Hat : No  
Wearing_Lipstick : No  
Wearing_Necklace : No  
Wearing_Necktie : No  
Young : Yes
```

<https://github.com/tornadomeet/mxnet-face>

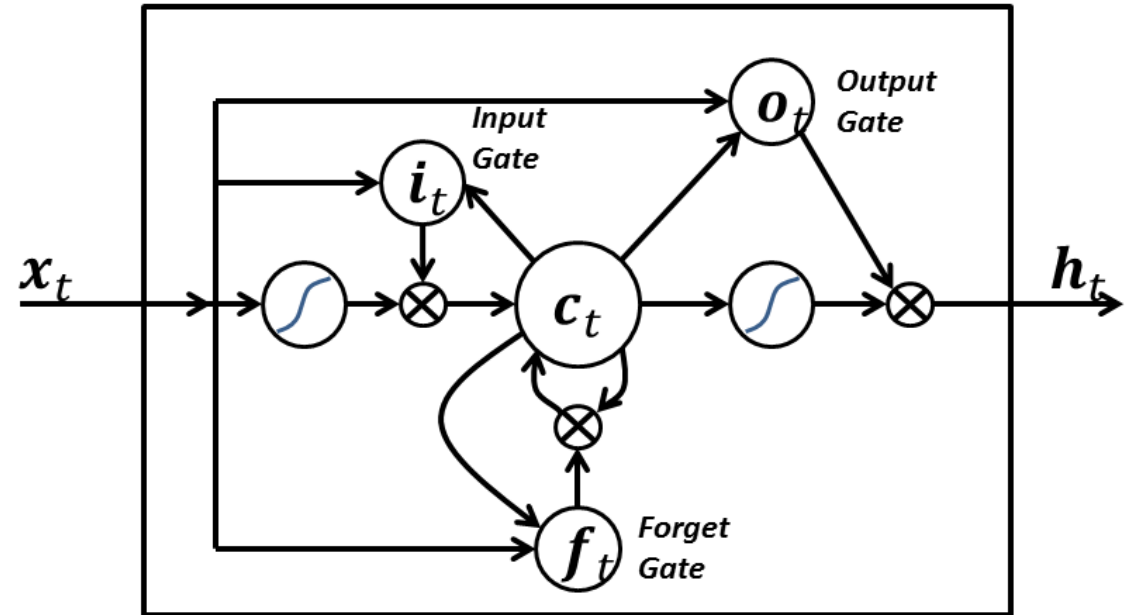
Real-Time Pose Estimation



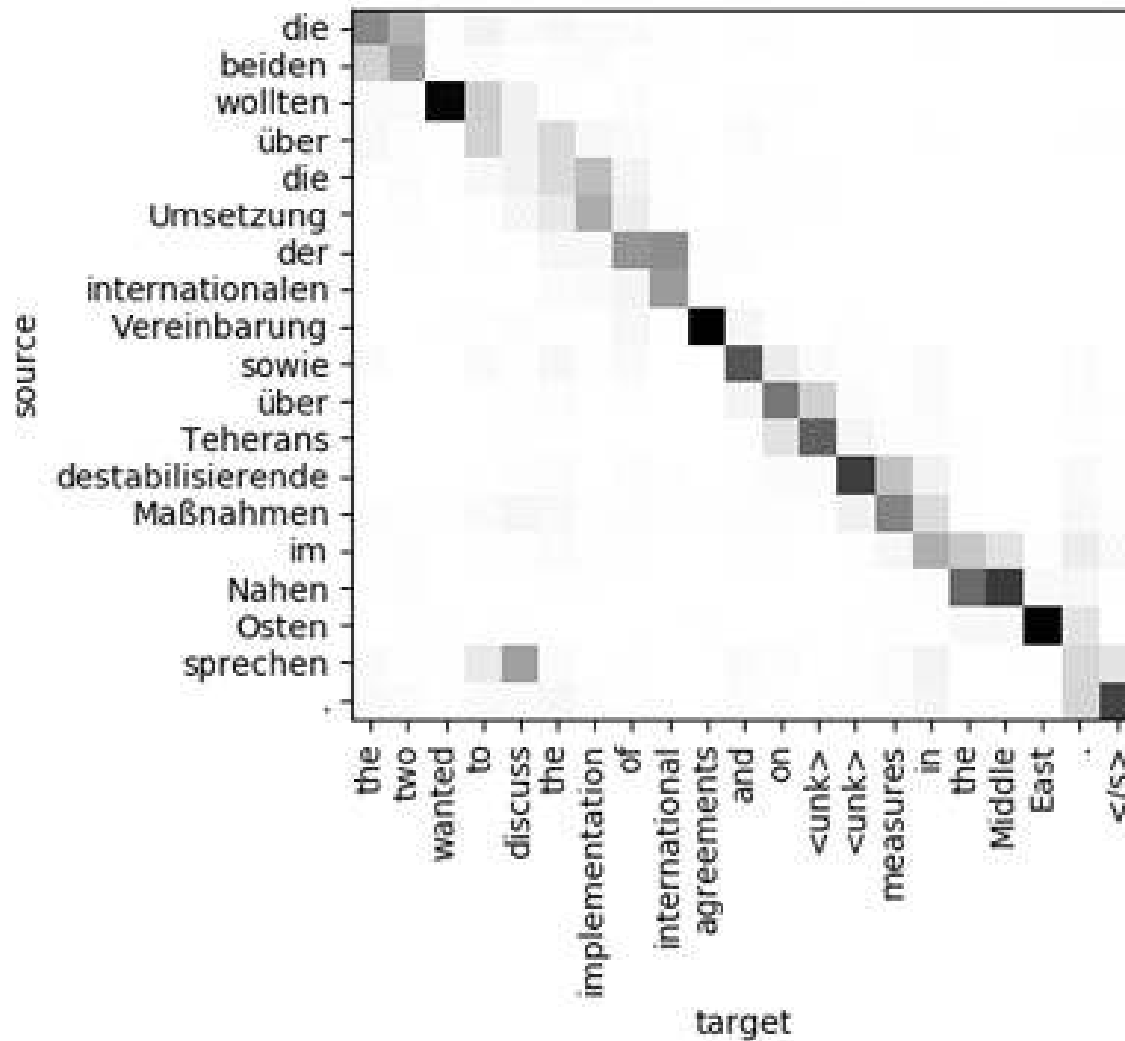
<https://github.com/dragonfly90/mxnet> Realtime Multi-Person Pose Estimation

Long Short Term Memory Networks (LSTM)

- A LSTM neuron computes the output based on the input and a **previous state**
- LSTM networks have **memory**
- They're great at predicting **sequences**, e.g. machine translation



LSTM: Machine Translation




<https://github.com/aws-labs/sockeye>


Generative Adversarial Networks (GAN)

- Goodfellow, 2014
- Dual-network architecture
- **Detector network**
 - Sees real samples
 - Learn how to **detect** real samples from fake ones created by the Generator network
- **Generator network**
 - Doesn't see real samples
 - Creates images from **random data**
 - Applies the **same weight updates** as the Generator
 - Learns gradually how to generate **better fake samples**

Adapted from <https://arxiv.org/pdf/1511.06434.pdf>




Adapted from <https://arxiv.org/pdf/1511.06434.pdf>



GAN: Welcome to the (un)real world, Neo



Generating new "celebrity" faces https://github.com/tkarras/progressive_growing_of_gans



From semantic map to 2048x1024 picture
<https://tcwang0509.github.io/pix2pixHD/>

Apache MXNet

Apache MXNet: Open Source library for Deep Learning




Programmable

Simple syntax,
multiple
languages




Most Open

Accepted into the
Apache Incubator




Portable

Highly efficient
models for
mobile
and IoT



Best On AWS

Optimized for
Deep Learning on AWS




High Performance

Near linear scaling
across hundreds of
GPUs



MXNet 1.0 released on December 4th



Image



Video



Speech




Events

"People Riding Bikes"


Text

Input $\begin{bmatrix} 1 \\ 3 \\ \dots \\ 4 \end{bmatrix}$ Weights $\begin{bmatrix} 0.2 \\ -0.1 \\ \dots \\ 0.7 \end{bmatrix} = 2$

```
mx.sym.FullyConnected(data, num_hidden=128)
```




```
mx.sym.Convolution(data, kernel=(5,5), num_filter=20)
```



```
mx.sym.Pooling(data, pool_type="max", kernel=(2,2),
```


stride=(2,2)




```
lstm.lstm_unroll(num_lstm_layer, seq_len, len, num_hidden, num_embed)
```

Queen $\rightarrow \begin{bmatrix} 0.2 \\ -0.1 \\ \dots \\ 0.7 \end{bmatrix}$ $\cos(w, \text{queen}) = \cos(w, \text{king}) - \cos(w, \text{man}) + \cos(w, \text{woman})$


```
mx.symbol.Embedding(data, input_dim, output_dim = k)
```



```
mx.sym.Activation(data, act_type="xxxx")
```



"sigmoid"



"tanh"



"relu"



"softrelu"



Image Segmentation



Face Search



Neural Art

"People Riding Bikes"

Image Caption

Bicycle, People, Road, Sport

Image Labels

"Οι άνθρωποι ιππασίας ποδήλατα"

Machine Translation

Model Server for Apache MXNet

ubuntu/python-2.7	ubuntu/python-3.5
 AWS CodeBuild Passing	 AWS CodeBuild Passing

Model Server for Apache MXNet (MMS) is a flexible and easy to use tool for serving Deep Learning models.

Use MMS Server CLI, or the pre-configured Docker images, to start a service that sets up HTTP endpoints to handle model inference requests.

<https://github.com/awslabs/mxnet-model-server/>

ONNX

OPEN NEURAL NETWORK EXCHANGE FORMAT

The new open ecosystem for interchangeable AI models

<https://aws.amazon.com/blogs/ai/announcing-onnx-support-for-apache-mxnet/>

The Apache MXNet API

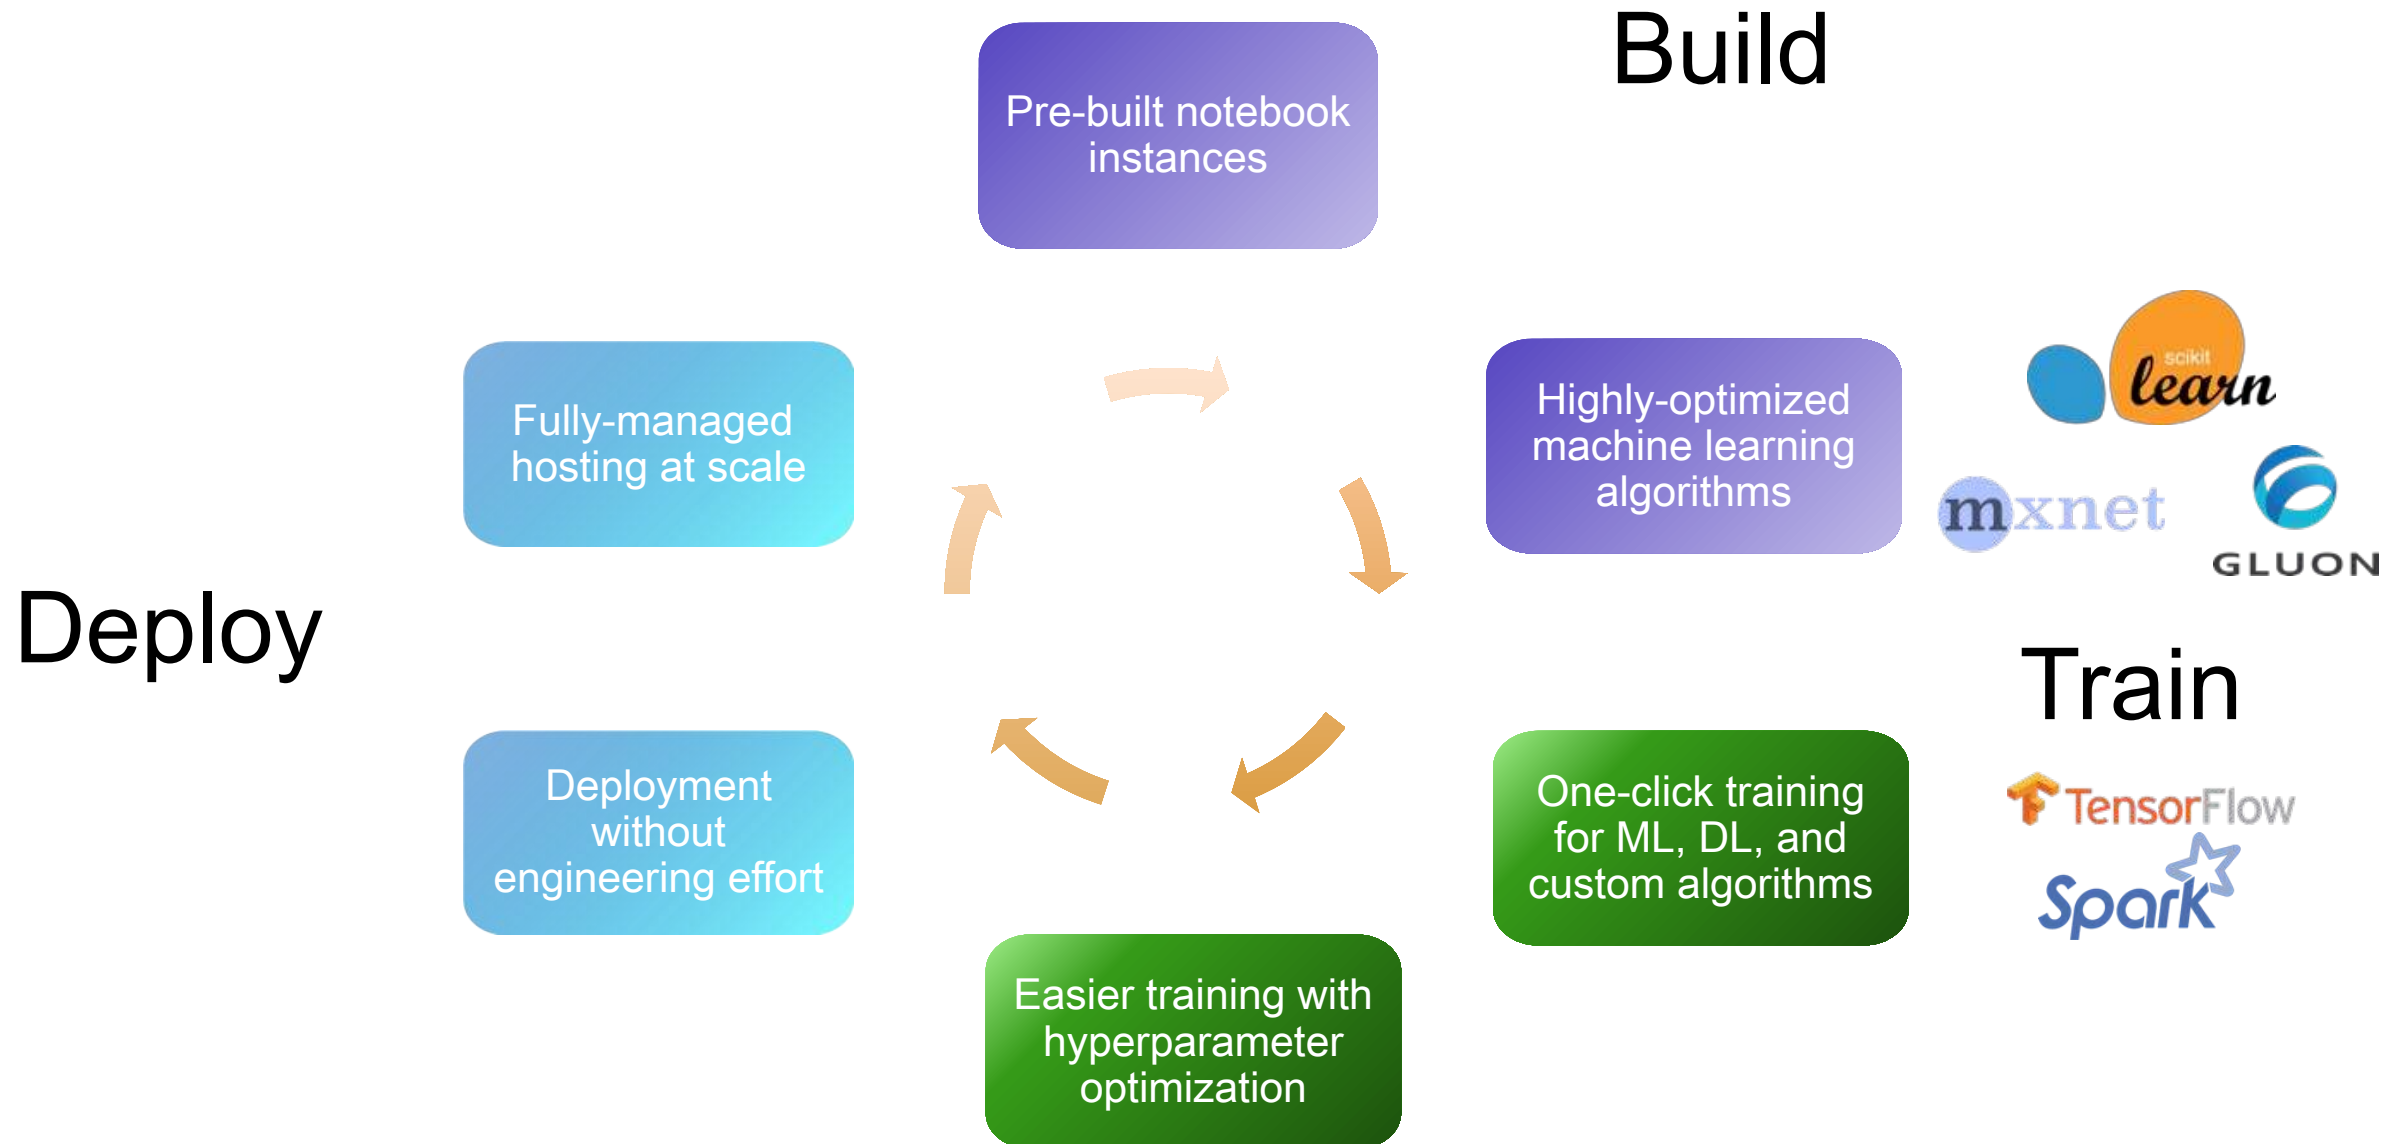
- Storing and accessing data in multi-dimensional arrays
→ *NDArray* API
- Building models (layers, weights, activation functions)
→ *Symbol* API
- Serving data during training and validation
→ *Iterators*
- Training and using models
→ *Module* API

Demos

<https://github.com/juliensimon/dlnotebooks>

- 1) Hello World: learn a synthetic data set
- 2) Classify images with pre-trained models
- 3) Learn and predict MNIST with a Multi-Layer Perceptron and the LeNet CNN
- 4) Train, host and deploy a model with Amazon SageMaker
- 5) Generate MNIST samples with a GAN

Amazon SageMaker



Resources

<https://aws.amazon.com/machine-learning>

<https://aws.amazon.com/sagemaker>

<https://aws.amazon.com/blogs/ai>

<https://mxnet.incubator.apache.org>

<https://github.com/apache/incubator-mxnet>

<https://github.com/gluon-api>

An overview of Amazon SageMaker <https://www.youtube.com/watch?v=ym7NEYEx9x4>

<https://medium.com/@julsimon>



Thank you!

**Julien Simon, AI Evangelist,
EMEA
@julsimon**