



Deep Learning for Developers

Julien Simon, AI Evangelist, EMEA
@julsimon

What to expect

- An introduction to Deep Learning
- An introduction to Apache MXNet
- Demos using Jupyter notebooks on Amazon SageMaker
- Resources

- **Artificial Intelligence**: design software applications which exhibit human-like behavior, e.g. speech, natural language processing, reasoning or intuition
- **Machine Learning**: teach machines to learn without being explicitly programmed
- **Deep Learning**: using neural networks, teach machines to learn from complex data where features cannot be explicitly expressed

Amazon Alexa is based on Deep Learning



Amazon AI is based on Deep Learning

Vision Services

Amazon Rekognition Image

Deep learning-based image analysis

[Learn more »](#)

Amazon Rekognition Video

Deep learning-based video analysis

[Learn more »](#)



Conversational chatbots

Amazon Lex

Build chatbots to engage customers

[Learn more »](#)

Language Services

Amazon Comprehend

Discover insights and relationships in text

[Learn more »](#)



Amazon Translate

Fluent translation of text

[Learn more »](#)



Amazon Transcribe

Automatic speech recognition

[Learn more »](#)



Amazon Polly

Natural sounding text to speech

[Learn more »](#)


图森 **tu** Simple



Last June, tuSimple drove an autonomous truck
for 200 miles from Yuma, AZ to San Diego,
CA

An introduction to Deep Learning

The neuron



$$\sum_{i=1}^l x_i * w_i = u$$

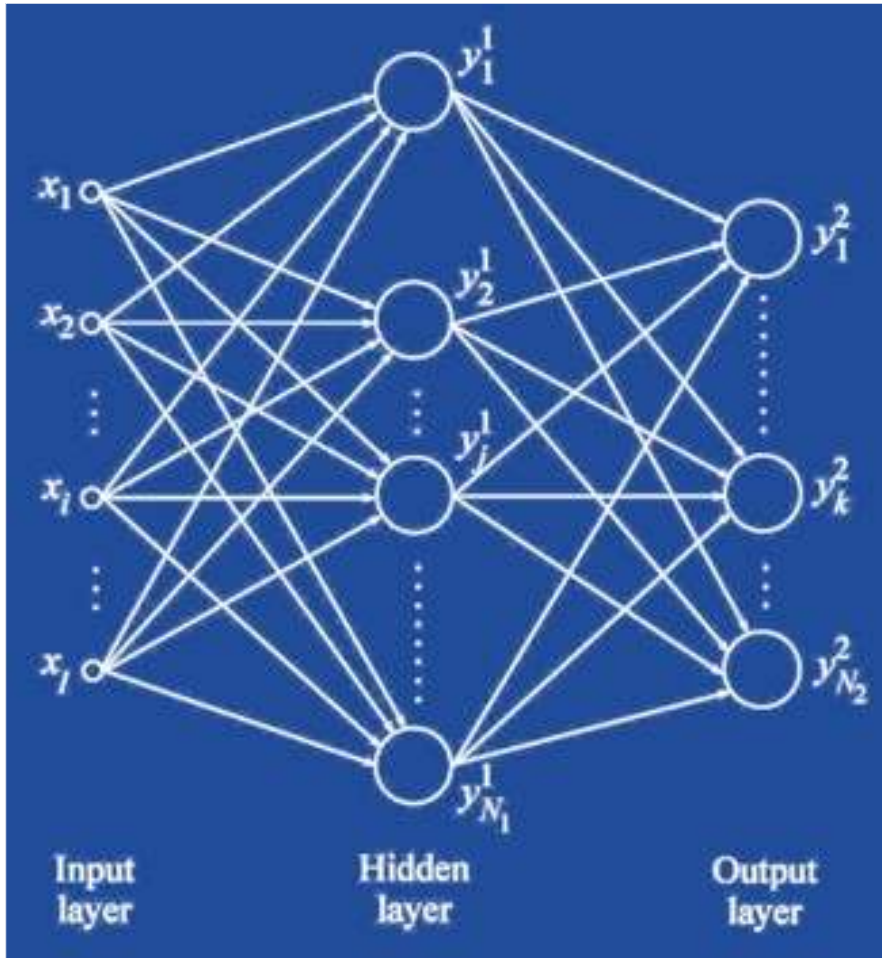
”Multiply and Accumulate”

Activation functions

Name	Plot	Equation
Identity		$f(x) = x$
Binary step		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Logistic (a.k.a. Soft step)		$f(x) = \frac{1}{1 + e^{-x}}$
TanH		$f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$
ArcTan		$f(x) = \tan^{-1}(x)$
Softsign ^{[7][8]}		$f(x) = \frac{x}{1 + x }$
Rectified linear unit (ReLU) ^[9]		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$

Source: Wikipedia

Neural networks



$$X = \begin{bmatrix} X_{11}, X_{12}, \dots, X_{1I} \\ X_{21}, X_{22}, \dots, X_{2I} \\ \dots \dots \dots \\ X_{m1}, X_{m2}, \dots, X_{mI} \end{bmatrix}$$

I features

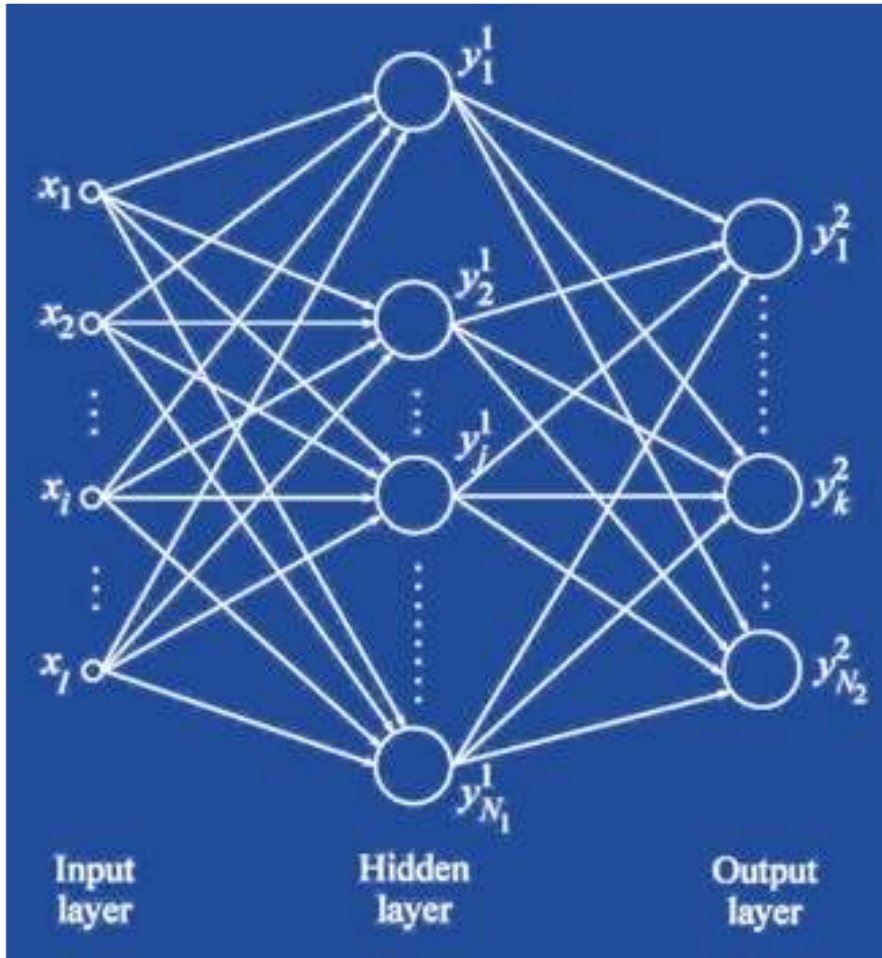
m samples

$$y = \begin{bmatrix} y_1 \\ y_2 \\ \dots \\ y_m \end{bmatrix} = \begin{bmatrix} 0,0,1,0,0,\dots,0 \\ 1,0,0,0,0,\dots,0 \\ \dots \\ 0,0,0,0,1,\dots,0 \end{bmatrix}$$

m labels,
 N_2 outputs

One-hot encoding

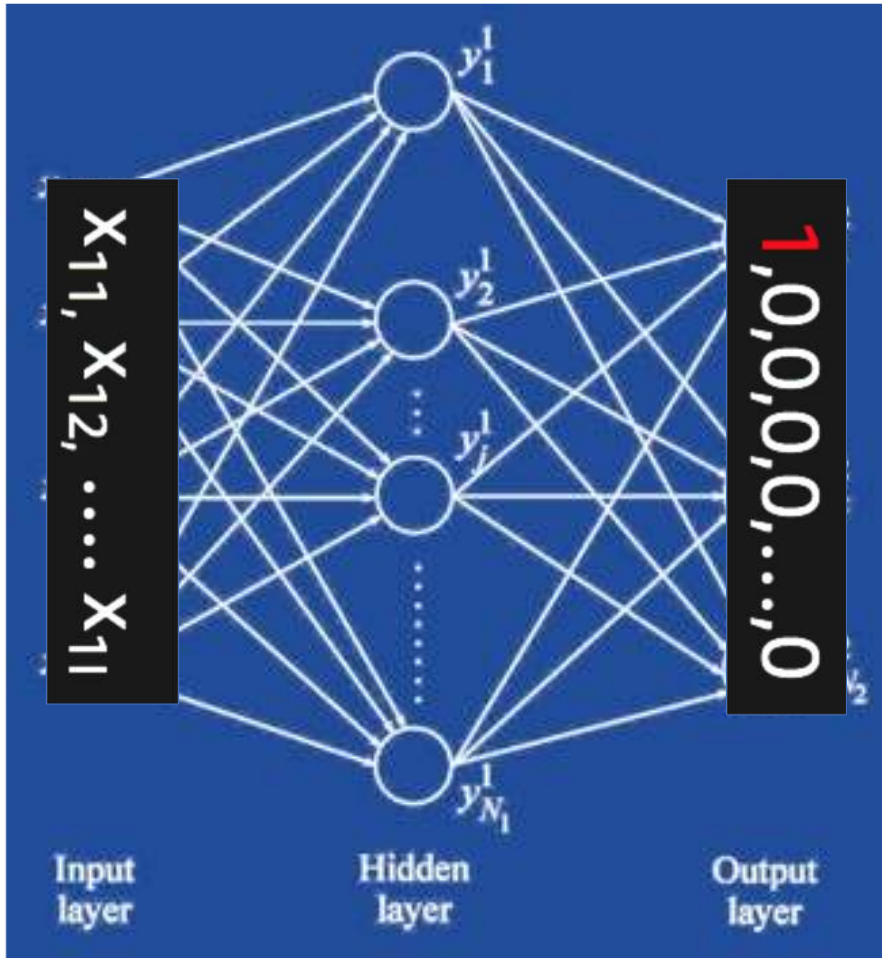
Neural networks



$$\begin{aligned}
 \mathbf{X} &= \begin{bmatrix} x_{11}, x_{12}, \dots, x_{1I} \\ x_{21}, x_{22}, \dots, x_{2I} \\ \vdots \\ x_{m1}, x_{m2}, \dots, x_{mI} \end{bmatrix} \quad \begin{matrix} I \text{ features} \\ m \text{ samples} \end{matrix} \\
 \mathbf{y} &= \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{bmatrix} = \begin{bmatrix} 0, 0, 1, \dots, 0 \\ 1, 0, 0, \dots, 0 \\ \vdots \\ 0, 0, 0, \dots, 0 \end{bmatrix} \quad \begin{matrix} m \text{ labels,} \\ N_2 \text{ categories} \end{matrix} \\
 &\quad \text{One-hot encoding}
 \end{aligned}$$

$$\text{Accuracy} = \frac{\text{Number of correct predictions}}{\text{Total number of predictions}}$$

Neural networks



Initially, the network will **not** predict correctly

$$f(X_1) = Y'_1$$

A **loss function** measures the difference between the **real label** Y_1 and the **predicted label** Y'_1

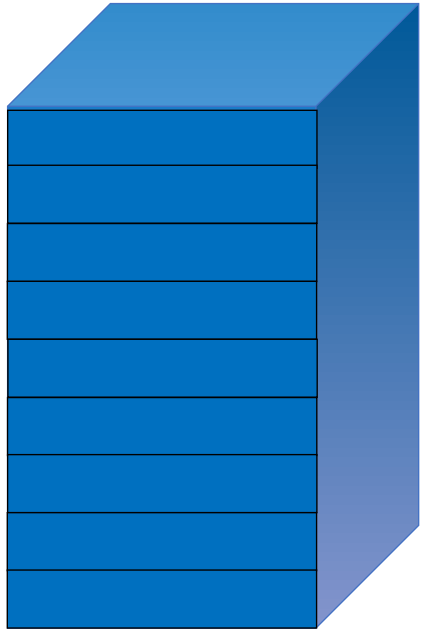
$$\text{error} = \text{loss}(Y_1, Y'_1)$$

For a **batch** of samples:


$$\sum_{i=1}^{\text{batch size}} \text{loss}(Y_i, Y'_i) = \text{batch error}$$

The purpose of the training process is to **minimize error** by gradually **adjusting weights**

Training



Training data set




Trained
neural network

Backpropagation


Batch size
Learning rate
Number of epochs

} Hyper parameters

Validation



Validation data set



Trained
neural network




Validation
accuracy

Prediction at
the end of
each epoch

Save the model at the end of each epoch

Early stopping



Apache MXNet

Apache MXNet: Open Source library for Deep Learning




Programmable

Simple syntax,
multiple
languages




Most Open

Accepted into the
Apache Incubator




Portable

Highly efficient
models for
mobile
and IoT



Best On AWS

Optimized for
Deep Learning on AWS




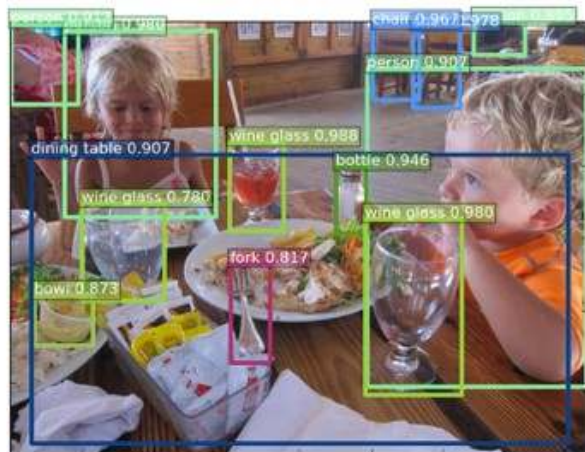

High Performance

Near linear scaling
across hundreds of
GPUs



MXNet 1.0 released on December 4th

Object Detection



<https://github.com/precedenceguo/mx-rcnn>

<https://github.com/zhreshold/mxnet-yolo>

Object Segmentation



<https://github.com/TuSimple/mx-maskrcnn>

Text Detection and Recognition



<https://github.com/Bartzi/stn-ocr>


Face Detection



```
attribution is:  
5_o_Clock_Shadow : No  
Arched_Eyebrows : No  
Attractive : No  
Bags_Under_Eyes : No  
Bald : No  
Bangs : No  
Big_Lips : No  
Big_Nose : No  
Black_Hair : No  
Blond_Hair : No  
Blurry : Yes  
Brown_Hair : No  
Bushy_Eyebrows : No  
Chubby : No  
Double_Chin : No  
Eyeglasses : No  
Goatee : No  
Gray_Hair : No  
Heavy_Makeup : No  
High_Cheekbones : No  
Male : Yes  
Mouth_Slightly_Open : No  
Mustache : No  
Narrow_Eyes : Yes  
No_Beard : Yes  
Oval_Face : No  
Pale_Skin : No  
Pointy_Nose : No  
Receding_Hairline : No  
Rosy_Cheeks : No  
Sideburns : No  
Smiling : No  
Straight_Hair : No  
Wavy_Hair : No  
Wearing_Earrings : No  
Wearing_Hat : No  
Wearing_Lipstick : No  
Wearing_Necklace : No  
Wearing_Necktie : No  
Young : Yes
```


<https://github.com/tornadomeet/mxnet-face>

Real-Time Pose Estimation



https://github.com/dragonfly90/mxnet_Realtime_Multi-Person_Pose_Estimation

Machine Translation



<https://github.com/awslabs/sockeye>

The Apache MXNet API

- Storing and accessing data in multi-dimensional arrays
→ *NDArray* API
- Building models (layers, weights, activation functions)
→ *Symbol* API
- Serving data during training and validation
→ *Iterators*
- Training and using models
→ *Module* API

Demos

<https://github.com/juliensimon/dlnotebooks>

- 1) Synthetic data set
- 2) Classify images with pre-trained models
- 3) Learn MNIST with a Multi-Layer Perceptron
- 4) Learn MNIST with the LeNet CNN
- 5) Predict handmade MNIST samples
- 6) Train and host a MNIST model with Amazon SageMaker

Resources

<https://aws.amazon.com/machine-learning>

<https://aws.amazon.com/blogs/ai>

<https://mxnet.incubator.apache.org>

<https://github.com/apache/incubator-mxnet>

<https://github.com/gluon-api>

<https://github.com/awslabs/sockeye>

An overview of Amazon SageMaker <https://www.youtube.com/watch?v=ym7NEYEx9x4>

<https://medium.com/@julsimon>



Thank you!

**Julien Simon, AI Evangelist,
EMEA
@julsimon**