

Advanced Task Scheduling with Amazon ECS and Blox

Julien Simon
Principal Technical Evangelist, AWS
@julsimon

Docker on Amazon Web Services

Amazon EC2 Container Service (ECS)

- <https://aws.amazon.com/ecs/>
- Launched in 04/2015
- No additional charge
- Since June: integration with **Spot instances**
- October 2nd: **per-second billing** for EC2 and EBS

Amazon EC2 Container Registry (ECR)

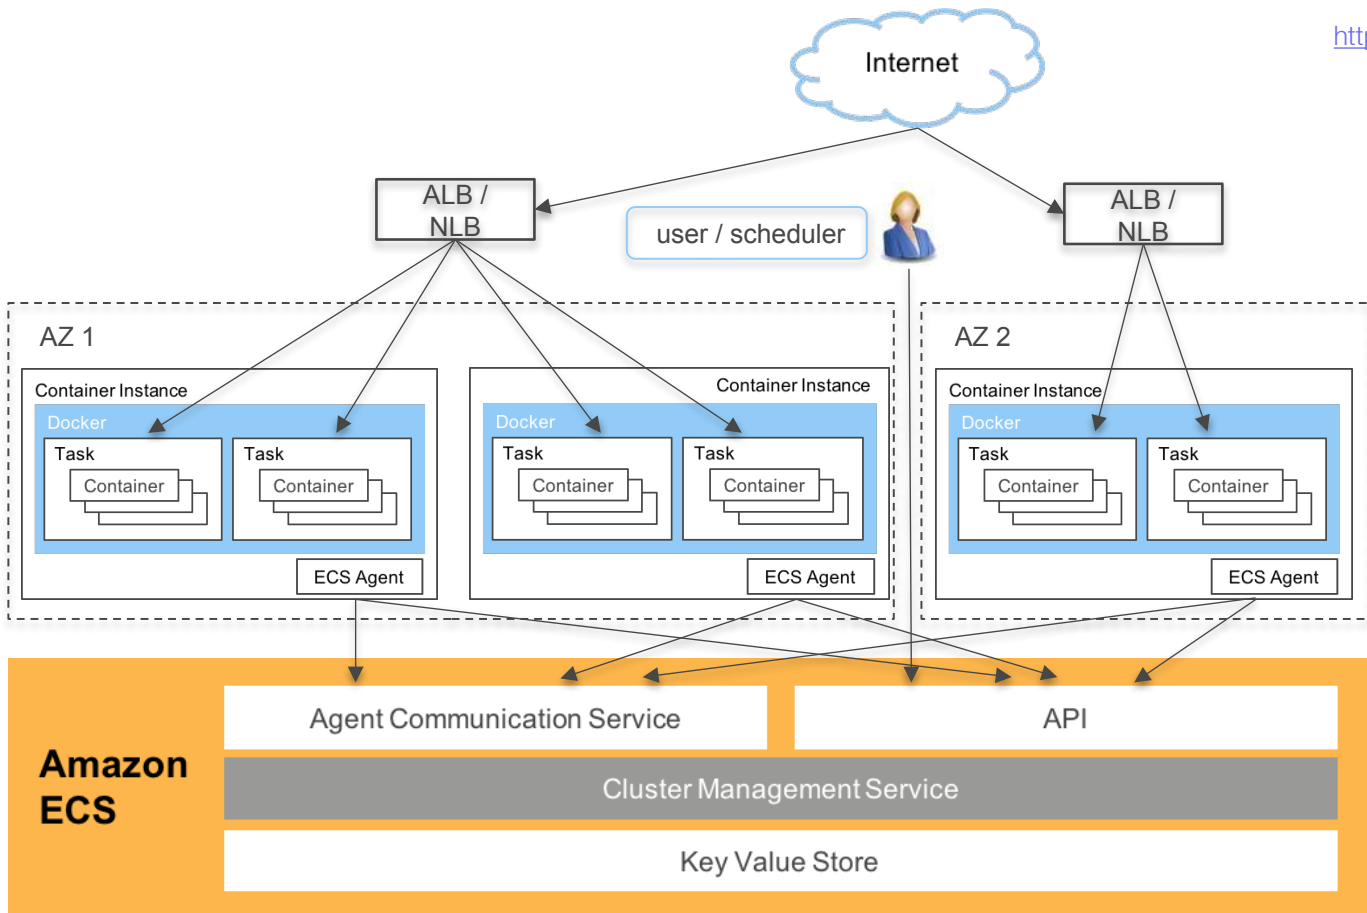
- <https://aws.amazon.com/ecr/>
- Launched in 12/2015
- Free tier: 500MB / month for a year
- \$0.10 / GB / month + outgoing traffic

ECS & ECR are available in 12 regions (US, EU, APAC, China)



Amazon ECS: Under the Hood

<https://github.com/aws/amazon-ecs-agent>



A close-up photograph of a red, circular engine start button. The button is set into a dark, metallic-looking dashboard. It has a silver-colored ring around its perimeter. The words "ENGINE" and "START" are printed in white, bold, sans-serif capital letters on the red surface. The button is slightly recessed, and there are concentric circular ridges on the dashboard around it. To the right, a portion of another circular control, possibly a gear shifter or a knob, is visible.

ENGINE
START

Selected ECS customers

Case studies and re:Invent videos: <https://aws.amazon.com/ecs/>

airtime



GILT



here



meetup

Aol.



okta



Blackboard



SUP
ERC
ELL



Container Partners



ECS Scheduling

The problem

Given a certain amount of
computing power and memory,

how can we best manage
an arbitrary number of apps
running in Docker containers?



Case study: Coursera



<https://www.youtube.com/watch?v=a45J6xAGUvA>

Coursera deliver **Massive Open Online Courses** (14 million students, 1000+ courses). Their platform runs a large number of batch jobs, notably to **grade programming assignments**. Grading jobs need to run in **near-real time** while preventing execution of **untrusted code** inside the Coursera platform.

After trying out some other Docker solutions, Coursera have picked **Amazon ECS** and have even written **their own scheduler**.

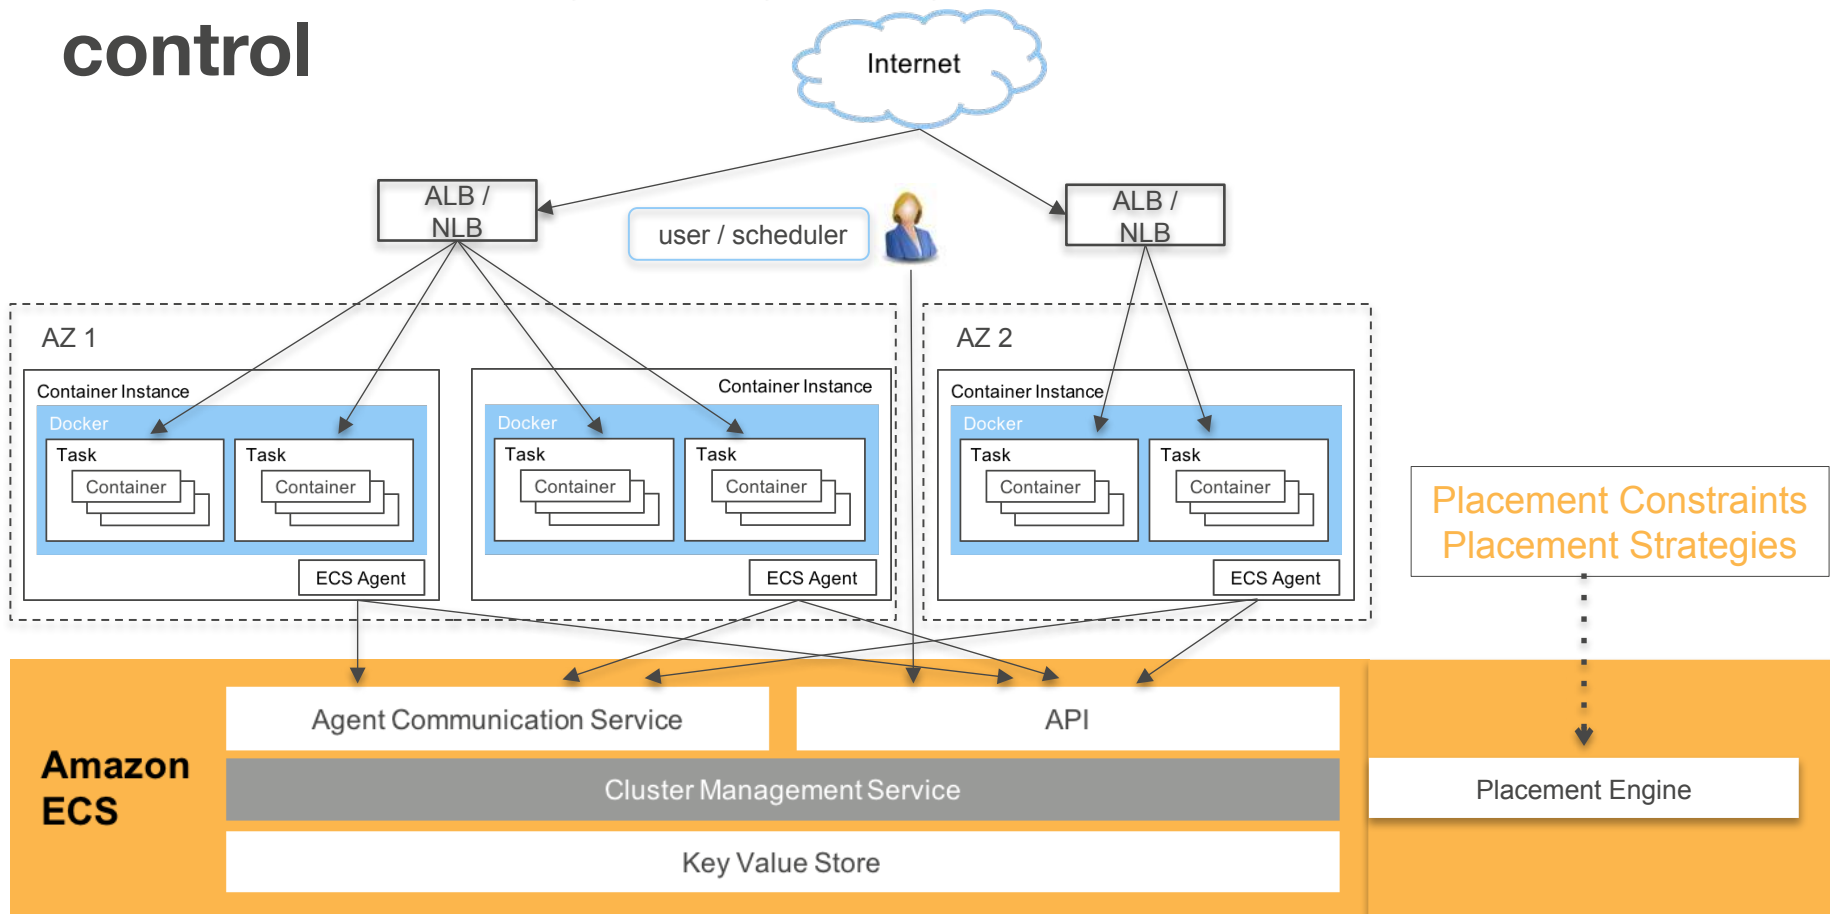
“Amazon ECS enabled Coursera to focus on releasing new software rather than spending time managing clusters” - Frank Chen, Software Engineer

Scheduling on ECS: two options so far


1. Let **ECS** handle scheduling through **Services**
 - Task Definition
 - ECS equivalent of the Docker Compose file
 - Versioned
 - `cpu_shares`, `mem_limit`
 - Number of containers
2. Implement a **custom scheduler** with the **ECS API**
 - Describe cluster state
 - Select a specific ECS instance according to custom logic
 - Run task on this instance

ECS Placement Engine

Placement Engine: giving developers more control



Placement Constraints



Name	Example
✓ AMI ID	<code>attribute:ecs.ami-id == ami-eca289fb</code>
✓ Availability Zone	<code>attribute:ecs.availability-zone == us-east-1a</code>
✓ Instance Type	<code>attribute:ecs.instance-type == t2.small</code>
✓ Distinct Instances	<code>type="distinctInstance"</code>
✓ Custom	<code>attribute:stack == prod</code>

Example: Constraint on Instance Family/Type

```
aws ecs list-container-instances --cluster ecs-demo --filter "attribute:ecs.instance-type matches t2.*"
{
  "containerInstanceArns": [
    "arn:aws:ecs:us-east-1:123456789000:container-instance/3ced5d42-537c-40b4-9551-b9022cc13b78",
    "arn:aws:ecs:us-east-1:123456789000:container-instance/442d988b-4b00-40bf-85ae-34e0819454f2",
    "arn:aws:ecs:us-east-1:123456789000:container-instance/5dell1ede-6c22-411e-a469-8301eeebae0f",
    "arn:aws:ecs:us-east-1:123456789000:container-instance/6bfb12c8-1c3c-4d4a-976c-ce3c2c79b031",
    "arn:aws:ecs:us-east-1:123456789000:container-instance/7eb87781-abab-4a6a-9a0d-602a4da59549",
    "arn:aws:ecs:us-east-1:123456789000:container-instance/af8d48ba-73c4-409a-b40f-66596aa86c5d",
    "arn:aws:ecs:us-east-1:123456789000:container-instance/b5c08e3e-bd25-4ec9-9a88-ce1d53640542",
  ]
}
```

```
aws ecs list-container-instances --cluster ecs-demo --filter "attribute:ecs.instance-type matches t2.small"
{
  "containerInstanceArns": [
    "arn:aws:ecs:us-east-1:123456789000:container-instance/af8d48ba-73c4-409a-b40f-66596aa86c5d",
    "arn:aws:ecs:us-east-1:123456789000:container-instance/b5c08e3e-bd25-4ec9-9a88-ce1d53640542",
  ]
}
```

Example: Constraint on Availability Zone

```
aws ecs list-container-instances --cluster ecs-demo --filter "attribute:ecs.availability-zone matches us-east-1.*"
{
  "containerInstanceArns": [
    "arn:aws:ecs:us-east-1:123456789000:container-instance/3ced5d42-537c-40b4-9551-b9022cc13b78",
    "arn:aws:ecs:us-east-1:123456789000:container-instance/442d988b-4b00-40bf-85ae-34e0819454f2",
    "arn:aws:ecs:us-east-1:123456789000:container-instance/5de11ede-6c22-411e-a469-8301eeebae0f",
    "arn:aws:ecs:us-east-1:123456789000:container-instance/6bfb12c8-1c3c-4d4a-976c-ce3c2c79b031",
    "arn:aws:ecs:us-east-1:123456789000:container-instance/7eb87781-abab-4a6a-9a0d-602a4da59549",
    "arn:aws:ecs:us-east-1:123456789000:container-instance/af8d48ba-73c4-409a-b40f-66596aa86c5d",
  ]
}
```

```
aws ecs list-container-instances --cluster ecs-demo --filter "attribute:ecs.availability-zone == us-east-1a"
{
  "containerInstanceArns": [
    "arn:aws:ecs:us-east-1:123456789000:container-instance/3ced5d42-537c-40b4-9551-b9022cc13b78",
    "arn:aws:ecs:us-east-1:123456789000:container-instance/442d988b-4b00-40bf-85ae-34e0819454f2",
    "arn:aws:ecs:us-east-1:123456789000:container-instance/5de11ede-6c22-411e-a469-8301eeebae0f",
  ]
}
```

Example: Combining Multiple Constraints


```
aws ecs list-container-instances --cluster ecs-demo --filter "attributes:ecs.instance-type matches t2.* and attribute:ecs.availability-zone == us-east-1a"
```

```
{
  "containerInstanceArns": [
    "arn:aws:ecs:us-east-1:123456789000:container-instance/3ced5d42-537c-40b4-9551-b9022cc13b78",
    "arn:aws:ecs:us-east-1:123456789000:container-instance/442d988b-4b00-40bf-85ae-34e0819454f2",
    "arn:aws:ecs:us-east-1:123456789000:container-instance/5dellede-6c22-411e-a469-8301eeebae0f",
  ]
}
```


```
aws ecs list-container-instances --cluster ecs-demo --filter "(attribute:ecs.instance-type in [t2.small, t2.medium] or attribute:ecs.instance-type matches g2.*) and attribute:ecs.availability-zone != us-east-1d"
```

```
{
  "containerInstanceArns": [
    "arn:aws:ecs:us-east-1:123456789000:container-instance/3ced5d42-537c-40b4-9551-b9022cc13b78",
    "arn:aws:ecs:us-east-1:123456789000:container-instance/442d988b-4b00-40bf-85ae-34e0819454f2",
    "arn:aws:ecs:us-east-1:123456789000:container-instance/5dellede-6c22-411e-a469-8301eeebae0f",
    "arn:aws:ecs:us-east-1:123456789000:container-instance/6bfb12c8-1c3c-4d4a-976c-ce3c2c79b031",
    "arn:aws:ecs:us-east-1:123456789000:container-instance/7eb87781-abab-4a6a-9a0d-602a4da59549",
    "arn:aws:ecs:us-east-1:123456789000:container-instance/d45f5b92-4faa-44a9-a9a7-2d744566e510",
    "arn:aws:ecs:us-east-1:123456789000:container-instance/f3d92b17-7d95-4cff-b623-390e871c6b60",
    "arn:aws:ecs:us-east-1:123456789000:container-instance/f7858158-5806-4d8d-82ea-f0eb4680e6cf",
    "arn:aws:ecs:us-east-1:123456789000:container-instance/fc751042-590a-440d-90a7-e8ebce02d234"
  ]
}
```



Placement Strategies




Binpacking



Spread




Affinity



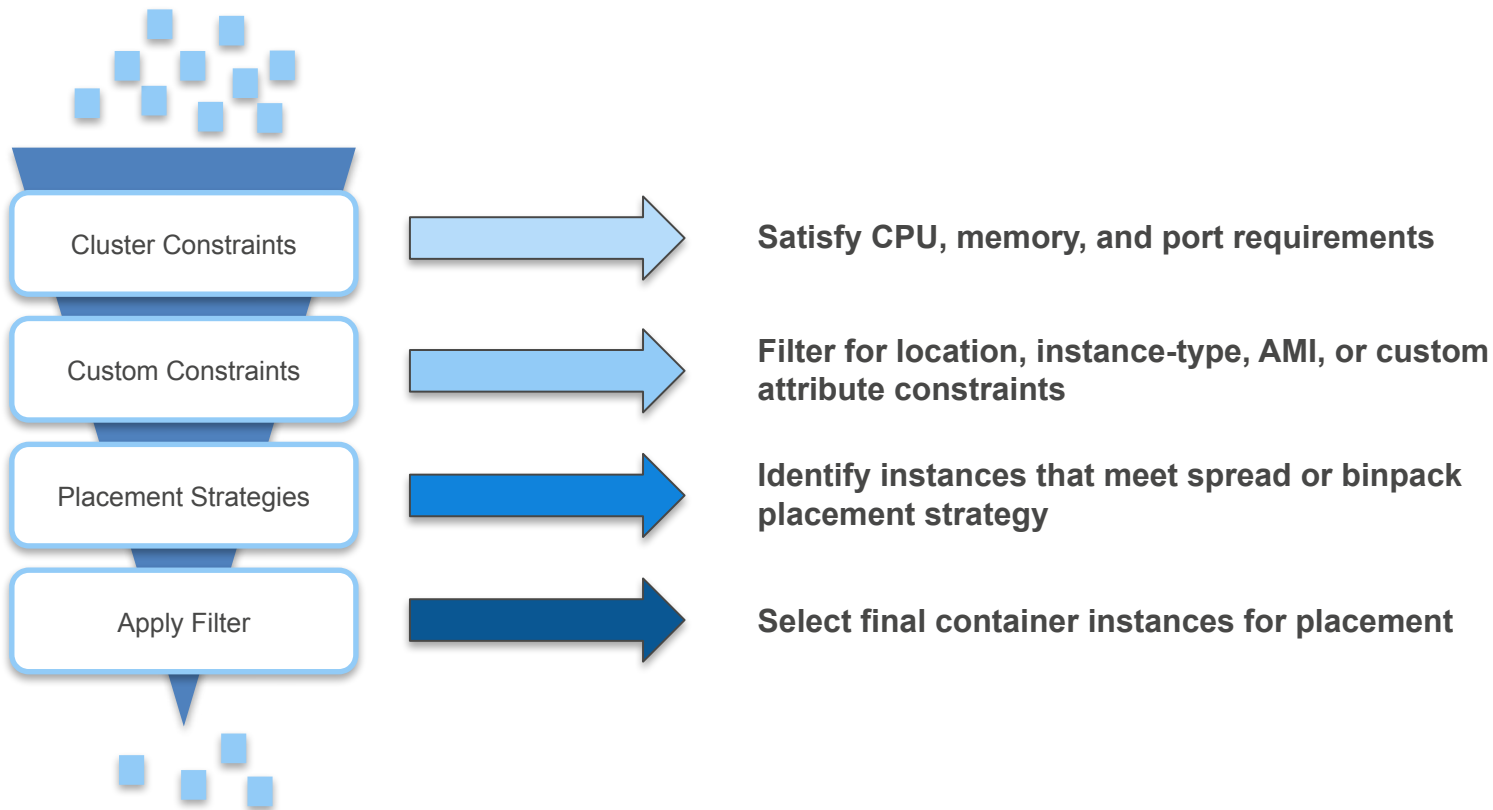
Distinct Instance

Placement Strategy Chaining




Placing Tasks

Anatomy of Task Placement




Placement: Targeting Instance Type


```
aws ecs run-task --cluster ecs-demo --task-definition myapp --count 5 --placement-constraints  
type="memberOf",expression="attribute:ecs.instance-type == g2.2xlarge"
```




g2.2xlarge



t2.small



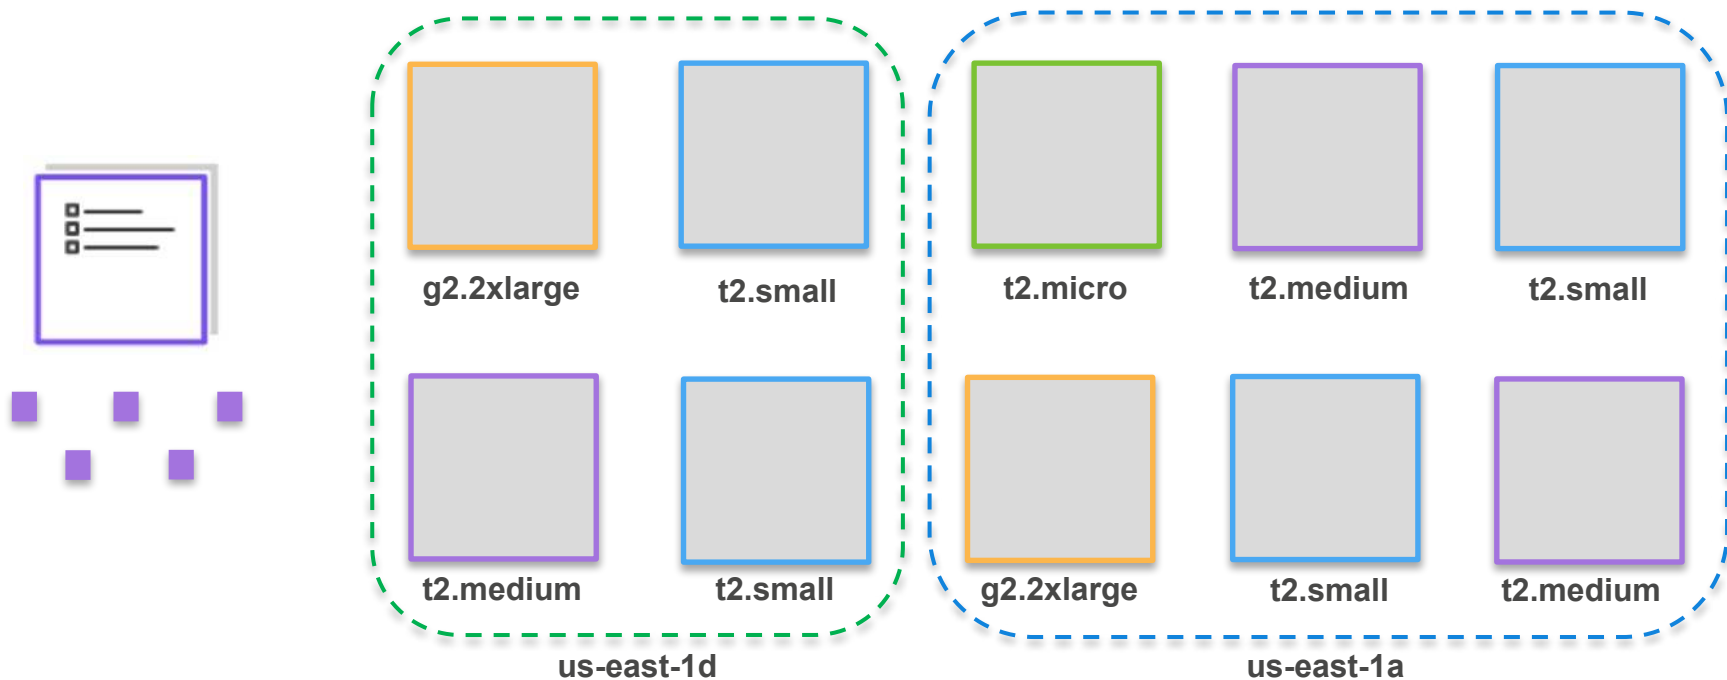
g2.2xlarge



g2.2xlarge


Placement: Targeting Instance Type & Zone

```
aws ecs run-task --cluster ecs-demo --task-definition myapp --count 5 --placement-constraints  
type="memberOf",expression="(attribute:ecs.instance-type == t2.small or  
attribute:ecs.instance-type == t2.medium) and attribute:ecs.availability-zone != us-east-1d"
```




Placement: Availability Zone Spread

```
aws ecs run-task --cluster ecs-demo --task-definition myapp --count 9 --placement-strategy  
type="spread",field="attribute:ecs.availability-zone"
```



Placement: Spread across Zone and Binpack

```
aws ecs run-task --cluster ecs-demo --task-definition myapp --count 9 --placement-strategy  
type="spread",field="attribute:ecs.availability-zone" type="binpack",field="memory"
```




Placement: Multiple Services on a Cluster



```
aws ecs create-service --service-name srvc-binpck --cluster ecs-demo --task-definition myapp-binpck  
--desired-count 5 --placement-strategy type="binpack",field="memory"
```





```
aws ecs create-service --service-name srvc-spread --cluster ecs-demo --task-definition myapp-spread  
--desired-count 6 --placement-strategy type="spread",field="attribute:ecs.availability-zone"
```



Placement: Services – Distinct Instances

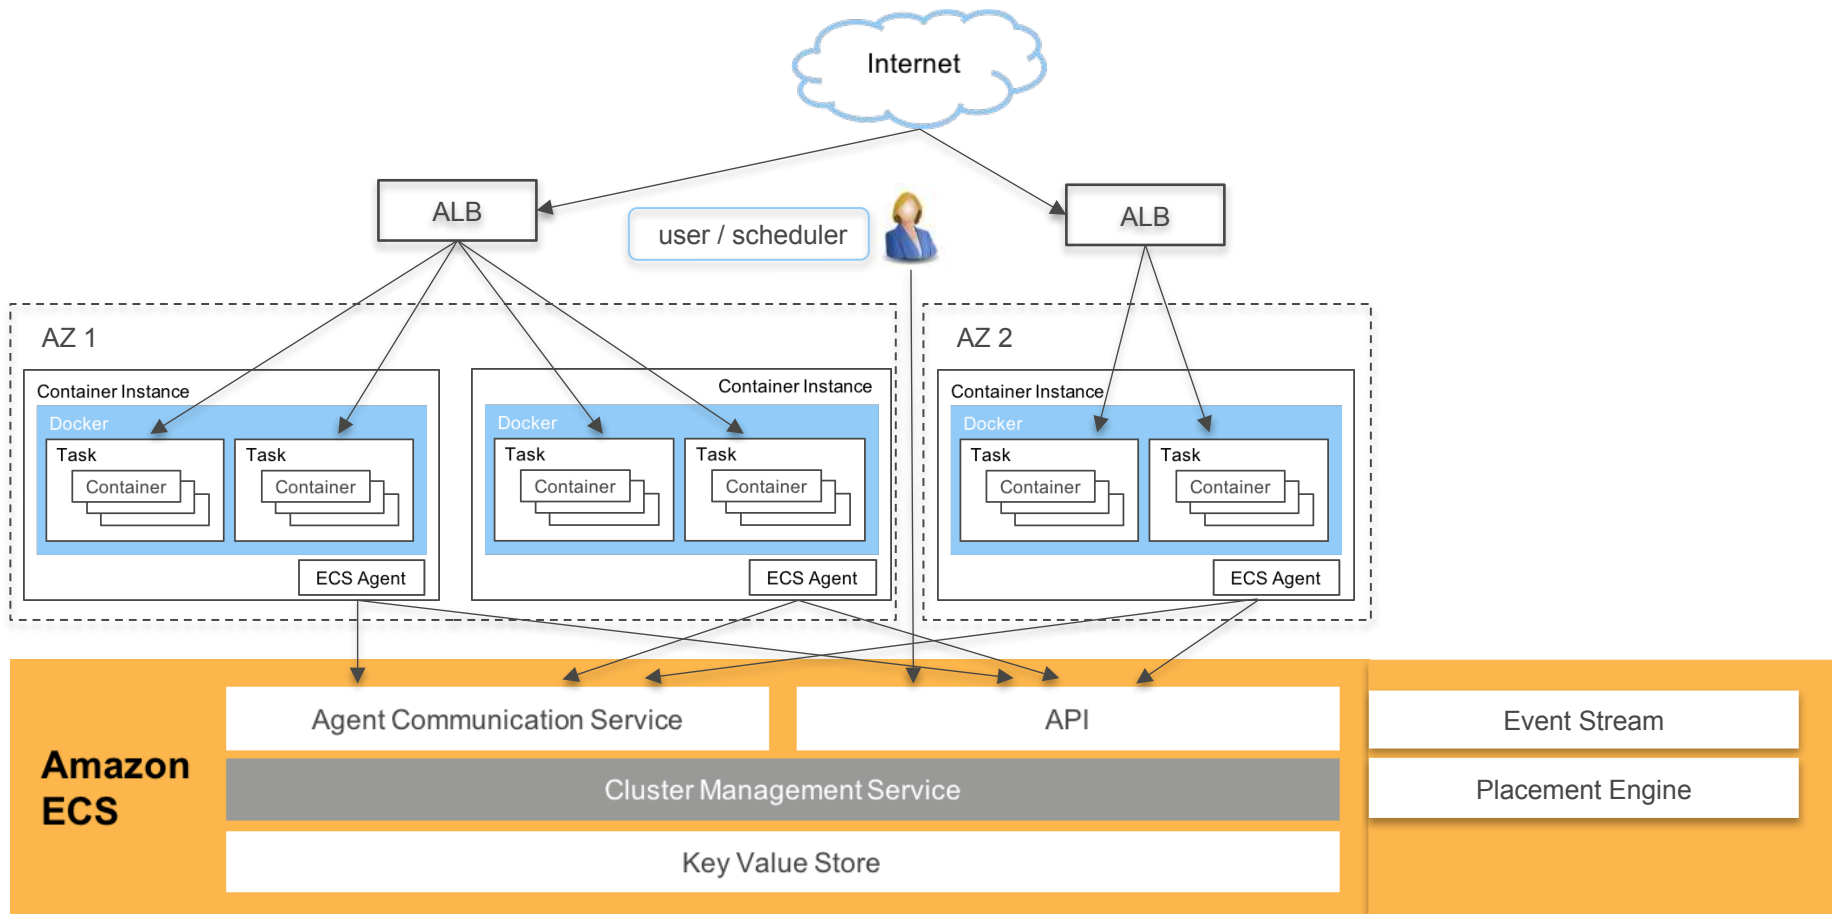
```
aws ecs create-service --service-name myapp-gpu --cluster ecs-demo --task-definition myapp-gpu --desired-count 3 --placement-constraints type="memberOf",expression="attribute:ecs.instance-type =~ g2.*" type="distinctInstance"
```

```
aws ecs create-service --service-name myapp --cluster ecs-demo --task-definition myapp --desired-count 9 --placement-constraints type="memberOf",expression="(attribute:ecs.instance-type == t2.small or attribute:ecs.instance-type == t2.medium)" type="distinctInstance"
```




Event Stream & Blox

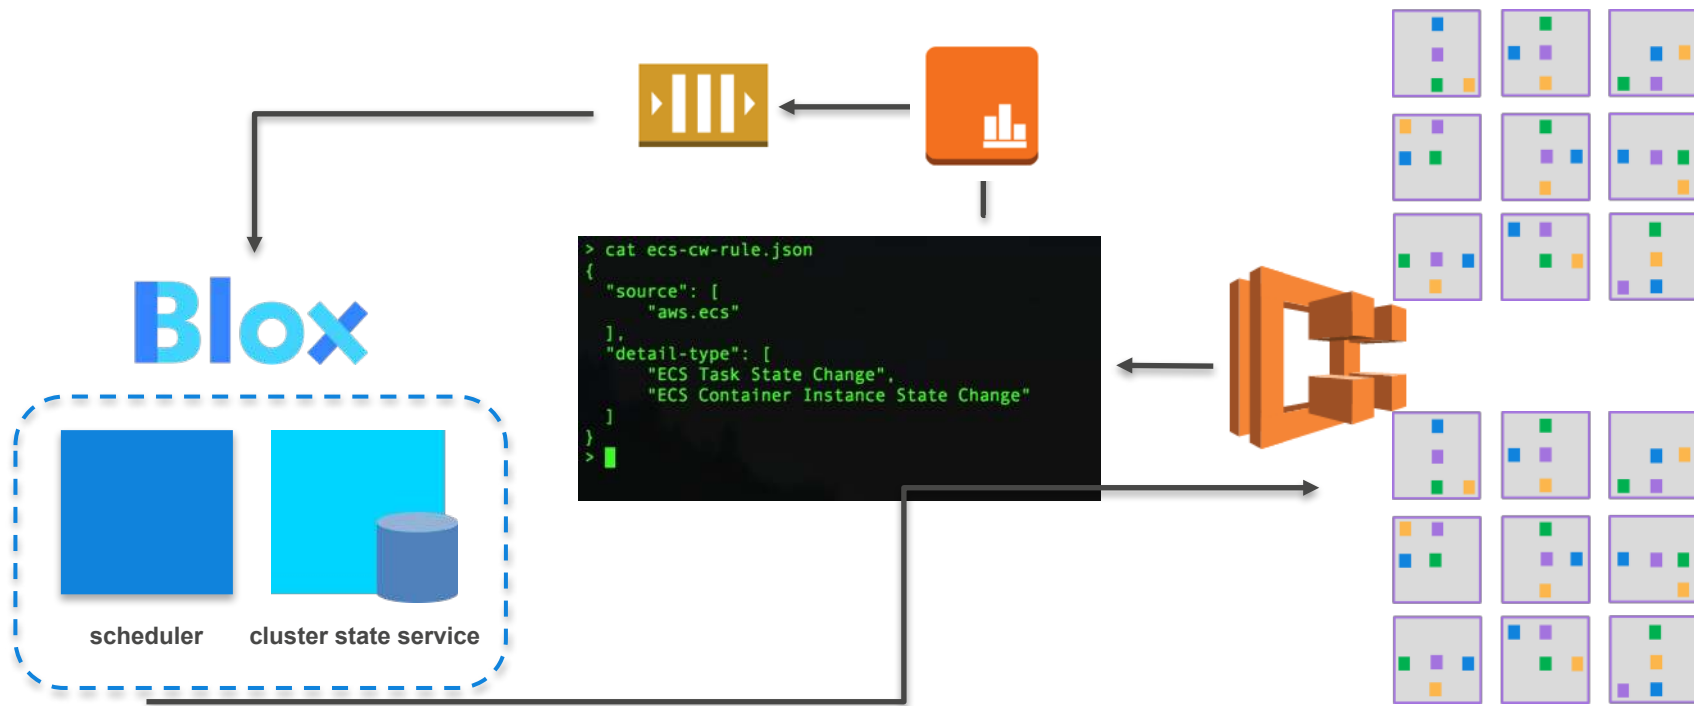
Amazon ECS: Under the Hood



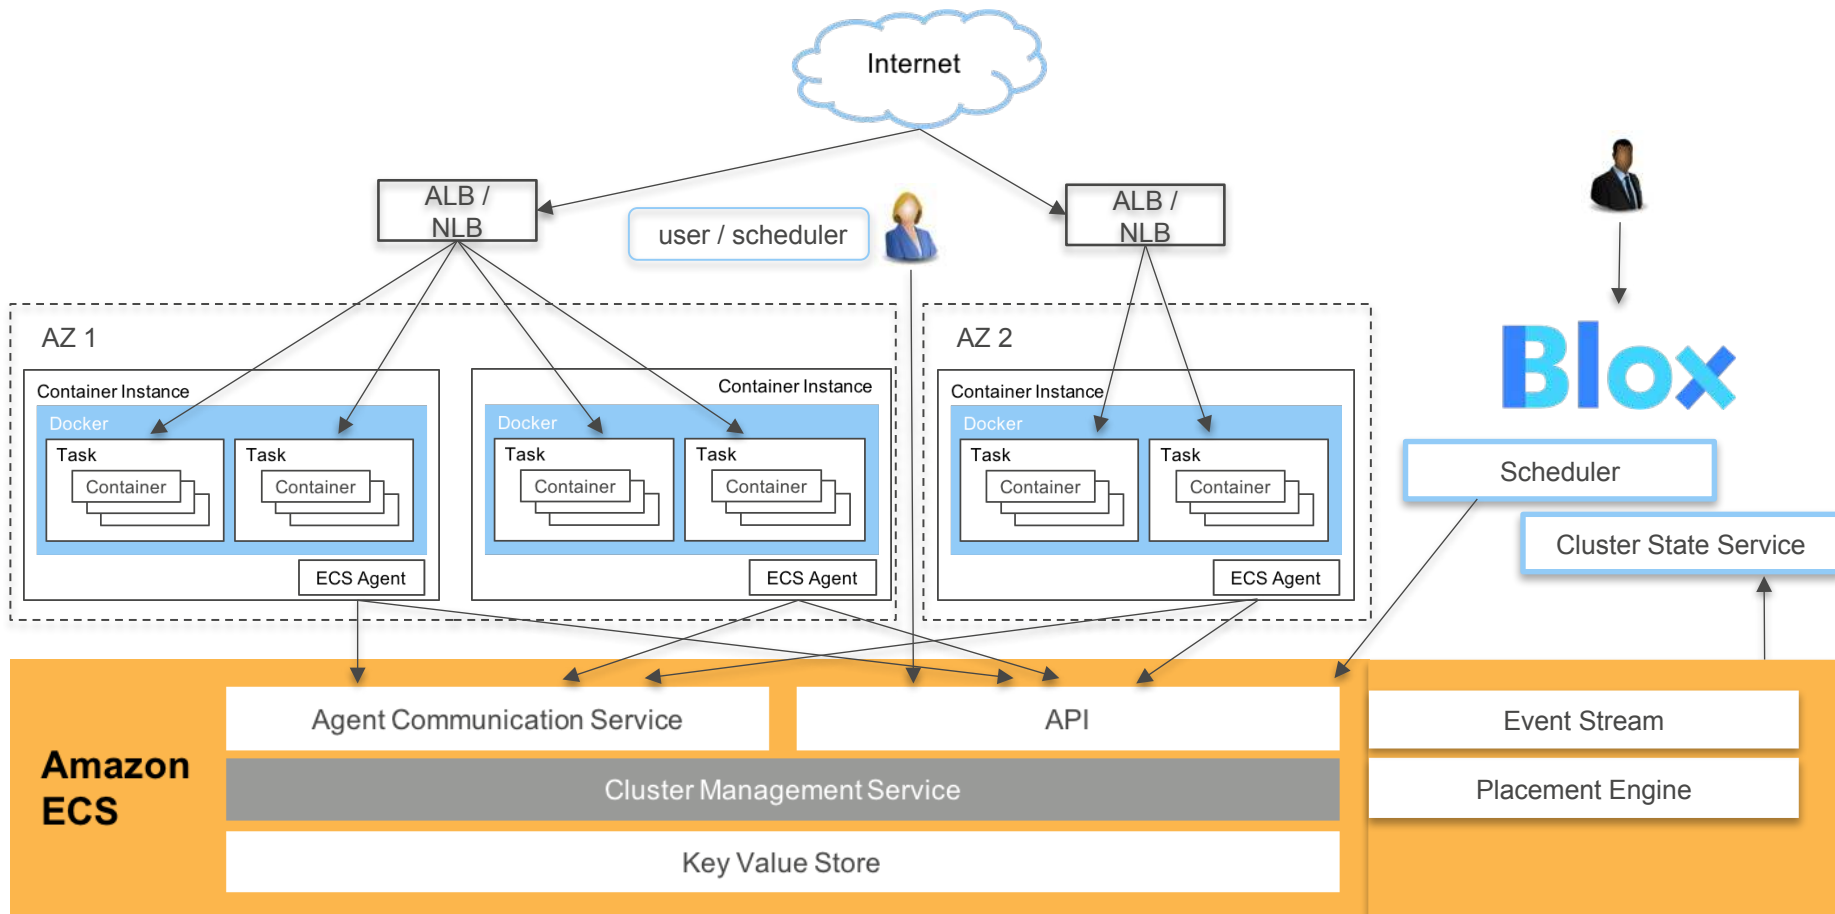
Consuming Real-time Events



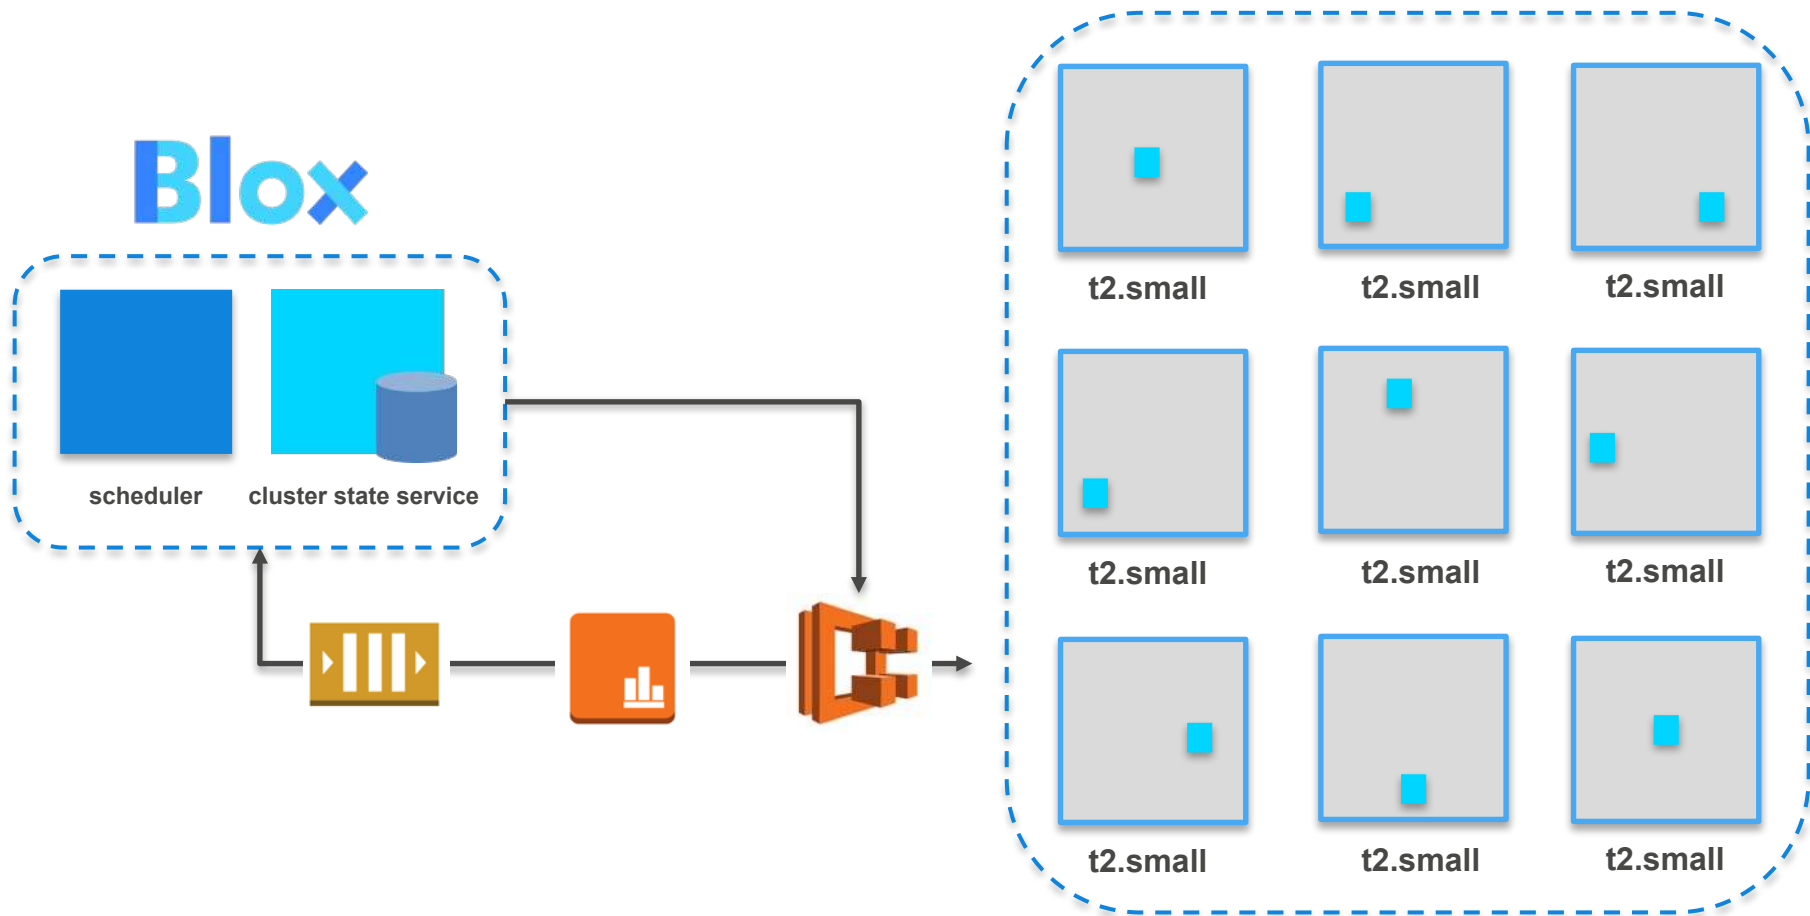
Handling ECS events with Blox



Amazon ECS: Under the Hood



Blox: Daemon Scheduler



Demo: Deploying Blox on AWS

<https://github.com/juliensimon/aws/tree/master/blox>

Creating Clusters

Create an ECS cluster for Blox

CF template: https://github.com/blox/blox/blob/dev/deploy/aws/conf/cloudformation_template.json

- CloudWatch Event Rule + SQS queue
- Daemon Scheduler + Cluster State Service + etcd
- REST API exposing the Daemon Scheduler API

Create another ECS cluster managed by Blox

```
$ ecs-cli configure --cluster WebCluster --region ap-southeast-1
```

```
$ ecs-cli up --keypair admin --capability-iam --size 3 --instance-type t2.micro
```

Invoke the scheduler API

‘demo-cli’ tool: <https://github.com/blox/blox/tree/dev/deploy/demo-cli>

Listing Task Definitions

Grab the ARN for an *nginx* Task Definition, which the Daemon Scheduler will manage on 'WebCluster'.

```
$ ./list-task-definitions.py --region ap-southeast-1
```

```
== Blox Demo CLI - List Task Definitions ==
```

```
{  
  "taskDefinitionArns": [  
    "arn:aws:ecs:ap-southeast-1:ACCOUNT:task-definition/BloxFramework:2",  
    "arn:aws:ecs:ap-southeast-1:ACCOUNT:task-definition/nginx:1",  
    "arn:aws:ecs:ap-southeast-1:ACCOUNT:task-definition/nginx:2"  
  ]  
}
```

Creating an Environment

```
$ ./blox-create-environment.py --environment WebEnvironment --cluster  
WebCluster --task-definition "arn:aws:ecs:ap-southeast-1:ACCOUNT:task-  
definition/nginx:2" --stack Blox --apigateway --region ap-southeast-1
```

== Blox Demo CLI - Create Blox Environment ==

HTTP Response Code: 200

```
{  
  "taskDefinition":  
    "arn:aws:ecs:ap-southeast-1:ACCOUNT:task-definition/nginx:2",  
    "deploymentToken": "17248257-08ec-4438-888f-e0ac28397653",  
    "health": "healthy",  
    "name": "WebEnvironment",  
    "instanceGroup": {  
      "cluster": "arn:aws:ecs:ap-southeast-1:ACCOUNT:cluster/WebCluster"  
    }  
}
```

Listing Environments

```
$ ./blox-list-environments.py --stack Blox --apigateway --region ap-southeast-1
```

```
== Blox Demo CLI - List Blox Environments ==
```

```
HTTP Response Code: 200
```

```
{
  "items": [
    {
      "taskDefinition":
"arn:aws:ecs:ap-southeast-1:ACCOUNT:task-definition/nginx:2",
      "deploymentToken": "17248257-08ec-4438-888f-e0ac28397653",
      "health": "healthy",
      "name": "WebEnvironment",
      "instanceGroup": {
        "cluster": "arn:aws:ecs:ap-southeast-1:ACCOUNT:cluster/WebCluster"
      }
    }
  ]
}
```

Creating a Deployment

```
$ ./blox-create-deployment.py --environment WebEnvironment --deployment-token  
"17248257-08ec-4438-888f-e0ac28397653" --stack Blox --apigateway --region ap-  
southeast-1
```

== Blox Demo CLI - Create Blox Deployment ==

HTTP Response Code: 200

```
{  
  "status": "pending",  
  "environmentName": "WebEnvironment",  
  "id": "7a05ea99-27a9-4339-a7a6-f4120065aea3",  
  "failedInstances": [],  
  "taskDefinition":  
    "arn:aws:ecs:ap-southeast-1:613904931467:task-definition/nginx:2"  
}
```

Listing Deployments

```
$ ./blox-list-deployments.py --environment WebEnvironment --stack Blox --  
apigateway --region ap-southeast-1
```

```
== Blox Demo CLI - List Blox Deployments ==
```

```
HTTP Response Code: 200
```

```
{  
  "items": [  
    {  
      "status": "completed",  
      "environmentName": "WebEnvironment",  
      "id": "7a05ea99-27a9-4339-a7a6-f4120065aea3",  
      "failedInstances": [],  
      "taskDefinition":  
"arn:aws:ecs:ap-southeast-1:ACCOUNT:task-definition/nginx:2"  
    }  
  ]  
}
```

Scaling a Deployment

```
$ ecs-cli ps
```

Name	State	Ports	TaskDefinition
26313cbe-d929-49de-9cc3-873bf5f32a91/nginx	RUNNING		nginx:2
98442432-fd5c-434d-b93c-0737bd06aaab/nginx	RUNNING		nginx:2
ce9bf217-4b34-4f31-9c7b-a8c3402f1ffd/nginx	RUNNING		nginx:2

```
$ ecs-cli scale --size 4 --capability-iam
```

```
$ ecs-cli ps
```

Name	State	Ports	TaskDefinition
26313cbe-d929-49de-9cc3-873bf5f32a91/nginx	RUNNING		nginx:2
98442432-fd5c-434d-b93c-0737bd06aaab/nginx	RUNNING		nginx:2
c404ac9a-0948-4cc8-b5b0-2238ccdf4035/nginx	RUNNING		nginx:2
ce9bf217-4b34-4f31-9c7b-a8c3402f1ffd/nginx	RUNNING		nginx:2

Additional resources

Tech articles by Werner Vogels, CTO, Amazon.com

<http://www.allthingsdistributed.com/2014/11/amazon-ec2-container-service.html>

<http://www.allthingsdistributed.com/2015/04/state-management-and-scheduling-with-ecs.html>

<http://www.allthingsdistributed.com/2015/07/under-the-hood-of-the-amazon-ec2-container-service.html>

Blox

<https://blox.github.io/>

Amazon ECS videos @ AWS re:Invent 2016

<https://aws.amazon.com/blogs/compute/amazon-ec2-container-service-at-aws-reinvent-2016-wrap-up/>

What's new on Amazon ECS

<https://aws.amazon.com/ecs/release-notes/>

<https://aws.amazon.com/blogs/compute/powering-your-amazon-ecs-cluster-with-amazon-ec2-spot-instances/>

Thank you!

@julsimon

<http://aws.amazon.com/evangelists/julien-simon>