



Machine Learning sur AWS

Julien Simon

Principal Technical Evangelist, AI and Machine Learning

@julsimon

A Long History of ML at Amazon



Personalized
recommendations



Fulfillment
automation and
inventory
management



Drones



Voice driven
interactions



Inventing entirely
new customer
experiences

Our mission

Put Machine Learning in the hands
of every developer and data scientist

More ML is built on AWS than anywhere else



More ML is built on AWS than anywhere else





Le marketing du futur ? Le marketing individualisé

Olivier Marc, co-fondateur et CEO, Advalo

Le marketing de masse est polluant, intrusif, peu efficace

**En moyenne,
99%
des personnes
ciblées ne
réagissent pas**




**Organisées en silo, les entreprises sont incapables
d'accompagner le consommateur au long de son parcours**

**En moyenne,
1%
des leads
sont convertis**



Le marketing de demain ? Un marketing individualisé qui intervient uniquement au bon moment



**En moyenne,
70% des
consommateurs
ont pris leur
décision d'achat
AVANT d'aller en
magasin**

La combinaison data / IA est la seule solution pour répondre à ces nouveaux enjeux





- Création 2014, Paris/Rennes, 40+ personnes
- Editeur de la première plateforme de marketing individualisé
- Hyper croissance
- Série A de 5 millions d'euros fin 2017

CACHE CACHE.

Cuisines
ixina

MORGAN
MORGAN DE TOI



DEVREO 1902

ŠKODA



CÔTÈSLISHI
NIKKEI FROM PERU

Notre plate-forme technique

- Sur AWS depuis le premier jour
- 5 milliards de données traitées par jour
Amazon EMR, Apache Spark
- 14 Teraoctets
Amazon S3, Amazon ElastiCache, Amazon RDS
- 250 sources de données
- 150 modèles prédictifs en production



Amazon
EC2




Amazon
S3



Amazon
RDS



Amazon
ElastiCache



Amazon
EMR

Use case #1 : SEAT

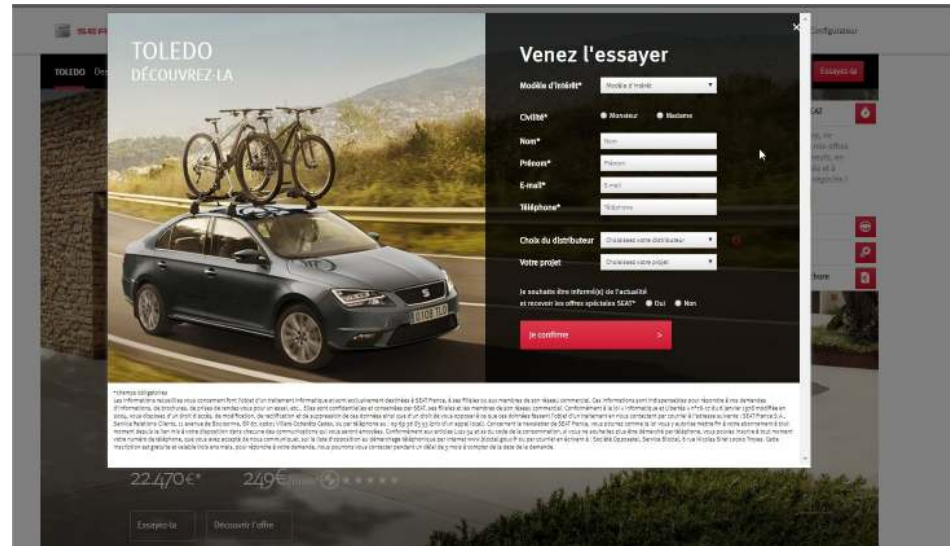


Enjeu : améliorer la **conversion** de *leads*



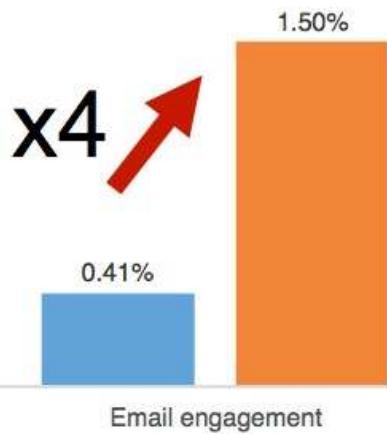
Challenge : exploiter l'ensemble des données pour scorer les **intentionnistes** actifs

Problème : le site génère beaucoup de bruit avec des amateurs de voiture qui n'ont pas particulièrement d'intentions d'achats

Solution : un modèle **prédicatif** en temps réel identifie les intentionnistes



Une efficacité prometteuse



Use case #2 : Bessec

BESSEC

CHAUSSEUR DEPUIS 1862

Enjeu : générer du **trafic** en magasin

Challenge : **individualiser** et **personnaliser** la prise de parole dans le digital comme le ferait un *personal shopper*

Problème : les algorithmes de recommandation classiques ne fonctionnent pas lors des changements de collections

Solution: le **Deep Learning** au service de la détection de styles



Des résultats encourageants

**5% des communications
sont individualisées**

**60% des KPIs
d'engagement**

BESSEC
CHAUSSEUR DEPUIS 1842

[FEMME](#) | [HOMME](#) | [ENFANT](#) | [MAGASINS](#)



Bonjour
Parce-que votre satisfaction est au cœur de nos préoccupations, ce mois-ci nous vous
avons concocté une sélection personnalisée de modèles qui devrait vous plaire.

NOTRE SÉLECTION RIEN QUE POUR VOUS :


MEPHISTO
RAINBOW
185 €
[JE DÉCOUVRE ►](#)


MEPHISTO
VINCENZO
165 €
[JE L'ADORE ►](#)


MEPHISTO
HARRY PERF
160 €
[JE CRAQUE ►](#)

Prochaines étapes

Doublément de l'équipe

Ouverture d'un bureau aux USA début 2019



<http://advalo.com>

Merci !

AWS ML Stack

Application Services

API-driven services: Vision, Language & Speech Services, Chatbots

Platform Services

Deploy machine learning models with high-performance machine learning algorithms, broad framework support, and one-click training, tuning, and inference.


Frameworks & Infrastructure

Develop sophisticated models with any framework, create managed, auto-scaling clusters of GPUs for large scale training, or run prediction on trained models.

Application Services

Amazon Rekognition

Deep Learning-based image analysis service



Object & Scene Detection




City	91.4%
Downtown	91.4%
Metropolis	91.4%
Urban	91.4%
Building	56.2%
High Rise	56.2%
Dock	55.6%
Pier	55.6%
Dawn	52.4%
Dusk	52.4%

Facial Analysis





looks like a face	99.8%
appears to be female	100%
age range	23 - 38 years old
smiling	99.4%
appears to be happy	93.2%
wearing eyeglasses	99.9%
wearing sunglasses	97.6%

Crowd Detection – up to 100 faces




Choose a sample image



Use your own image


Image must be .jpg or .png format and no larger than 5MB. Your image isn't stored.

 **Upload** or drag and drop

Go

Done with the demo?
[Learn more](#)

▼ Results



looks like a face	99.9 %
appears to be male	99.9 %
age range	26 - 43 years old
smiling	98.5 %
appears to be happy	99.7 %
not wearing glasses	99.9 %

[Show more](#)

Facial Search



=



Similarity

97%



≠



≠



Image Moderation



▼ Results

Suggestive	83.5%
Female Swimwear Or Underwear	83.5%


Explicit Nudity

- Nudity
- Graphic Male Nudity
- Graphic Female Nudity
- Sexual Activity
- Partial Nudity

Suggestive

- Female Swimwear or Underwear
- Male Swimwear or Underwear
- Revealing Clothes

Celebrity Recognition



Done with the demo?

[Learn more](#)

▼ Results



Demis Roussos

[Learn More](#)

Match confidence

99 %

► Request

► Response

Text in Image



▼ Results

US English only

| 5T6E652 |


► Request

▼ Response

```
{
  "TextDetections": [
    {
      "Confidence": 96.38984680175781,
      "DetectedText": "5T6E652",
      "Geometry": {
        "BoundingBox": {
          "Height": 0.05606506019830704,
          "Left": 0.38905566930770874,
          "Top": 0.7383729219436646,
          "Width": 0.10844030976295471
        },
        "Polygon": [
          {
            "X": 0.38905566930770874,
            "Y": 0.7383729219436646
          },
          {
            "X": 0.49749597907066345,
            "Y": 0.7726403474807739
          },
          {
            "X": 0.4896218776702881,
            "Y": 0.8287054300308228
          },
          {
            "X": 0.38905566930770874,
            "Y": 0.7383729219436646
          }
        ]
      }
    }
  ]
}
```

Amazon Rekognition Video

Deep Learning-based video analysis service



Video Analysis



Rekognition

Context


Spatial

Temporal

Auditory

Features

Objects, Scenes and Concepts



Person: 98.9%
Plant: 91.4%
Sports: 60.0%

Facial Analysis



Female: 99.9%
Eyes are open: 98.4%
Calm: 86.0%

Face Recognition




Person 1

Activity Detection

Sports, Running

Object Tracking





Sky News will use **Amazon Rekognition** to perform **real-time identification of guests** as they enter St. George's Chapel

Amazon Polly

Deep Learning-based text-to-speech service




Amazon Polly: Text In, Life-like Speech Out

“Salut, je
m’appelle Léa. Je
suis la nouvelle
voix française de
Polly.”



Amazon Polly



A Focus On Voice Quality & Pronunciation


Support for Speech Synthesis Markup Language (SSML) Version 1.0

<https://www.w3.org/TR/speech-synthesis>

<speaking xml:lang="en-US">

The price of this book is <prosody rate="60%">45€</prosody>

</speaking>



Breathing



<https://aws.amazon.com/blogs/machine-learning/amazon-polly-releases-new-ssml-breath-feature/>

<amazon:auto-breaths>

Amazon Translate

Neural Machine Translation Service

Now generally available



FREE TIER

Amazon Translate

Natural and fluent language translation



Real-time
translation



Batch analysis



Automatic language
recognition



Low cost

Languages

Available now

English to/from:

- Spanish
- Portuguese
- German
- French
- Arabic
- Simplified Chinese

Coming soon

- Japanese
- Russian
- Italian
- Traditional Chinese
- Turkish
- Czech

Amazon Transcribe

Automatic speech recognition service

Now generally available



FREE TIER

Amazon Transcribe

Automatic conversion of speech into accurate, grammatically correct text




Multiple
languages



Intelligent punctuation
and formatting



Timestamp
generation



Support for
telephony audio




Recognize multiple
speakers



Custom
vocabulary

Amazon Comprehend


Natural Language Processing



FREE TIER

Fully managed natural language processing

Discover valuable insights from text



API explorer

Paste the text that you would like to analyze with natural language processing.

[Clear text](#)

The accompanying shots from the **Newcastle Herald**'s award-winning team of photographers are worth **more than 1000 words** apiece. But that's the maximum number of words you will require to have a shot at winning the **Herald's sixth annual short story competition**. It's time to let your inner author off the leash and **allow** your imagination to run wild. Stories can be of any genre but must be your own work. The plot must draw inspiration in some way from **one** of these **four images**. Herald photographers **Simone De Peak**, **Marina Neil**, **Jonathan Carroll** and **Max Mason-Hubers** captured the evocative pictures. Stories submitted for the competition must be between **900** and **1000 words**. Judges include **Newcastle Writers Festival** director **Rosemarie Milsom**, **Herald** deputy editor **Matt Carr**, **Hunter Writers Centre** director **Karen Crofts** and president **Megan Buxton**, and **Fairfax Media's** **Newcastle-Hunter** group managing editor **Chad Watson**.

4

916 of 1000 characters used

Language English

Analyze

Entity

This API returns the named entities ("Person", "Organization", "Locations", etc.) within the text you analyzed.

List Tiles JSON Filter Show all categories

Entity	Category	Count	Confidence
Newcastle Herald	Organization	1	0.89
more than 1000 words	Quantity	1	0.91
Herald	Organization	2	0.46
sixth annual short story competition	Quantity	1	0.79
one	Quantity	1	0.4
four images	Quantity	1	0.96
Simone De Peak	Person	1	0.99
Marina Neil	Person	1	0.99+
Jonathan Carroll	Person	1	0.99+
Max Mason-Hubers	Person	1	0.99
900	Quantity	1	0.99

Show all

Topic modeling



**Amazon
Comprehend**

STORM



WORLD
SERIES



STOCK
MARKET



WASHINGTON



AUSTRALIA HEALTH



CRISIS



MACHINE
LEARNING



Platform Services

ML is still too complicated for everyday developers



Collect and prepare
training data



Choose and
optimize your ML
algorithm



Set up and manage
environments for
training



Train and tune model
(trial and error)



Deploy model
in production



Scale and manage
the production
environment

Amazon SageMaker



Pre-built
notebooks for
common
problems



Built-in, high-
performance
algorithms



One-click
training



Hyperparameter
optimization



One-click
deployment



Fully managed
hosting with auto-
scaling

Build

Train

Deploy

FREE TIER

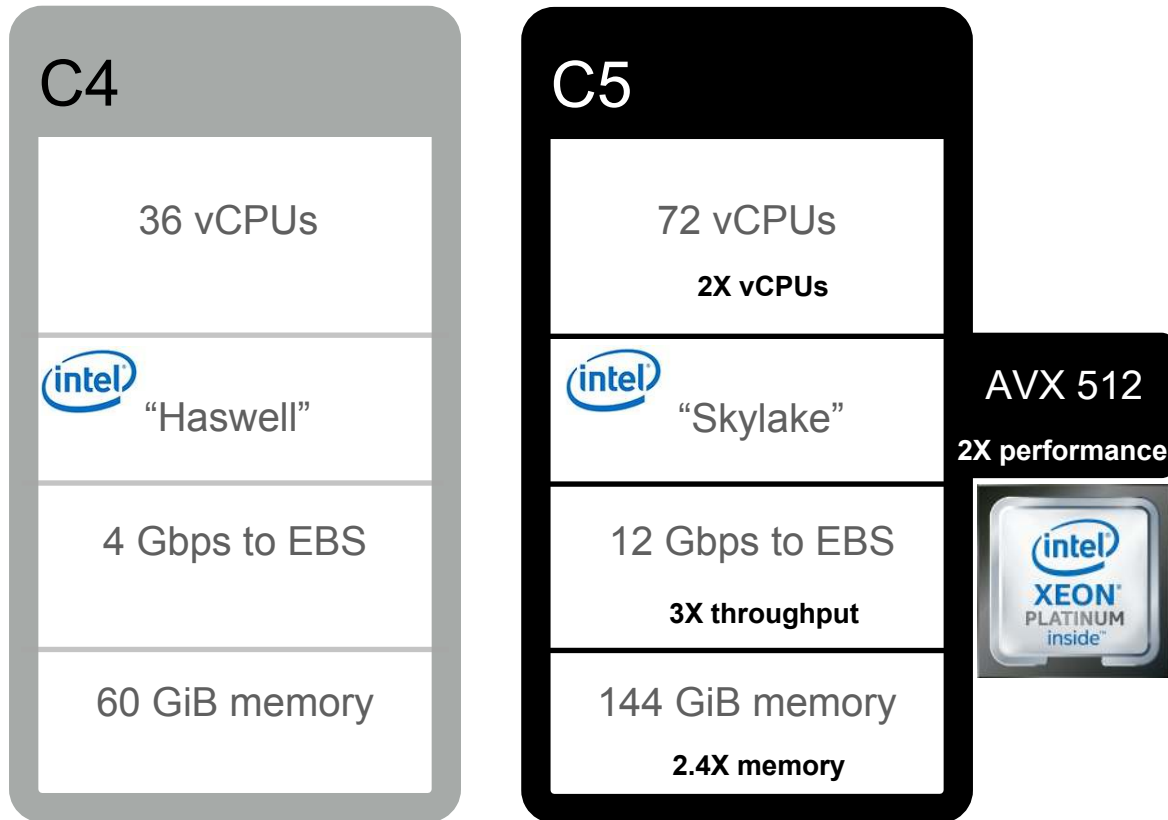
Frameworks & Infrastructure

Amazon EC2 C5 instances

C5: Next Generation
Compute-Optimized
Instances with
Intel® Xeon® Scalable
Processor

AWS Compute optimized instances support the new Intel® AVX-512 advanced instruction set, enabling you to more efficiently run vector processing workloads with single and double floating point precision, such as AI/machine learning or video processing.

*25% improvement in
price/performance over C4*



Amazon EC2 P3 Instances

The fastest, most powerful GPU instances in the cloud

- P3.2xlarge, P3.8xlarge, P3.16xlarge
- Up to eight NVIDIA Tesla V100 GPUs in a single instance
 - 40,960 CUDA cores, 5120 Tensor cores
 - 128GB of GPU memory
- 1 PetaFLOPs of computational performance – *14x better than P2*
- 300 GB/s GPU-to-GPU communication (NVLink) – *9x better than P2*

AWS Deep Learning AMIs

Preconfigured environments to quickly build Deep Learning applications

Conda AMI

For developers who want pre-installed pip packages of DL frameworks in separate virtual environments.

Base AMI

For developers who want a clean slate to set up private DL engine repositories or custom builds of DL engines.

AMI with source code

For developers who want preinstalled DL frameworks and their source code in a shared Python environment.




GLUON

PYTORCH

<https://aws.amazon.com/machine-learning/amis/>

Faster TensorFlow training on C5



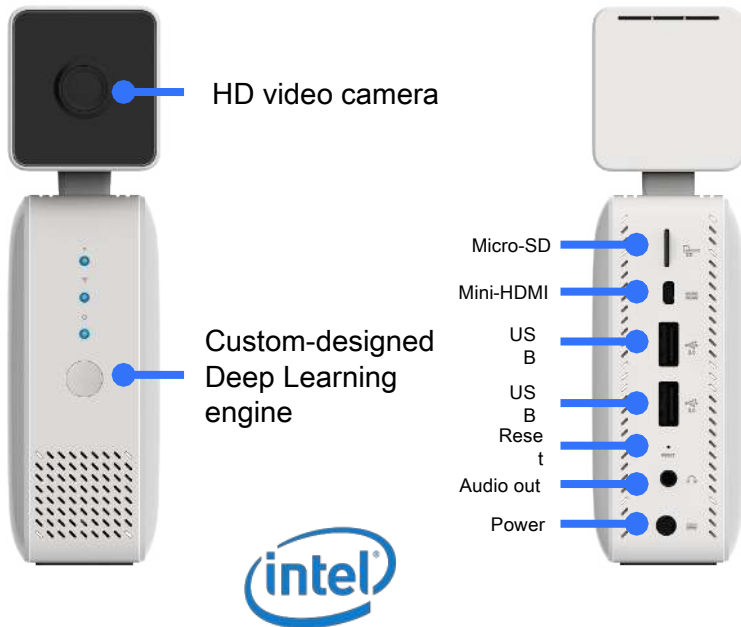
<https://aws.amazon.com/blogs/machine-learning/faster-training-with-optimized-tensorflow-1-6-on-amazon-ec2-c5-and-p3-instances/>



Can we do more to put ML in the
hands of all developers (literally)?

AWS DeepLens

The world's first Deep Learning-enabled video camera for developers



HD video camera
with on-board
compute optimized
for Deep Learning



Integrates with Amazon
SageMaker and AWS
Lambda



From unboxing
to prediction in
<10 minutes



Tutorials, examples,
demos, and pre-built
models

Object detection with AWS DeepLens



Democratization of ML

APPLICATION SERVICES



**Amazon
Rekognition**



**Amazon
Polly**



**Amazon
Lex**

**Amazon
Comprehend**



**Amazon
Rekognition
Video**



**Amazon
Transcribe**



**Amazon
Translate**

PLATFORM SERVICES

**Amazon
SageMaker**

AWS DeepLens

FRAMEWORKS AND INTERFACES

**Deep
Learning AMI**

Apache
MXNet

Caffe2

CNTK

PyTorch

TensorFlow

Theano

Torch

Keras

Gluon

Thank you!



<https://aws.amazon.com/machine-learning>

<https://aws.amazon.com/blogs/ai>

<https://medium.com/@julsimon>

<https://youtube.com/juliensimonfr>

Julien Simon

Principal Technical Evangelist, AI and Machine Learning

@**julsimon**