AWS re:Invent

ENT 321

Build, Train, and Deploy Machine Learning for the Enterprise with Amazon SageMaker

Julien Simon
Principal Technical Evangelist, AI & Machine Learning
Amazon Web Services
@julsimon

AWS re:Invent

Agenda

- Welcome & housekeeping
- Slides: quick overview of Amazon SageMaker
- Labs

- What we'll cover today:
  - Loading data from Amazon S3
  - Training and deploying with built-in algorithms,
  - Finding optimal hyper parameters with Automatic Model Tuning,
  - Running HTTPS predictions and batch predictions,
  - Beyond built-in algorithms: a peek at Deep Learning.

Amazon SageMaker

AWS re:Invent

Amazon SageMaker

Easily build, train, and deploy Machine Learning models

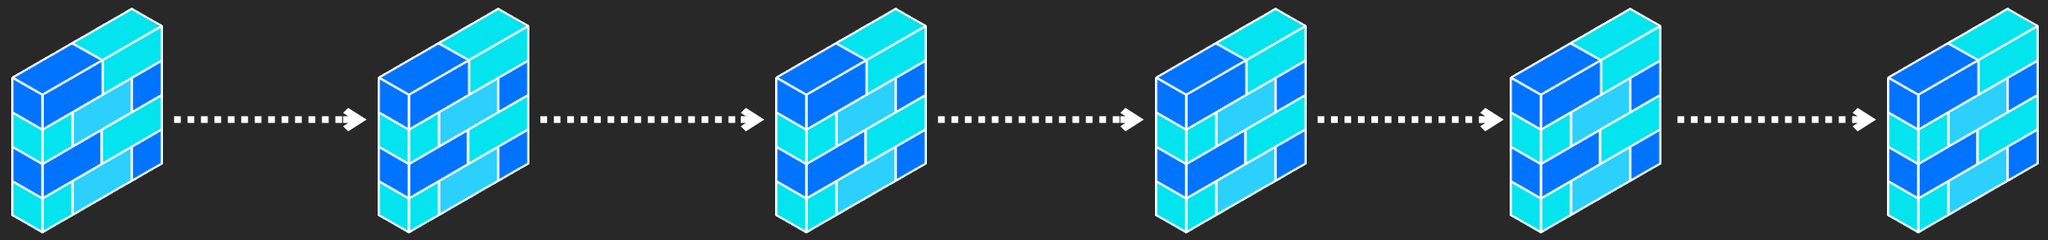

Collect and prepare training data

Choose and optimize your ML algorithm

Set up and manage environments for training

Train and tune model (trial and error)

Deploy model in production

Scale and manage the production environment

FREE TIER

Amazon SageMaker

Notebook instances

Built-in, high-performance algorithms

ALGORITHMS	K-Means Clustering Principal Component Analysis Neural Topic Modelling Factorization Machines Linear Learner	XGBoost Latent Dirichlet Allocation Image Classification Seq2Seq, And more!
FRAMEWORKS	Apache MXNet, Chainer TensorFlow, PyTorch	Caffe2, CNTK, Torch

Set up and manage environments for training

Train and tune model (trial and error)

Deploy model in production

Scale and manage the production environment

Build

AWS re:Invent

Amazon SageMaker

Notebook instances

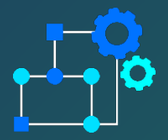

Built-in, high-performance algorithms

Build

One-click training

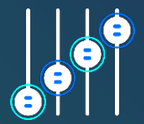

Automatic Model Tuning

Train

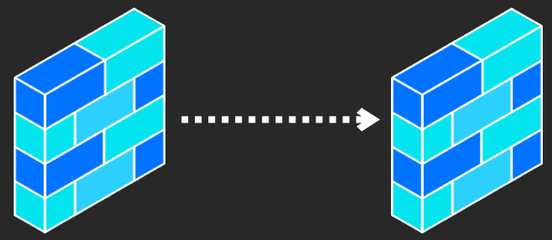

Deploy model in production

Scale and manage the production environment

AWS re:Invent

Amazon SageMaker

Notebook instances

Built-in, high-performance algorithms

Build

One-click training

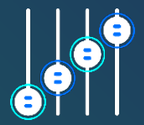

Automatic Model Tuning

Train

One-click deployment

Fully managed hosting with auto-scaling

Deploy

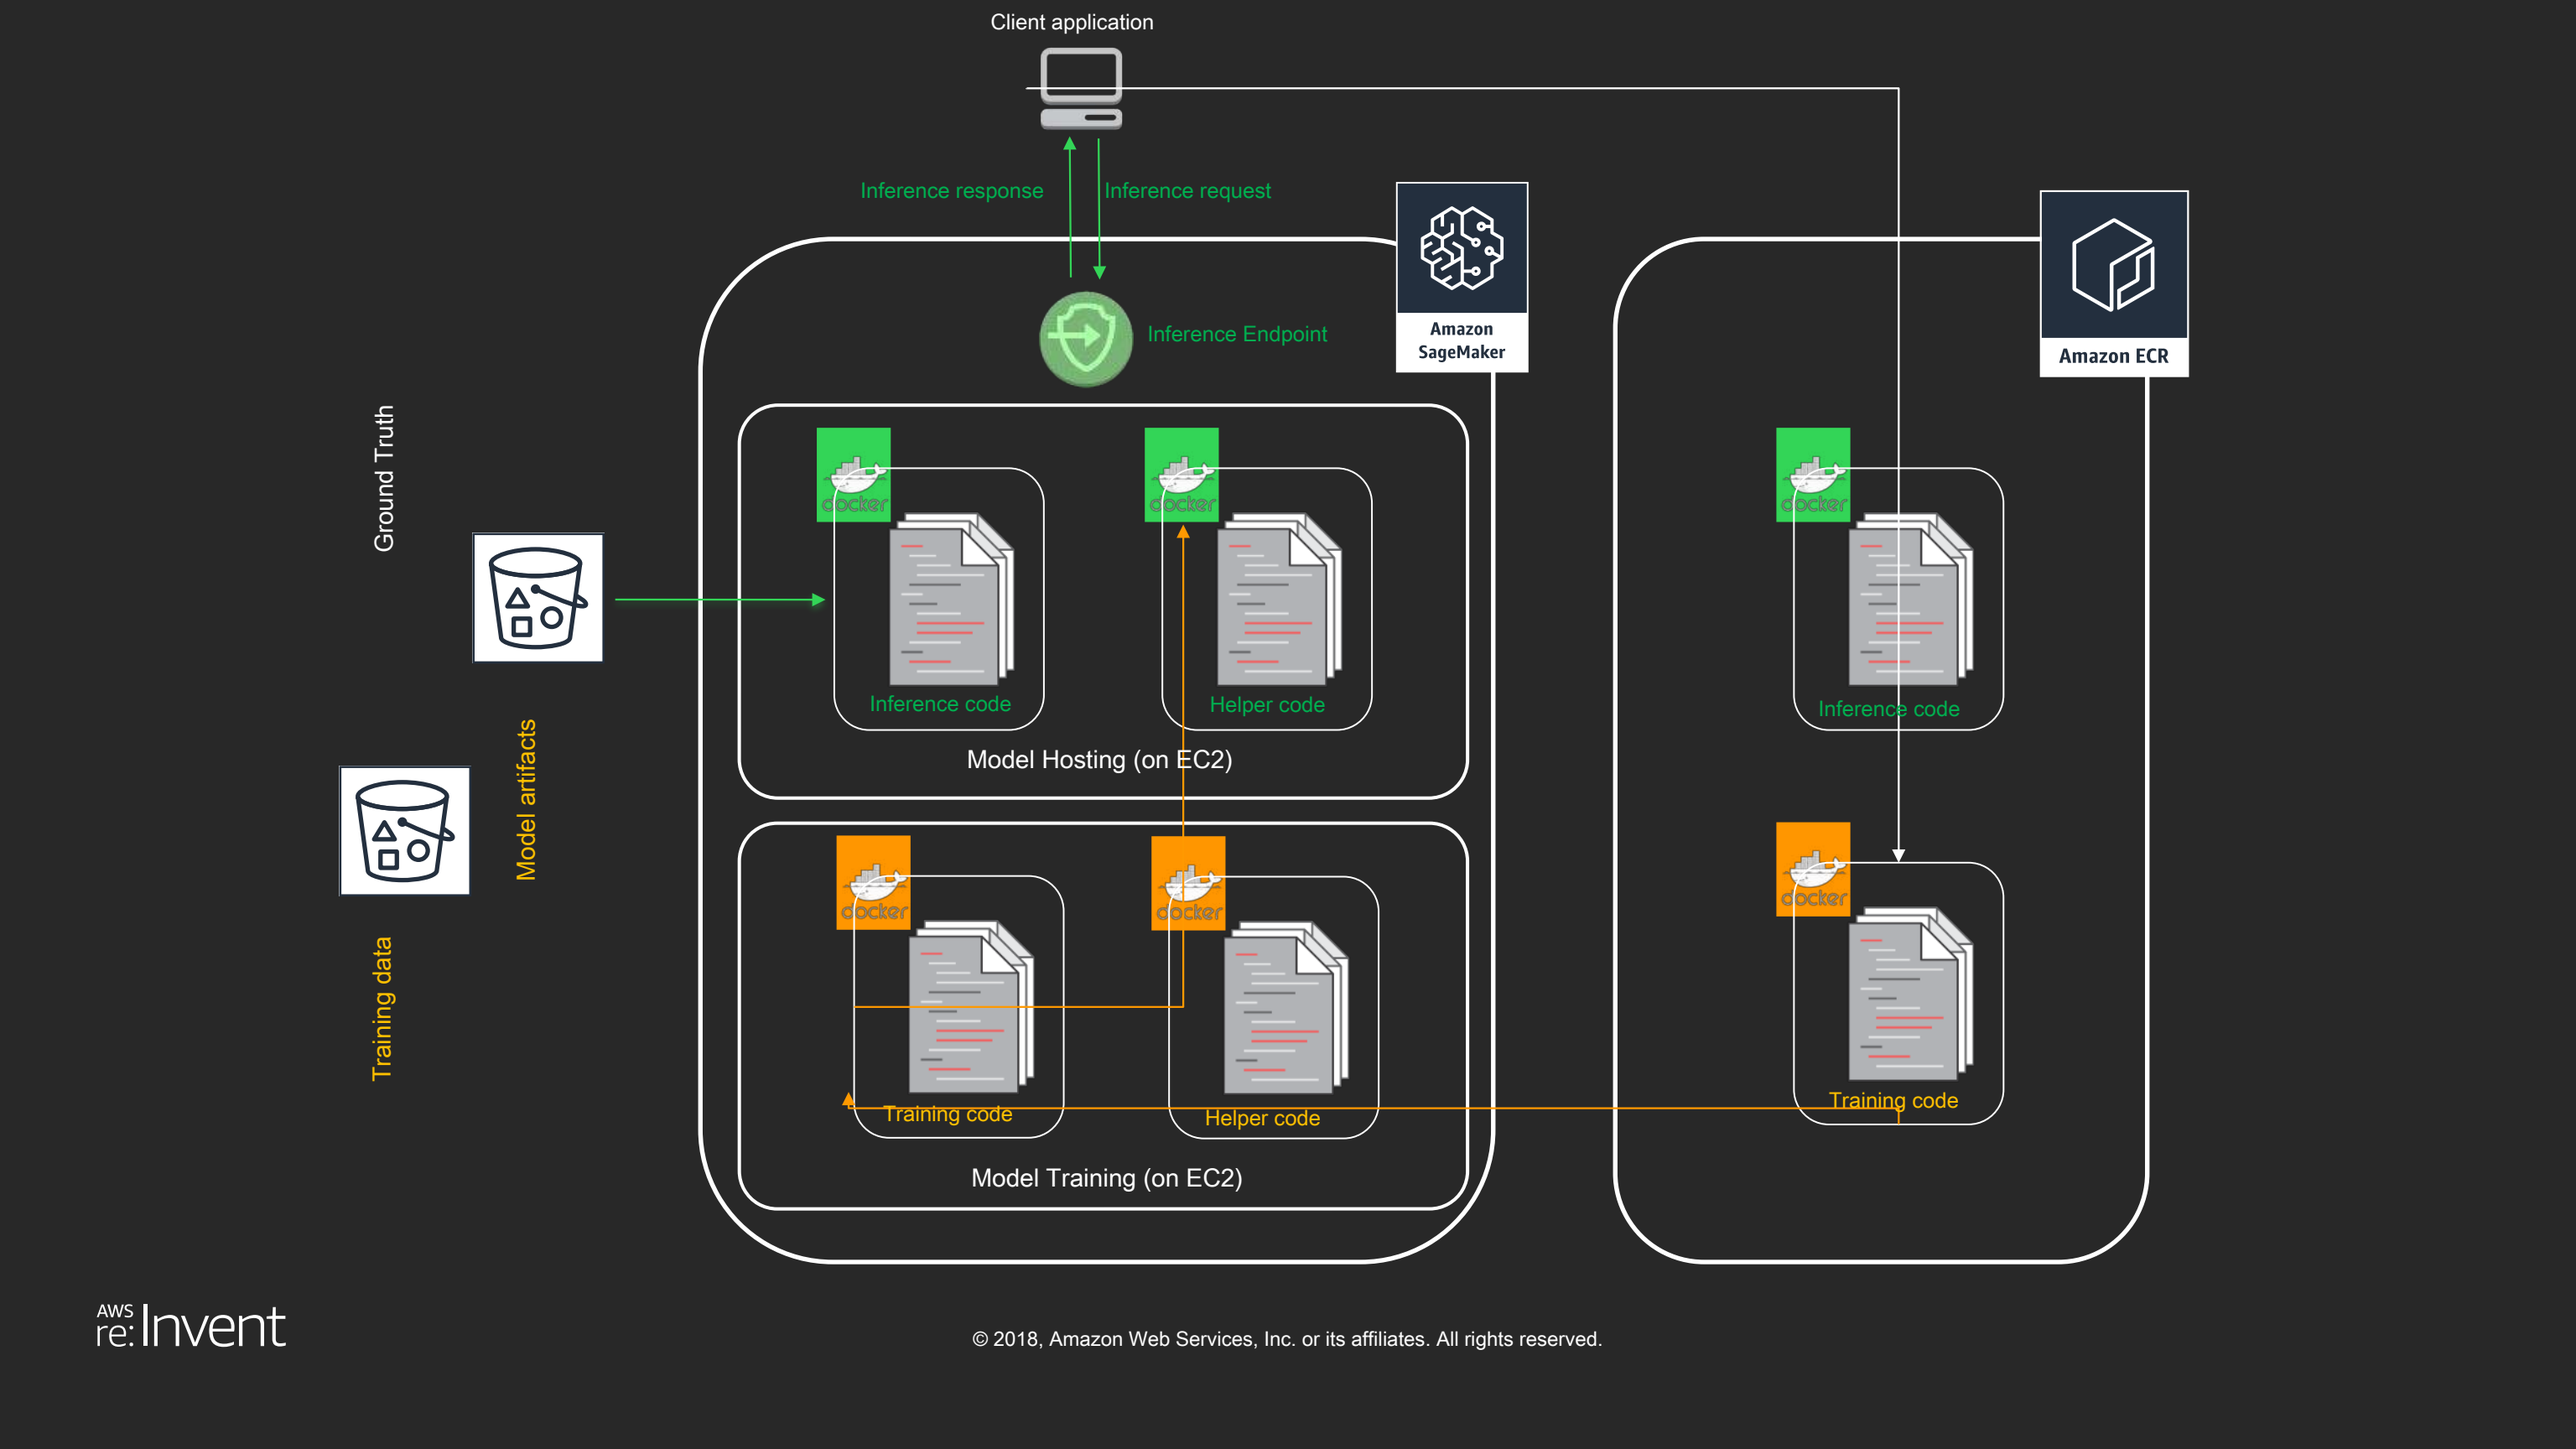

Model options

Training code

Factorization Machines
Linear Learner
Principal Component Analysis
K-Means
XGBoost
And more

Built-in Algorithms

mxnet
TensorFlow
PyTorch

Bring Your Own Script

Bring Your Own Container

The Amazon SageMaker SDK

- Python SDK orchestrating all Amazon SageMaker activity
  - Algorithm selection, training, deploying, hyper parameter optimization, etc.
  - There's also a Spark SDK (Python and Scala) which we won't cover today.

- High-level objects for:
  - Some built-in algos: Kmeans, PCA, etc.
  - Deep Learning libraries: TensorFlow, MXNet, PyTorch, Chainer.
  - Sagemaker.estimator.estimator for everything else.

https://github.com/aws/sagemaker-python-sdk
https://sagemaker.readthedocs.io/en/latest/

Built-in algorithms

pink: supervised, blue: unsupervised

Linear Learner: regression, classification	**Image Classification**: Deep Learning (ResNet)
Factorization Machines: regression, classification, recommendation	**Object Detection**: Deep Learning (VGG or ResNet)
K-Nearest Neighbors: non-parametric regression and classification	**Neural Topic Model**: topic modeling
XGBoost: regression, classification, ranking https://github.com/dmlc/xgboost	**Latent Dirichlet Allocation**: topic modeling (mostly)
K-Means: clustering	**Blazing Text**: GPU-based Word2Vec, and text classification
Principal Component Analysis: dimensionality reduction	**Sequence to Sequence**: machine translation, speech to text and more
Random Cut Forest: anomaly detection	**DeepAR**: time-series forecasting (RNN)
Object2Vec: general-purpose embeddings	**IP Insights**: usage patterns for IP addresses

XGBoost

- Open Source project
- Popular tree-based algorithm for regression, classification and ranking
- Builds a collection of trees.
- Handles missing values and sparse data
- Supports distributed training
- Can work with data sets larger than RAM

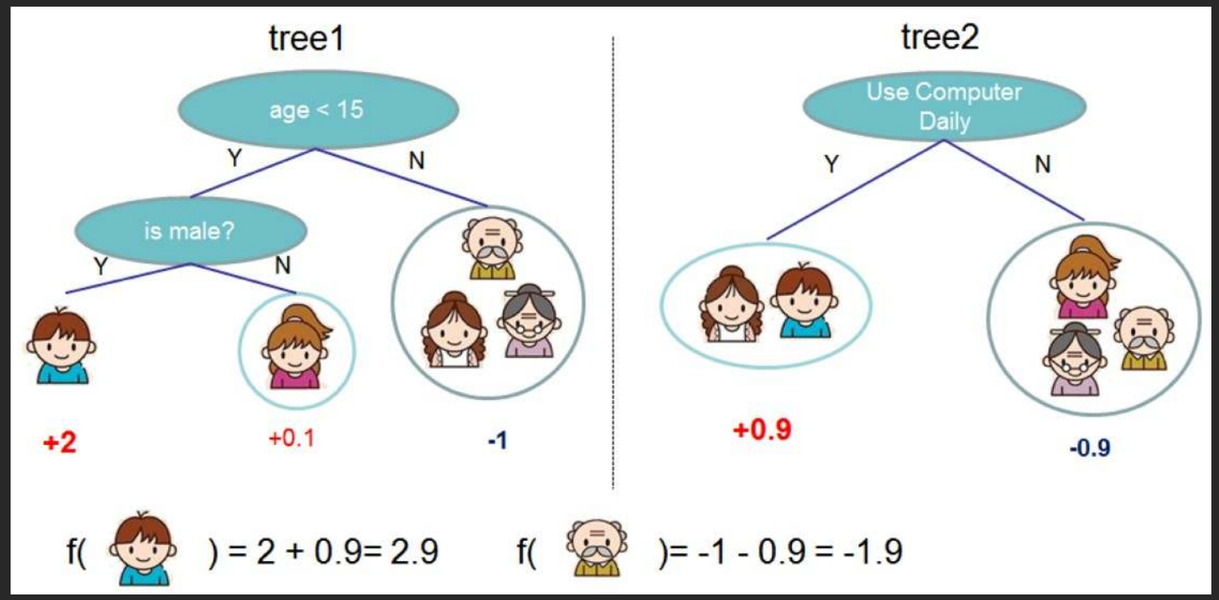

https://github.com/dmlc/xgboost
https://xgboost.readthedocs.io/en/latest/
https://arxiv.org/abs/1603.02754

Loading training data from Amazon S3

- Two modes: File Mode and Pipe Mode.
  - *input_mode* parameter in *sagemaker.estimator.Estimator.*
- File Mode copies the data set to training instances.
  - You need to provision enough storage.
  - *S3DataSource* object.
  - *S3DataDistributionType* : *FullyReplicated* | *ShardedByS3Key*
  - Differerent data formats are supported: CSV, protobuf, JSON, libsvm (check algo docs!).
- Pipe Mode streams the data set to training instances.
  - This allows you to process infinitely-large data sets.
  - Training starts faster.
  - This mode is supported by some built-in algos as well as Tensorflow.
  - Your data set must be in recordIO-encoded protobuf format.

Walkthrough:

- AWS credits
- SageMaker console
- Notebook instance setup

AWS re:Invent

Labs

AWS re:Invent

Labs

1. Training, deploying and predicting with XGBoost
2. Finding optimal hyper parameters with Automatic Model Tuning
3. Running HTTPS predictions and batch predictions,
4. Beyond built-in algorithms: a peek at TensorFlow.

Resources

AWS re:Invent

Resources

https://ml.aws

https://aws.amazon.com/sagemaker
https://github.com/awslabs/amazon-sagemaker-examples
https://github.com/aws/sagemaker-python-sdk

https://medium.com/@julsimon

Thank you!

Julien Simon
Principal Technical Evangelist, AI & Machine Learning
Amazon Web Services
@julsimon

AWS re:Invent