

Big Data Architectural Patterns and Best Practices on AWS

Julien Simon
Principal Technical Evangelist
[@julsimon](#)

What to Expect from the Session

Big data challenges

Architectural principles


How to simplify big data processing

What technologies should you use?


- Why?
- How?

Reference architecture

Design patterns



Ever Increasing Big Data



Big Data Evolution

Batch
processing



Stream
processing



Artificial
intelligence



Cloud Services Evolution


Virtual
machines

Managed
services

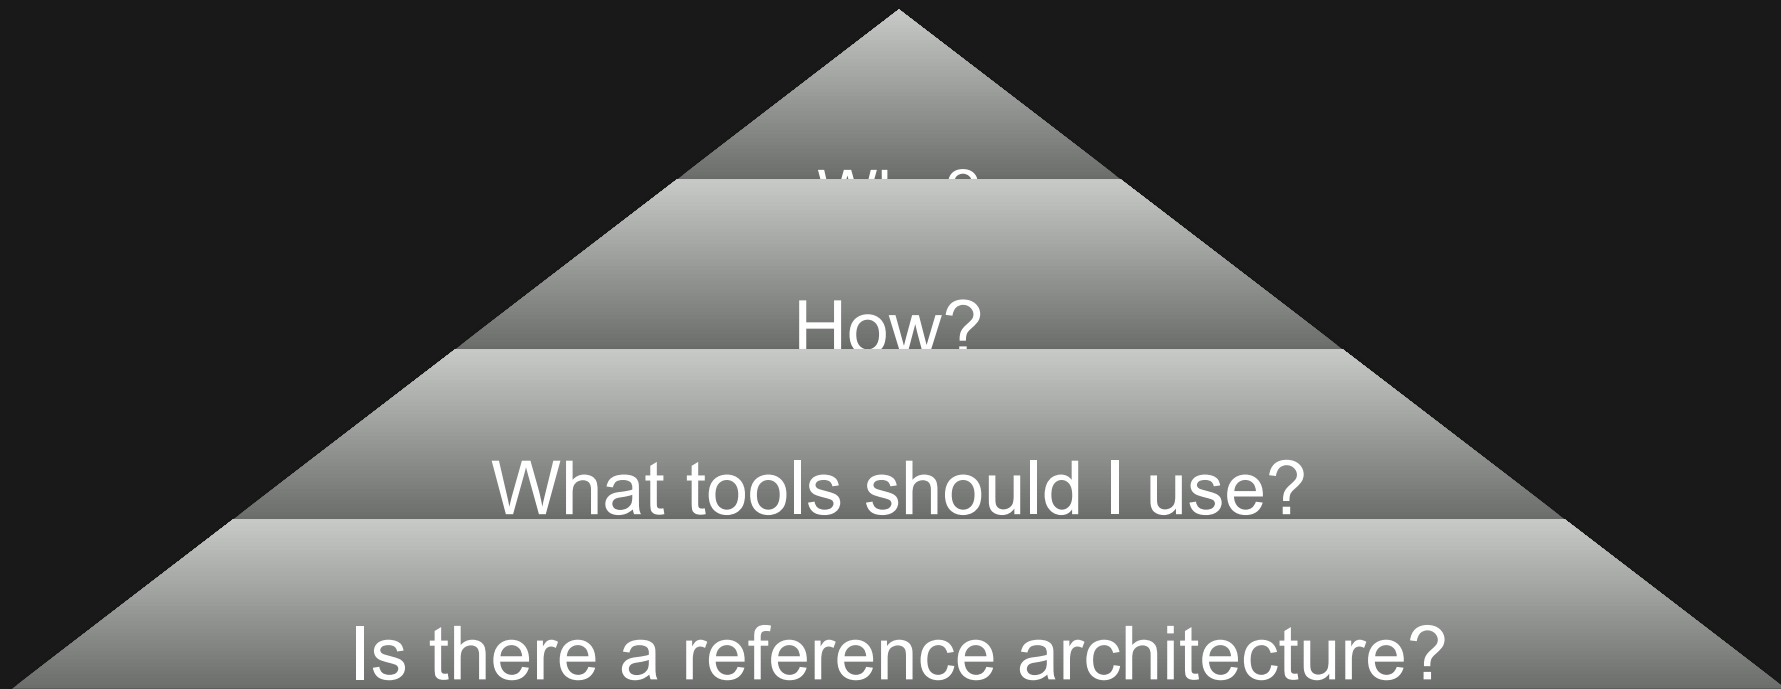
Serverless



Plethora of Tools



Big Data Challenges



Architectural Principles



Build decoupled systems

- Data → Store → Process → Store → Analyze → Answers

Use the right tool for the job

- Data structure, latency, throughput, access patterns

Leverage managed and serverless services

- Scalable/elastic, available, reliable, secure, no/low admin

Use log-centric design patterns

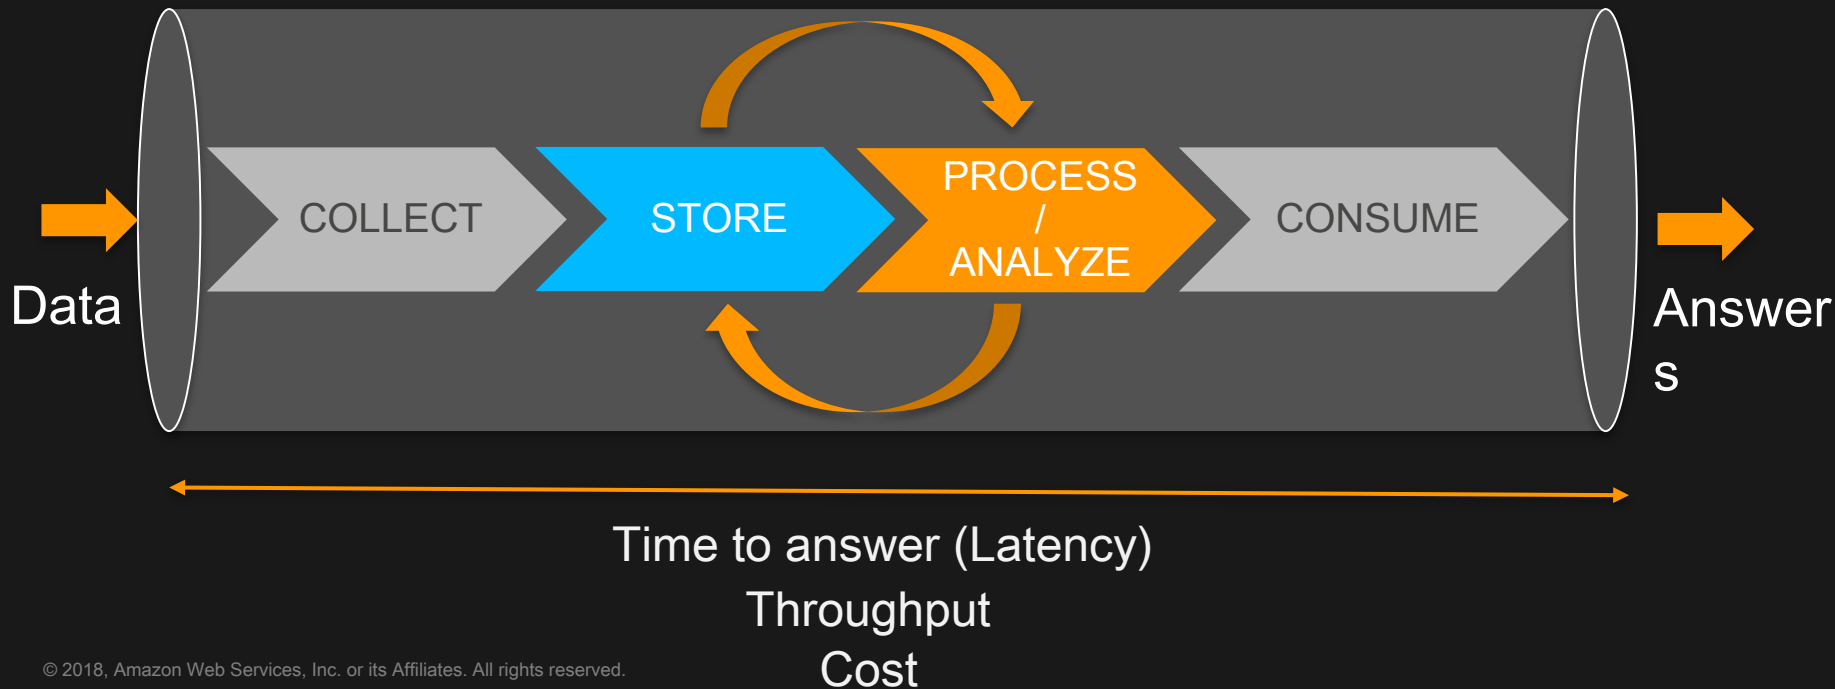
- Immutable logs (data lake), materialized views

Be cost-conscious


- Big data ≠ big cost

AI/ML enable your applications

Simplify Big Data Processing



What Is the Temperature of Your Data?



Data Characteristics: Hot, Warm, Cold

| | Hot | Warm | Cold |
|---------------------------------------|-----------|----------|-----------|
| Volume | MB–GB | GB–TB | PB–EB |
| Item size | B–KB | KB–MB | KB–TB |
| Latency | ms | ms, sec | min, hrs |
| Durability | Low–high | High | Very high |
| Request rate | Very high | High | Low |
| Cost/GB | \$\$\$ | \$ \$ \$ | \$ |
| <div>Hot dataWarm dataCold data</div> | | | |

A light gray arrow pointing to the right, centered on a dark gray background. The word "COLLECT" is written in a bold, dark gray, sans-serif font in the center of the arrow.

COLLECT

COLLECT

Type of Data

Data structures
Database records

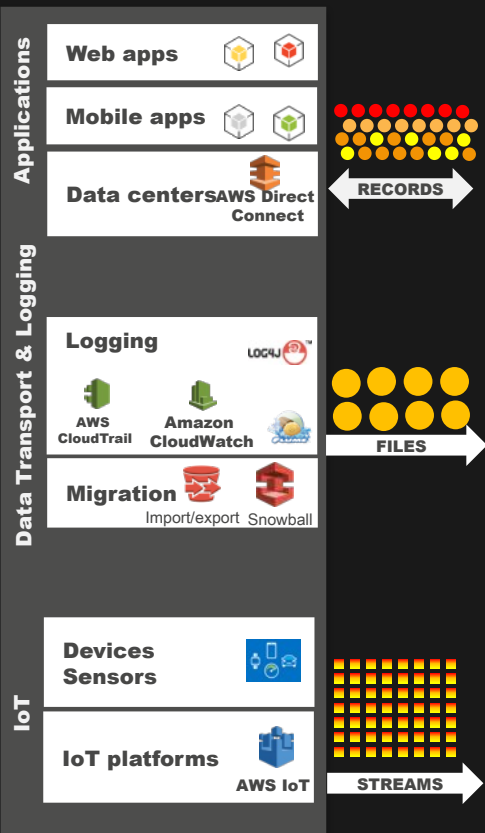
Transactions

Media files
Log files

Files

Data streams

Events



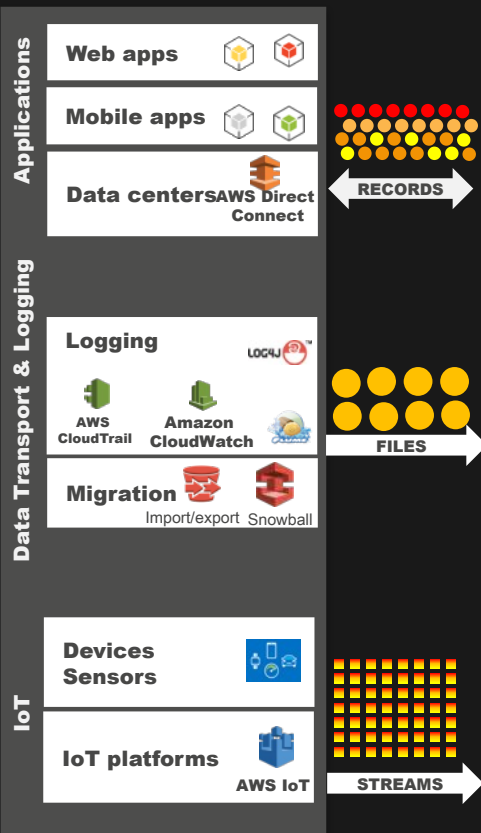
A blue arrow pointing right, centered on a dark background. The word "STORE" is written in white, uppercase letters inside the arrow.

STORE

COLLECT

Type of Data

STORE



Data structures
Database records

Media files
Log files

Data streams

In-memory

NoSQL

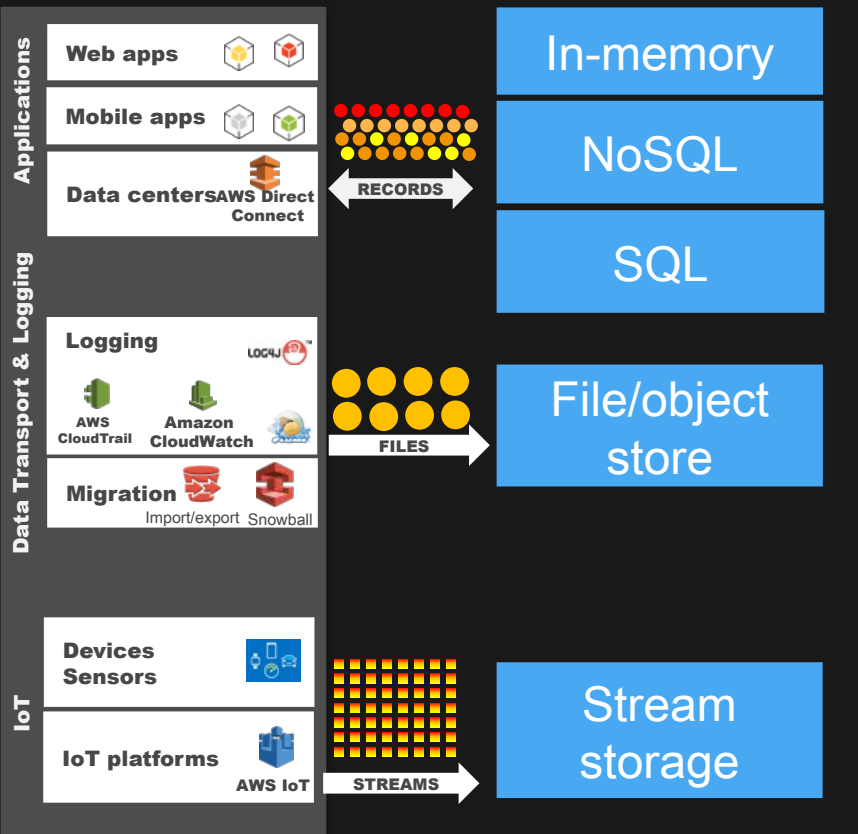
SQL

File/object
store

Stream
storage

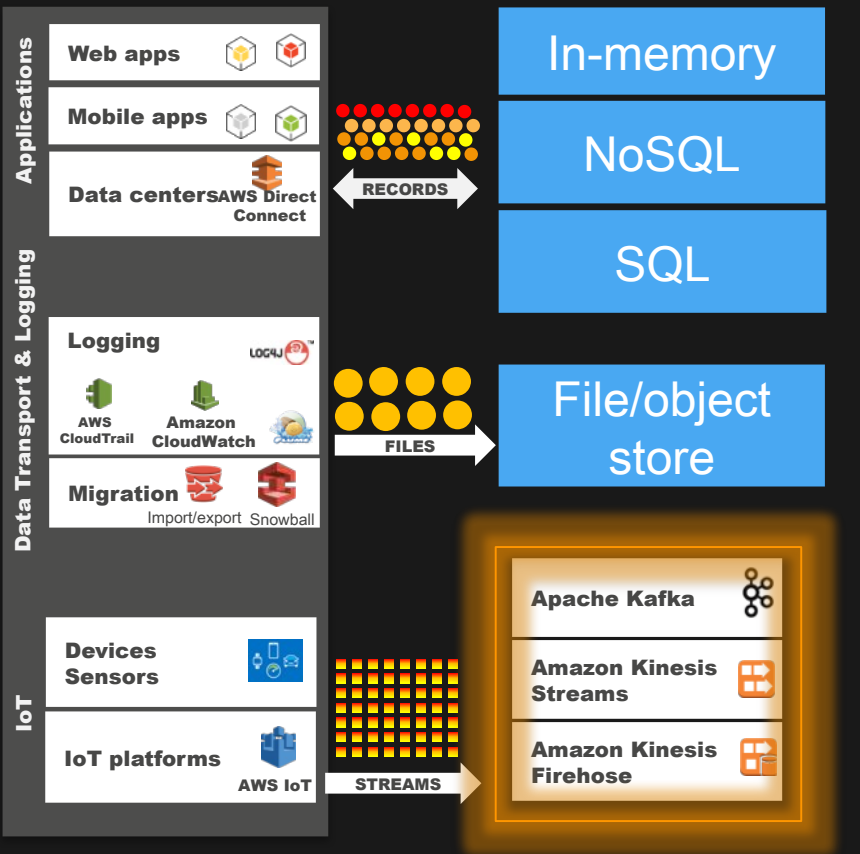
COLLECT

STORE



COLLECT

STORE



Stream Storage

Apache Kafka

- High throughput distributed streaming platform

Amazon Kinesis Streams

- Managed stream storage

Amazon Kinesis Firehose

- Managed data delivery

Why Stream Storage?

Decouple producers & consumers

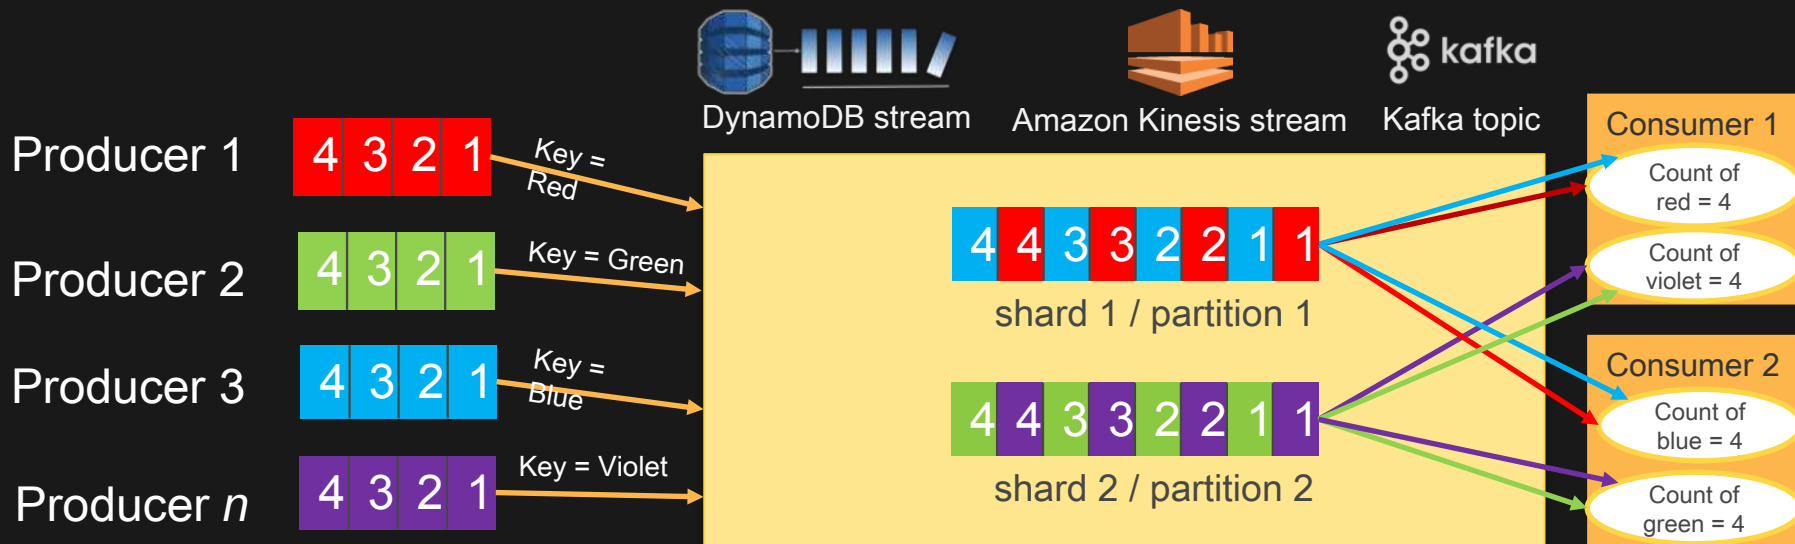
Persistent buffer

Collect multiple streams

Preserve client ordering

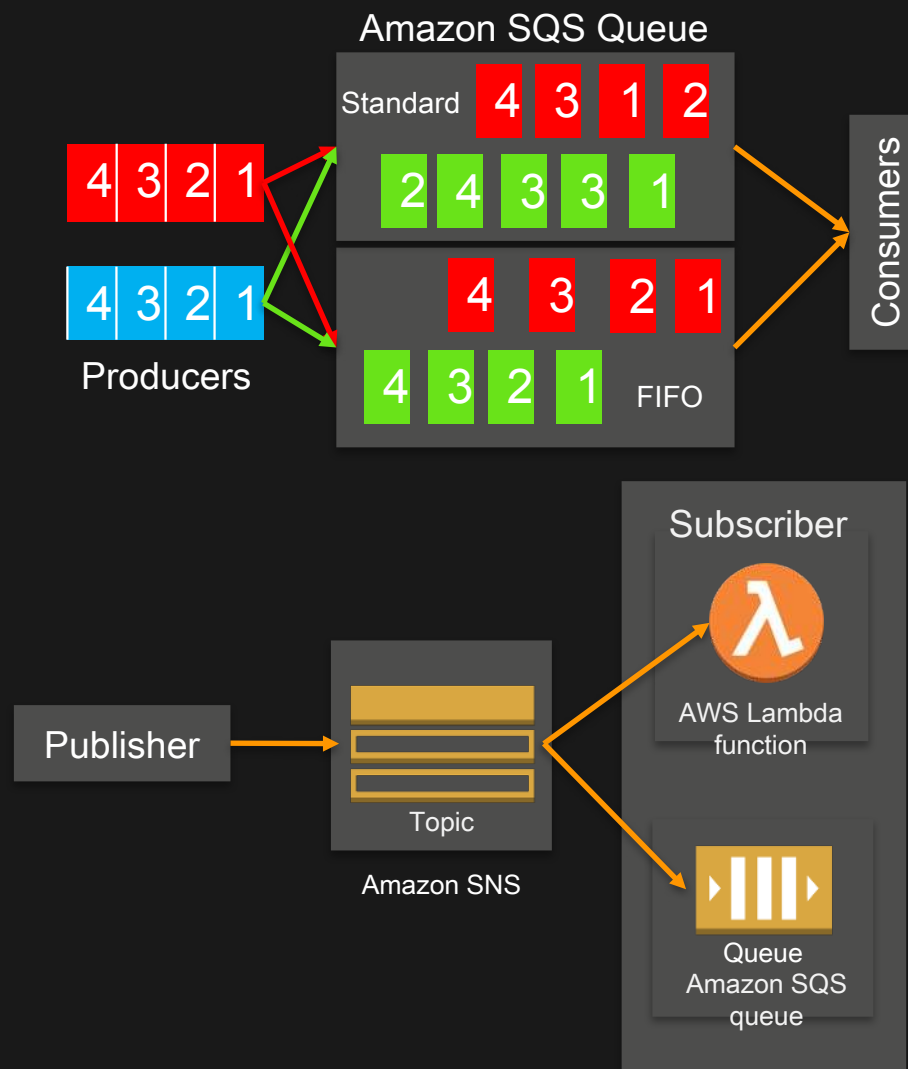
Parallel consumption

Streaming MapReduce



What About Amazon SQS?

- Decouple producers & consumers
- Persistent buffer
- Collect multiple streams
- **No** client ordering (standard)
 - FIFO queue preserves client ordering
- **No** streaming MapReduce
- **No** parallel consumption
 - Amazon SNS can publish to multiple SNS subscribers (queues or Lambda functions)



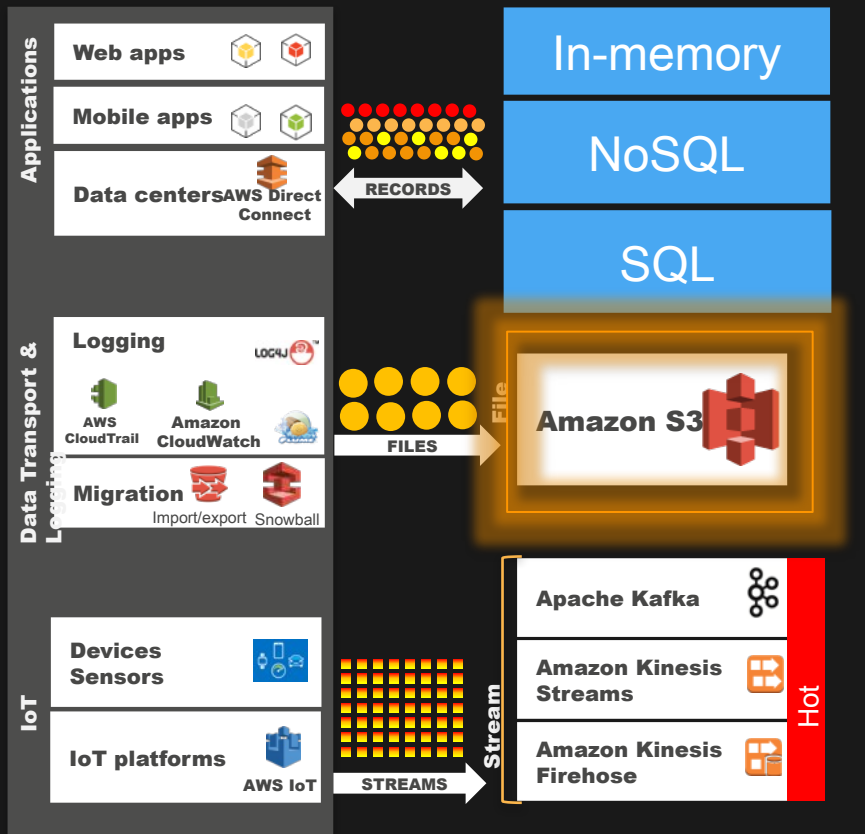
Which Stream/Message Storage Should I Use?

| | Amazon Kinesis Streams | Amazon Kinesis Firehose | Apache Kafka (on Amazon EC2) | Amazon SQS (Standard) | Amazon SQS (FIFO) |
|------------------------------|------------------------|-----------------------------|-------------------------------|-----------------------|-------------------|
| AWS managed | Yes | Yes | No | Yes | Yes |
| Guaranteed ordering | Yes | No | Yes | No | Yes |
| Delivery (deduping) | At least once | At least once | At least/At most/exactly once | At least once | Exactly once |
| Data retention period | 7 days | N/A | Configurable | 14 days | 14 days |
| Availability | 3 AZ | 3 AZ | Configurable | 3 AZ | 3 AZ |
| Scale / throughput | No limit / ~ shards | No limit / automatic | No limit / ~ nodes | No limits / automatic | 300 TPS / queue |
| Parallel consumption | Yes | No | Yes | No | No |
| Stream MapReduce | Yes | N/A | Yes | N/A | N/A |
| Row/object size | 1 MB | Destination row/object size | Configurable | 256 KB | 256 KB |
| Cost | Low | Low | Low (+admin) | Low-medium | Low-medium |

COLLECT

STORE

File/Object Storage



Amazon S3


Managed object storage service built to store and retrieve any amount of data

Use **Amazon S3** as Your Persistent File Store

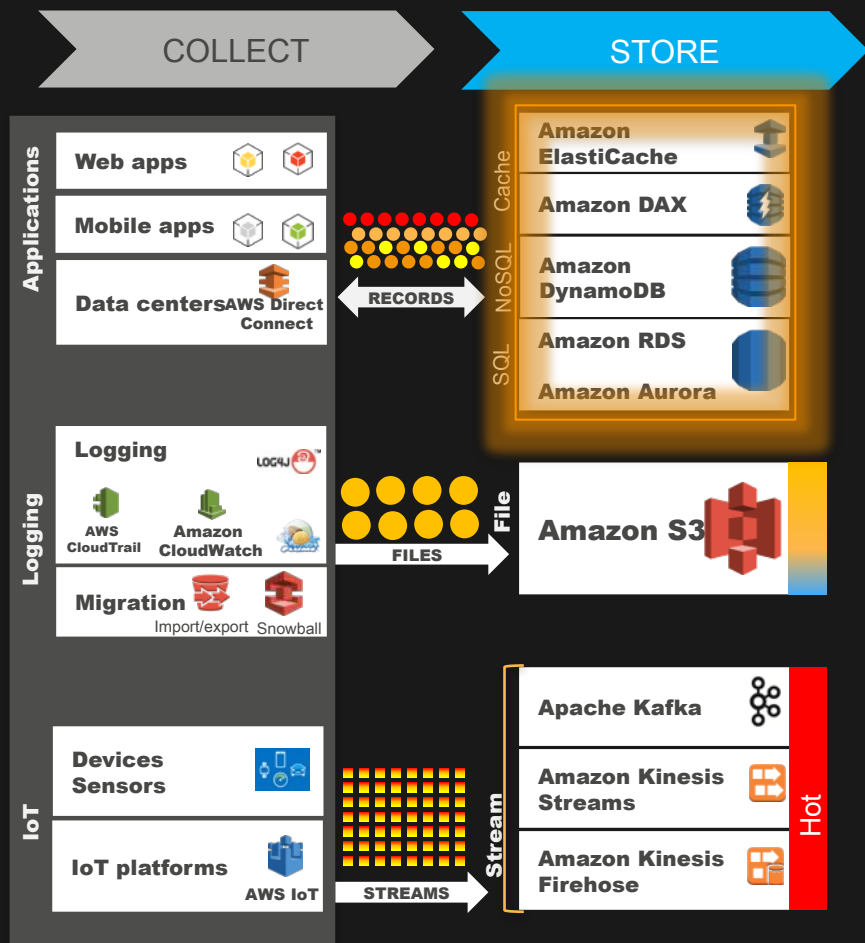
- Natively supported by big data frameworks (Spark, Hive, Presto, etc.)
- Decouple storage and compute
 - No need to run compute clusters for storage (unlike HDFS)
 - Can run transient Amazon EMR clusters with Amazon EC2 Spot Instances
 - Multiple & heterogeneous analysis clusters and services can use the same data
- Designed for 99.999999999% durability
- No need to pay for data replication within a region
- Secure: SSL, client/server-side encryption at rest
- Low cost

What About HDFS & Data Tiering?

- Use **HDFS** for hottest datasets (e.g. iterative read on the same datasets)
- Use **Amazon S3 Standard** for frequently accessed data
- Use **Amazon S3 Standard – IA** for less frequently accessed data
- Use **Amazon Glacier** for archiving cold data
- Use S3 Analytics to optimize tiering strategy



S3 Analytics



Cache & Database

Amazon ElastiCache

- Managed Memcached or Redis service

Amazon DynamoDB Accelerator (DAX)

- Managed in-memory cache for DynamoDB

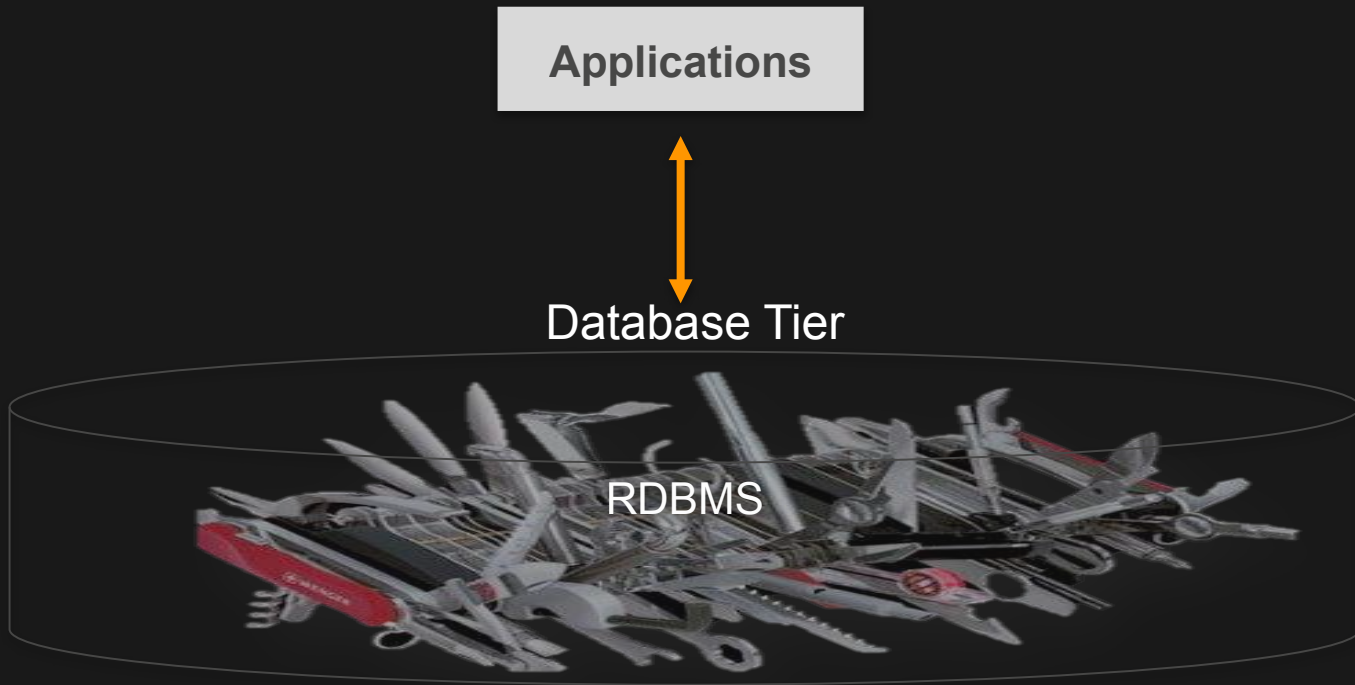
Amazon DynamoDB

- Managed NoSQL database service

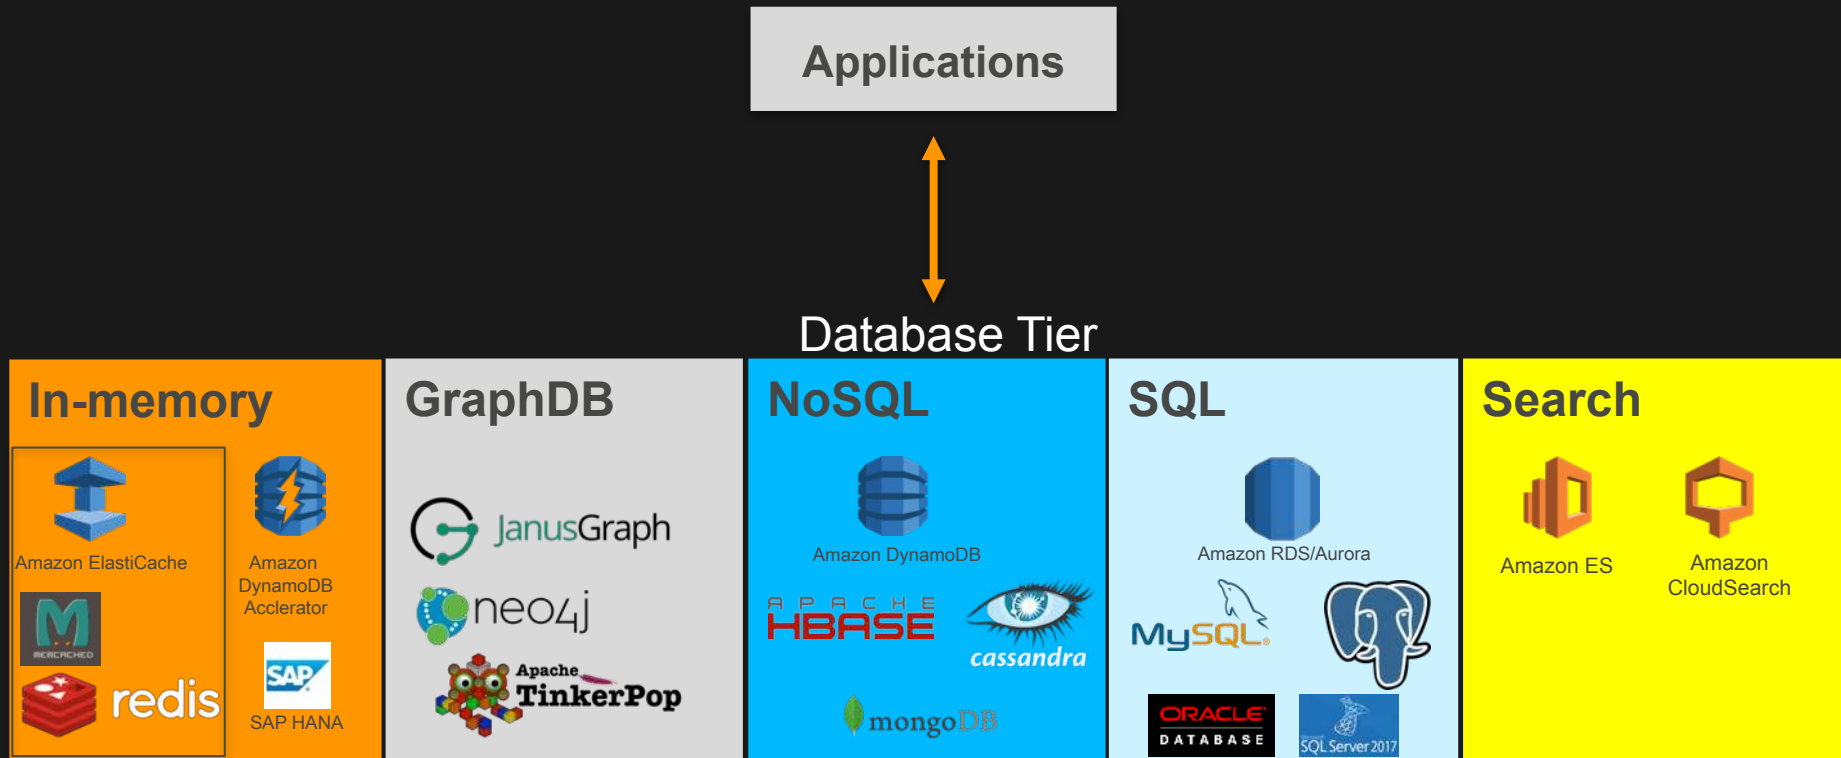
Amazon RDS

- Managed relational database service

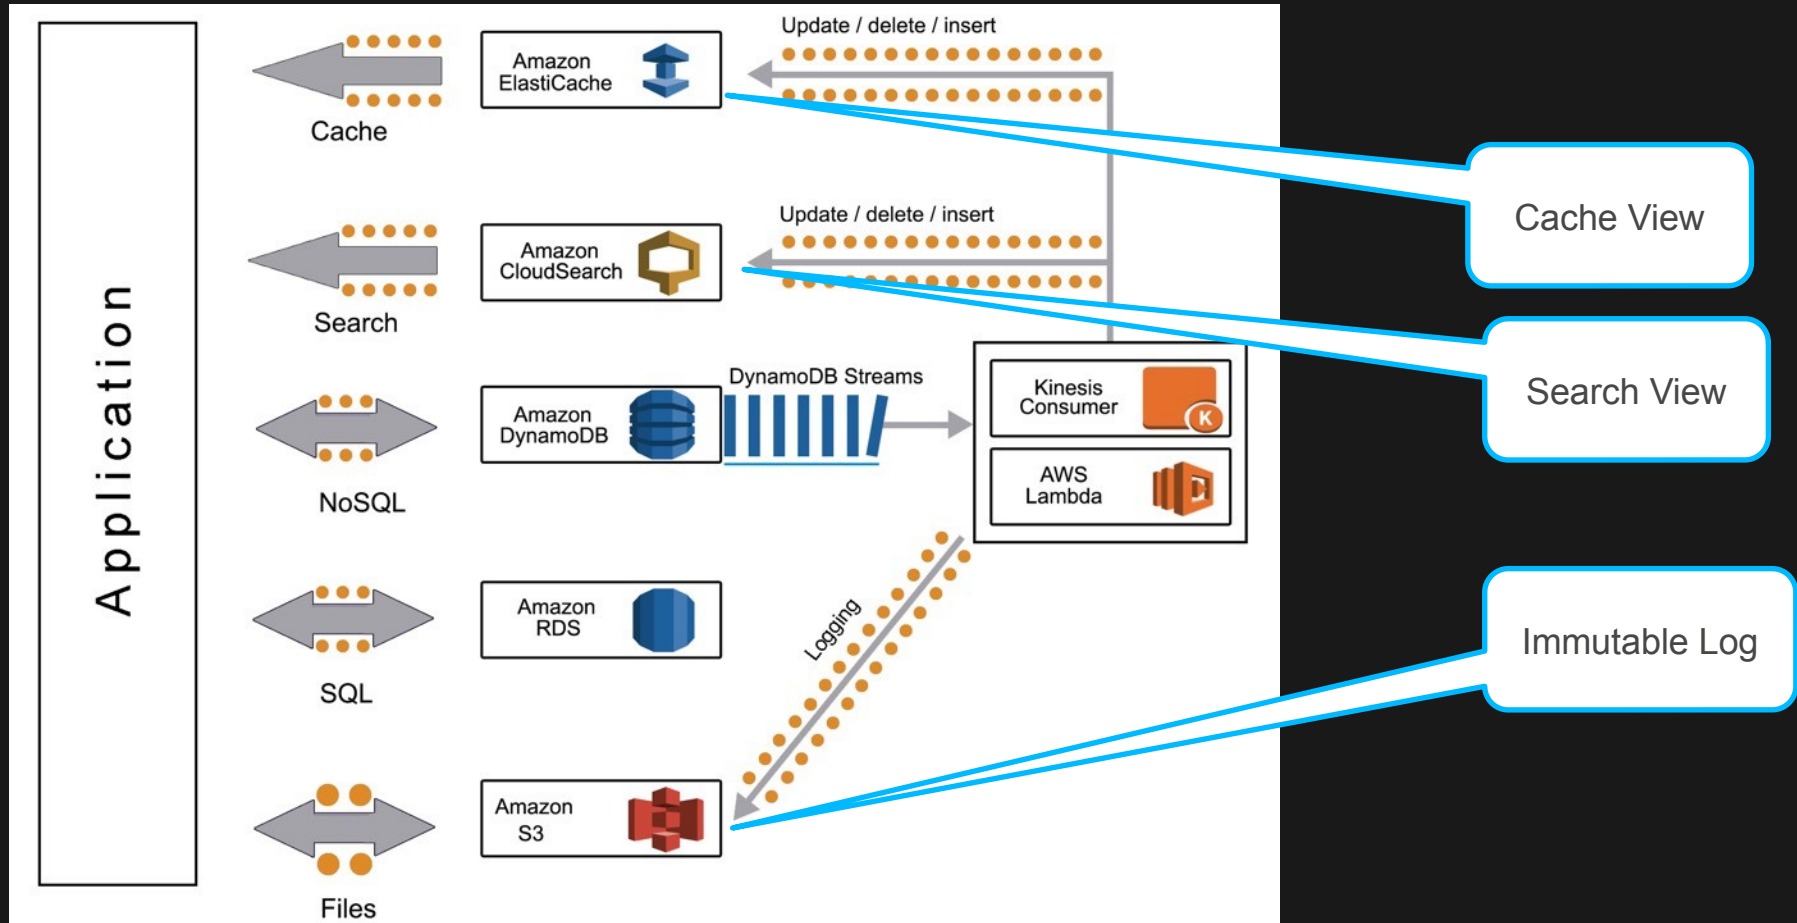
Anti-Pattern



Best Practice: Use the Right Tool for the Job ✓



Materialized Views and Immutable Log



Which Data Store Should I Use?

Data structure → Fixed-schema, JSON, Key/Value,

Access patterns → Store data in the format you will access it

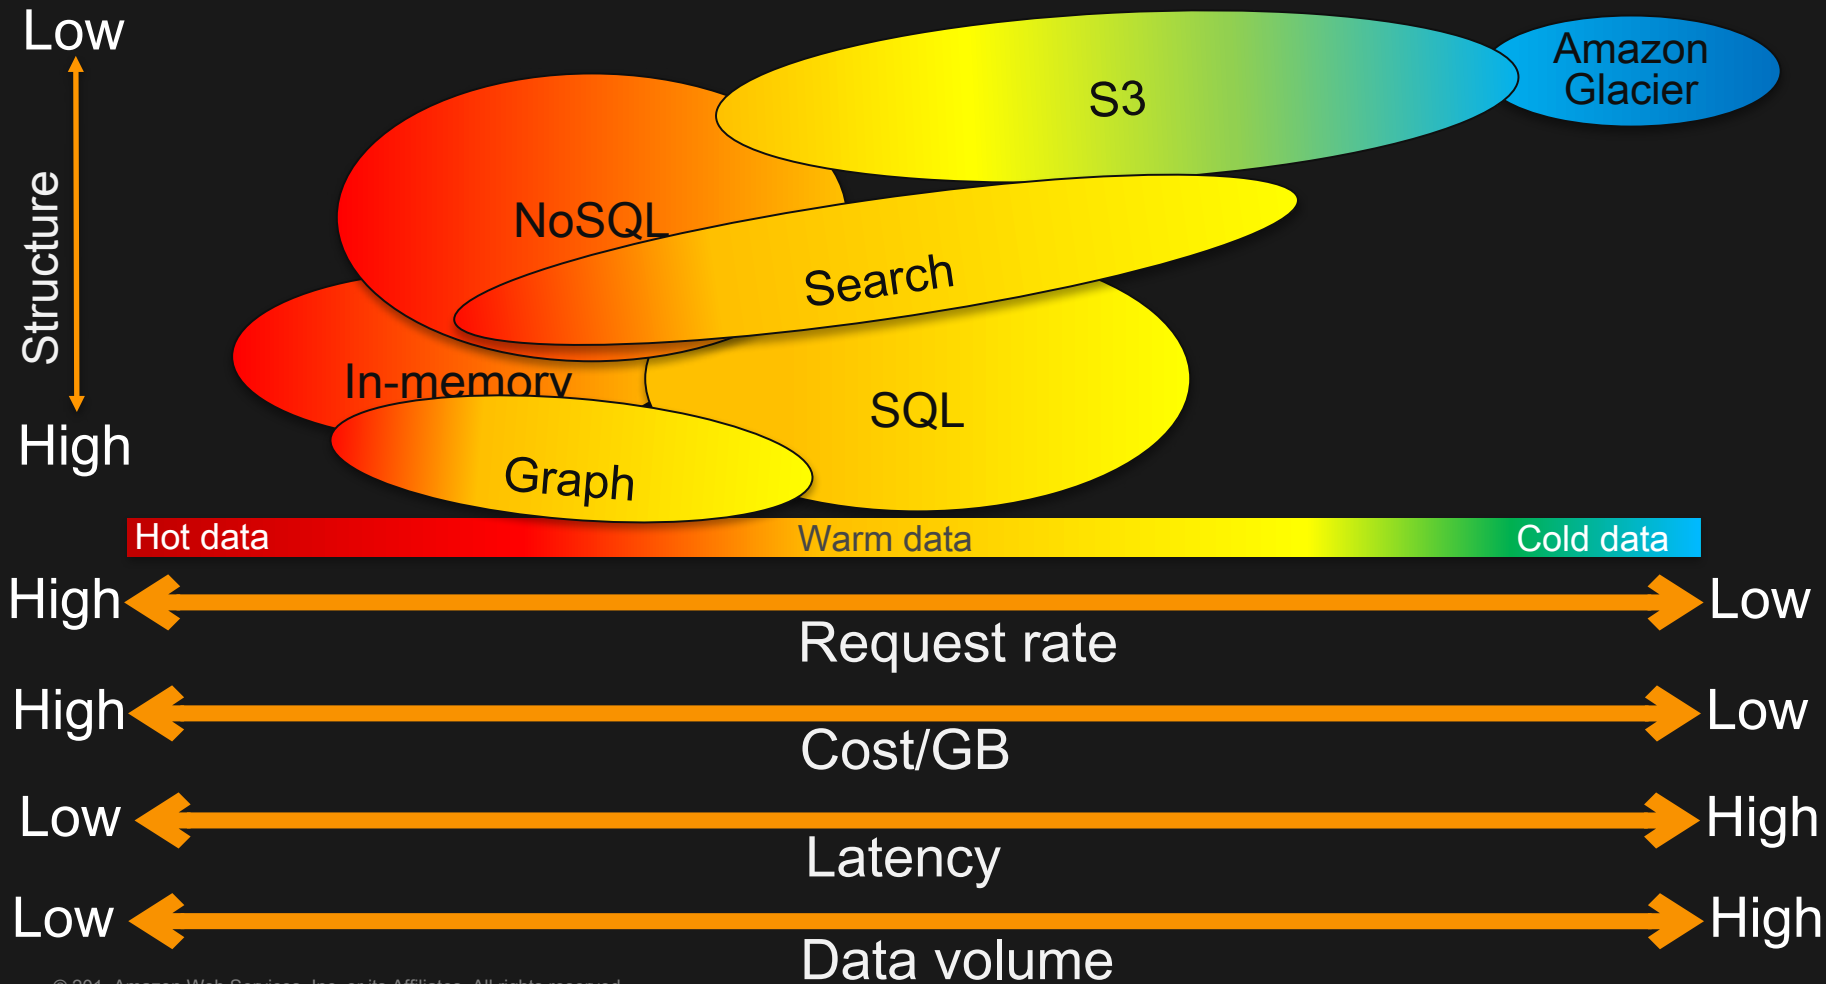
Data characteristics → Hot, warm, cold

Cost → Right cost

Data Structure and Access Patterns

| Data Structure | What to use? |
|--------------------|------------------|
| Fixed schema | SQL, NoSQL |
| Schema-free (JSON) | NoSQL, Search |
| Key/Value | In-memory, NoSQL |
| Graph | GraphDB |

| Access Patterns | What to use? |
|-------------------------------------|------------------|
| Put/Get (key, value) | In-memory, NoSQL |
| Simple relationships → 1:N, M:N | NoSQL |
| Multi-table joins, transaction, SQL | SQL |
| Faceting, Search | Search |
| Graph traversal | GraphDB |



Which Data Store Should I Use?

| | Amazon ElastiCache | Amazon DAX | Amazon DynamoDB | Amazon RDS (Aurora) | Amazon ES | Amazon S3 | Amazon Glacier |
|-----------------------|-----------------------|--------------------|-------------------------|---------------------------|--------------------|--------------------------|---------------------|
| Average Latency | µs-ms | µs-ms | ms | ms, sec | ms,sec | ms,sec,min (~ size) | hrs |
| Typical Data Volume | GB | GB | GB-TBs (no limit) | GB-TB (64 TB max) | GB-TB | MB-PB (no limit) | GB-PB (no limit) |
| Typical Item Size | B-KB | KB (400 KB max) | KB (400 KB max) | KB (64 KB max) | B-KB (2 GB max) | KB-TB (5 TB max) | GB (40 TB max) |
| Request Rate | High – very high | High – very high | Very high (no limit) | High | High | Low – high (no limit) | Very low |
| Storage Cost GB/Month | \$\$ | \$\$ | ¢¢ | ¢¢ | ¢¢ | ¢ | ¢4/10 |
| Durability | Low - moderate | NA | Very high | Very high | High | Very high | Very high |
| Availability | Hot data | | | Warm data | | | Cold data |
| | High 2 AZ | High 3 AZ | Very high 3 AZ | Very high 3 AZ | High 2 AZ | Very high 3 AZ | Very high 3 AZ |

“I’m currently scoping out a project. The design calls for **many small files**, perhaps up to a **billion during peak**. The **total size** would be on the order of **1.5 TB per month...**”

| Request rate (Writes/sec) | Object size (Bytes) | Total size (GB/month) | Objects per month |
|------------------------------|------------------------|--------------------------|-------------------|
| 300 | 2048 | 1483 | 777,600,000 |

Amazon S3 or Amazon DynamoDB?



<https://calculator.s3.amazonaws.com/index.html>

| Request rate (writes/sec) | Object size (bytes) | Total size (GB/month) | Objects per month |
|------------------------------|------------------------|--------------------------|----------------------|
| 300 | 2,048 | 1,483 | 777,600,000 |



Amazon DynamoDB is a high performance non-relational database service that is easy to set up, operate, and scale. It is designed to address the core problems of database management, performance, scalability, and reliability. It also provides predictable high performance and low latency at scale.

Indexed Data Storage:

Dataset Size:

1483 GB

Provisioned Throughput Capacity *:

Item Size (All attributes):

2

KB

Number of items read per second:

0

Reads/Second

Read Consistency:



Strongly
Consistent



Eventually Consistent
(cheaper)

Number of items written per second:

300

Writes/Second



Amazon S3 is storage for the Internet. It is designed to make web-scale computing easier for developers.

Storage:

Storage:

1483

GB

Reduced Redundancy Storage:

0

GB

Requests:

PUT/COPY/POST/LIST Requests:

777600000

Requests

GET and Other Requests:

0

Requests

Amazon S3 or Amazon DynamoDB?

aws


SIMPLE MONTHLY CALCULATOR

Amazon S3 Standard

| | |
|-------------------|----------------|
| Storage | \$34 |
| Put/list requests | \$3,888 |
| Total | \$3,922 |

Amazon DynamoDB

| | |
|------------------------|--------------|
| Provisioned throughput | \$273 |
| Indexed data storage | \$383 |
| Total | \$656 |



[Scenario 1](#)

Amazon DynamoDB Wins!

Request rate (writes/sec)

300

Object size (bytes)

2,048

Total (GB/month)

1,483

Object per month

777,600,000


[Scenario 2](#)

300

32,768

23,730

777,600,000



Amazon S3 Wins!

Amazon S3 Standard

| | |
|-------------------|----------------|
| Storage | \$545 |
| Put/List Requests | \$3,888 |
| Total | \$4,433 |

Amazon DynamoDB

| | |
|------------------------|-----------------|
| Provisioned Throughput | \$4,556 |
| Indexed Data Storage | \$5,944 |
| Total | \$10,500 |

A large, solid orange arrow pointing to the right, centered on a dark gray background. The arrow has a black outline and contains the text "PROCESS / ANALYZE" in white, bold, sans-serif capital letters.

PROCESS /
ANALYZE

Predictive Analytics

Amazon AI

API-driven Services

- Amazon Lex – Chatbots
- Amazon Polly & Transcribe – Speech
- Amazon Rekognition – Image analysis
- Amazon Translate & Comprehend -- NLP

Managed ML Platforms

- Amazon SageMaker
- Apache Spark ML on EMR

AWS Deep Learning AMI


- Pre-installed with Apache MXNet, TensorFlow, Caffe, Theano, Torch, Microsoft Cognitive Toolkit, and Keras; plus DL tools/drivers

Developers

Data
scientists

Deep
learning
experts

PROCESS / ANALYZE



Interactive and Batch Analytics

Amazon ES

- Managed Service for Elasticsearch

Amazon Redshift and Amazon Redshift Spectrum

- Managed Data Warehouse
- Spectrum enables querying Amazon S3

Amazon Athena

- Serverless Interactive Query Service

Amazon EMR

- Managed Hadoop Framework for running Apache Spark, Flink, Presto, Tez, Hive, Pig, HBase, etc.



Stream/Real-time Analytics

Spark Streaming on Amazon EMR

Amazon Kinesis Analytics


- Managed Service for running SQL on Streaming data

Amazon KCL

- Amazon Kinesis Client Library

AWS Lambda

- Run code Serverless (without provisioning or managing servers)
- Services such as S3 can publish events to Lambda
- Lambda can pool event from a Kinesis



Which Analytics Should I Use?

Batch

Takes minutes to hours

Example: Daily/weekly/monthly reports

Amazon EMR (MapReduce, Hive, Pig, Spark)

Interactive

Takes seconds

Example: Self-service dashboards

Amazon Redshift, Amazon Athena, Amazon EMR (Presto, Spark)

Stream

Takes milliseconds to seconds

Example: Fraud alerts, 1 minute metrics

Amazon EMR (Spark Streaming), Amazon Kinesis Analytics, KCL, AWS Lambda, etc.

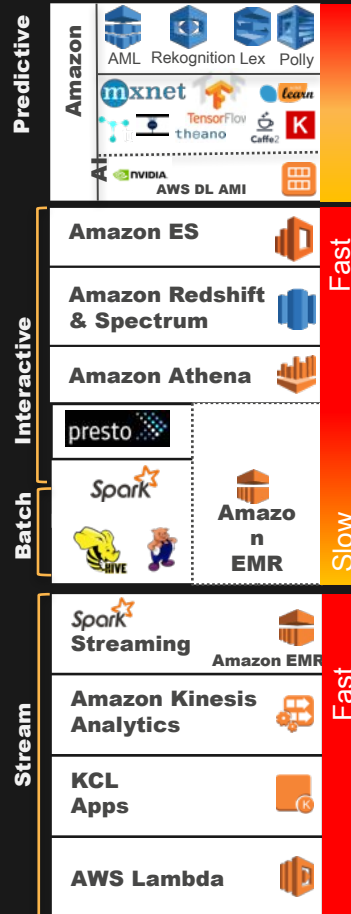
Predictive

Takes milliseconds (real-time) to minutes (batch)

Example: Fraud detection, Forecasting demand, Speech recognition

Amazon AI (Lex, Polly, ML, Amazon Rekognition), Amazon EMR (Spark ML), Deep Learning AMI (MXNet, TensorFlow, Theano, Torch, CNTK, and Caffe)

PROCESS / ANALYZE



Which Stream Processing Technology Should I Use?

| | Amazon EMR (Spark Streaming) | KCL Application | Amazon Kinesis Analytics | AWS Lambda |
|--------------------------|---------------------------------|----------------------------------|-------------------------------------|--------------------------|
| Managed Service | Yes | No (EC2 + Auto Scaling) | Yes | Yes |
| Serverless | No | No | Yes | Yes |
| Scale/Throughput | No limits / ~ nodes | No limits / ~ nodes | No limits / automatic | No limits / automatic |
| Availability | Single AZ | Multi-AZ | Multi-AZ | Multi-AZ |
| Programming Languages | Java, Python, Scala | Java, others via MultiLangDaemon | ANSI SQL with extensions | Node.js, Java, Python |
| Sliding Window Functions | Build-in | App needs to implement | Built-in | No |
| Reliability | KCL and Spark checkpoints | Managed by KCL | Managed by Amazon Kinesis Analytics | Managed by AWS Lambda |

Fast

Which Analytics Tool Should I Use?

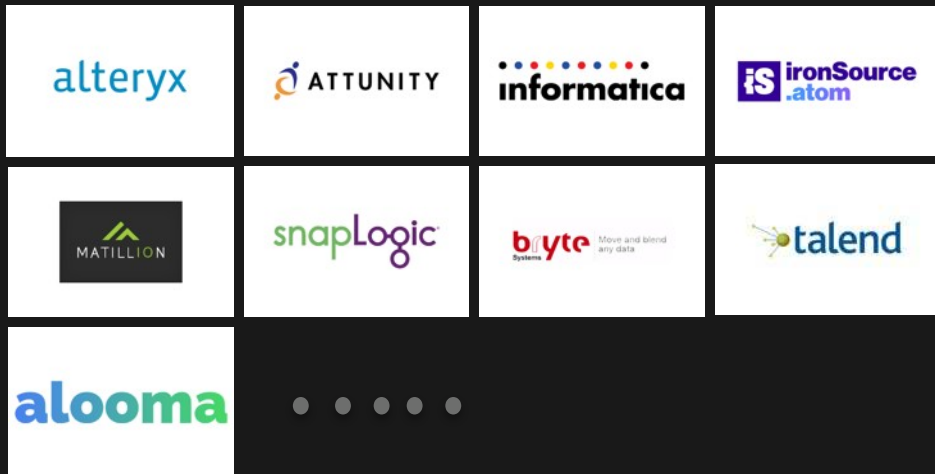
| | Amazon Redshift | Amazon Redshift Spectrum | Amazon Athena | Amazon EMR | | |
|----------------------|---|---|---|---------------------------------|-----------------|-------|
| | | | | Presto | Spark | Hive |
| Use case | Optimized for data warehousing | Query S3 data from Amazon Redshift | Interactive Queries over S3 data | Interactive Query | General purpose | Batch |
| Scale/Throughput | ~Nodes | ~Nodes | Automatic | ~ Nodes | | |
| Managed Service | Yes | Yes | Yes, Serverless | Yes | | |
| Storage | Local storage | Amazon S3 | Amazon S3 | Amazon S3, HDFS | | |
| Optimization | Columnar storage, data compression, and zone maps | AVRO, PARQUET TEXT, SEQ RCFILE, ORC, etc. | AVRO, PARQUET TEXT, SEQ RCFILE, ORC, etc. | Framework dependent | | |
| Metadata | Amazon Redshift Catalog | Glue Catalog | Glue Catalog | Glue Catalog or Hive Meta-store | | |
| Auth/Access controls | IAM, Users, groups, and access controls | IAM, Users, groups, and access controls | IAM | IAM, LDAP & Kerberos | | |
| UDF support | Fastest | | Fast | Slow | | |

What About Extract Transform and Load?



Data Integration Partners

Reduce the effort to move, cleanse, synchronize, manage, and automatize data-related processes.



AWS Glue



Data Catalog



Job Authoring



Job Execution

AWS Glue is a fully managed (serverless) ETL service that makes it simple and cost-effective to categorize your data, clean it, enrich it, and move it reliably between various data stores.

COLLECT

STORE

ETL

PROCESS/ANALYZE

CONSUME

Applications

Web apps



Mobile apps

Data centers
AWS Direct
Connect

RECORDS

Cache

Amazon
ElastiCache

Hot

Amazon DAX

Amazon
DynamoDB

SQL

Amazon RDS



Amazon Aurora



Warm

Data Transport & Logging

Logging

AWS
CloudTrailAmazon
CloudWatch

Migration



FILES

File

Amazon S3



IoT

Devices
Sensors

IoT platforms



AWS IoT



STREAMS


Stream

Apache Kafka

Amazon Kinesis
StreamsAmazon Kinesis
Firehose

Hot

Predictive



Interactive



Batch

Stream

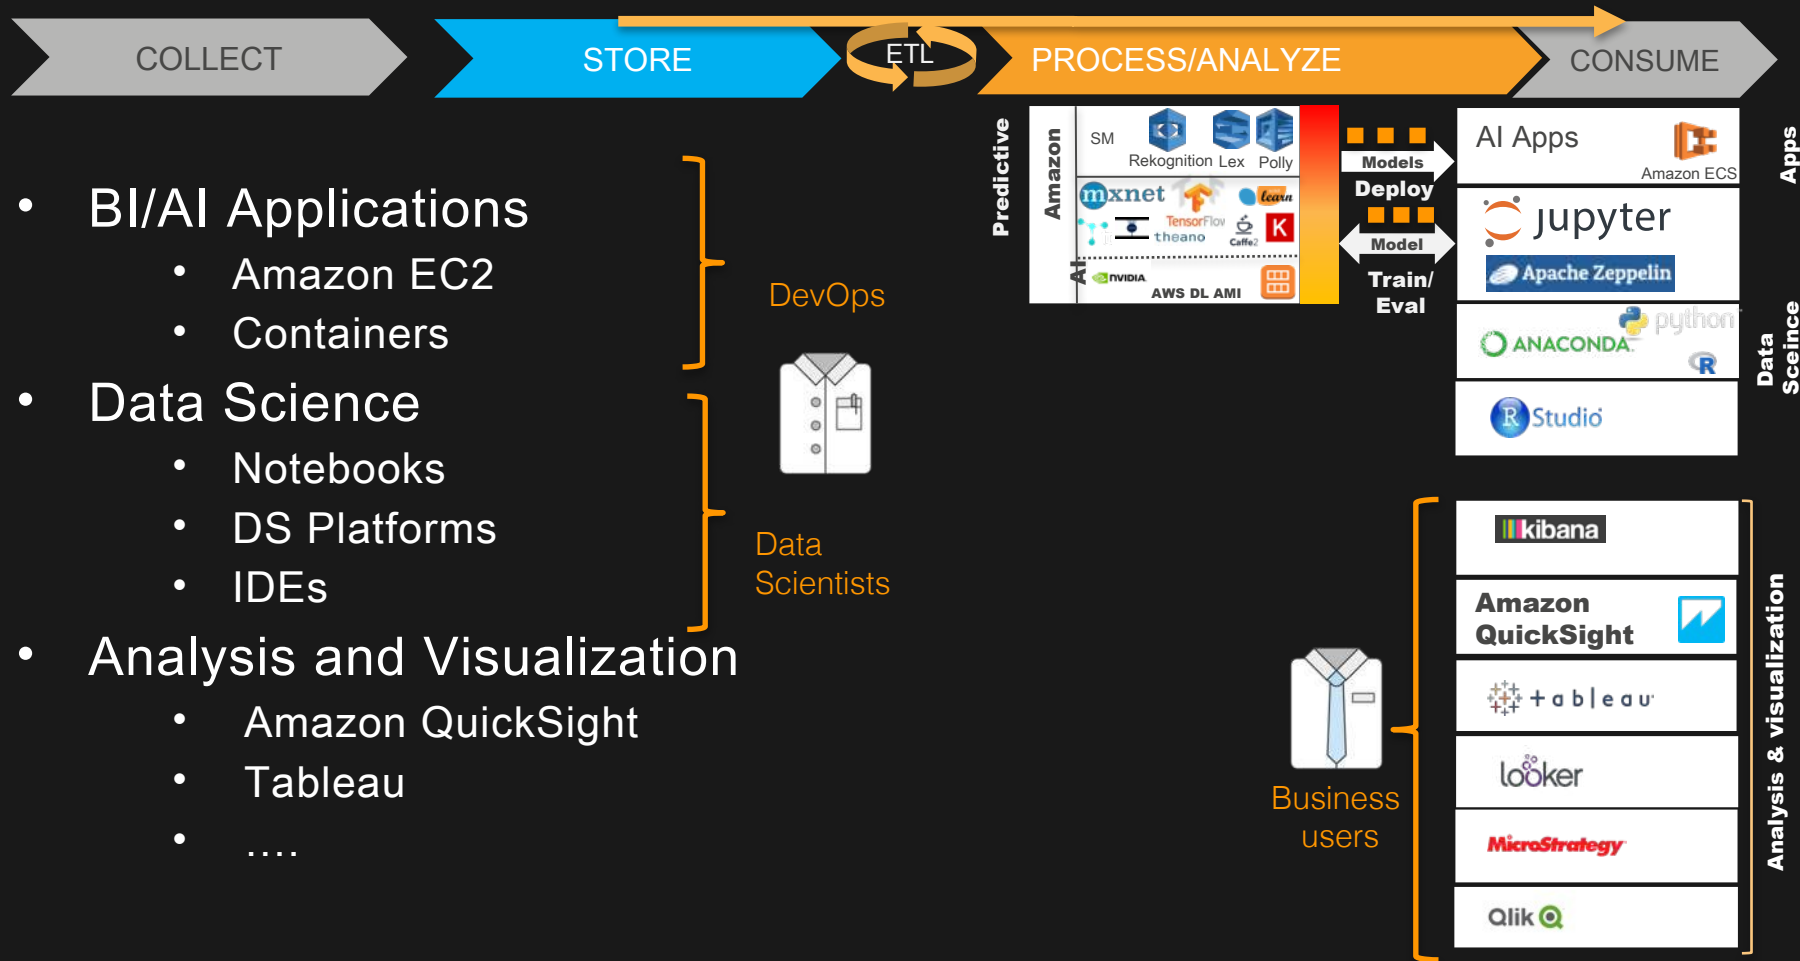


Fast

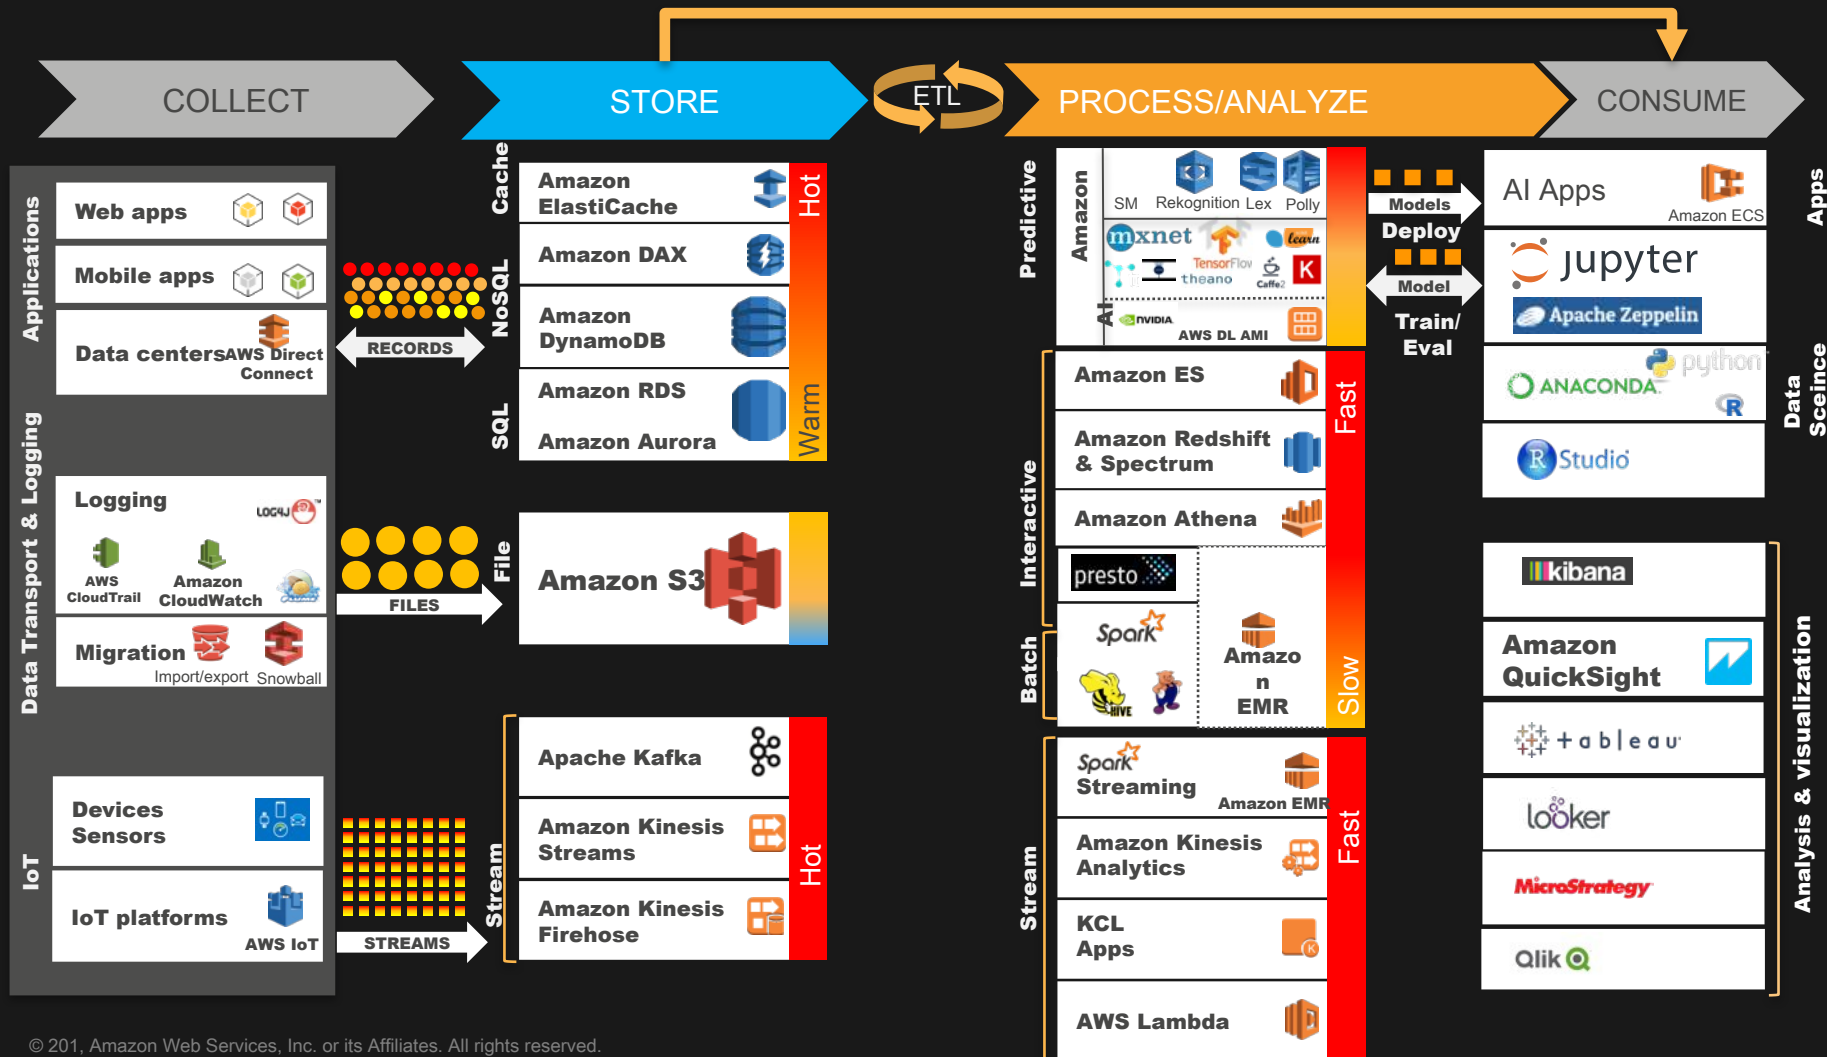
Fast



CONSUME

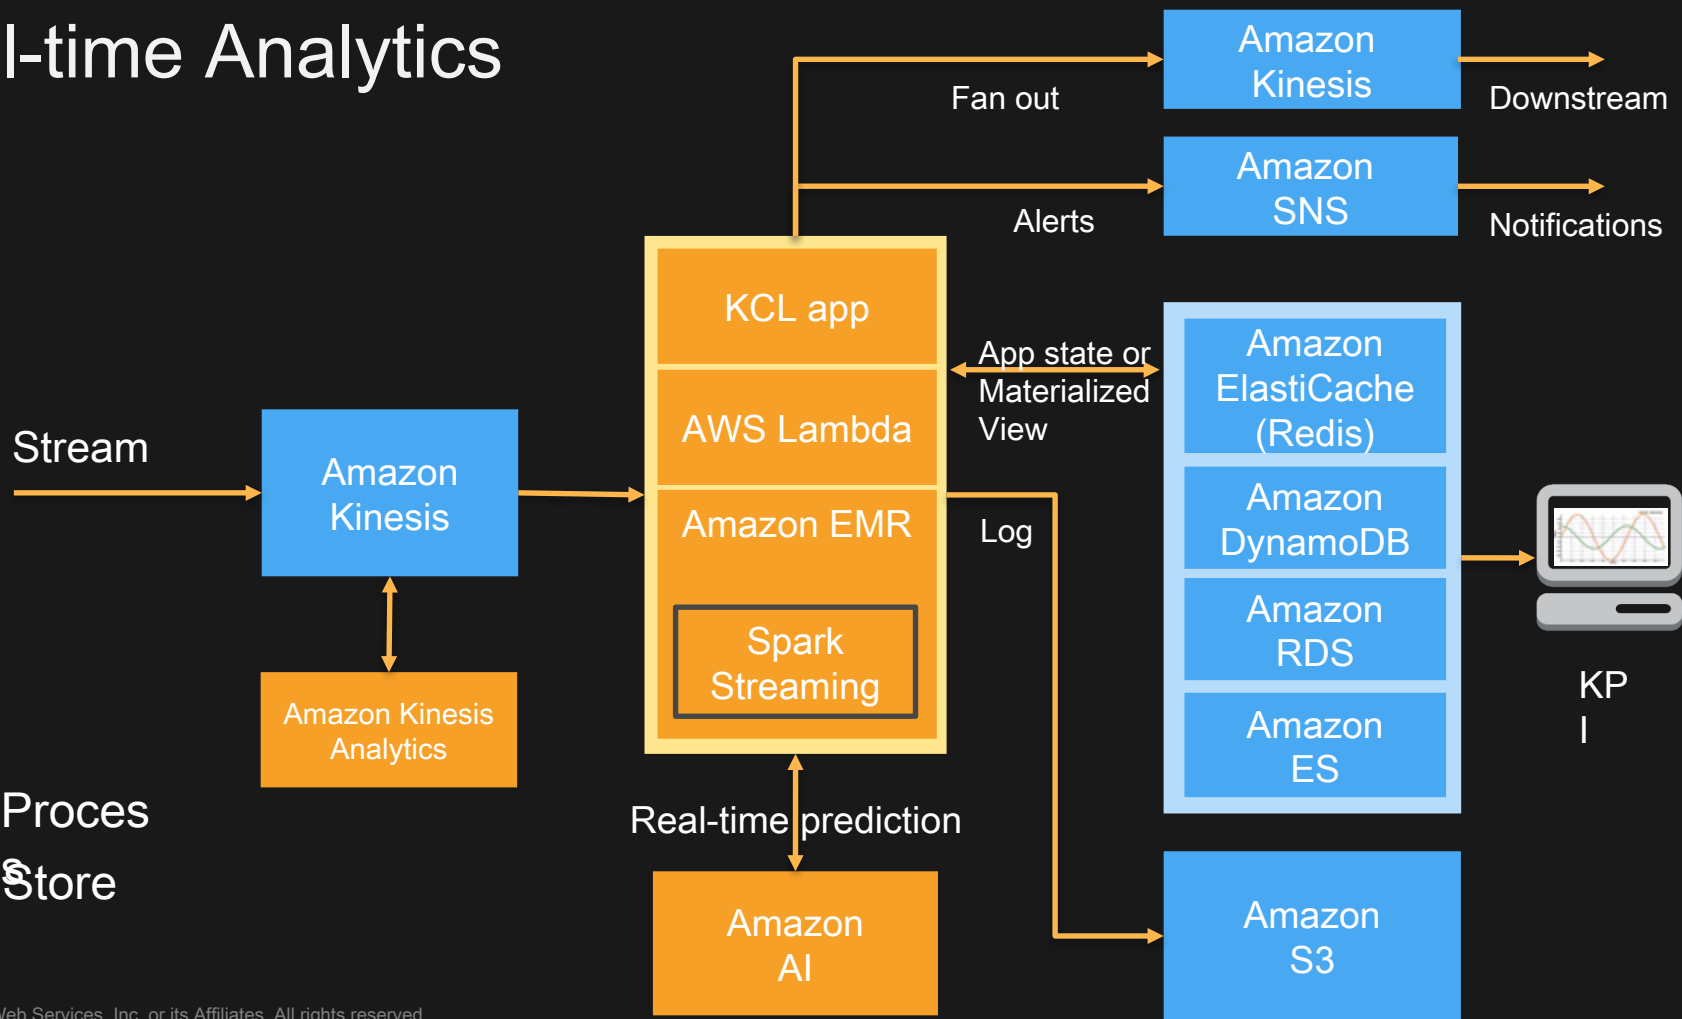


Putting It All Together

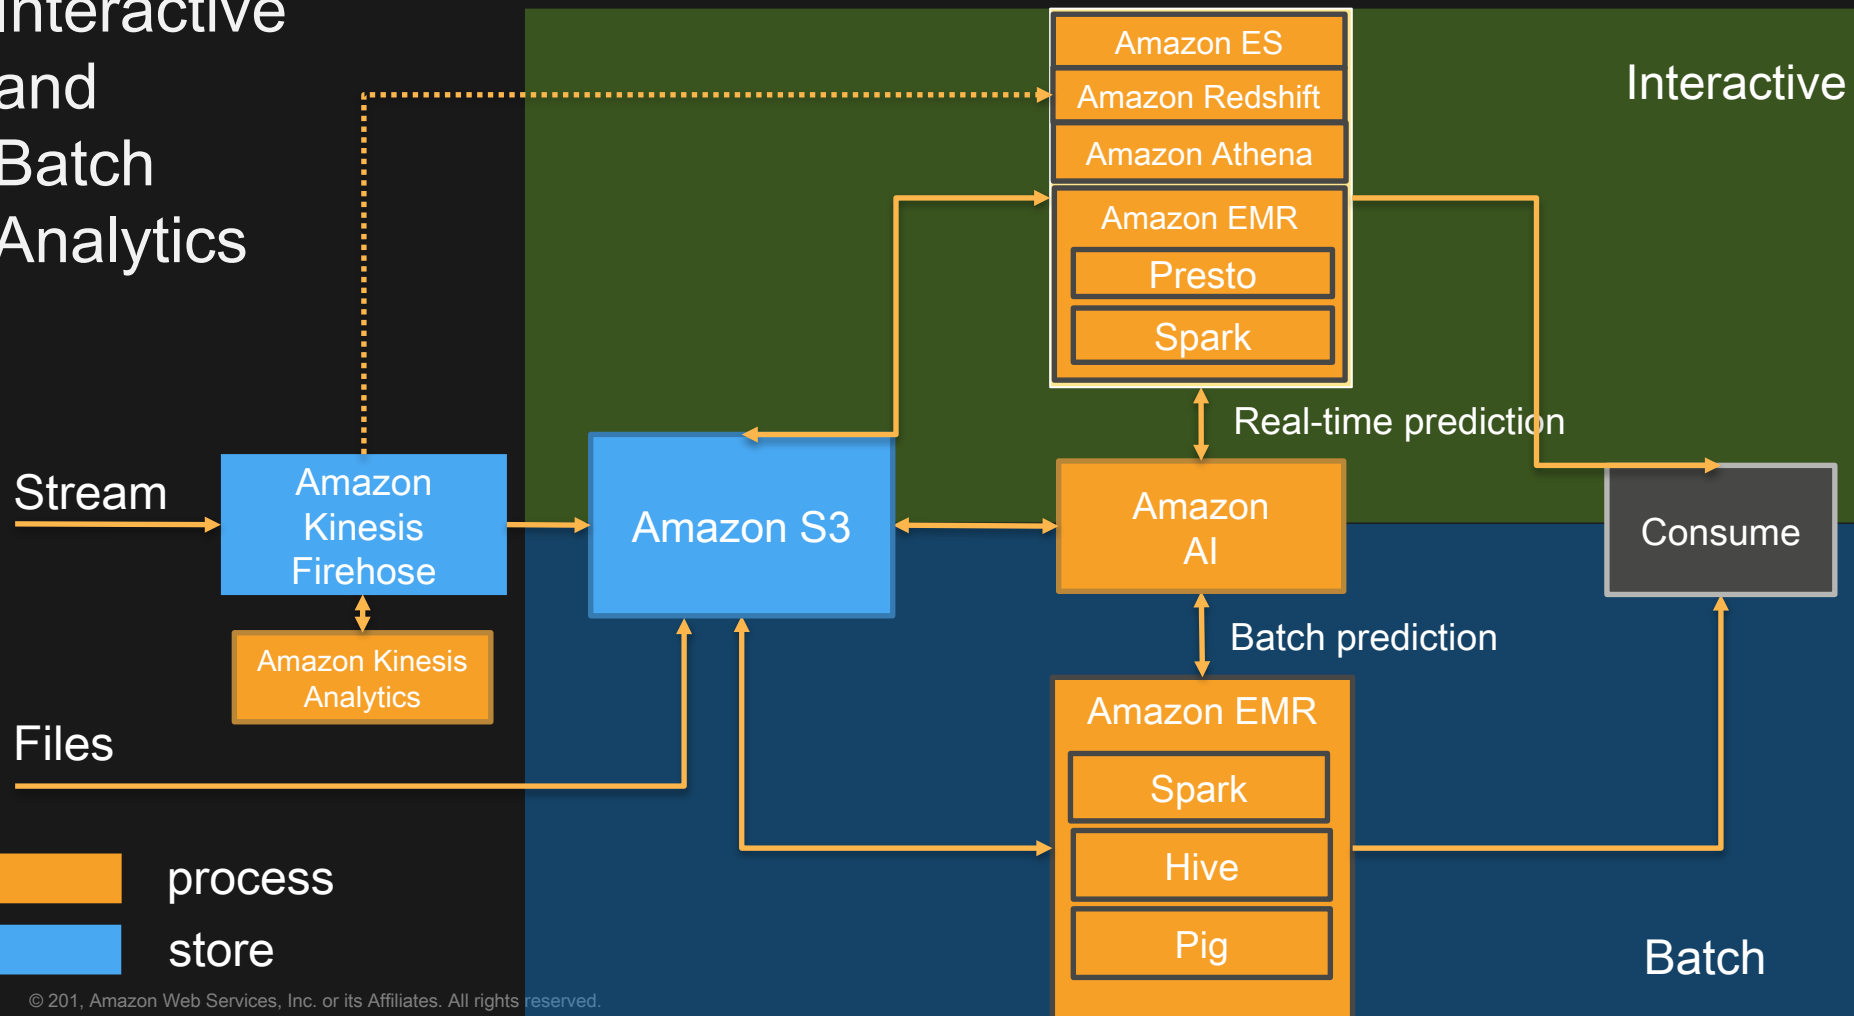


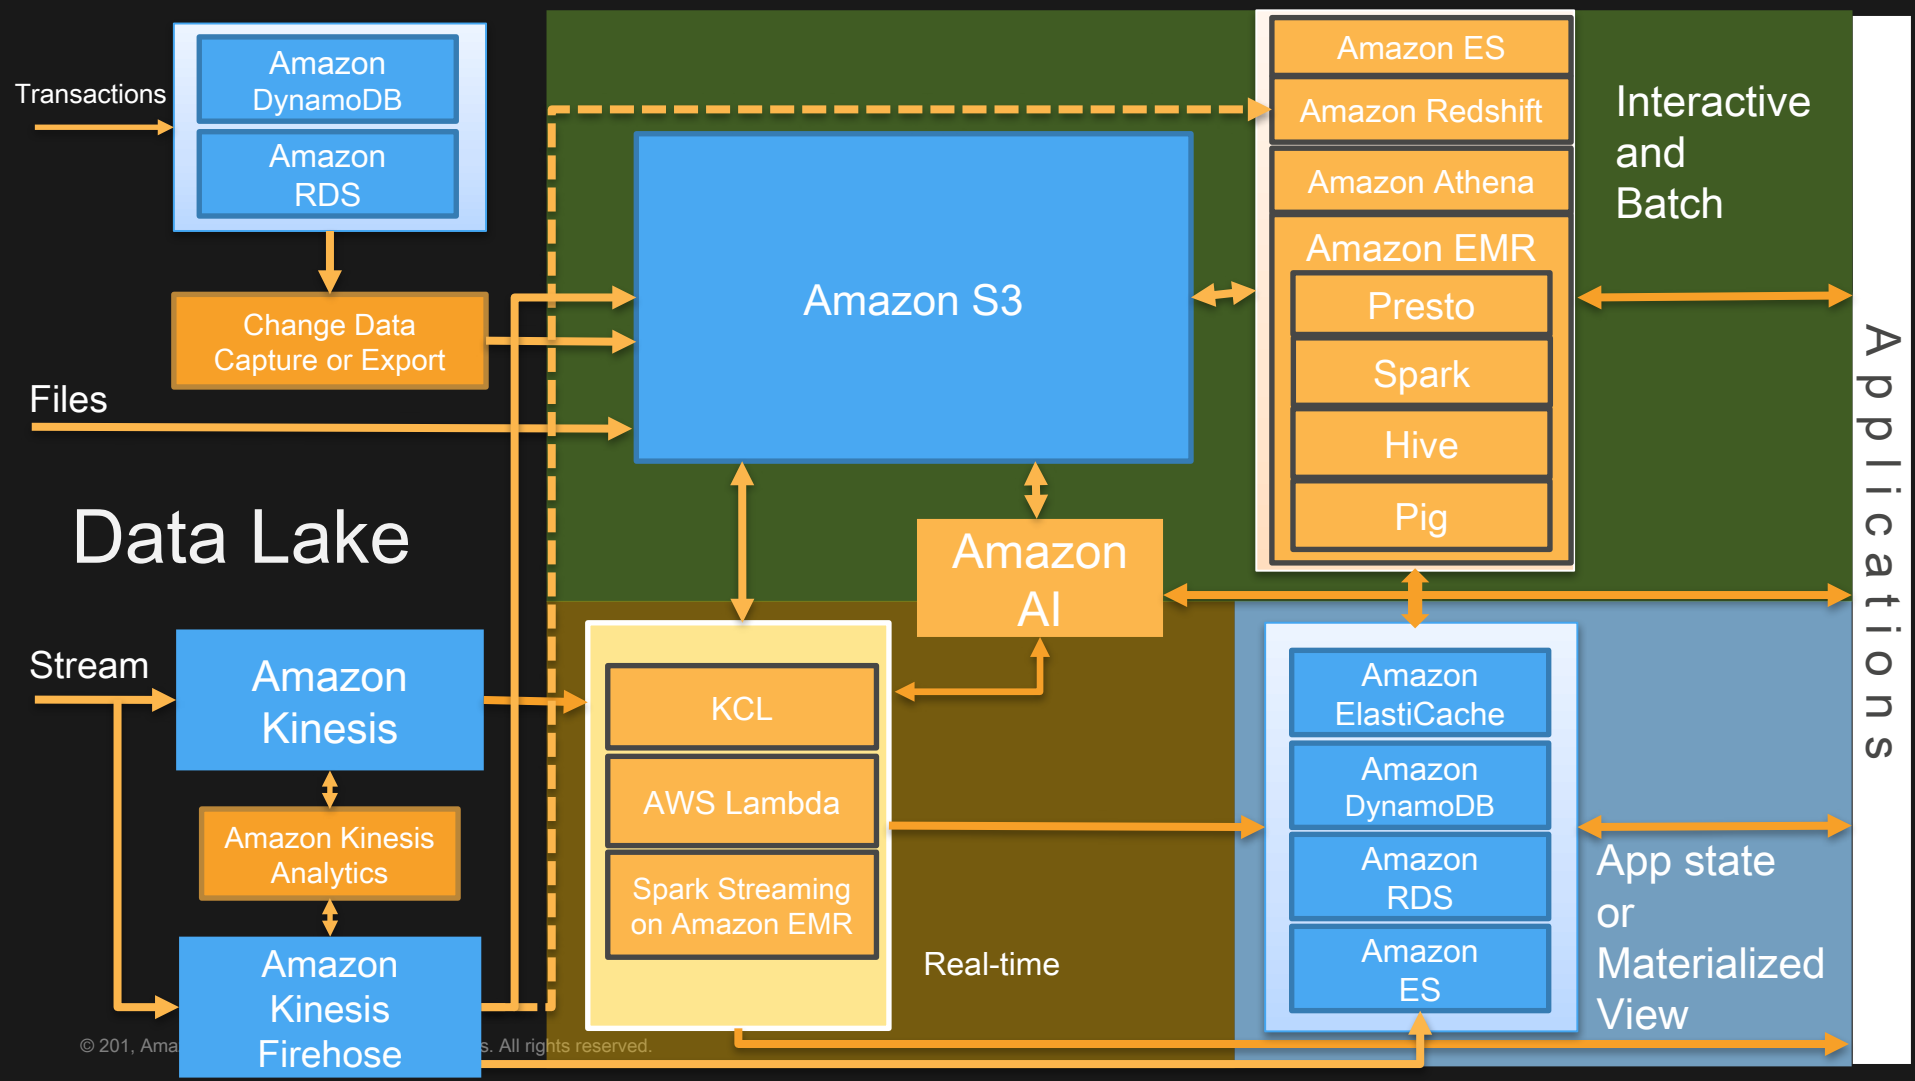
Design Patterns

Real-time Analytics



Interactive and Batch Analytics





What about Metadata?

- Glue Catalog
 - Hive Metastore compliant
 - Crawlers - Detect new data, schema, partitions
 - Search - Metadata discovery
 - Amazon Athena, Amazon EMR, and Amazon Redshift Spectrum compatible
- Hive Metastore (Presto, Spark, Hive, Pig)
 - Can be hosted on Amazon RDS



Security & Governance

- AWS Identity and Access Management (IAM)
- Amazon Cognito
- Amazon CloudWatch & AWS CloudTrail
- AWS KMS
- AWS Directory Service
- Apache Ranger

Security &
Governance



IAM



Amazon
Cognito



Amazon
CloudWatch



AWS
CloudTrail



AWS
KMS



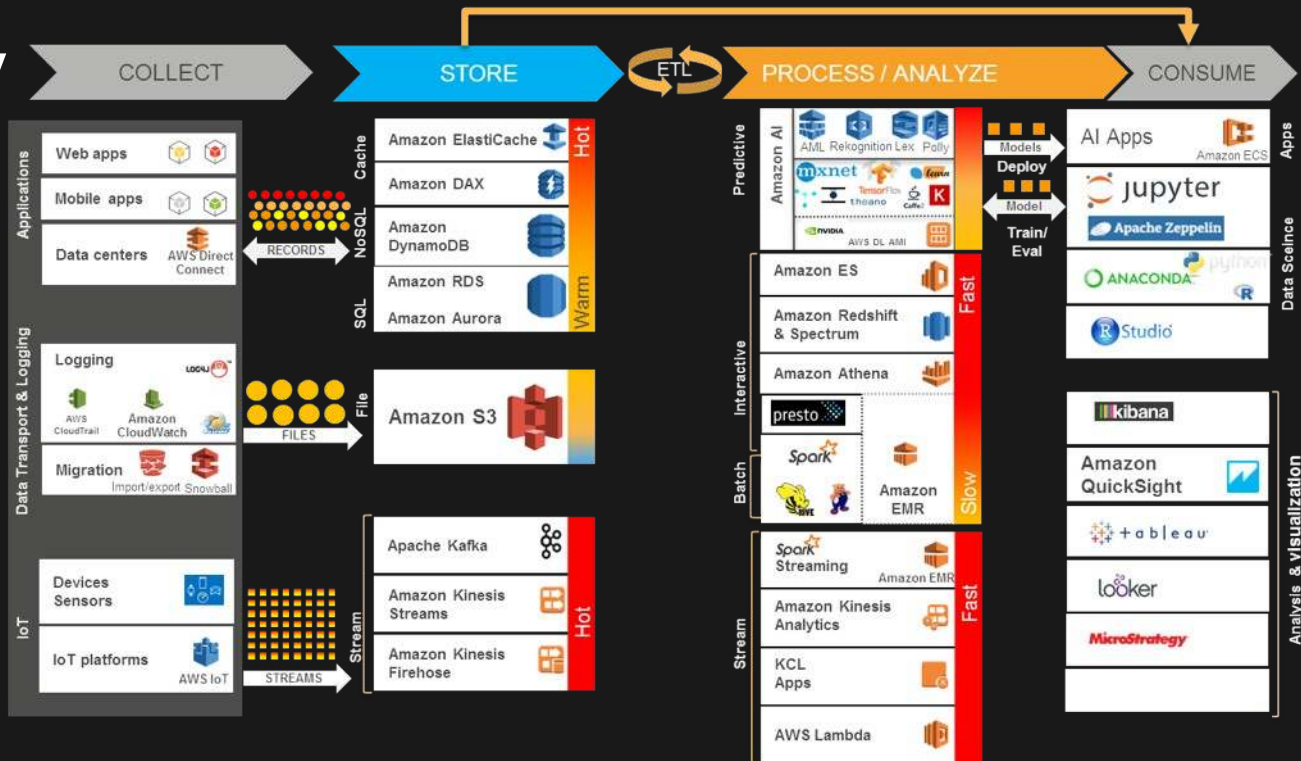
AWS
CloudHSM



AWS Directory
Service

Apache Ranger

Summary



Data lake Reference Architecture



Summary



Build decoupled systems

- Data → Store → Process → Store → Analyze → Answers

Use the right tool for the job

- Data structure, latency, throughput, access patterns

Leverage AWS managed and serverless services

- Scalable/elastic, available, reliable, secure, no/low admin

Use log-centric design patterns

- Immutable logs, data lake, materialized views

Be cost-conscious

- Big data ≠ Big cost

AI/ML enable your applications

Resources

<https://aws.amazon.com/big-data/>

<https://aws.amazon.com/whitepapers/>

<https://aws.amazon.com/blogs/big-data/>

<https://aws.amazon.com/blogs/big-data/aws-big-data-analytics-sessions-at-reinvent-2017/>



Thank you!

Julien Simon
Principal Technical Evangelist
@julsimon