

CITY NAME

# DEV DAY



© 2020, Amazon Web Services, Inc. or its affiliates. All rights reserved.

**DEV** DAY  
CITY NAME

MLT3

# Natural Language Processing: concepts, algorithms & use cases

Speaker Name  
Job Title  
Company/Org Name



# What to expect

## 1 – Word Vectors

## 2 – Algorithms

- Word2Vec
- GloVe
- FastText, BlazingText
- ELMo
- BERT
- XLNet

## 3 – Use cases & demos

- Word vectors
- Word similarity
- Word analogy
- Text classification
- Sentiment analysis

## 4 – Getting started

# Problem statement

- NLP is a **major** field in AI
  - Text classification, machine translation, text generation, chat bots, vocal assistants, etc.
  - You could even say that strong AI requires efficient NLP
- NLP apps require a **language model** in order to predict the next word
  - Given a sequence of words  $(w_1, \dots, w_n)$ , predict  $w_{n+1}$  that has the highest probability
- Vocabulary size can be **hundreds of thousands** of words  
... in **millions of documents**
- Can we build a compact **mathematical representation** of language, that will help with a variety of downstream NLP tasks?

« *You shall know a word by the company it keeps* », Firth (1957)

- **Word vectors** are built from co-occurrence counts
  - Also called **word embeddings**
  - High dimensional: at least 50, up to 300
- Words with **similar meanings** should have **similar vectors**
  - "car"  $\approx$  "automobile"  $\approx$  "sedan"
- The distance between vectors for the same concepts should be similar
  - distance("Paris", "France")  $\approx$  distance("Berlin", "Germany")
  - distance("hot", "hotter")  $\approx$  distance("cold", "colder")

# High-level view

1. Start from a **large text corpus** (100s of millions of words, even billions)
  2. Preprocess the corpus
    - **Tokenize**: « hello, world! » → « <BOS>hello<SP>world<SP>!<EOS> »
    - **Multi-word entities**: « Rio de Janeiro » → « rio\_de\_janeiro »
  3. Build the **vocabulary**
    - Remove very rare words?
  4. Learn **vector representations** for all words
- ... or simply use (or fine-tune) **pre-trained vectors** (more on this later)



DEV DAY

# Algorithms





# Word2Vec (2013)

<https://arxiv.org/abs/1301.3781>

<https://code.google.com/archive/p/word2vec/>

- **Continuous bag of words (CBOW):**
  - the model **predicts the current word** from surrounding context words
  - Word **order** doesn't matter (hence the 'bag of words')
- **Skipgram**
  - the model **uses the current word to predict** the surrounding window of context words
  - This may work better when little data is available
- CBOW trains faster, skipgram is more accurate
- C code, based on shallow neural network





# Global Vectors aka GloVe (2014)

<https://nlp.stanford.edu/projects/glove/>

<https://github.com/stanfordnlp/GloVe>

<https://www.quora.com/How-is-GloVe-different-from-word2vec>

- Performance generally similar to Word2Vec
- Pre-trained models: up to 840 billion tokens, 2.2 million vocabulary, 300 dimensions
- C code, based on matrix factorization

# FastText (2016)

<https://arxiv.org/abs/1607.04606>

<https://arxiv.org/abs/1802.06893>

<https://fasttext.cc/>

<https://www.quora.com/What-is-the-main-difference-between-word2vec-and-fastText>

- Extension of Word2Vec: each word is treated as a **set of subwords** aka character n-grams
  - « Computer », n=5 : <START>Comp , compu, omput, mpute, puter, uter<END>
  - A word vector is the sum or average of its **subword vectors**
- Subwords help with **rare/unknown/mispelled words**, as they share **subwords** with known words
  - « Computerization » and « Cmputer » should be close to « Computer »
- Unsupervised learning: **compute word vectors**, with pre-trained vectors for 294 languages
- Supervised learning: **use word vectors** for multi-label, multi-class text classification
- Also language detection for 170 languages
- Multithreaded C++ code, with Python API

# BlazingText (2017)

<https://dl.acm.org/citation.cfm?id=3146354>

<https://aws.amazon.com/blogs/machine-learning/enhanced-text-classification-and-word-vectors-using-amazon-sagemaker-blazingtext/>

- Amazon-invented algorithm, available in Amazon SageMaker
- Extends FastText with GPU capabilities
- Unsupervised learning: word vectors
  - 20x faster
  - CBOW and skip-gram with subword support
  - Batch skip-gram for distributed training
- Supervised learning: text classification
  - 100x faster
  - Models are compatible with FastText

	Word2Vec (unsupervised learning)			Text Classification (supervised learning)
Modes	Skip-gram (supports subwords)	CBOW (supports subwords)	batch_skipgram	supervised
Single CPU instance	✓	✓	✓	✓
Single GPU instance (with 1 or more GPUs)	✓	✓		✓*
Multiple CPU instances			✓	

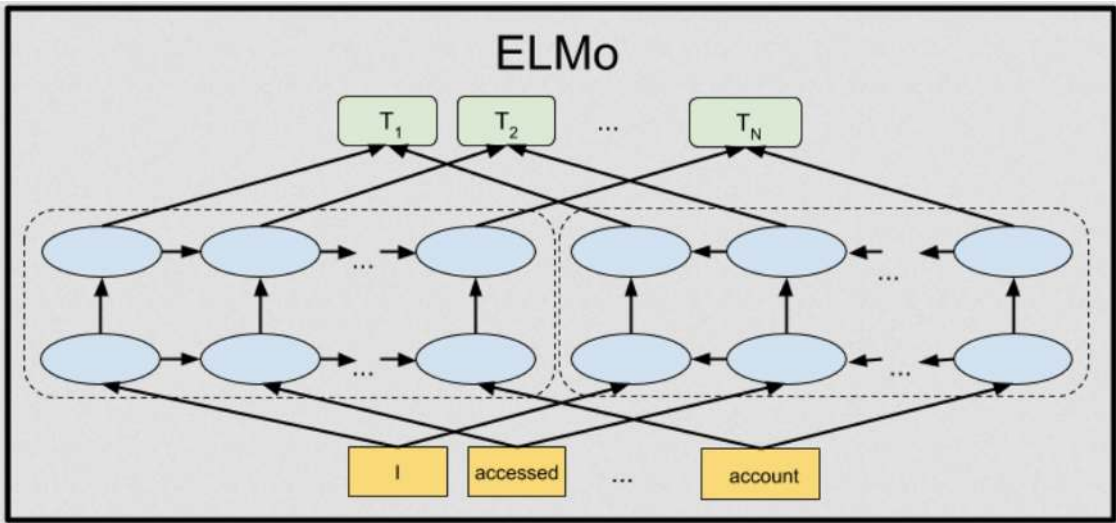
# Limitations of Word2Vec (and family)

- Some words have different meanings (aka **polysemy**)
  - « *Kevin, stop throwing **rocks!*** » vs. « *Machine Learning **rocks*** »
  - Word2Vec encodes the **different meanings** of a word into the **same vector**
- Bidirectional context is not taken into account
  - Previous words (**left-to-right**) and next words (**right-to-left**)

# Embeddings from Language Models aka ELMo (02/2018)

<https://arxiv.org/abs/1802.05365>  
<https://allennlp.org/elmo>

- ELMo generates a context-aware vector for each word
  - Character-level CNN
  - Bidirectional context, with two unidirectional LSTMs → No cheating possible (can't peek at the future)
  - “Deep” embeddings, reflecting output from all layers
- No vocabulary, no vector file:  
you need to use the model itself
- Reference implementation with TensorFlow



Source:  
Google

Source		Nearest Neighbors
GloVe	play	playing, game, games, played, players, plays, player, Play, football, multiplayer
	Chico Ruiz made a spectacular <u>play</u> on Alusik 's grounder {...}	Kieffer , the only junior in the group , was commended for his ability to hit in the clutch , as well as his all-round excellent play .
biLM	Olivia De Havilland signed to do a Broadway <u>play</u> for Garson {...}	{...} they were actors who had been handed fat roles in a successful <u>play</u> , and had talent enough to fill the roles competently , with nice understatement .

Source: ELMo  
paper


# Bidirectional Encoder Representations from Transformers aka BERT (10/2018)

<https://arxiv.org/abs/1810.04805>

<https://github.com/google-research/bert>

<https://www.quora.com/What-are-the-main-differences-between-the-word-embeddings-of-ELMo-BERT-Word2vec-and-GloVe>

- BERT improves on ELMo
  - Replace LSTM with **Transformers**, which deal better with long-term dependencies
  - Truly bidirectional architecture: left-to-right and right-to-left contexts are **learned by the same network**
  - Words are **randomly masked** during training to prevent cheating
- Pre-trained models: **BERT Base** and **BERT Large**
  - Masked word prediction
  - Next sentence prediction
- Reference implementation with TensorFlow



Source:  
Google



# Limitations of BERT

- BERT cannot handle more than 512 input tokens
- BERT masks words during training, but not during fine-tuning (aka training/fine-tuning discrepancy)
- BERT isn't trained to predict the next word, so it's not great at text generation
- BERT doesn't learn dependencies for masked words
  - Train « I am going to <MASK> my <MASK> » on « walk » / « dog », « eat » / « breakfast », and « debug » / « code ».
  - BERT could legitimately predict « I am going to eat my code » or « I am going to debug my dog » :-/

# XLNet (06/2019)

<https://arxiv.org/abs/1906.08237>

<https://github.com/zihangdai/xlnet>

- XLNet beats BERT at 20 tasks
- XLNet uses **bidirectional context**, but words are **randomly permuted**
  - No cheating possible
  - No masking required
- **XLNet Base** and **XLNet Large**
- Reference implementation with TensorFlow

07/2019: ERNIE 2.0 (Baidu)

beats BERT & XLNet

<https://github.com/PaddlePaddle/ERNIE/>

# Train yourself or not?

- Word2Vec and friends
  - Try **pre-trained embeddings** first
    - Check that the training corpus is similar to your own data
    - Same language, similar vocabulary
  - Remember that subword models will help with unknown / misspelled words
  - If you have exotic requirements AND lots of data, training is not expensive
- EIMo, BERT, XLNet
  - Training is very **expensive**: several days using several GPUs
  - Fine-tuning is **cheap**: just a few GPU hours for SOTA results
  - Fine-tuning scripts and pre-trained models are available: start there!
- In both cases, you'll still have to pre-process data (yeaaaaah)

DEV DAY

# Use cases & demos



# Demo: training Word2Vec subword vectors with BlazingText on Amazon SageMaker

[https://github.com/aws-labs/amazon-sagemaker-examples/tree/master/introduction to amazon algorithms/blazingtext word2vec subwords text8](https://github.com/aws-labs/amazon-sagemaker-examples/tree/master/introduction%20to%20amazon%20algorithms/blazingtext_word2vec_subwords_text8)

# Word similarity

- Words with a **similar meaning** are expected to have **similar vectors**


- ‘cosine similarity’, i.e. normalized dot product

-1 → words are not similar

+1 → words are similar

$$\text{similarity} = \cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\| \|\mathbf{B}\|} = \frac{\sum_{i=1}^n A_i B_i}{\sqrt{\sum_{i=1}^n A_i^2} \sqrt{\sum_{i=1}^n B_i^2}},$$

- Using word vectors for your vocabulary:
  - Pick a word
  - Compute the cosine similarity of its vector with respect to all other vectors
  - Keep the top ‘k’ cosines similarities
  - Return the corresponding words

nearest neighbors of <i>frog</i>	Litoria	Leptodactylidae	Rana	Eleutherodactylus
Pictures				




Source:  
GloVe





# Word analogy

- The **distance** between two word vectors defines the **relationship** between the two words.
- A **similar distance** between two other vectors reflects a **similar relationship**
- « King » - « Man »  $\approx$  « Queen » - « Woman »
- « King » - « Man » + « Woman »  $\approx$  « Queen »
- Meaning: « Man » is to « King » what « Woman » is to « Queen »
- Now we can ask :« Paris » is to « France » what « Rome » is to... ?
- Answer: vector closest to « France » - « Paris » + « Rome », hopefully « Italy » ☺



Source:  
GloVe

# Demo: finding similarities and analogies with Gluon NLP and pre-trained GloVe embeddings

<https://gitlab.com/juliensimon/dlnotebooks/gluonnlp/>

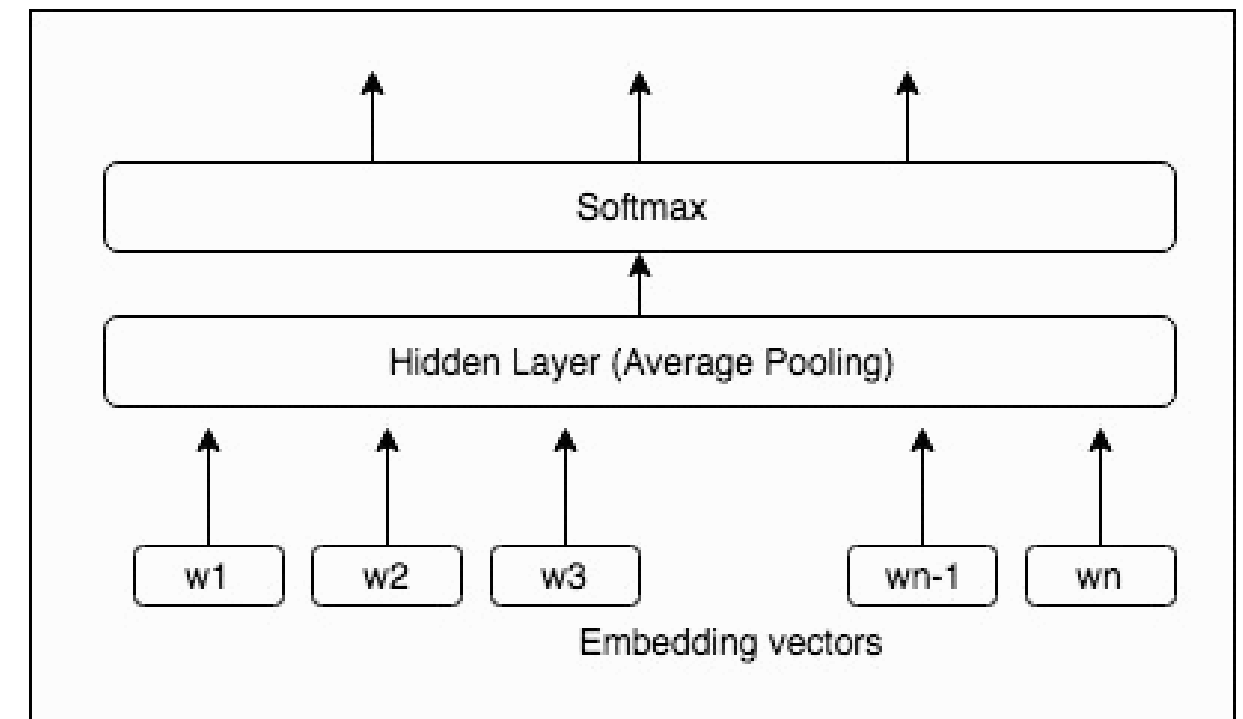
# Demo: embeddings with ELMo on TensorFlow

<https://gitlab.com/juliensimon/dlnotebooks/blob/master/nlp/ELMO%20TensorFlow.ipynb>

# Text classification

Sentiment analysis, spam detection, sentence pair comparison, etc.

1. Build a dataset of **labeled sentences**
2. Grab a pre-trained model, and add a classification layer
3. Convert each sentence to a list of vectors
4. Train or fine-tune the model to predict the correct class



Source:  
Wikipedia

# Demo: sentiment analysis on movie review with ktrain and pre-trained BERT

<https://gitlab.com/juliensimon/dlnotebooks/ktrain/>

# Getting started on AWS

<https://ml.aws>

<https://aws.amazon.com/marketplace/solutions/machine-learning/natural-language-processing>

<https://aws.amazon.com/sagemaker>

<https://github.com/aws-labs/amazon-sagemaker-examples>



DEV DAY

# Thank you!

