

DEV DAY

Optimize your Machine Learning workloads

Julien Simon
Global Evangelist, AI & Machine Learning
@julsimon



Amazon SageMaker

Build



Pre-built
notebooks for
common
problems



Built-in, high-
performance
algorithms

Train



One-click
training



Hyperparameter
optimization


Deploy




One-click
deployment



Fully managed
hosting with auto-
scaling

 P3DN, C5N
TensorFlow on 256 GPUs
Dynamic Training on MXNet
Automatic Model Tuning

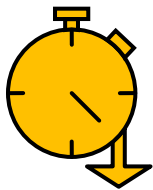
 Model compilation
Elastic inference
Inference pipelines

Optimizing infrastructure

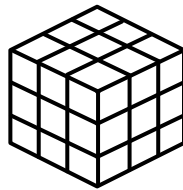


Amazon EC2 P3dn

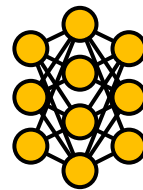
<https://aws.amazon.com/blogs/aws/new-ec2-p3dn-gpu-instances-with-100-gbps-networking-local-nvme-storage-for-faster-machine-learning-p3-price-reduction/>



Reduce machine
learning training time



Better GPU
utilization



Support larger, more
complex models

KEY FEATURES

100Gbps of
networking
bandwidth

8 NVIDIA Tesla
V100 GPUs

32GB of
memory per
GPU
(2x more P3)

96 Intel
Skylake vCPUs
(50% more than P3)
with AVX-512

Amazon EC2 C5n

<https://aws.amazon.com/blogs/aws/new-c5n-instances-with-100-gbps-networking/>

Intel Xeon Platinum 8000

Up to 3.5GHz single core speed

Up to 100Gbit networking

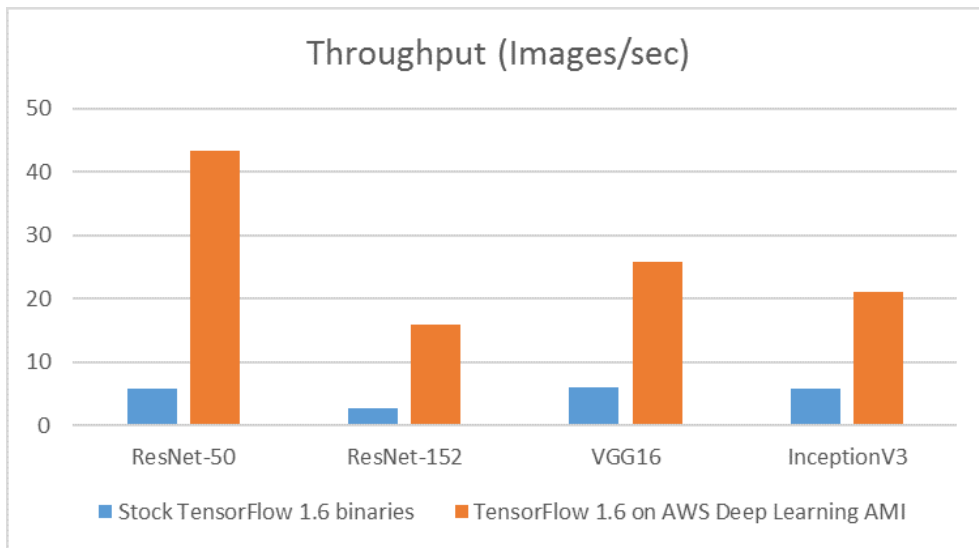
Based on Nitro hypervisor for bare metal-like performance

Instance Name	vCPUs	RAM	EBS Bandwidth	Network Bandwidth
c5n.large	2	5.25 GiB	Up to 3.5 Gbps	Up to 25 Gbps
c5n.xlarge	4	10.5 GiB	Up to 3.5 Gbps	Up to 25 Gbps
c5n.2xlarge	8	21 GiB	Up to 3.5 Gbps	Up to 25 Gbps
c5n.4xlarge	16	42 GiB	3.5 Gbps	Up to 25 Gbps
c5n.9xlarge	36	96 GiB	7 Gbps	50 Gbps
c5n.18xlarge	72	192 GiB	14 Gbps	100 Gbps

Optimizing frameworks



Making TensorFlow faster



<https://aws.amazon.com/blogs/machine-learning/faster-training-with-optimized-tensorflow-1-6-on-amazon-ec2-c5-and-p3-instances/>
(March 2018)

Training a ResNet-50 benchmark with the synthetic ImageNet dataset using our optimized build of TensorFlow 1.11 on a c5.18xlarge instance type is **11x faster** than training on the stock binaries.

<https://aws.amazon.com/about-aws/whats-new/2018/10/chainer4-4-theano-1-0-2-launch-deep-learning-ami/> (October 2018)

Scaling TensorFlow near-linearly to 256 GPUs

<https://aws.amazon.com/about-aws/whats-new/2018/11/tensorflow-scalability-to-256-gpus/>

Stock
TensorFlow

65%

scaling efficiency
with 256 GPUs



AWS-Optimized
TensorFlow

90%

scaling efficiency
with 256 GPUs

Available with
Amazon SageMaker
and the AWS Deep
Learning AMIs

30m

training time



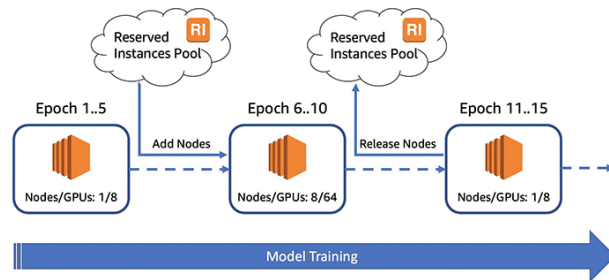
14m

training time

Dynamic training with Apache MXNet and RIs

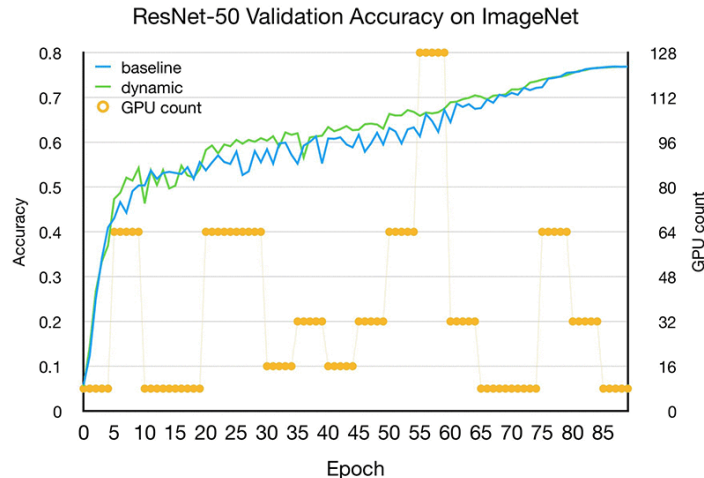
<https://aws.amazon.com/blogs/machine-learning/introducing-dynamic-training-for-deep-learning-with-amazon-ec2/>

Use a variable number of instances for distributed training



No loss of accuracy

Coming soon
spot instances,
additional frameworks



Optimizing models



Examples of hyperparameters

XGBoost

Tree depth
Max leaf nodes
Gamma
Eta
Lambda
Alpha
...

Neural Networks

Number of layers
Hidden layer width
Learning rate
Embedding
dimensions
Dropout
...

Automatic Model Tuning

Finding the optimal set of hyper parameters

1. **Manual Search** ("I know what I'm doing")
2. **Grid Search** ("X marks the spot")
 - Typically training hundreds of models
 - Slow and expensive
3. **Random Search** ("Spray and pray")
 - Works better and faster than Grid Search
 - But... but... but... it's random!
4. **HPO**: use Machine Learning
 - Training fewer models
 - **Gaussian Process Regression** and **Bayesian Optimization**
 - You can now resume from a previous tuning job



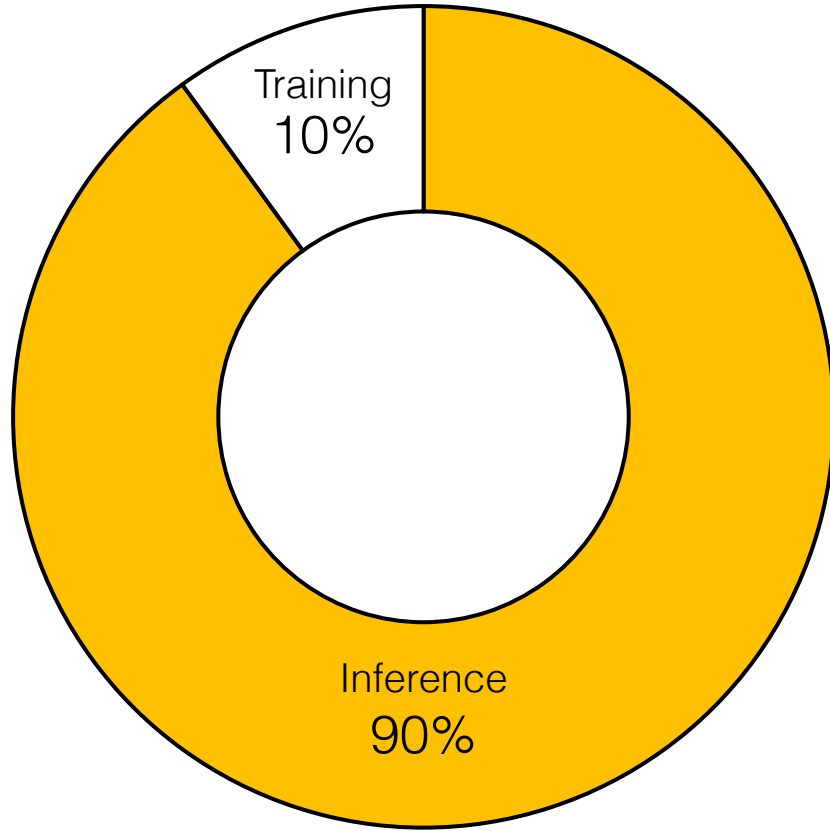
Demo: HPO with Apache MXNet

https://github.com/aws-labs/amazon-sagemaker-examples/tree/master/hyperparameter_tuning/hpo_mxnet_mnist.ipynb

Optimizing inference



Predictions drive
complexity and
cost in production



Model optimization is extremely complex

mxnet

TensorFlow

PYTORCH

intel

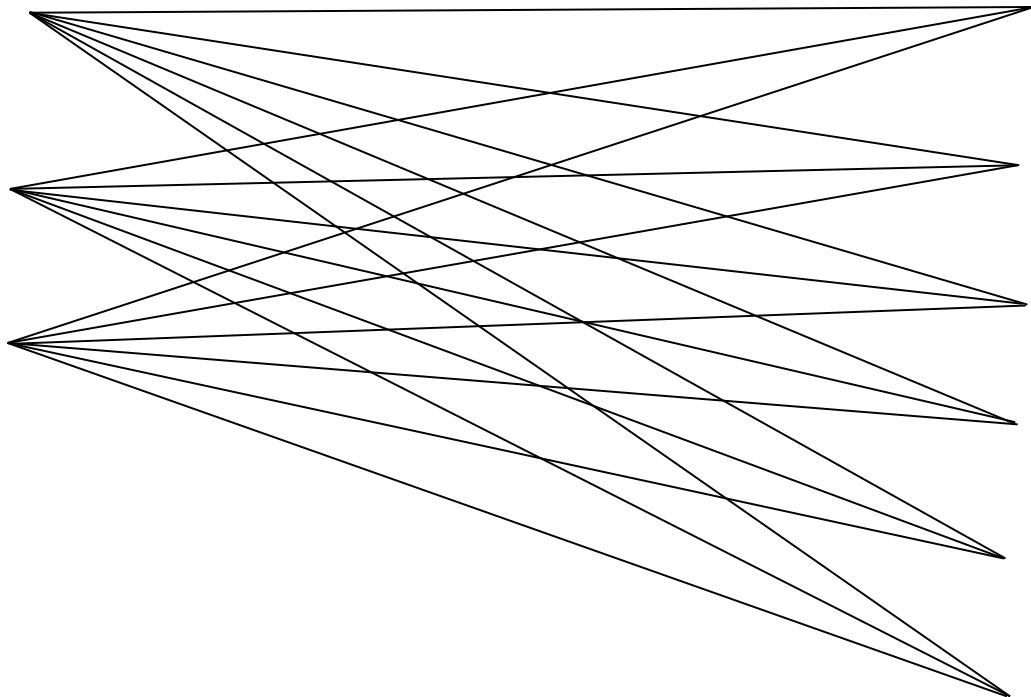
nvidia

Qualcomm

XILINX

cadence

arm



Amazon Neo: compiling models

<https://aws.amazon.com/blogs/aws/amazon-sagemaker-neo-train-your-machine-learning-models-once-run-them-anywhere/>

- Train once, run anywhere
- Frameworks and algorithms
 - TensorFlow, Apache MXNet, PyTorch, ONNX, and XGBoost
- Hardware architectures
 - ARM, Intel, and NVIDIA starting today
 - Cadence, Qualcomm, and Xilinx hardware coming soon
- Amazon SageMaker Neo is open source, enabling hardware vendors to customize it for their processors and devices:
<https://github.com/neo-ai/>

Compiling ResNet-50 for the Raspberry Pi

Configure the compilation job

```
{
  "RoleArn": $ROLE_ARN,
  "InputConfig": {
    "S3Uri": "s3://jsimon-neo/model.tar.gz",
    "DataInputConfig": "{\"data\": [1, 3, 224, 224]}",
    "Framework": "MXNET"
  },
  "OutputConfig": {
    "S3OutputLocation": "s3://jsimon-neo/",
    "TargetDevice": "rasp3b"
  },
  "StoppingCondition": {
    "MaxRuntimeInSeconds": 300
  }
}
```

Compile the model

```
$ aws sagemaker create-compilation-job
--cli-input-json file://config.json
--compilation-job-name resnet50-mxnet-pi

$ aws s3 cp s3://jsimon-neo/model-
rasp3b.tar.gz .

$ gtar tfz model-rasp3b.tar.gz
compiled.params
compiled_model.json
compiled.so
```

Predict with the compiled model

```
from dlr import DLRModel
model = DLRModel('resnet50', input_shape,
output_shape, device)
out = model.run(input_data)
```

Demo: compiling a pre-trained PyTorch model with Neo

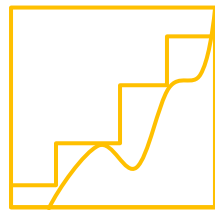
[https://github.com/aws-labs/amazon-sagemaker-examples/blob/master/advanced_functionality/
pytorch_torchvision_neo/pytorch_torchvision_neo.ipynb](https://github.com/aws-labs/amazon-sagemaker-examples/blob/master/advanced_functionality/pytorch_torchvision_neo/pytorch_torchvision_neo.ipynb)

Amazon Elastic Inference

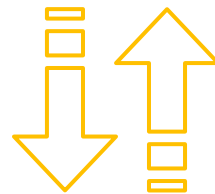
<https://aws.amazon.com/blogs/aws/amazon-elastic-inference-gpu-powered-deep-learning-inference-acceleration/>



Lower inference costs
up to 75%



Match capacity
to demand



Available between 1 to 32
TFLOPS

Integrated with
Amazon EC2,
Amazon SageMaker,
and Amazon DL
AMIs

KEY
FEATURES
Support for TensorFlow,
Apache MXNet, and
ONNX
with PyTorch coming soon

Single and
mixed-precision
operations

Demo:

Elastic Inference with TensorFlow

https://github.com/aws-labs/amazon-sagemaker-examples/blob/master/sagemaker-python-sdk/tensorflow/iris_dnn_classifier_using_estimators/

Train & predict faster

Save time

Save money

**Save your sanity (no
plumbing!)**

Getting started

<https://ml.aws>

<https://aws.amazon.com/sagemaker>

<https://github.com/awslabs/amazon-sagemaker-examples>

DEV DAY

Thank you!

Julien Simon

Global Evangelist, AI & Machine Learning

@julsimon

<https://medium.com/@julsimon>

