



Enabling Deep Learning in IoT Applications with Apache MXNet

Julien Simon

Principal AI/ML Evangelist, EMEA

@julsimon

March 2018

Agenda

- Deep Learning at the Edge?
- Apache MXNet
- Predicting in the Cloud or at the Edge?
- AWS DeepLens
- Getting started

Services covered: Apache MXNet, Deep Learning AMI, Amazon SageMaker, AWS IoT, AWS Greengrass (ML), AWS DeepLens

Deep Learning at the Edge?


Most machine data never reaches the cloud



Medical equipment



Industrial machinery

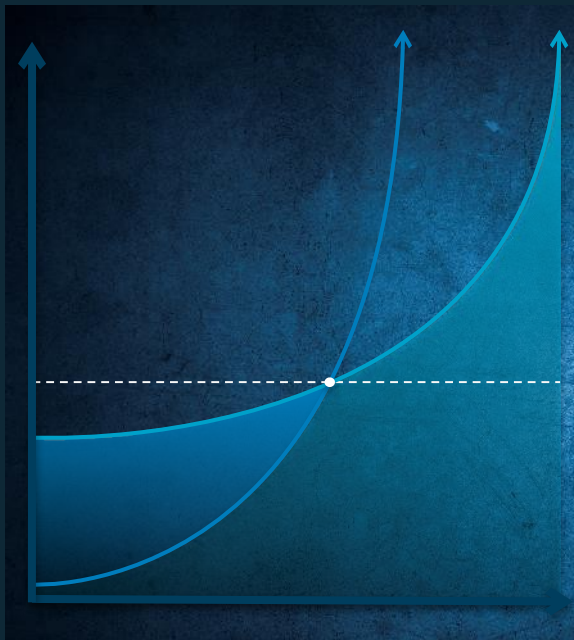


Extreme environments

Why this problem isn't going away



Law of physics



Law of economics



Law of the land

Deep Learning at the Edge

- Capturing data at the Edge and sending it to the Cloud is a good start.
- We can use it for analytics, model training, etc.
- Let's see how we could close the loop and use predictive models at the Edge.

Deep Learning challenges at the Edge

- **Resource-constrained devices**
 - CPU, memory, storage, power consumption.
- **Network connectivity**
 - Availability, cost, bandwidth, latency.
 - On-device prediction may be the only option.
- **Deployment**
 - Updating code and models on a fleet of devices is not easy.



Deep Learning wishlist at the Edge

- Rely on cloud-based services for seamless **training** and **deployment**.
- Have the option to use **cloud-based prediction**.
- Be able to run **device-based prediction** with good performance.
- Support different **technical environments** (CPUs, languages).

Apache MXNet

Apache MXNet: Open Source library for Deep Learning



Programmable

Simple syntax,
multiple
languages




Most Open

Accepted into the
Apache Incubator



Portable

Highly efficient
models for
mobile
and IoT



High Performance

Near linear scaling
across hundreds of
GPUs



Best On AWS

Optimized for
Deep Learning on AWS


1. Flexible experimentation in the Cloud.
2. Scalable training in the Cloud.
3. Good prediction performance at the Edge.
4. Prediction in the Cloud or at the Edge.

1 - Flexible experimentation in the Cloud

- API for Python, R, Perl, Matlab, Scala, C++.
- Gluon
 - Imperative programming aka 'define-by-run'.
 - Inspect, debug and modify models during training.
- Extensive model zoo
 - Pre-trained computer vision models.
 - DenseNet, SqueezeNet for resource-constrained devices.

2 - Scalable training in the Cloud

Amazon SageMaker



AWS Deep Learning AMI



Amazon EC2

c5



p3



3 - Good prediction performance at the Edge

- MXNet is written in C++.
- Gluon networks can be 'hybridized' for additional speed.
- Two libraries boost performance on CPU-only devices
 - Fast implementation of math primitives
 - Hardware-specific instructions, e.g. Intel AVX or ARM NEON
 - Intel Math Kernel Library <https://software.intel.com/en-us/mkl>
 - NNPACK <https://github.com/Maratyszczka/NNPACK>
- Mixed precision training on GPUs
 - Use float16 instead of float32 for weights and activations
 - Almost 2x reduction in model size, no loss of accuracy
 - <https://devblogs.nvidia.com/parallelforall/mixed-precision-training-deep-neural-networks/>



4 - Predicting in the Cloud or at the Edge

- Cloud-based: **invoke a Lambda function with AWS IoT.**
- Cloud-based: **invoke a SageMaker endpoint with HTTP.**
- Device-based: **bring your own code and model.**
- Device-based: **deploy your code and model with AWS Greengrass.**

Invoking a Lambda function with AWS IoT

- Train a model in **SageMaker** (or bring your own).
- Host it in **S3** (or embed it in a Lambda function).
- Write a **Lambda** function performing prediction.
- Invoke it through **AWS IoT**.



Best when

Devices can support neither HTTP nor local inference (e.g. Arduino).

Costs must be kept as low as possible.

Requirements

Network is available and reliable (MQTT is less demanding than HTTP).

Devices are provisioned in AWS IoT (certificate, keys).

<https://aws.amazon.com/blogs/compute/seamlessly-scale-predictions-with-aws-lambda-and-mxnet/>

Invoking a SageMaker endpoint with HTTP

- Train a model in **SageMaker** (or bring your own).
- Deploy it to a prediction endpoint.
- Invoke the HTTP endpoint from your devices.

Best when

Devices are not powerful enough for local inference.

Models can't be easily deployed to devices.

Additional cloud-based data is required for prediction.

Prediction activity must be centralized.

Requirements

Network is available and reliable.

Devices support HTTP.

Bring your own code and model

- Train a model in **SageMaker** (or bring your own).
- Bring your own application code.
- Provision devices at manufacturing time (or use your own update mechanism).

Best when

You don't want to or can't rely on cloud services
(no network connectivity?)

Requirements

Devices are powerful enough for local inference.

Models don't need to be updated, if ever.

DIY!

Deploy your code and model with AWS Greengrass

- Train a model in **SageMaker** (or bring your own).
- Write a **Lambda** function performing prediction.
- Add both as resources in your **Greengrass** group.
- Let **Greengrass** handle deployment and updates.



Best when
You want the same programming model in the Cloud and at the Edge.
Code and models need to be updated, even if network connectivity is infrequent or unreliable.
One device in the group should be able to perform prediction on behalf on other devices.

Requirements
Devices are powerful enough to run Greengrass (XXX HW requirements)
Devices are provisioned in AWS IoT (certificate, keys).

AWS Greengrass ML

The image displays three overlapping screenshots of the AWS Greengrass console, illustrating the configuration of a **GGOBJECTClassificationGroup**.

Leftmost screenshot: Shows the **GREENGRASS GROUP** header with the group name **GGOBJECTClassificationGroup** and version **Version d762d135-a5fd-**. A sidebar menu on the left lists **Deployments**, **Subscriptions**, **Cores**, **Devices**, **Lambdas**, **Resources**, and **Settings**. The **Resources** option is highlighted.

Middle screenshot: Shows the same header and sidebar. The **Resources** section is expanded, showing a list of resources. The **squeezenet_model** resource is highlighted with an orange box.

Rightmost screenshot: Shows the **Resources** section in detail. It includes a table with the following data:

Name	Resource Type	Source
videoCoreInterface	Device	/dev/vchiq
videoCoreShareMemory	Device	/dev/vcsm
squeezenet_model	Model	https://jsimon-greengrass-demo.s3...

AWS DeepLens

AWS DeepLens

World's first Deep Learning enabled video camera for developers



A new way to learn

Custom built for Deep Learning

Broad Framework Support

Deploy models from Amazon SageMaker

Integrated with AWS

Fully programmable with AWS Lambda

AWS DeepLens

DeepLens > Projects > Face-detection

Face-detection

Delete Deploy to device

Project

Copy Edit

Name	Description	Version
Face-detection	Detect all faces in your surroundings	-


ARN

arn:aws:deeplens:us-east-1:[redacted]:project/Face-detection

Project content

Type	Name
Function	arn:aws:lambda:us-east-1:[redacted]:function:deeplens-face-detection:1
Model	deeplens-face-detection

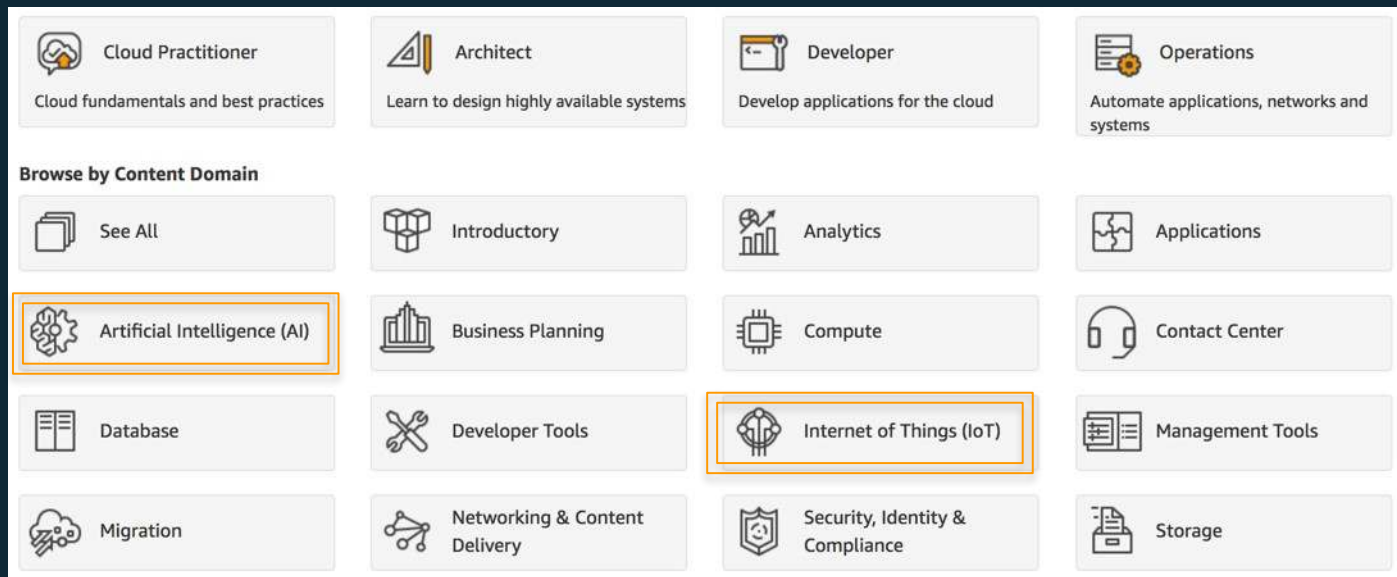
Object detection with AWS DeepLens



Getting started

Digital Training

AWS Training and Certification released **free digital training courses** that will make it easier for you to build your cloud skills.



Amazon Machine Learning Lab



Lots of companies doing
Machine Learning



Lack ML
expertise



Unable to unlock
business potential

Amazon ML Lab
provides the missing ML
expertise



Leverage Amazon experts with decades of ML
experience with technologies like Amazon Echo,
Amazon Alexa, Prime Air and Amazon Go



Brainstorming



Modeling



Teaching

<https://aws.amazon.com/ml-solutions-lab/>

Resources

<https://aws.amazon.com/machine-learning>

<https://aws.amazon.com/sagemaker>

<https://aws.amazon.com/greengrass>

<https://aws.amazon.com/greengrass/ml>

<https://aws.amazon.com/deeplens>

<https://aws.amazon.com/machine-learning/amis/>

<https://mxnet.incubator.apache.org>

<http://gluon.mxnet.io/>

<https://medium.com/@julsimon>

Thank you!

Julien Simon
Principal AI/ML Evangelist, EMEA
@julsimon