

Deep Learning at the Edge

Julien Simon, Principal Evangelist, AI & Machine Learning
@julsimon

Deep Learning at the Edge

1. Flexible experimentation in the Cloud.
2. Scalable training in the Cloud.
3. Good prediction performance at the Edge.
4. Simple deployment of code and model

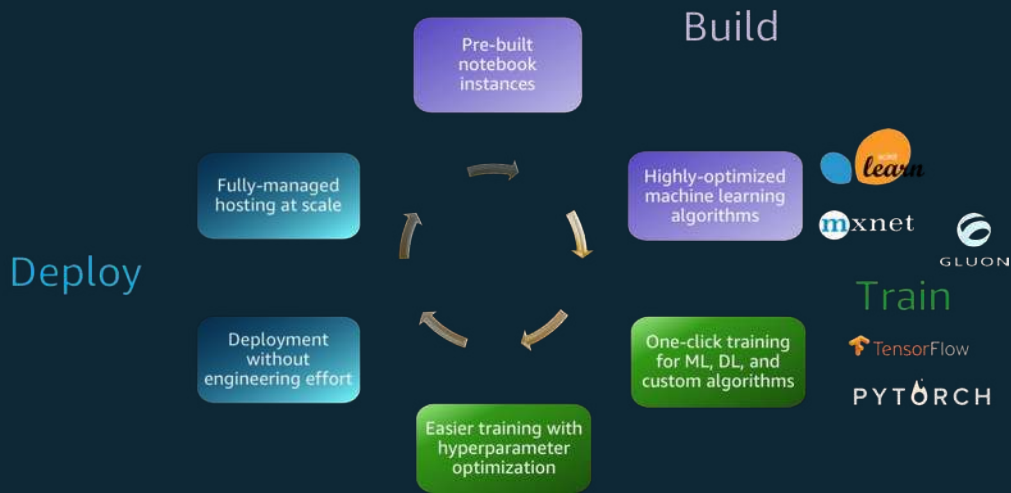
Flexible experimentation in the Cloud



- Apache MXNet: Python, R, Perl, Matlab, Scala, C++.
- Gluon
 - Imperative programming aka 'define-by-run'.
 - Inspect, debug and modify models during training.
- Extensive model zoo
 - Pre-trained computer vision models.
 - MobileNet, SqueezeNet for resource-constrained devices.

Scalable training in the Cloud

Amazon SageMaker



AWS Deep Learning AMI



Amazon
EC2

c5



p3



Good prediction performance at the Edge

- MXNet is written in C++.
- Gluon networks can be 'hybridized' for additional speed.
- Two libraries boost performance on CPU-only devices
 - Fast implementation of math primitives
 - Hardware-specific instructions, e.g. Intel AVX or ARM NEON
 - Intel Math Kernel Library <https://software.intel.com/en-us/mkl>
 - NNPACK <https://github.com/Maratyszczka/NNPACK>
- Mixed precision training
 - Use float16 instead of float32 for weights and activations
 - Almost 2x reduction in model size, no loss of accuracy, faster inference
 - <https://devblogs.nvidia.com/parallelforall/mixed-precision-training-deep-neural-networks/>



Simple deployment of code and model

- Train a model in **SageMaker** (or bring your own).
- Write a **Lambda** function performing prediction.
- Add both as resources in your **Greengrass** group.
- Let **Greengrass** handle deployment and updates.

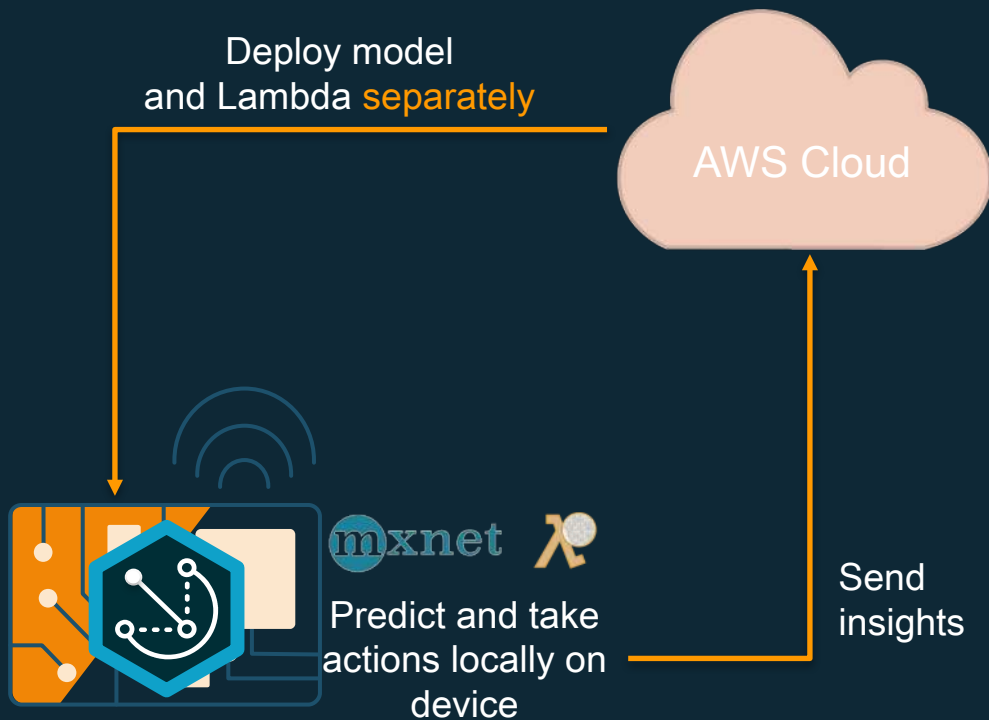


Best when
You want the same programming model in the Cloud and at the Edge.
Code and models need to be updated, even if network connectivity is infrequent or unreliable.
One device in the group should be able to perform prediction on behalf on other devices.

Requirements
Devices are powerful enough to run Greengrass (XXX HW requirements)
Devices are provisioned in AWS IoT (certificate, keys).

ML Inference using AWS Greengrass

PREVIEW
AVAILABLE

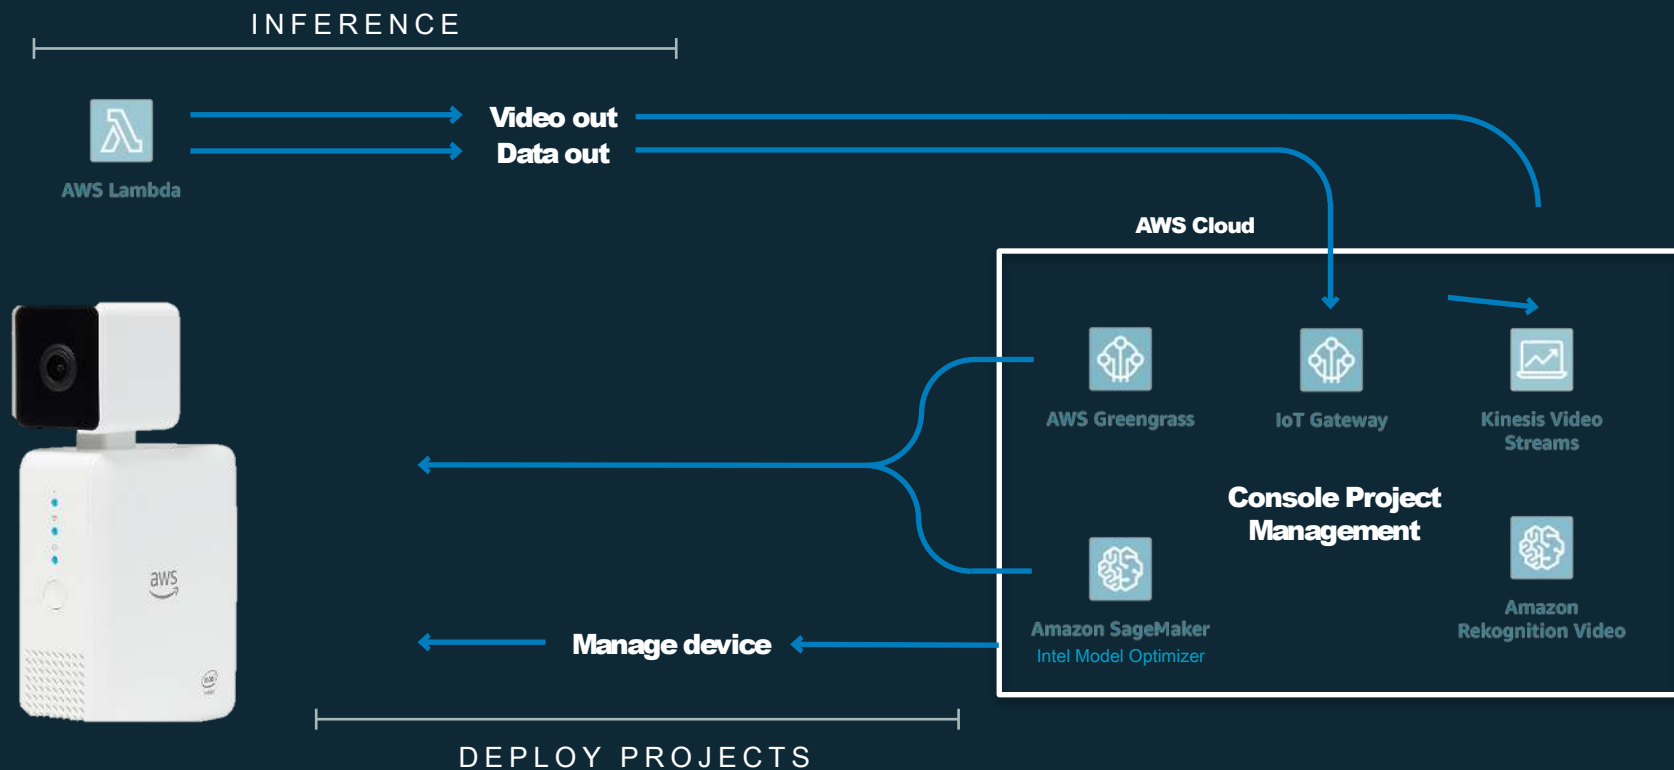


AWS DeepLens



- Intel Atom CPU
- Gen9 graphics
- Ubuntu 16.04 LTS
- 100 GFLOPS performance
- Dual band Wi-Fi
- 8 GB RAM
- 16 GB Storage (eMMC)
- 32 GB SD card
- 4 MP camera with MJPEG
- H.264 encoding at 1080p resolution
- 2 USB ports
- Micro HDMI
- Audio out
- AWS Greengrass preconfigured
- Intel c1DNN for Apache MXNet

AWS DeepLens Architecture



AWS DeepLens

Object-detection

DeleteDeploy to device

Project

CopyEdit

Name	Description	Version
Object-detection	Detect 20 popular objects	-


ARN

arn:aws:deeplens:us-east-1:[redacted]:project/Object-detection

Project content

Type	Name
Function	arn:aws:lambda:us-east-1:[redacted]:function:deeplens-object-detection:1
Model	deeplens-object-detection

Object detection with AWS DeepLens



Resources

Resources

<https://mxnet.incubator.apache.org>

<http://gluon.mxnet.io>

<https://aws.amazon.com/sagemaker> (free tier available)

An overview of Amazon SageMaker: <https://www.youtube.com/watch?v=ym7NEYEx9x4>

<https://github.com/aws-labs/amazon-sagemaker-examples>

<https://aws.amazon.com/greengrass> (free tier available)

<https://aws.amazon.com/deeplens>

<https://github.com/intel/clDNN>

<https://medium.com/@julsimon>



Thank you!

Julien Simon, Principal Evangelist, AI & Machine Learning
@julsimon