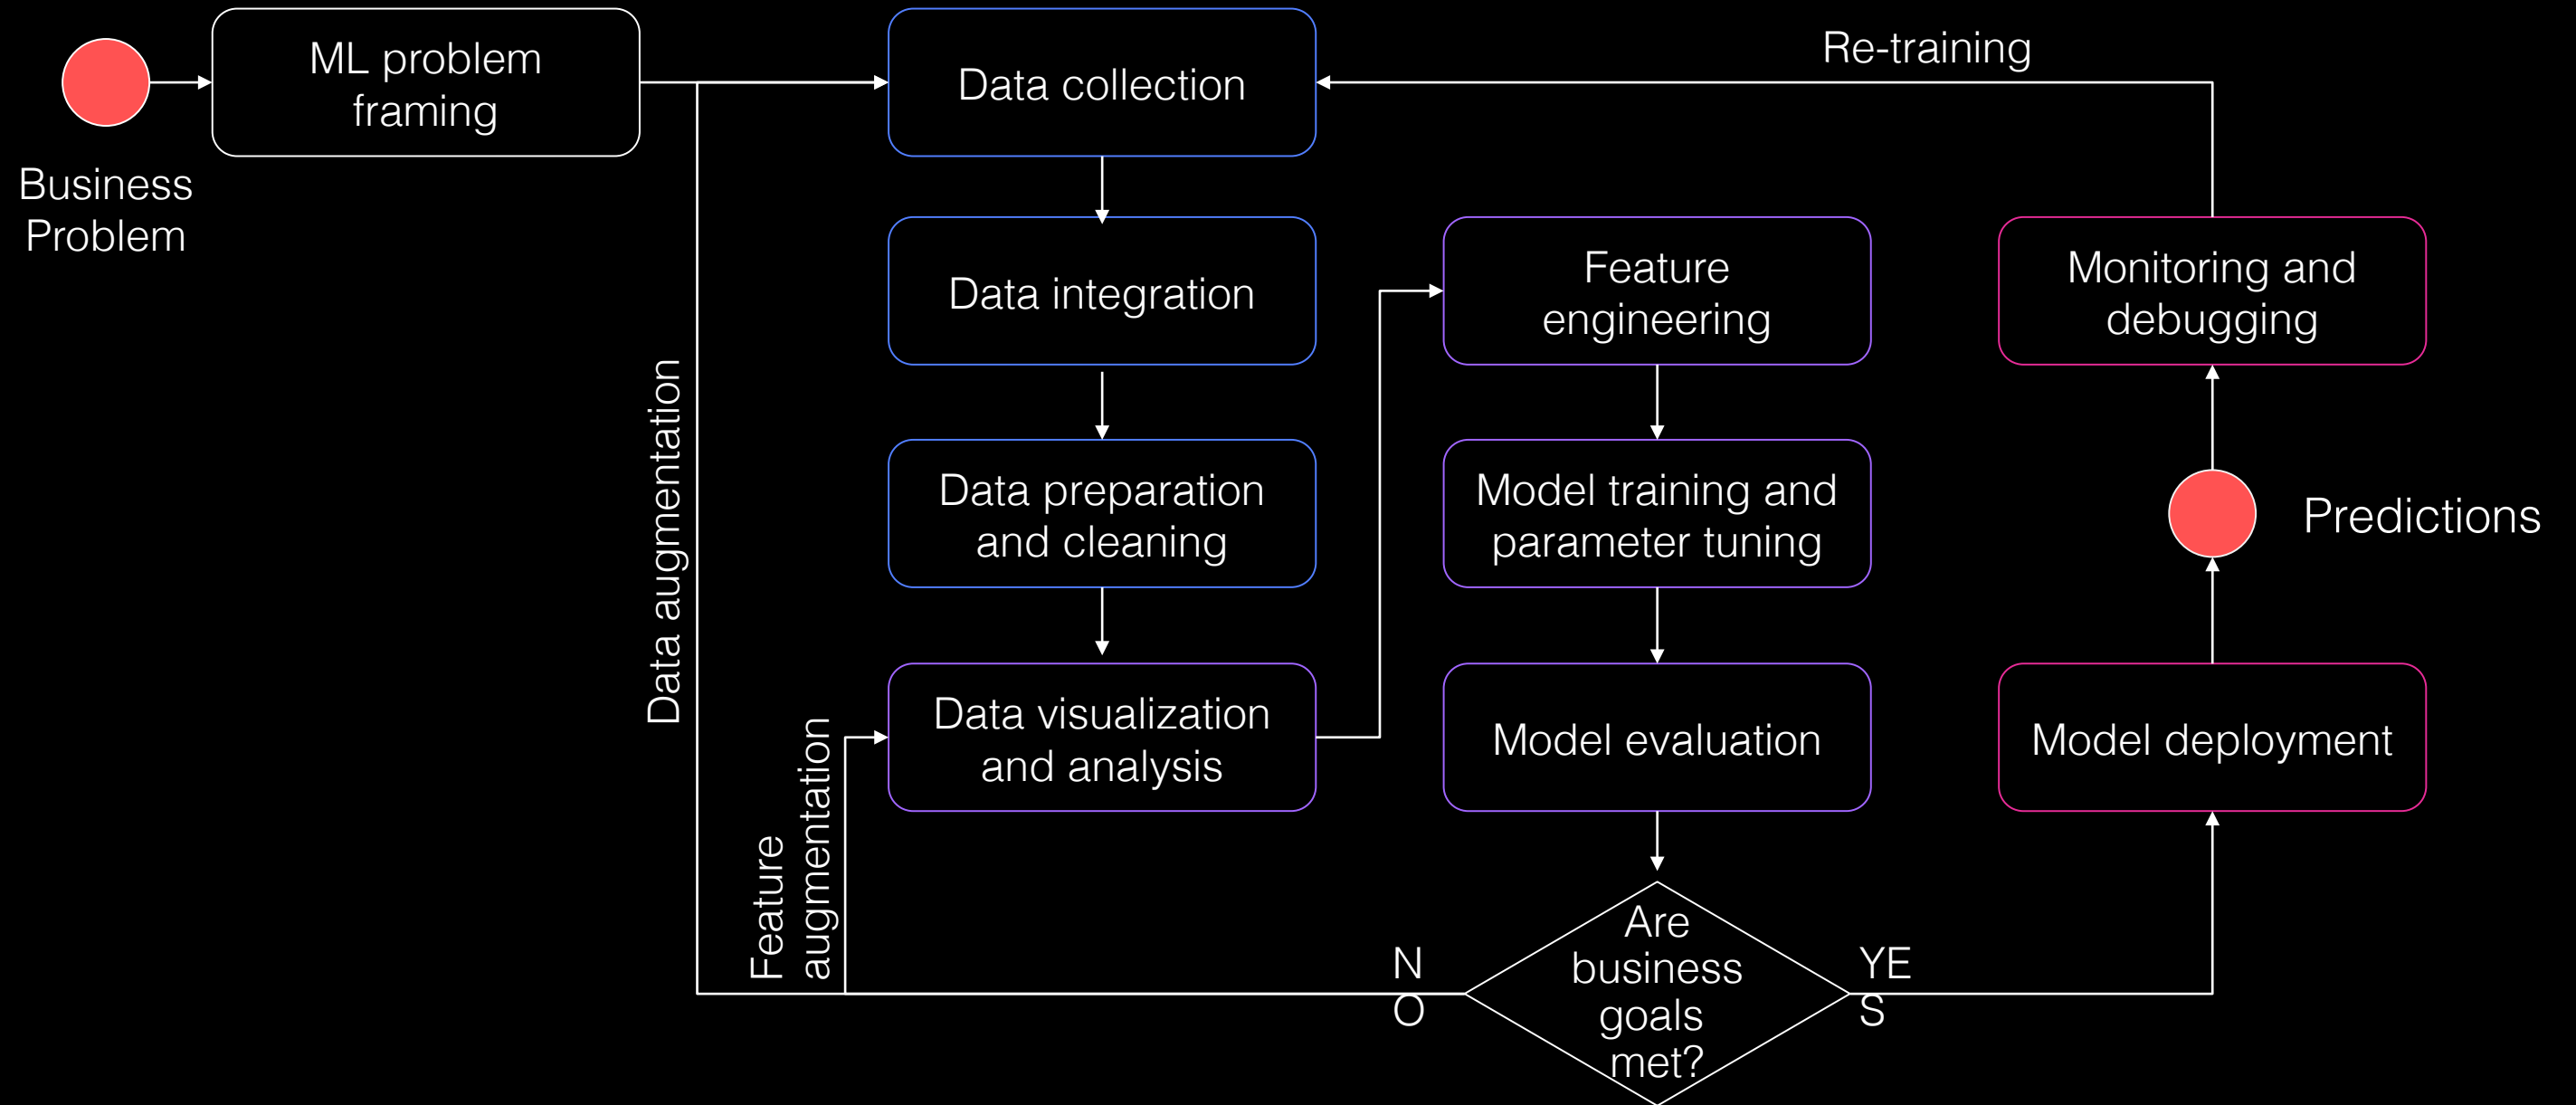


Become a Machine Learning developer with AWS services

Julien Simon
Global Evangelist, AI & Machine Learning, AWS
@julsimon

Breght
Boschker
CTO, SkinVision

Machine learning cycle



SkinVision

Applying Machine Learning to skin cancer detection

Breght Boschker, CTO
www.skinvision.com

SkinVision

Together we care

Early detection of skin cancer

1 in 5 people get skin cancer

Act when it's most treatable

Awareness is key

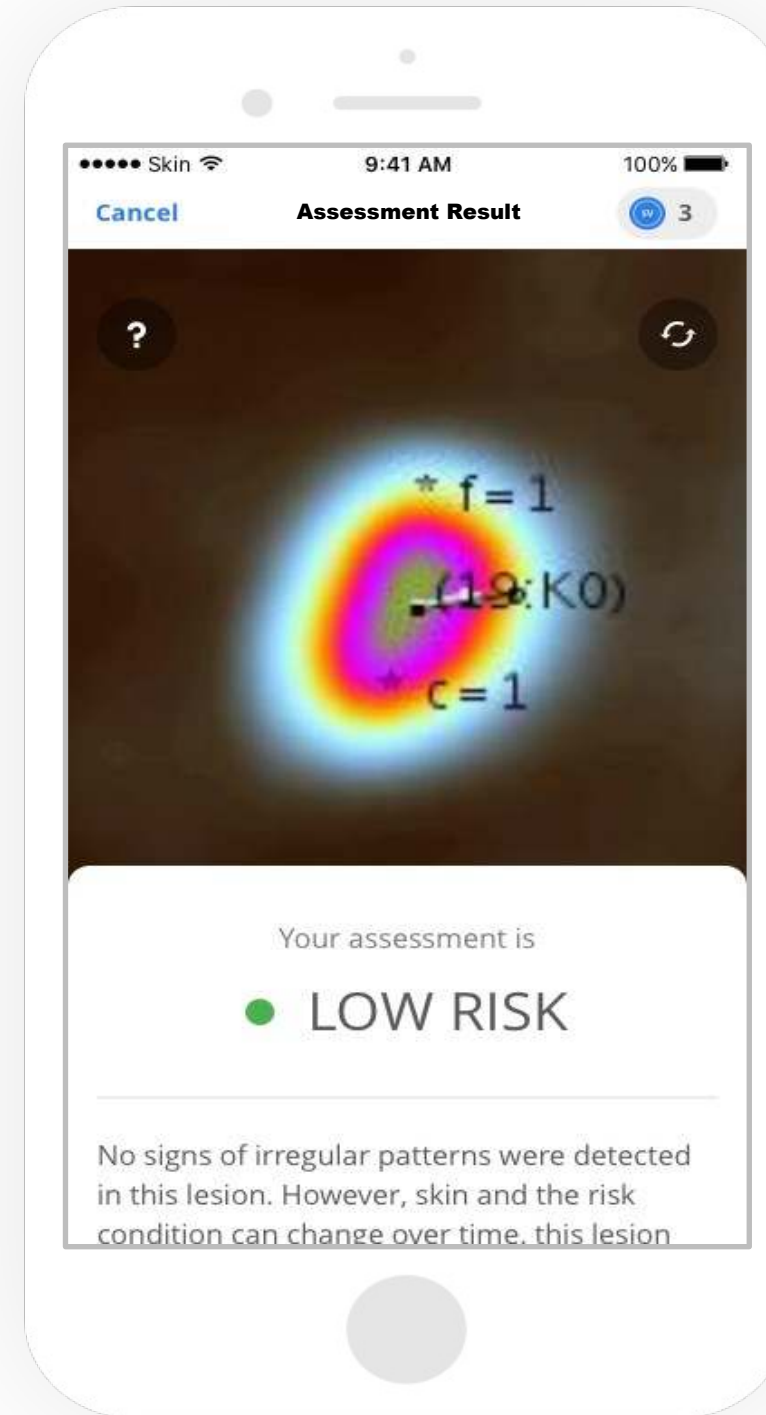
Using your smartphone

Scientifically proven technology

At the level of a specialized dermatologist

Available globally*

* Coming to USA & Canada



SkinVision

Together we care

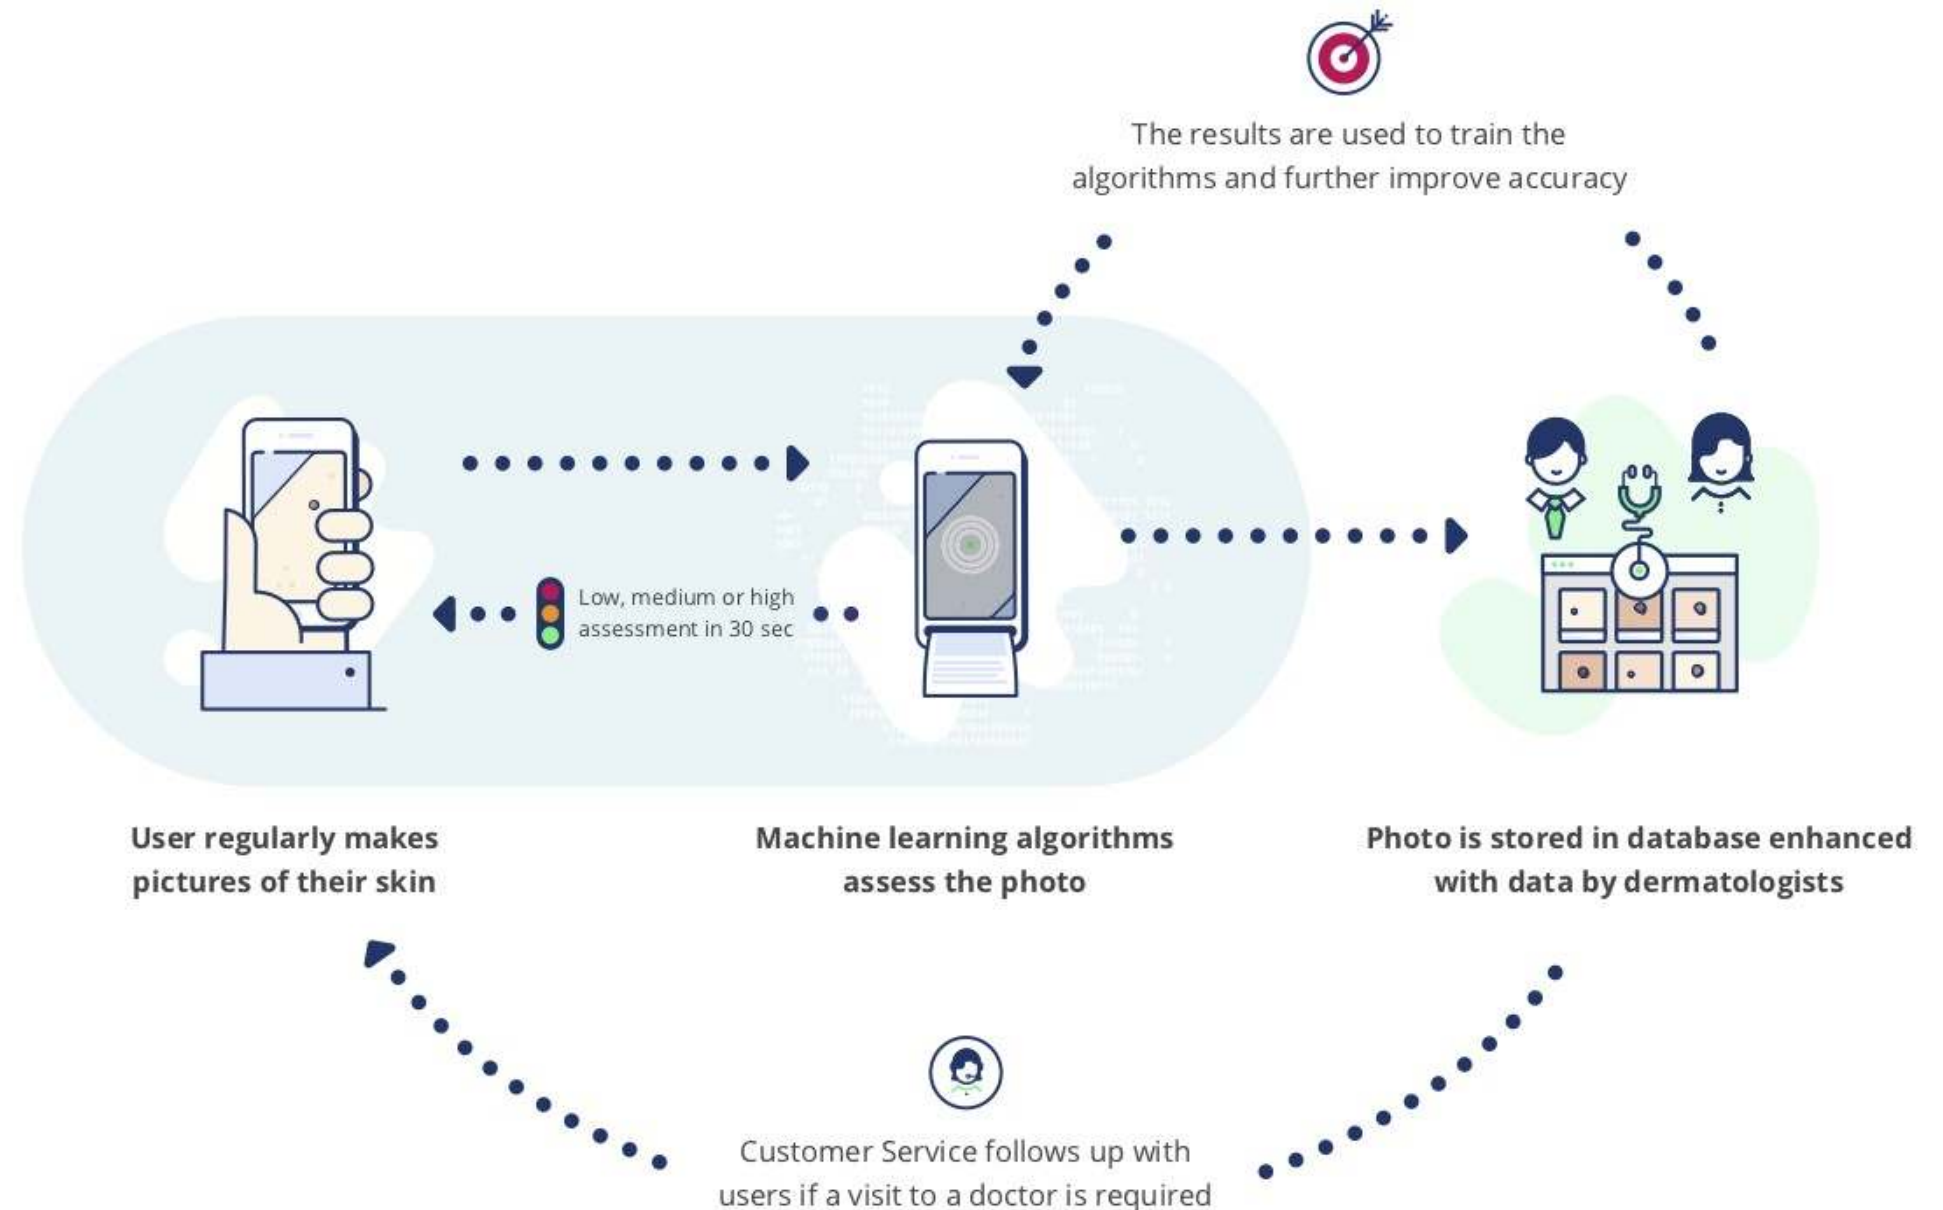
Guide a user in their health journey

- Should you see a Healthcare Professional?
- Continuous Monitoring
- Follow-up

Machine Learning Risk Assessment

Continuous data enrichment

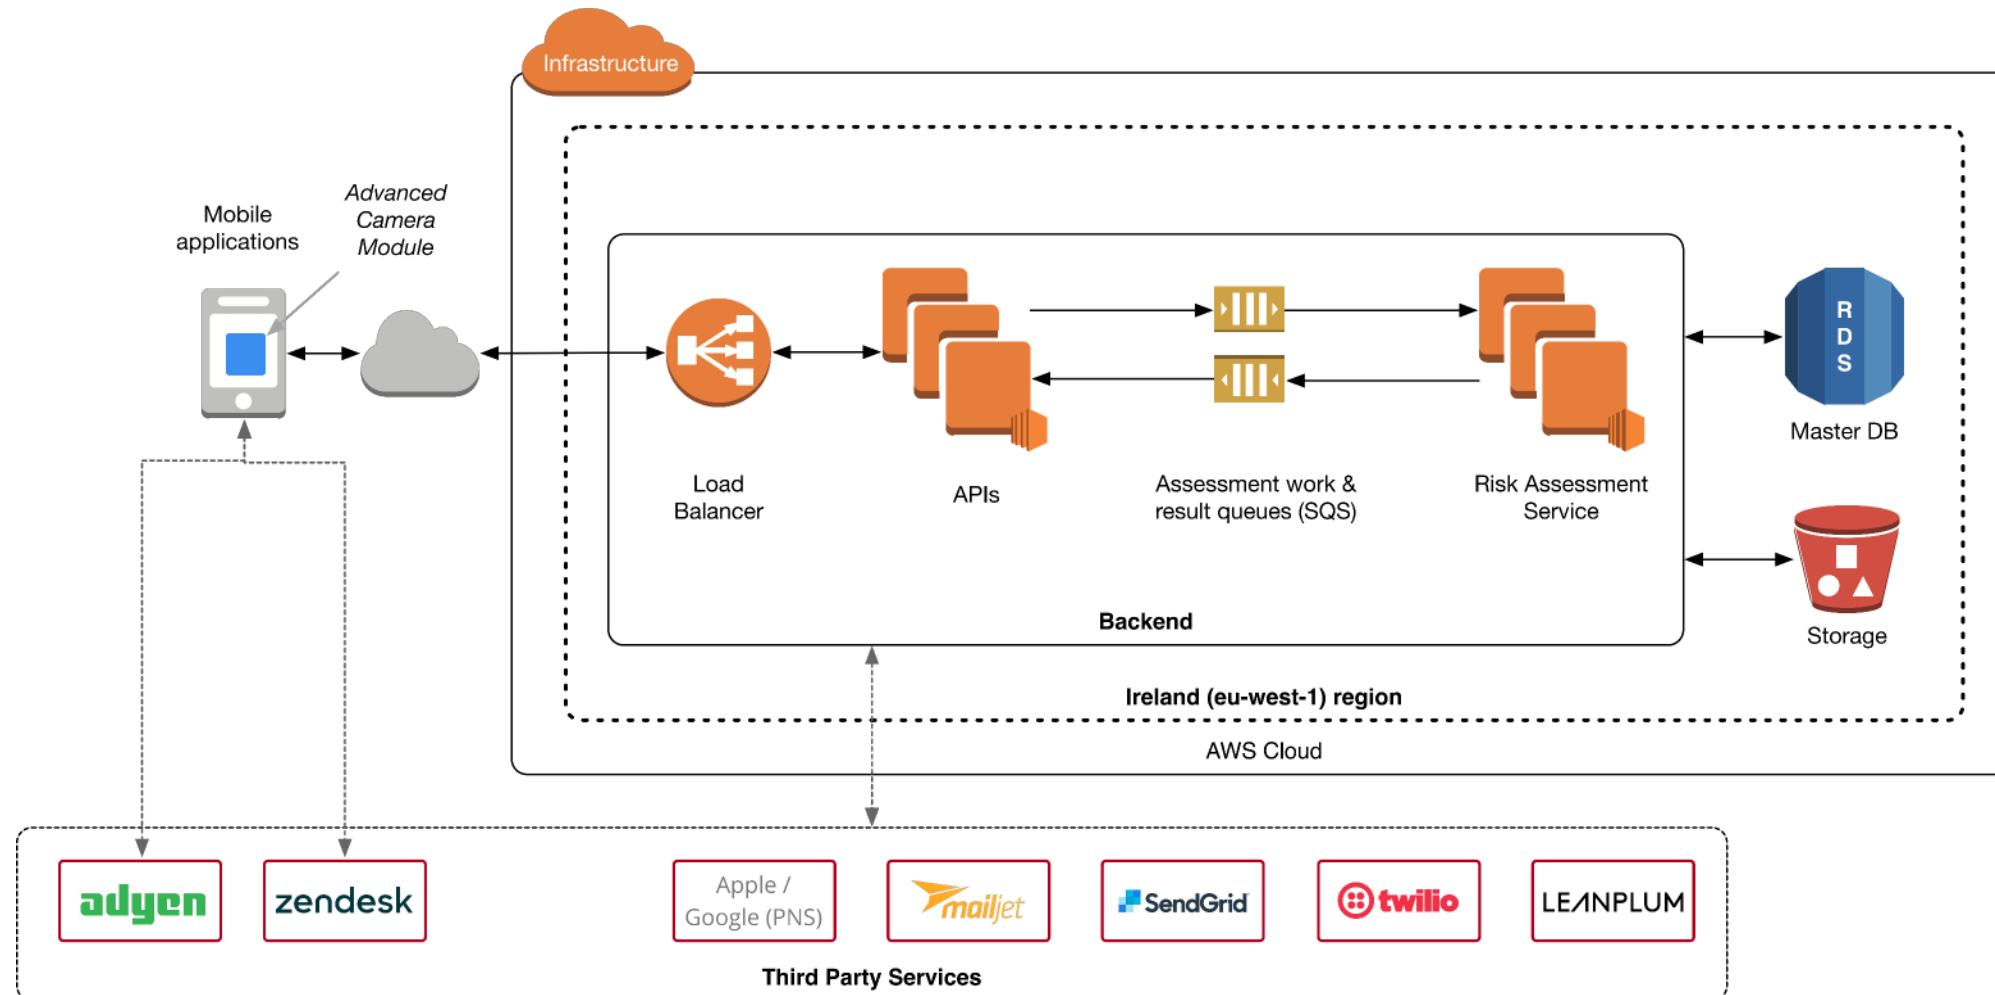
- Clinical Validation
- Customer Validation



SkinVision // AWS

- AWS has been SkinVision's first choice from the beginning
- Security
- Scalability
- Global availability
- Strong support & growth model
- Innovation driver
- Continuous focus on ease of use, automation

SkinVision Architecture



- Mobile applications
- Backend Systems
- ML Risk Assessment
- Internal Health Quality Management Systems
- Data Analysis & Business Intelligence
- Messaging


TensorFlow

 Keras


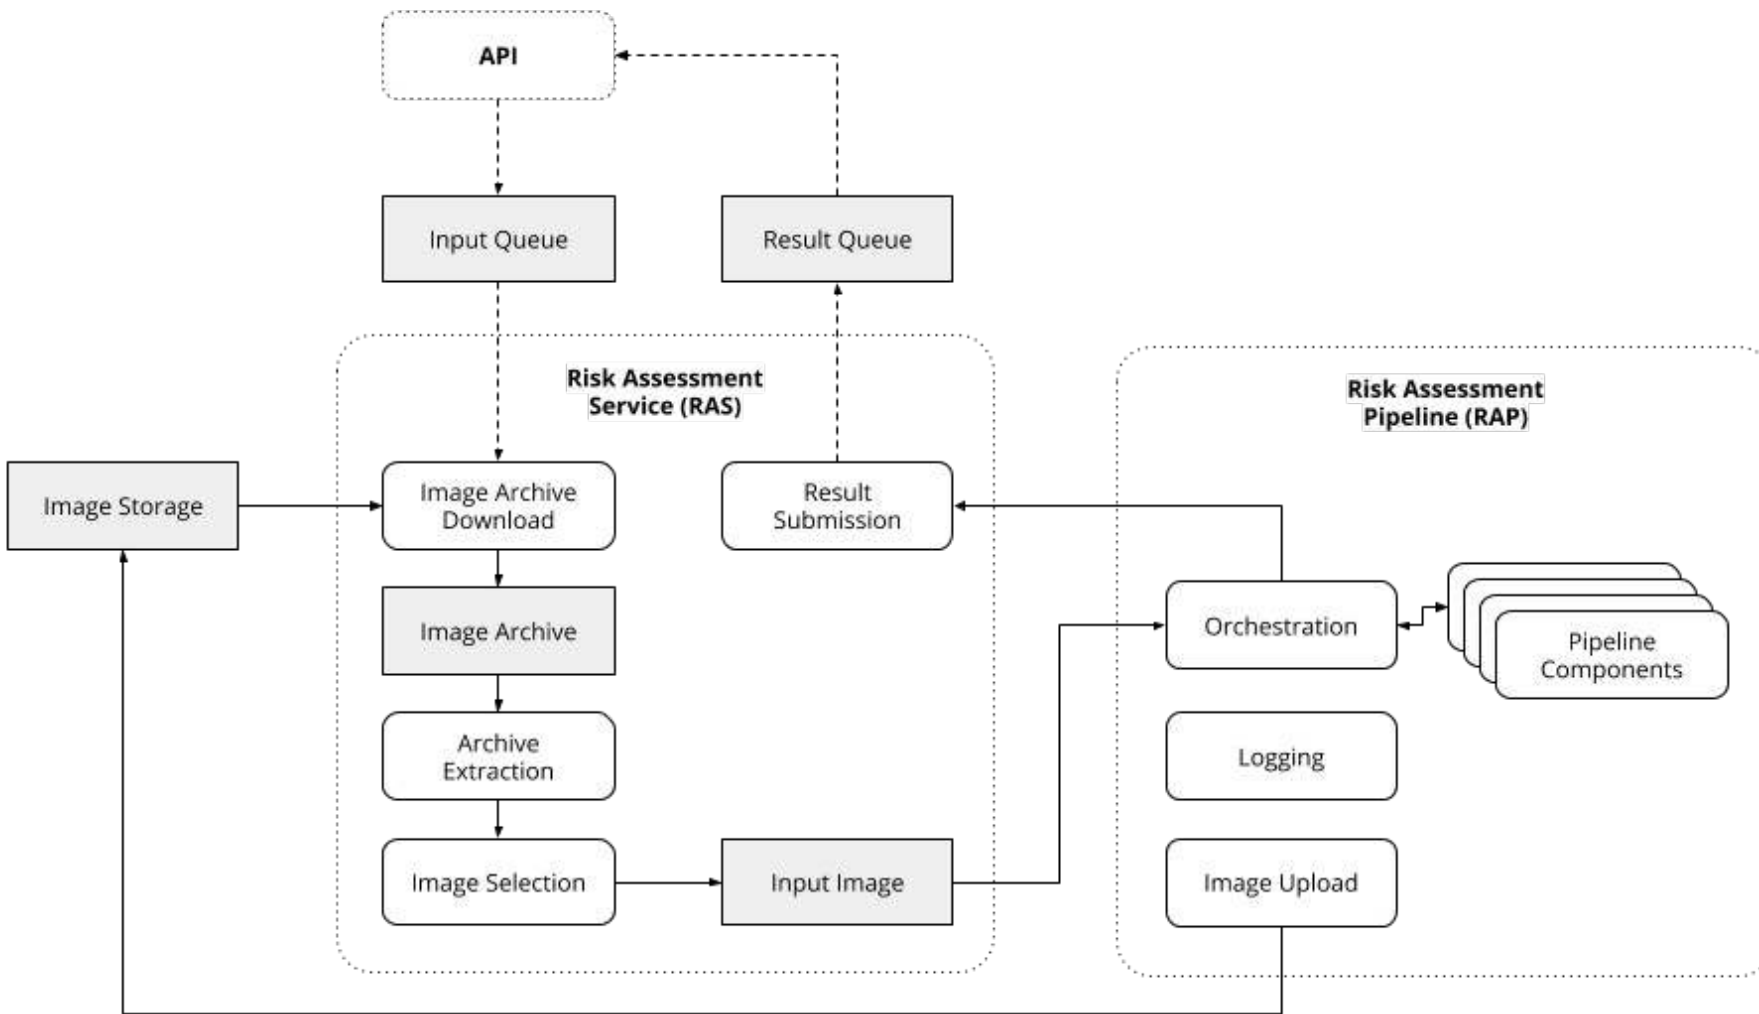
 looker

 amazon
REDSHIFT

SkinVision Architecture

Machine Learning at Scale


- Mobile apps for iOS & Android
- API
- Amazon SQS
- Service Workers: Docker
- Risk Assessment Pipeline: Docker Orchestration
- Amazon S3



SkinVision Approach

- Machine Learning as an **Engineering** Problem
 - Repeatable
 - Traceable
 - Measurable
 - Automated
 - Infrastructure as Code
 - Cost
- Proof through **data**
 - Scientific
 - 'Shadow' pipelines on AWS

Active: 1.2 million users, 1,500+ daily assessments
Effective: 27,000+ skin cancers found, 5,000+ melanoma found



Ongoing work

- Multiple disease areas
- Compliance Attestation & Compliance Automation
 - AWS Landing Zone
 - AWS Config
 - AWS GuardDuty
- Cost-down & scalability
 - Amazon SageMaker
 - Optimized frameworks
 - Amazon EKS / Container solutions
 - Amazon RedShift Spectrum


When was the last
time *you* checked
your skin?

<http://www.skinvision.com/download>




AWS services for Machine Learning

Build your dataset




Annotating data at scale is time-consuming and expensive




Amazon SageMaker Ground Truth


Build scalable and cost-effective labeling workflows



Quickly label
training data



Easily integrate
human labelers



Get accurate
results

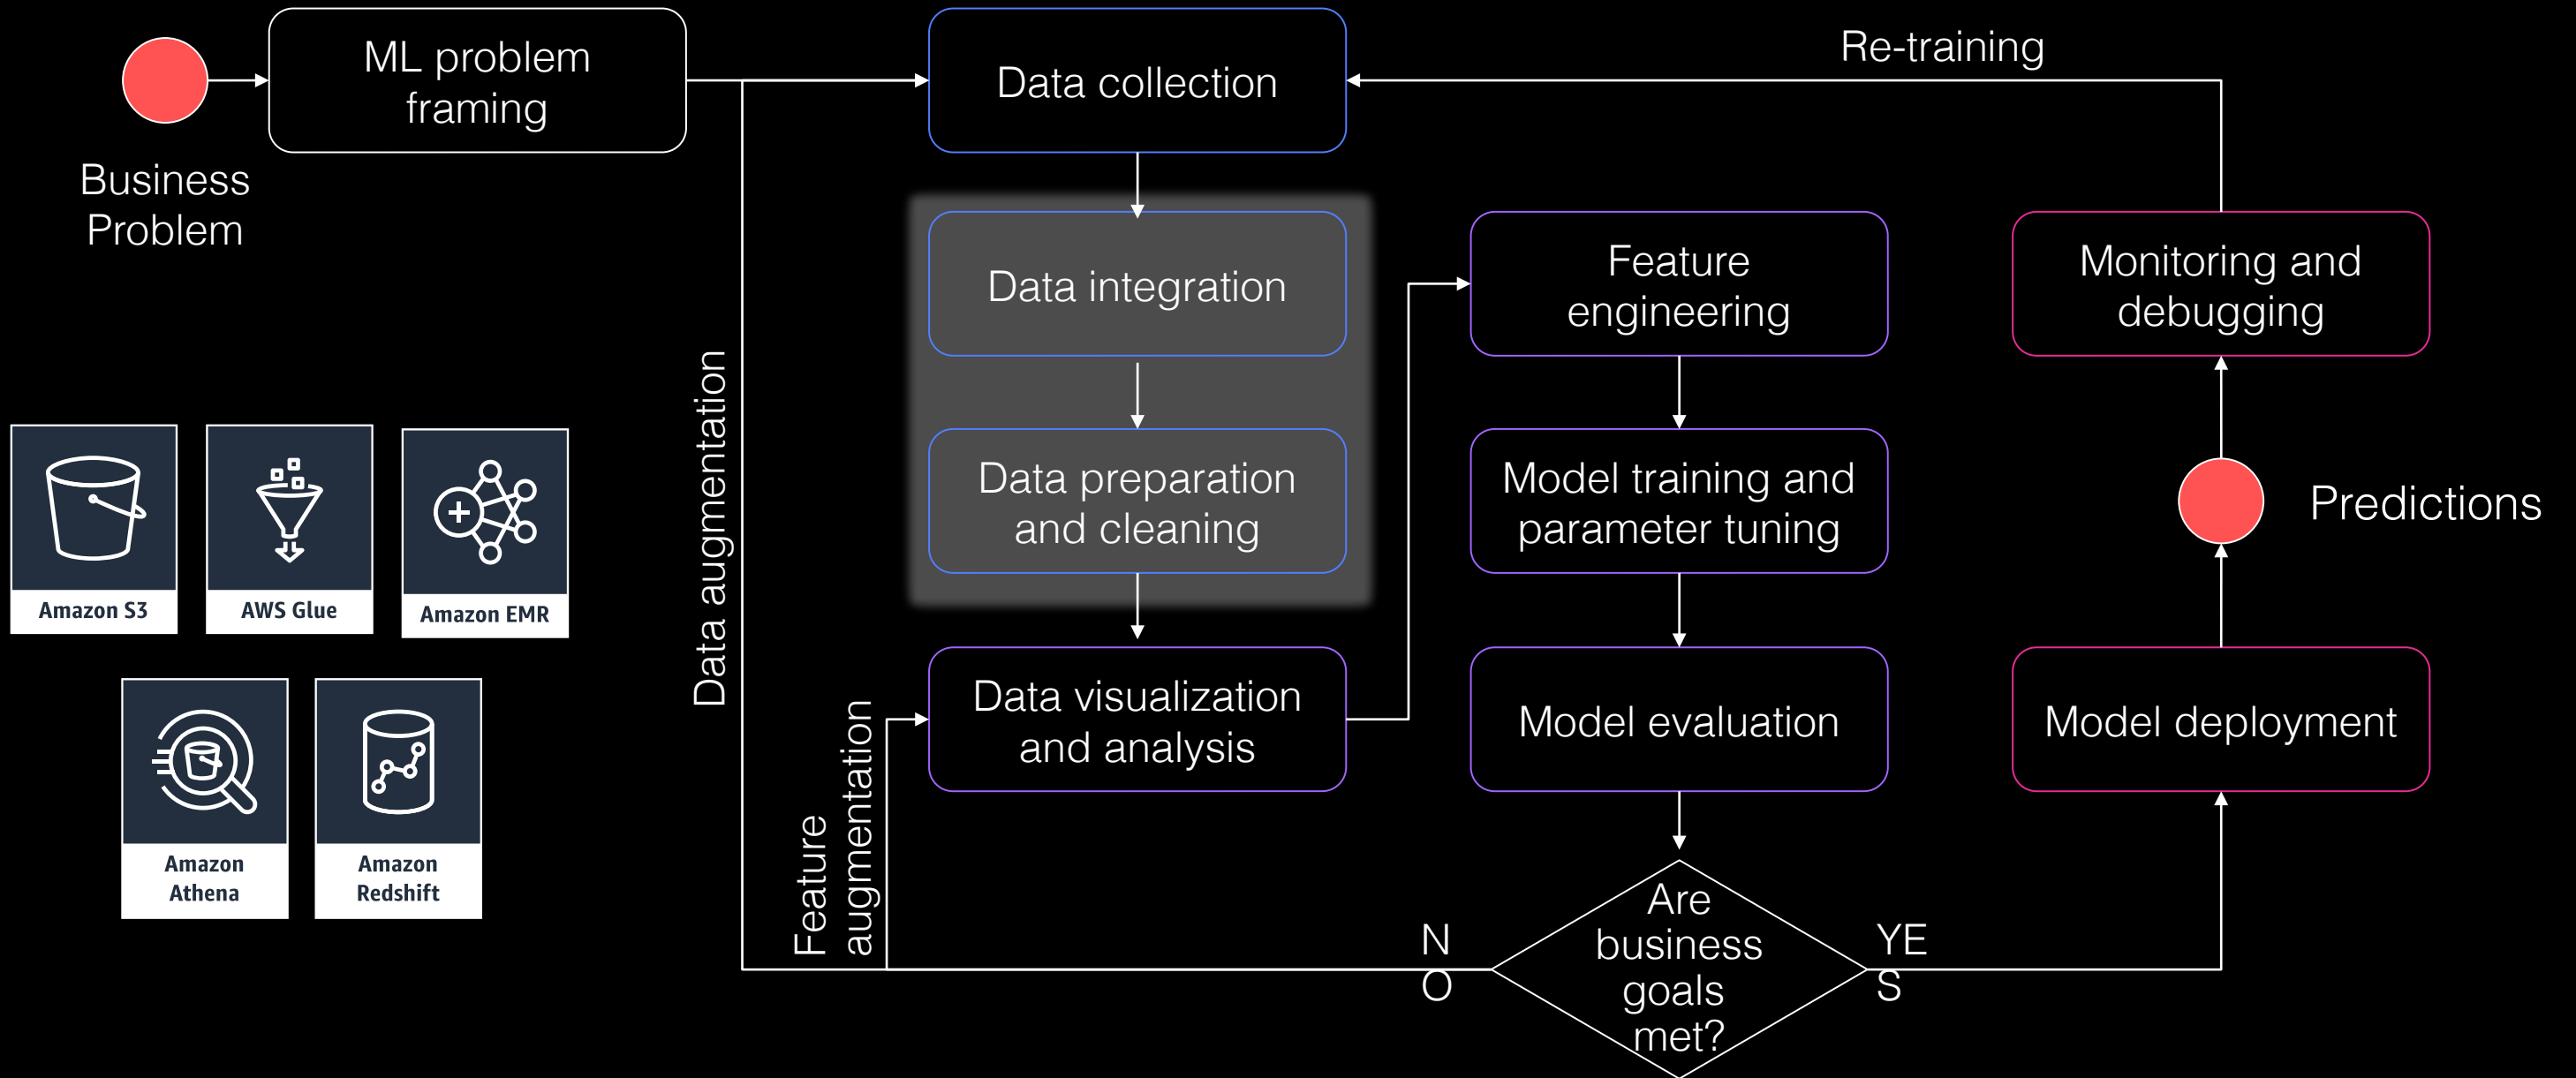
KEY FEATURES

Automatic labeling via
machine learning


Ready-made and
custom workflows for
image bounding box,
segmentation, and text

Private and public
human workforce

Prepare your dataset for Machine Learning



Build, train and deploy models using compute services



AWS Deep Learning AMIs

Preconfigured environments on Amazon Linux or Ubuntu

**NEW (March
27th)
Deep Learning
containers**

Conda AMI

For developers who want pre-installed pip packages of DL frameworks in separate virtual environments.

Base AMI

For developers who want a clean slate to set up private DL engine repositories or custom builds of DL engines.


AMI with source code

For developers who want preinstalled DL frameworks and their source code in a shared Python environment.




Amazon SageMaker


Amazon SageMaker




Collect and
prepare training
data



Choose and
optimize your
ML algorithm




Set up and
manage
environments
for training



Train and
Tune ML Models



Deploy models
in production



Scale and manage
the production
environment

Same service and APIs from experimentation to production

intuit



tinder



CONVOY

SIEMENS



DOW JONES




SONY



aws SUMMIT

© 2019, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Build, train and deploy models using SageMaker



The Amazon SageMaker API

- Python SDK **orchestrating** all Amazon SageMaker activity
 - High-level objects for **algorithm selection, training, deploying, automatic model tuning**, etc.
 - **Spark SDK** (Python & Scala)
- AWS SDK
 - For scripting and automation
 - CLI : *'aws sagemaker'*
 - Language SDKs: boto3, etc.

Model options




Training code

Factorization Machines
Linear Learner
Principal Component
Analysis
K-Means Clustering
XGBoost
And more

Built-in Algorithms (17)

No ML coding required
No infrastructure work required
Distributed training
Pipe mode



Built-in Frameworks

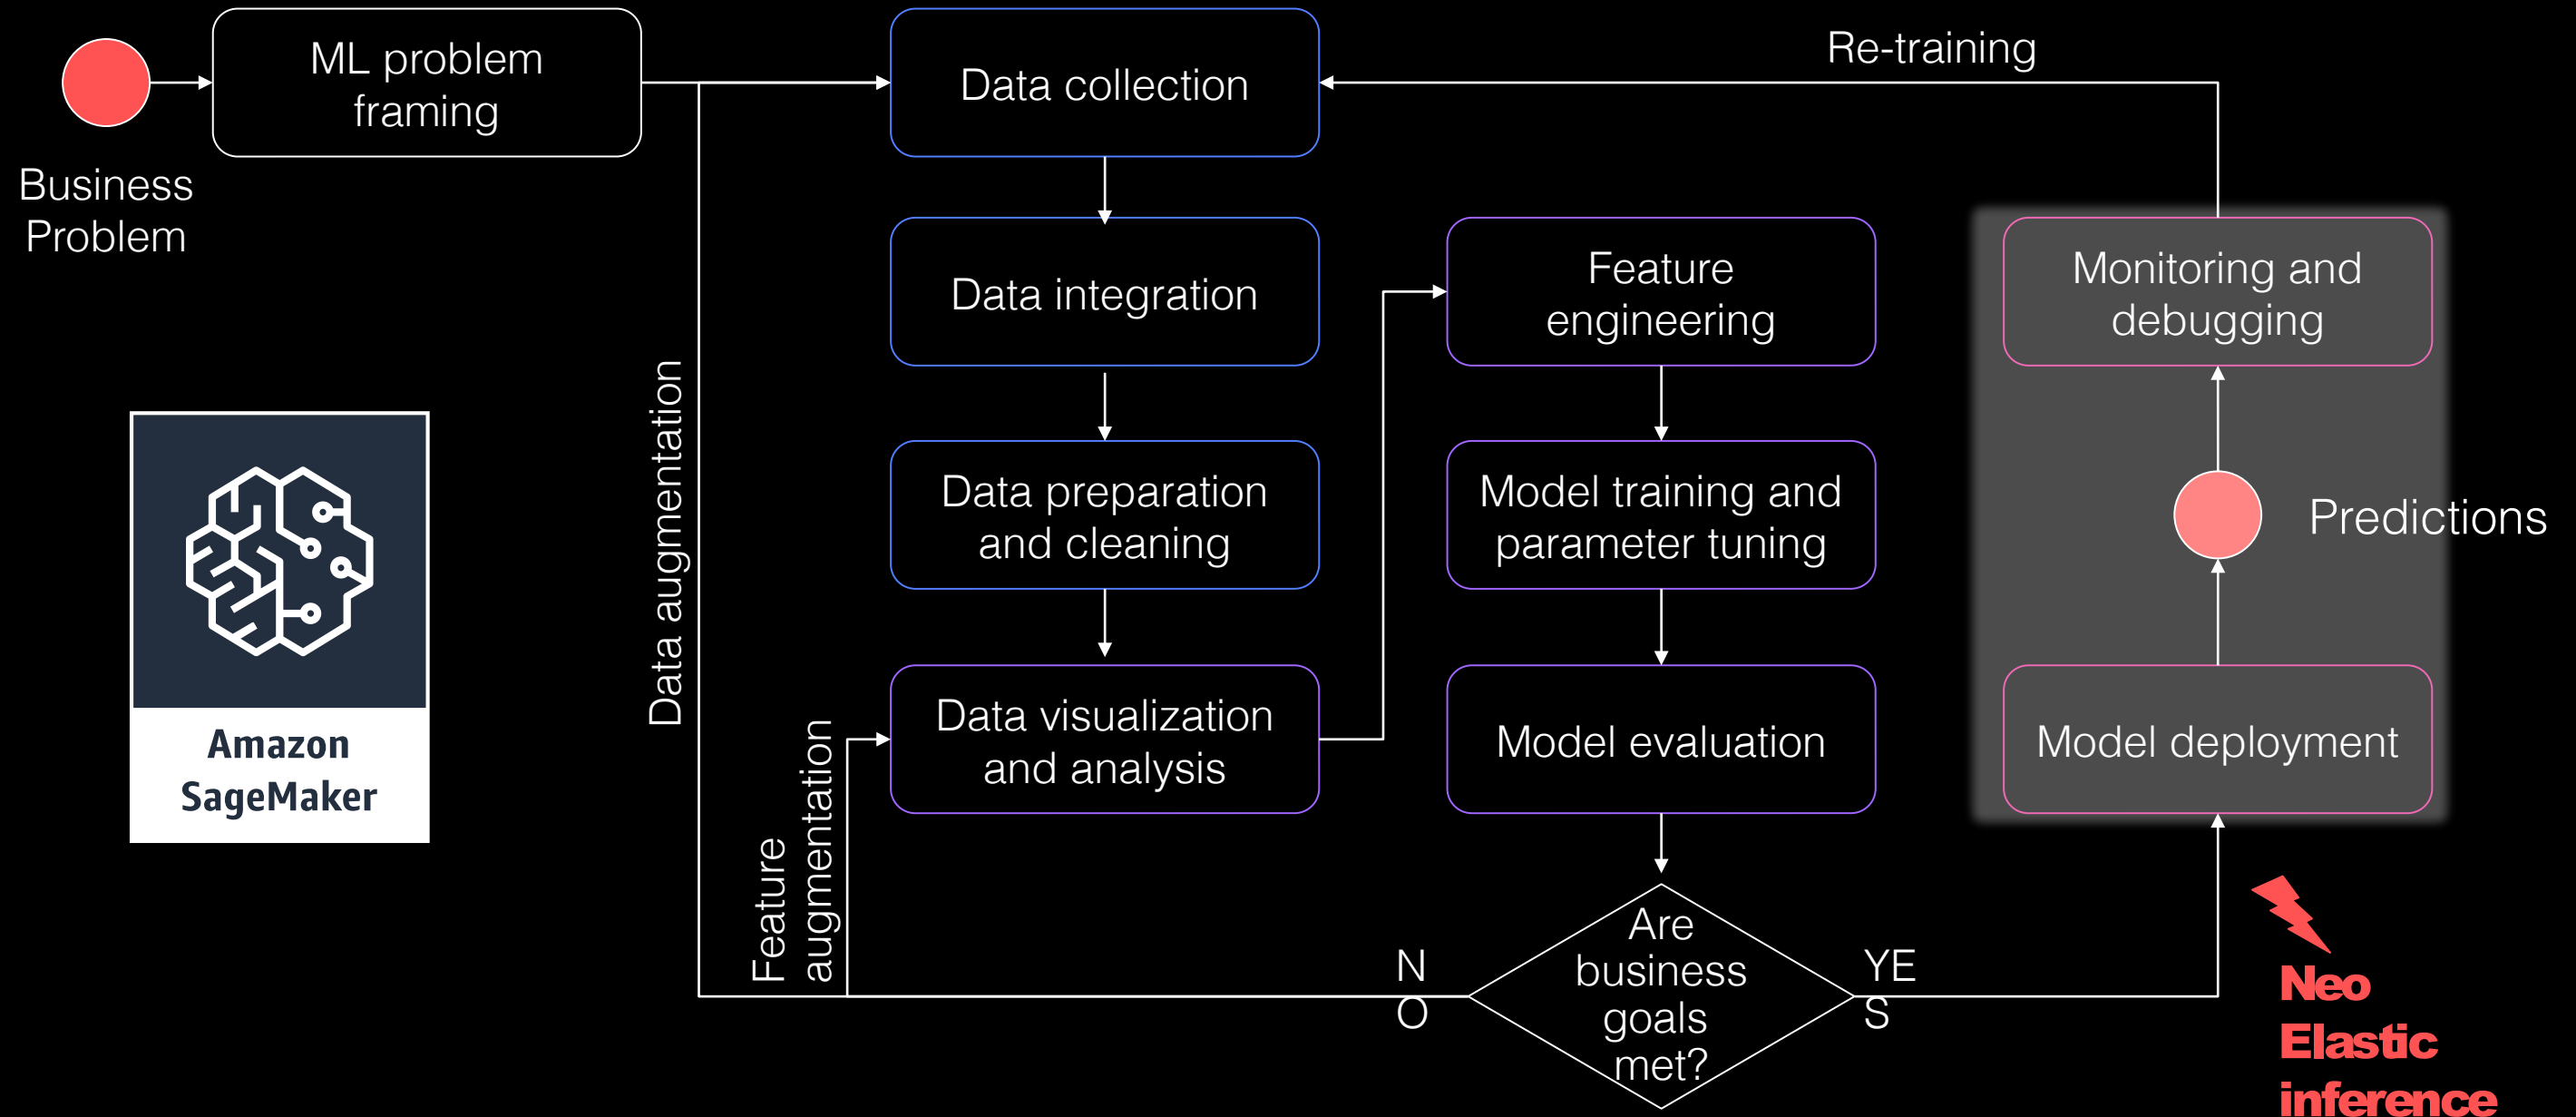
Bring your own code: script
mode
Open source containers
No infrastructure work required
Distributed training
Pipe mode



Bring Your Own
Container

Full control, run anything!
R, C++, etc.
No infrastructure work required

Optimize and deploy models using SageMaker



Amazon SageMaker Neo

Optimize models for the underlying hardware architecture

- Train once, run anywhere
- Supported **frameworks** and **algorithms**
 - TensorFlow, Apache MXNet, PyTorch, ONNX, and XGBoost
- Supported **hardware architectures**
 - ARM, Intel, and NVIDIA
 - Cadence, Qualcomm, and Xilinx hardware coming soon

The Neo compiler and runtime are **open source**, enabling hardware vendors to customize it for their processors and devices: <https://github.com/neo-ai>

Compiling ResNet-50 for the Raspberry Pi

Configure the compilation job

```
{
  "RoleArn": $ROLE_ARN,
  "InputConfig": {
    "S3Uri": "s3://jsimon-neo/model.tar.gz",
    "DataInputConfig": "{\"data\": [1, 3, 224, 224]}",
    "Framework": "MXNET"
  },
  "OutputConfig": {
    "S3OutputLocation": "s3://jsimon-neo/",
    "TargetDevice": "rasp3b"
  },
  "StoppingCondition": {
    "MaxRuntimeInSeconds": 300
  }
}
```

Compile the model

```
$ aws sagemaker create-compilation-job
--cli-input-json file://config.json
--compilation-job-name resnet50-mxnet-pi

$ aws s3 cp s3://jsimon-neo/model-
rasp3b.tar.gz .

$ gtar tfz model-rasp3b.tar.gz
compiled.params
compiled_model.json
compiled.so
```

Predict with the compiled model


```
from dlr import DLModel
model = DLModel('resnet50', input_shape,
output_shape, device)
out = model.run(input_data)
```

Amazon Elastic Inference

Attach fractional acceleration to any EC2 instance



Lower inference costs
up to 75%



Match capacity
to demand



Available between 1 to 32
TFLOPS

Integrated with
Amazon EC2,
Amazon SageMaker,
and Amazon DL
AMIs

Support for TensorFlow,
Apache MXNet, and
ONNX
with PyTorch coming soon

Single and
mixed-precision
operations

Demo:

Image classification on Caltech-256, with Automatic Model Tuning and Elastic Inference

<https://gitlab.com/juliensimon/dlnotebooks/blob/master/sagemaker/08-Image-classification-advanced.ipynb>

Getting started

<http://aws.amazon.com/free>

<https://ml.aws>

<https://aws.amazon.com/sagemaker>

<https://github.com/aws/sagemaker-python-sdk>

<https://github.com/aws/sagemaker-spark>

<https://github.com/aws-labs/amazon-sagemaker-examples>

<https://gitlab.com/juliensimon/ent321>

<https://medium.com/@julsimon>

<https://gitlab.com/juliensimon/dlnotebooks>

Dank u wel!

Julien Simon
Global Evangelist, AI & Machine Learning, AWS
@julsimon

Breght
Boschker
CTO, SkinVision



Please complete the
session survey.