


Get Started with Machine Learning and Computer Vision Using AWS DeepLens

Julien Simon
Global Evangelist, AI & Machine Learning
[@julsimon](#)

Computing is increasingly available at the edge



Machine Learning predictions at the edge would make devices smarter.
Could we simply invoke cloud-based models?

Most machine data can't reach the Cloud



Medical equipment



Industrial machinery




Extreme environments

Why this problem isn't going away



Law of physics



Law of economics




Law of the land

To-do list

- ❑ Build a data set
- ❑ Experiment with state-of-the art algorithms for computer vision
- ❑ Train in the Cloud at any scale
- ❑ Deploy inference code and model at the edge


Building an image dataset

Annotating large datasets is time-consuming




Amazon SageMaker Ground Truth


<https://aws.amazon.com/blogs/aws/amazon-sagemaker-ground-truth-build-highly-accurate-datasets-and-reduce-labeling-costs-by-up-to-70>



Quickly label
training data



Easily integrate
human labelers



Get accurate
results

KEY FEATURES

Automatic labeling via
machine learning

Ready-made and
custom workflows for
images and text


Private and public
human workforce

Label
management

Experimenting and training at any scale

Do it yourself or fully-managed: you decide!

Amazon SageMaker



AWS Deep Learning AMI

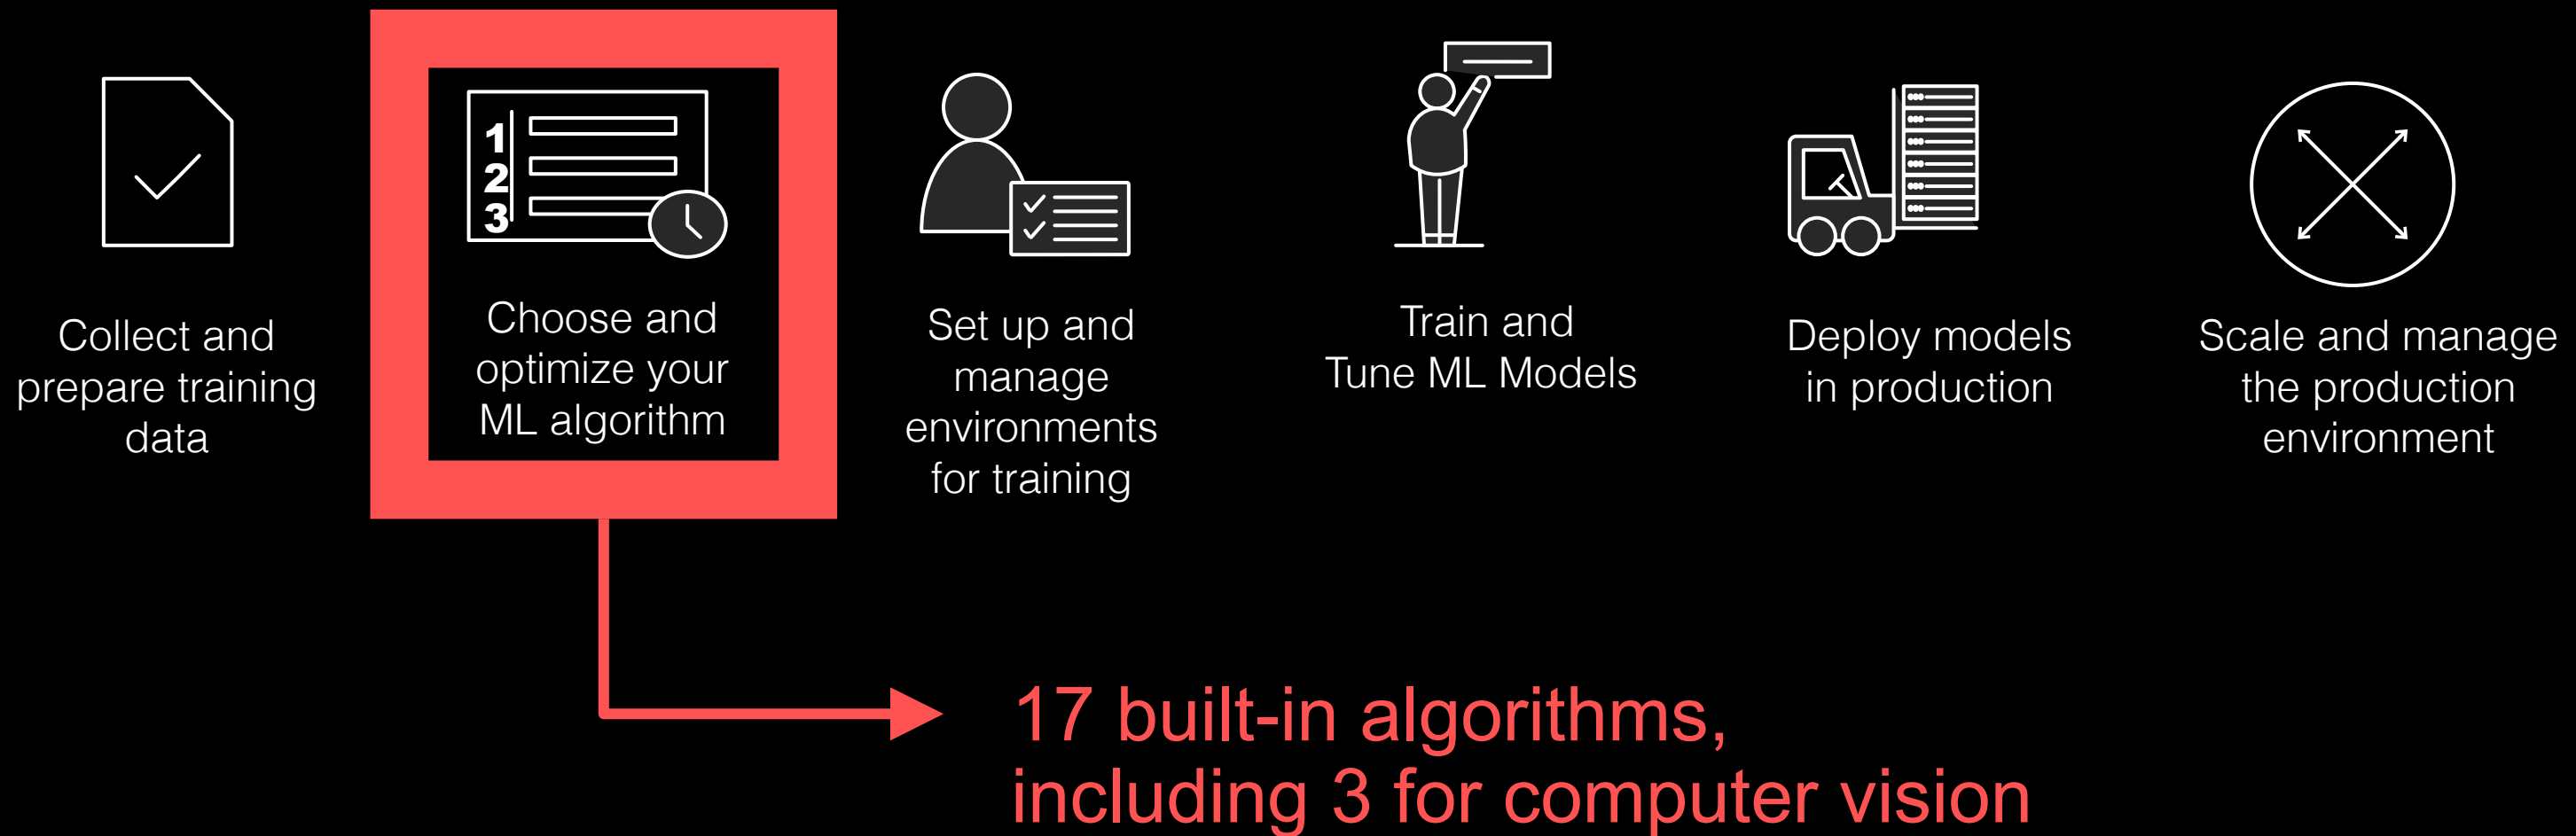


Amazon EC2



Amazon SageMaker

Build, train, and deploy ML Models at any scale





Deep Learning-based algorithms and pre-trained models

Classification, detection, segmentation




[electric_guitar],
with probability 0.671




Amazon SageMaker


Build, train, and deploy ML Models at any scale




Collect and
prepare training
data



Choose and
optimize your
ML algorithm




Set up and
manage
environments
for training



Train and
Tune ML Models



Deploy models
in production



Scale and manage
the production
environment

Deploying inference code and model at the edge

Deploying at the edge with AWS Greengrass



Install the
Greengrass
runtime



Deploy and run
Lambda functions



Manage from
AWS Console




Same programming
model as in the Cloud



Local
communication
and orchestration


Deploying models with AWS Greengrass ML Inference



Define models as
Greengrass
resources and
transfer them to
your devices

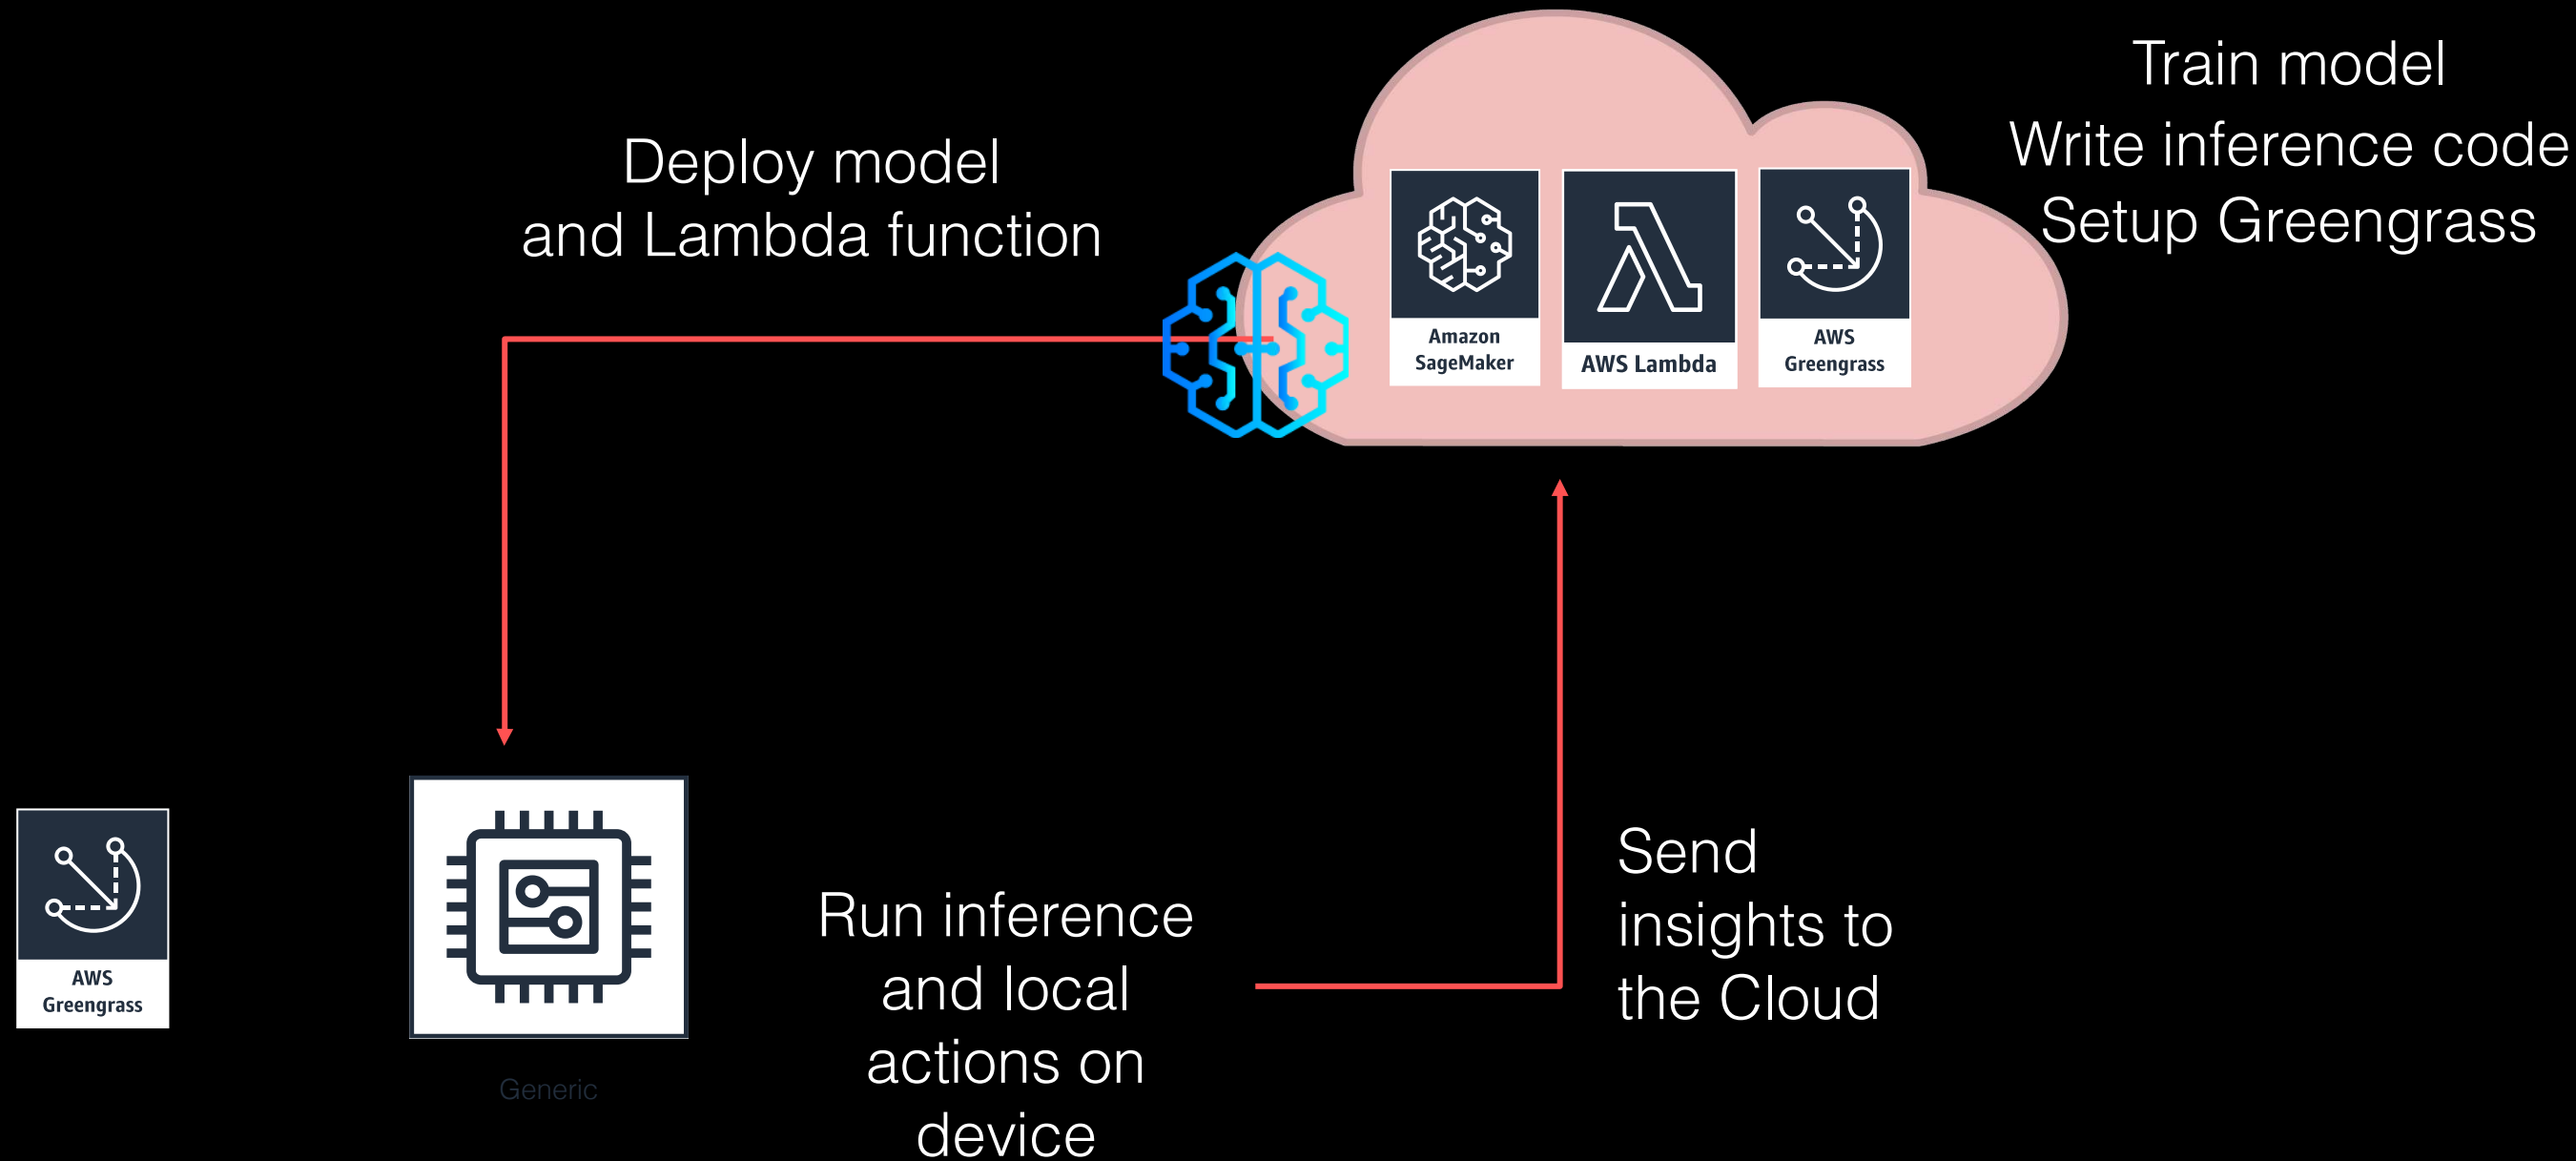


Inference
takes place
on devices



Devices take
action quickly
– even when
disconnected

Architecture



To-do list

- ✓ Build a data set
- ✓ Experiment with state-of-the art algorithms for computer vision
- ✓ Train in the Cloud at any scale
- ✓ Deploy inference code and model at the edge
- ❑ Put it all in practice with a fun device

AWS DeepLens




AWS DeepLens



HD video camera



Custom-designed deep learning inference engine



Micro-SD

Mini-HDMI


USB

USB

Reset

Audio out

Power



HD video camera with on-board compute optimised for deep learning



From unboxing to first inference in <10 minutes



Integrates with Amazon SageMaker and AWS Lambda



Tutorials, examples, demos, and pre-built models

Get started in minutes with sample projects



OBJECT DETECTION

Detect and recognise objects.



HOT DOG NOT HOT DOG

Classify your food.



CAT AND DOG

Detect a cat or dog.



ARTISTIC STYLE TRANSFER

Transfer a style onto video.



ACTIVITY RECOGNITION

Recognise common activities.



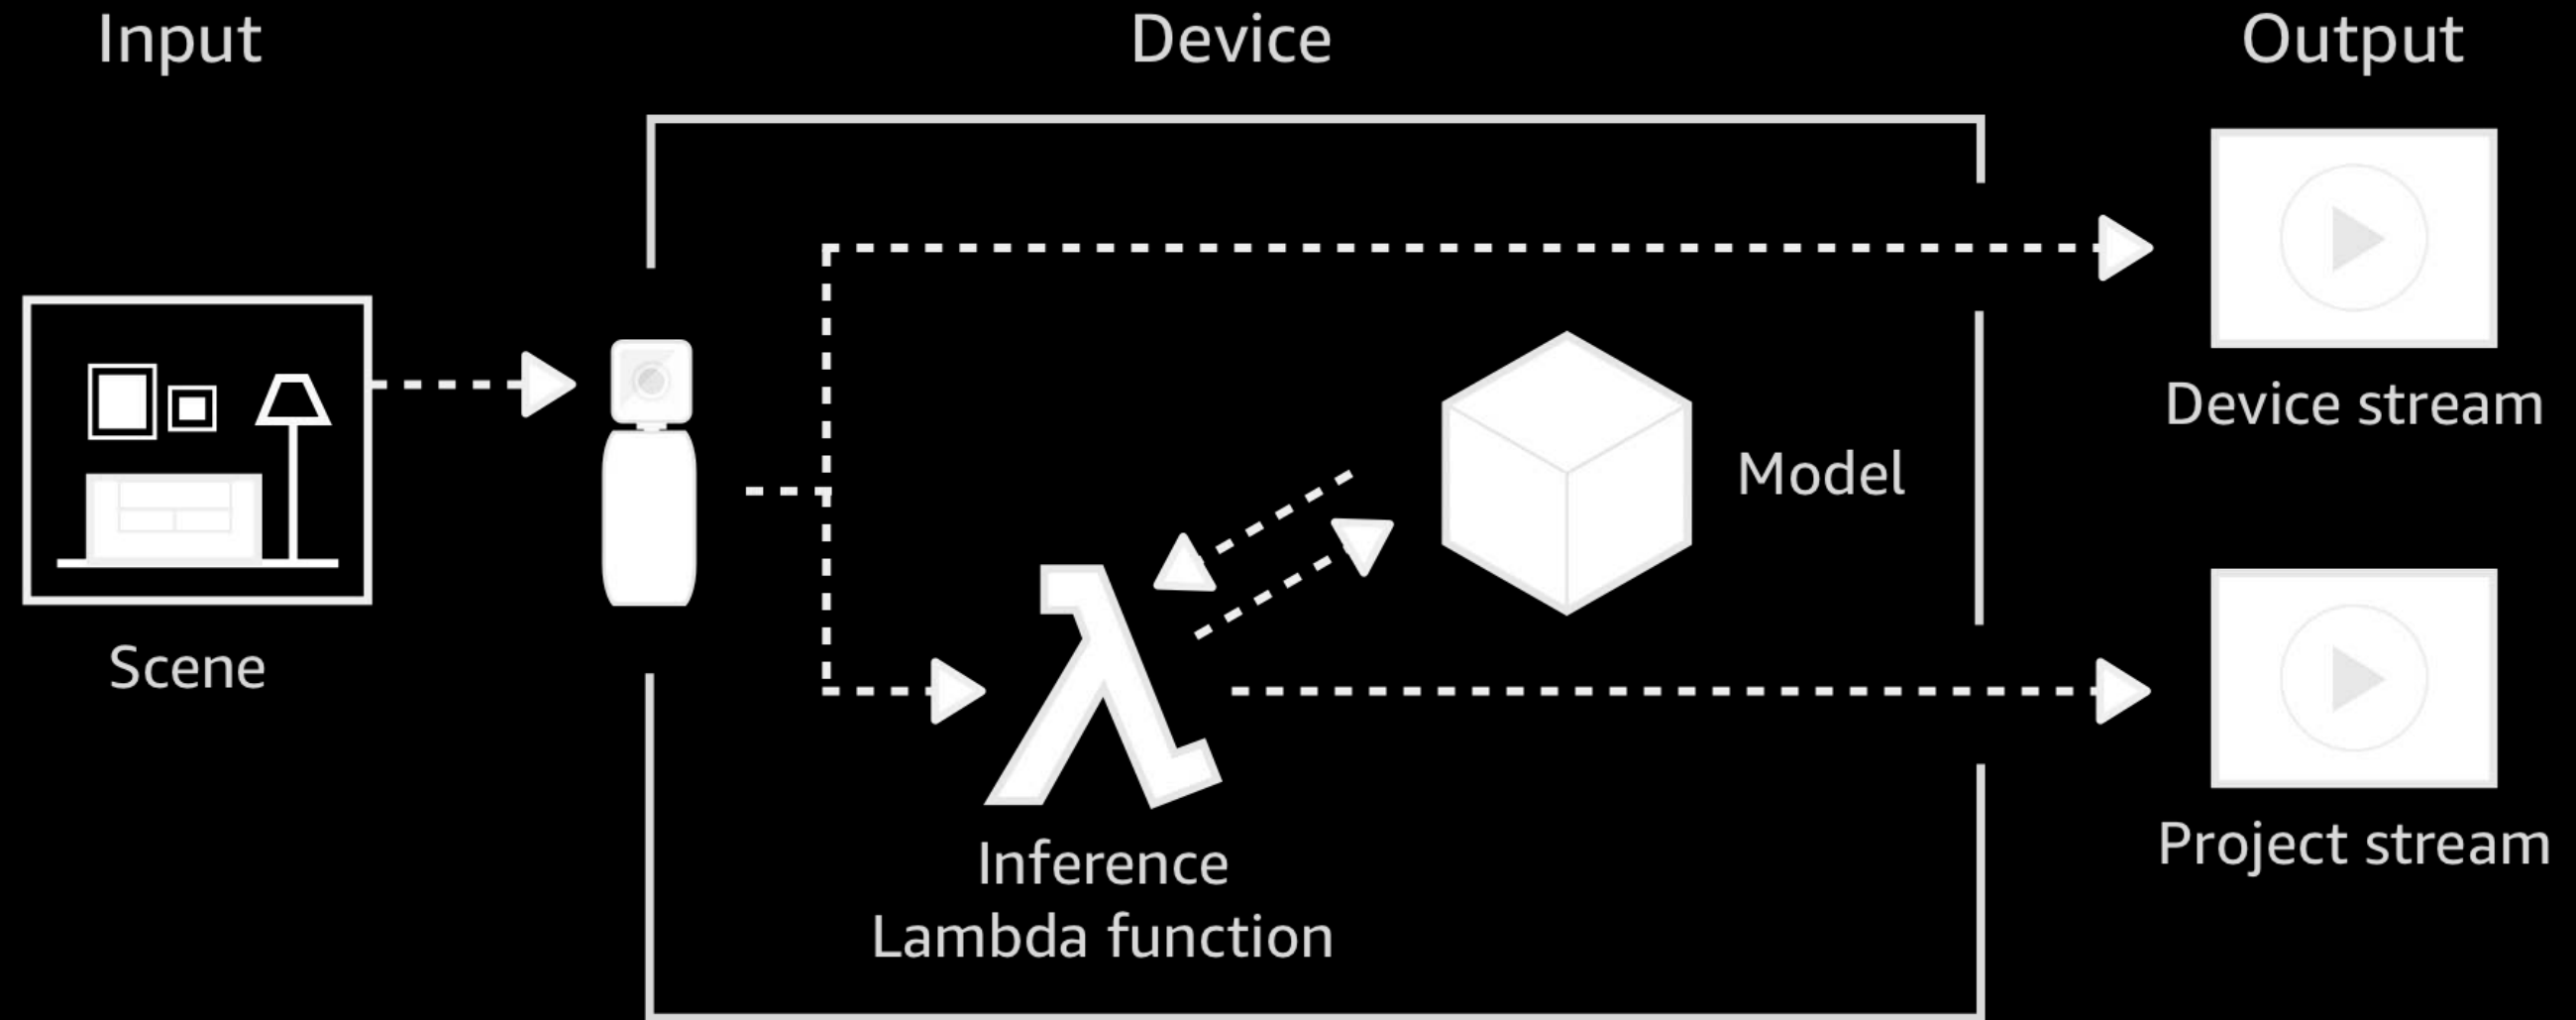
FACE DETECTION

Detect faces of people.

Use your own models with AWS DeepLens

- AWS DeepLens can run TensorFlow, Caffe and Apache MXNet models
 - Inception
 - MobileNet
 - NasNet
 - ResNet
 - Etc.
- Train or fine-tune your model on Amazon SageMaker
- Deploy to AWS DeepLens with AWS Greengrass

AWS DeepLens



Writing the Lambda function

not scary: it's mostly cut and paste ;)

Lambda function: load the model

```
# When deployed to a Greengrass core, this code will be executed immediately
# as a long-lived lambda function.

def greengrass_infinite_infer_run():
    try:
        modelPath = "/opt/awscam/artifacts/mxnet_squeezenet.xml"
        modelType = "classification"

        # Send a starting message to IoT console
        client.publish(topic=iotTopic, payload="Infinite inference starts now")

        # Load model to GPU (use {"GPU": 0} for CPU)
        mcfg = {"GPU": 1}
        model = awscam.Model(modelPath, mcfg)
        client.publish(topic=iotTopic, payload="Model loaded")
```

Lambda function: optimize a custom model

- Custom models need to be optimized for the on-board GPU.
- The first call optimizes the model, further calls do nothing.

```
error, model_path = mo.optimize(model_name, input_width, input_height)
```

Lambda function: get a video frame and predict

```
doInfer = True
while doInfer:
    # Get a frame from the video stream
    ret, frame = awscam.getLastFrame()
    numFrames += 1

    # Raise an exception if failing to get a frame
    if ret == False:
        raise Exception("Failed to get frame from the stream")

    # Resize frame to fit model input requirement
    frameResize = cv2.resize(frame, (224, 224))

    # Run model inference on the resized frame
    inferOutput = model.doInference(frameResize)
```

Lambda function: annotate live stream

```
# Output inference result to the fifo file so it can be viewed with mplayer
parsed_results = model.parseResult(modelType, inferOutput)['ssd']
label = '{}'
for obj in parsed_results:
    if obj['prob'] > max_threshold:
        xmin = int( xscale * obj['xmin'] ) + int((obj['xmin'] - input_width/2) + input_width/2)
        ymin = int( yscale * obj['ymin'] )
        xmax = int( xscale * obj['xmax'] ) + int((obj['xmax'] - input_width/2) + input_width/2)
        ymax = int( yscale * obj['ymax'] )
        cv2.rectangle(frame, (xmin, ymin), (xmax, ymax), (0, 255, 0))
        label += "{}: {:.2f}, ".format(outMap[obj['label']], obj['prob'])
        label_show = "{}: {:.2f}%".format(outMap[obj['label']], obj['prob'] * 100)
        cv2.putText(frame, label_show, (xmin, ymin-15), cv2.FONT_HERSHEY_SIMPLEX, 0.5, (0, 255, 0))
label += '"null": 0.0'
label += '}'
client.publish(topic=iotTopic, payload = label)
global jpeg
ret, jpeg = cv2.imencode('.jpg', frame)
```



Demo

- Use the built-in algorithm for image classification in Amazon SageMaker
- Fine tune a pre-trained model on the CIFAR-256 image dataset
- Write a simple Lambda function for inference
- Deploy function and model to AWS DeepLens

<https://gitlab.com/juliensimon/dlnotebooks/sagemaker/>

ml.aws
aws.amazon.com/deeplens
aws.training/
machinelearning

Thank you!

Julien Simon
Global Evangelist, AI and Machine Learning

@julsimon
<https://medium.com/julsimon>



Please complete the
session survey.