


DEV DAY

Build, train and deploy Machine Learning models at scale


Julien Simon
Global Evangelist, AI & Machine Learning
@julsimon




The Amazon ML Stack: Broadest & Deepest Set of Capabilities



ML SERVICES



ML FRAMEWORKS & INFRASTRUCTURE




Amazon SageMaker:

Build, Train, and Deploy ML Models at Scale




Collect and
prepare training
data



Choose and
optimize your
ML algorithm




Set up and
manage
environments
for training



Train and
Tune ML Models



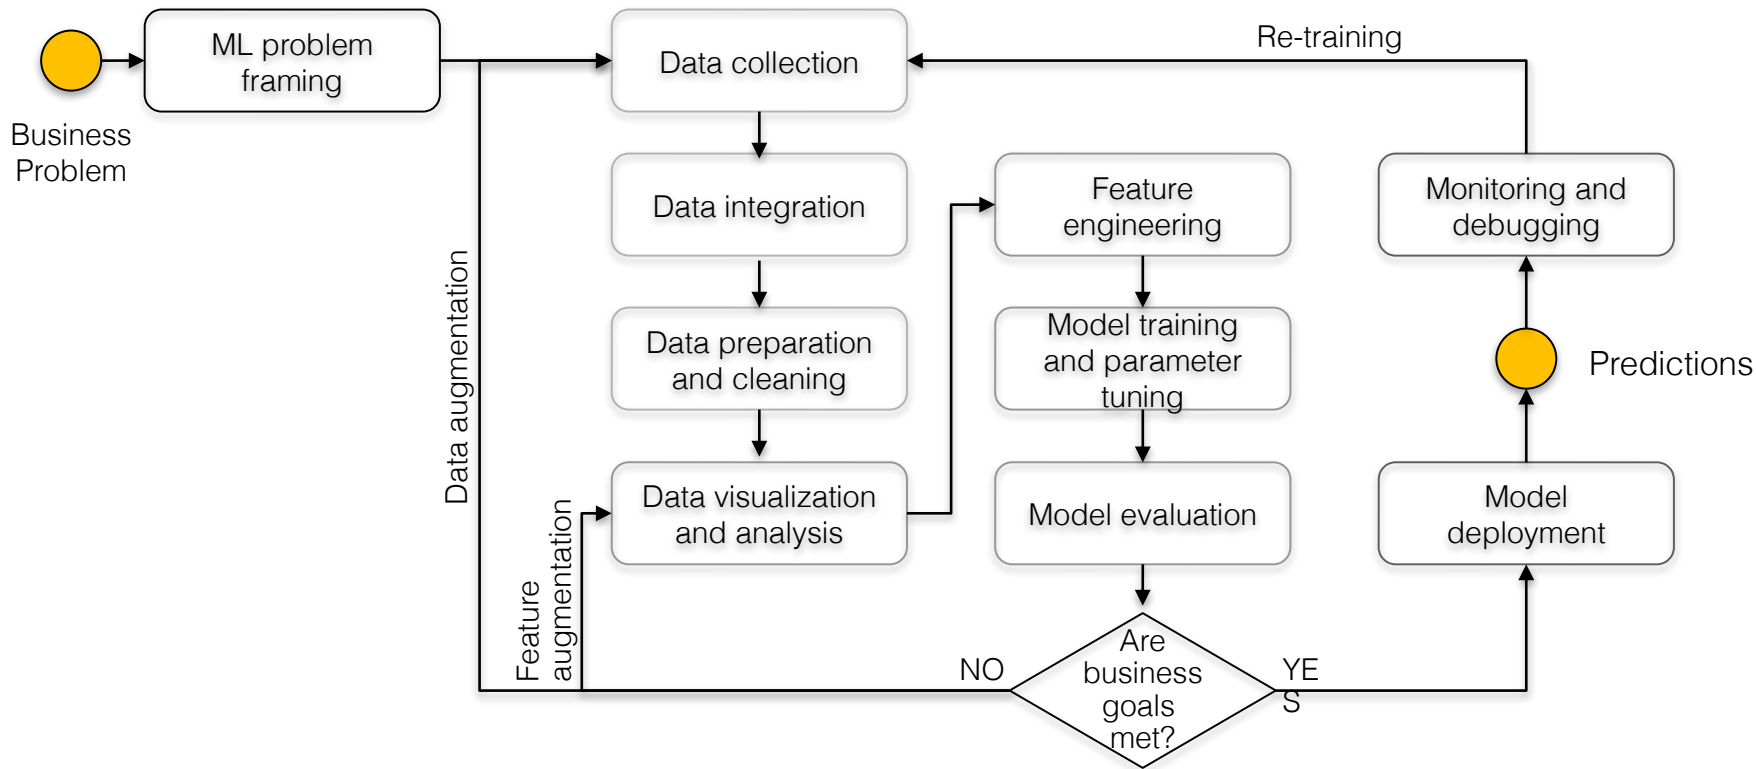
Deploy models
in production



Scale and manage
the production
environment




Machine learning cycle



Successful models require high-quality data




Successful models require high-quality data




Amazon SageMaker Ground Truth


<https://aws.amazon.com/blogs/aws/amazon-sagemaker-ground-truth-build-highly-accurate-datasets-and-reduce-labeling-costs-by-up-to-70>



Quickly label
training data



Easily integrate
human labelers



Get accurate
results

KEY FEATURES

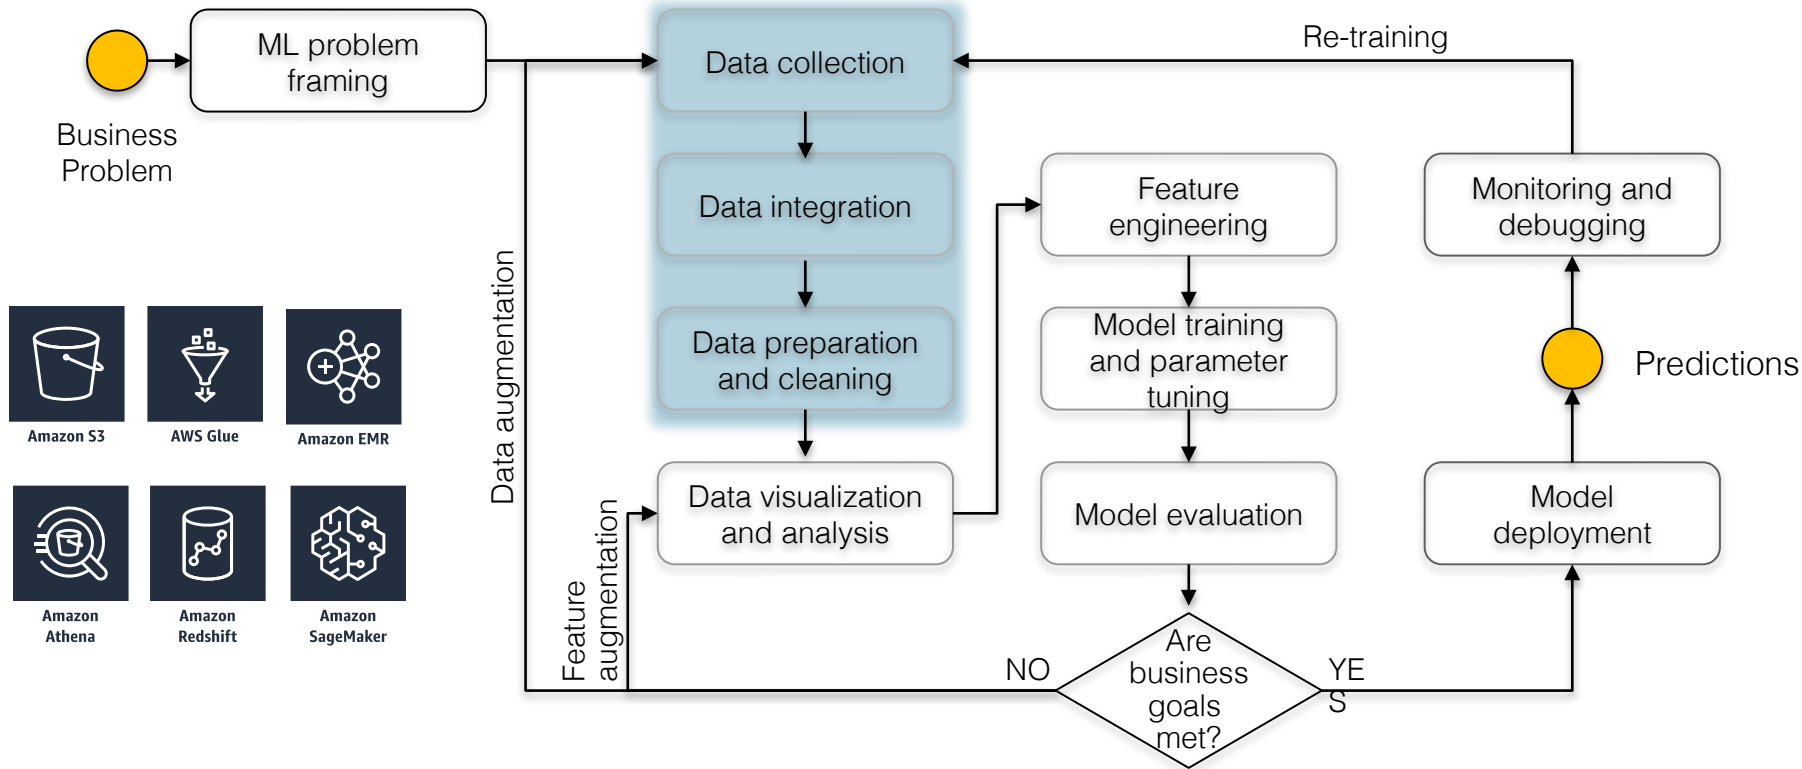
Automatic labeling via
machine learning

Ready-made and
custom workflows for
image bounding box,
segmentation, and text

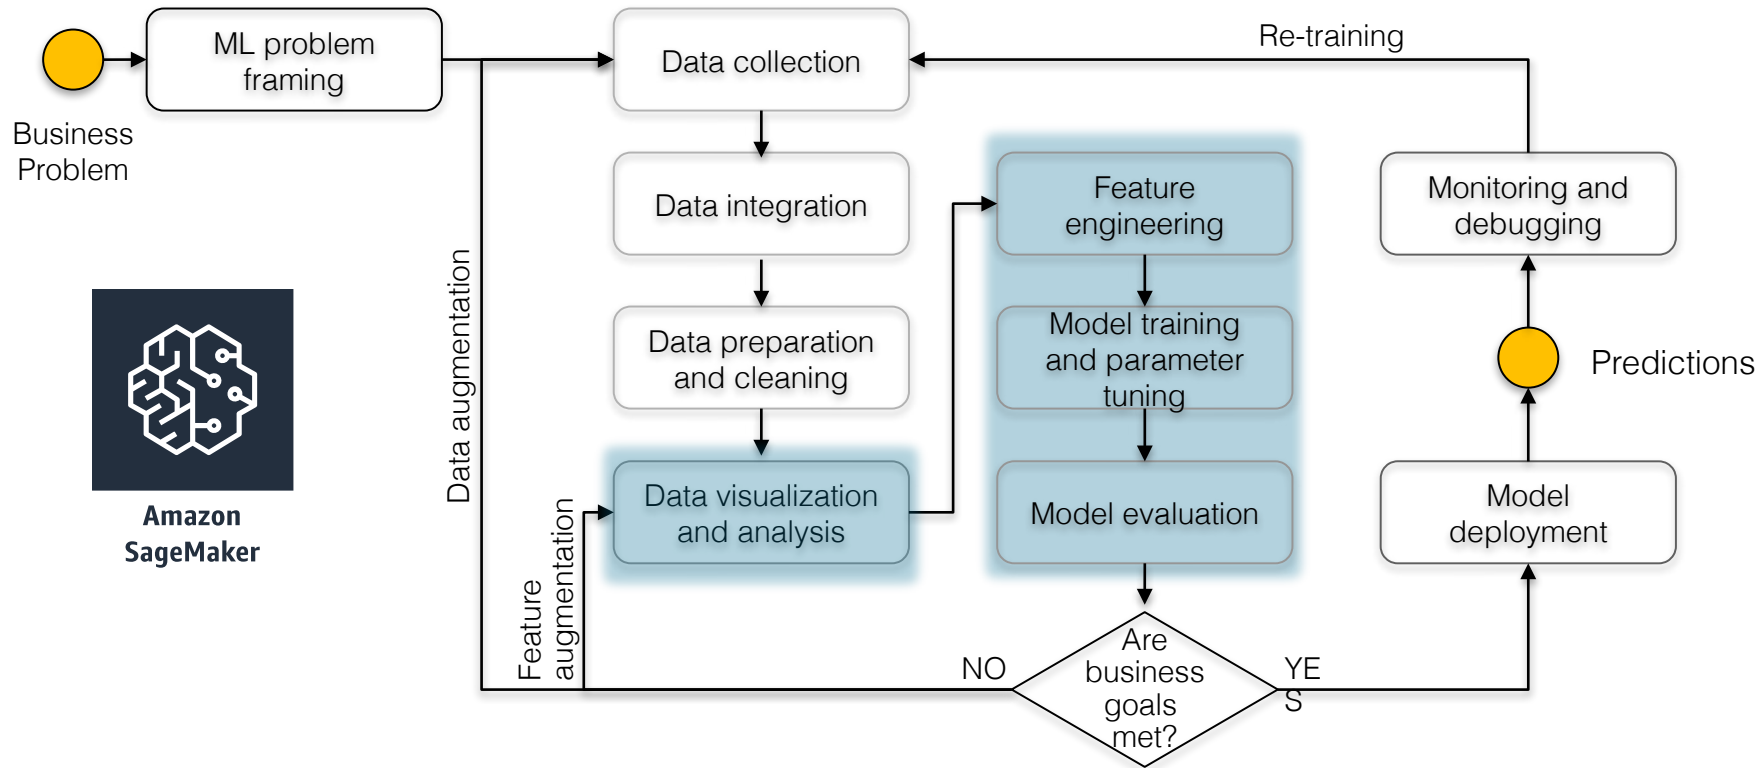
Private and public
human workforce

Label
management

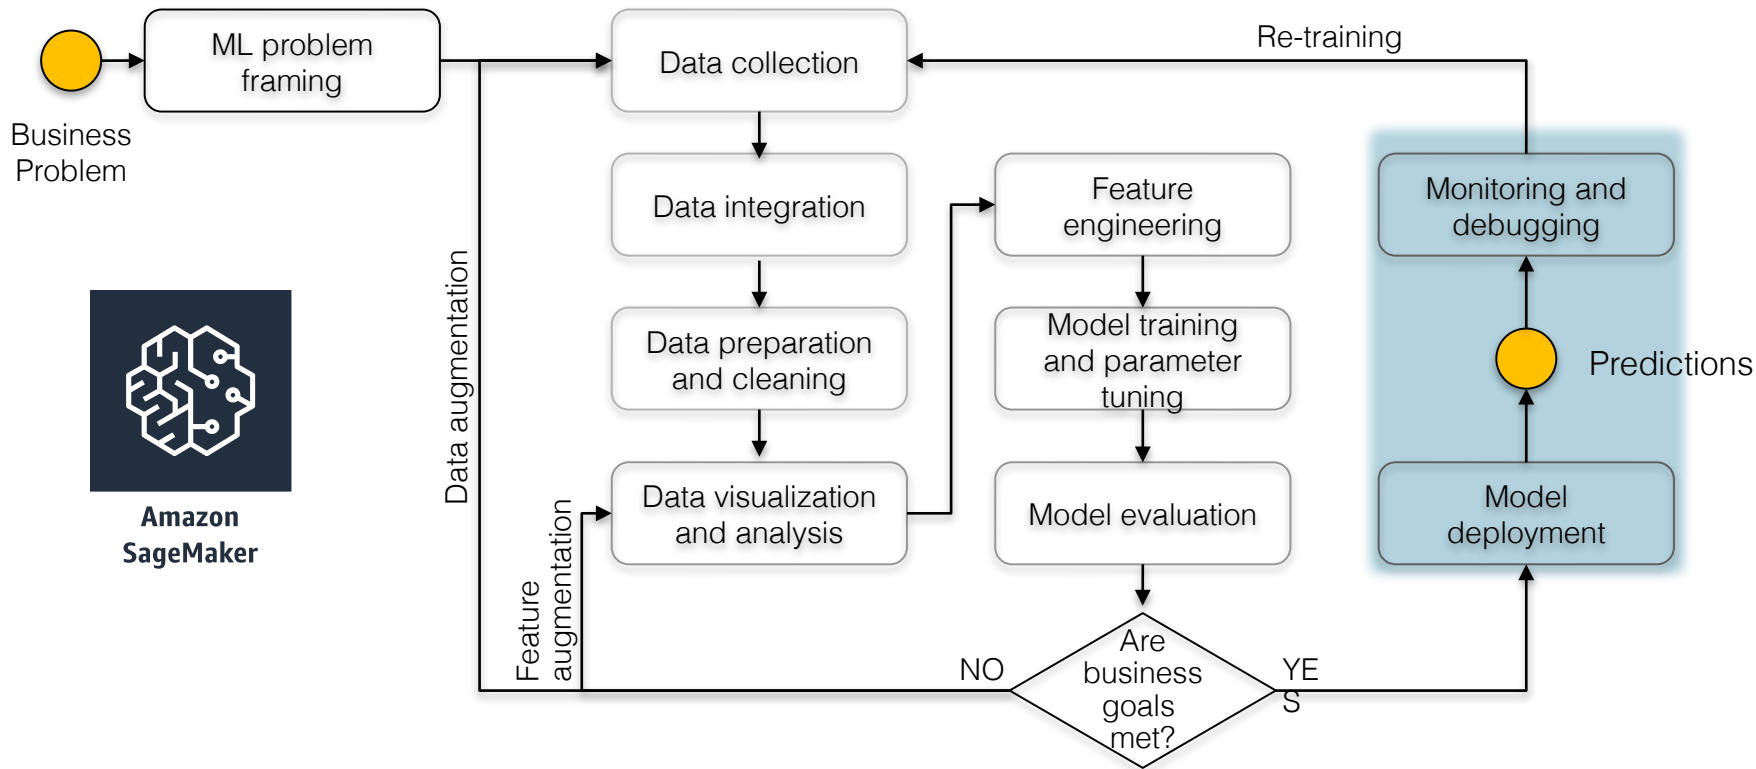
Manage data on AWS



Build and train models using SageMaker




Deploy models using SageMaker



Amazon SageMaker




Collect and
prepare training
data



Choose and
optimize your
ML algorithm




Set up and
manage
environments
for training



Train and
Tune ML Models



Deploy models
in production



Scale and manage
the production
environment



Git integration
Elastic inference



New built-in algorithms
scikit-learn environment
Model marketplace
Search




P3DN, C5N
TensorFlow on 256 GPUs
Resume HPO tuning job



Model compilation
Elastic inference
Inference pipelines

Machine Learning Marketplace



Hello, julien

Categories ▾Delivery Methods ▾Solutions ▾Migration Mapping AssistantYour Saved ListPartnersSell in AWS MarketplaceAmazon Web Services HomeHelp

Categories

All Categories

[Infrastructure Software \(27\)](#)

[Business Software \(7\)](#)

[Machine Learning \(61\)](#)

Filters

[Clear all filters](#)

Vendors

☐ RocketML (15)

☐ Sensifai (9)

☐ Intel® AI (7)

☐ Peak (5)

☐ Outpace Systems (5)

☐ improve.ai (4)

☐ H2O.ai (4)

☐ Dimensional Mechanics (4)

☐ TIBCO Software Inc. (3)

☐ bigfinite (1)

[Show more](#)

All Categories (61 results) showing 1 - 10

1

2

3


4

5

6

7


▶



Text Similarity Analyzer

★★★★★ (0) | Sold by TIBCO Software Inc.


Engineers word/document features on a corpus with NLP methods, and uses these features to compare new text to the corpus.



Automatic Image Tagging

★★★★★ (0) | Sold by Sensifai

Automatic Image Tagging and Recognition



Demand Forecasting for Intermittent Data

★★★★★ (0) | Sold by Peak

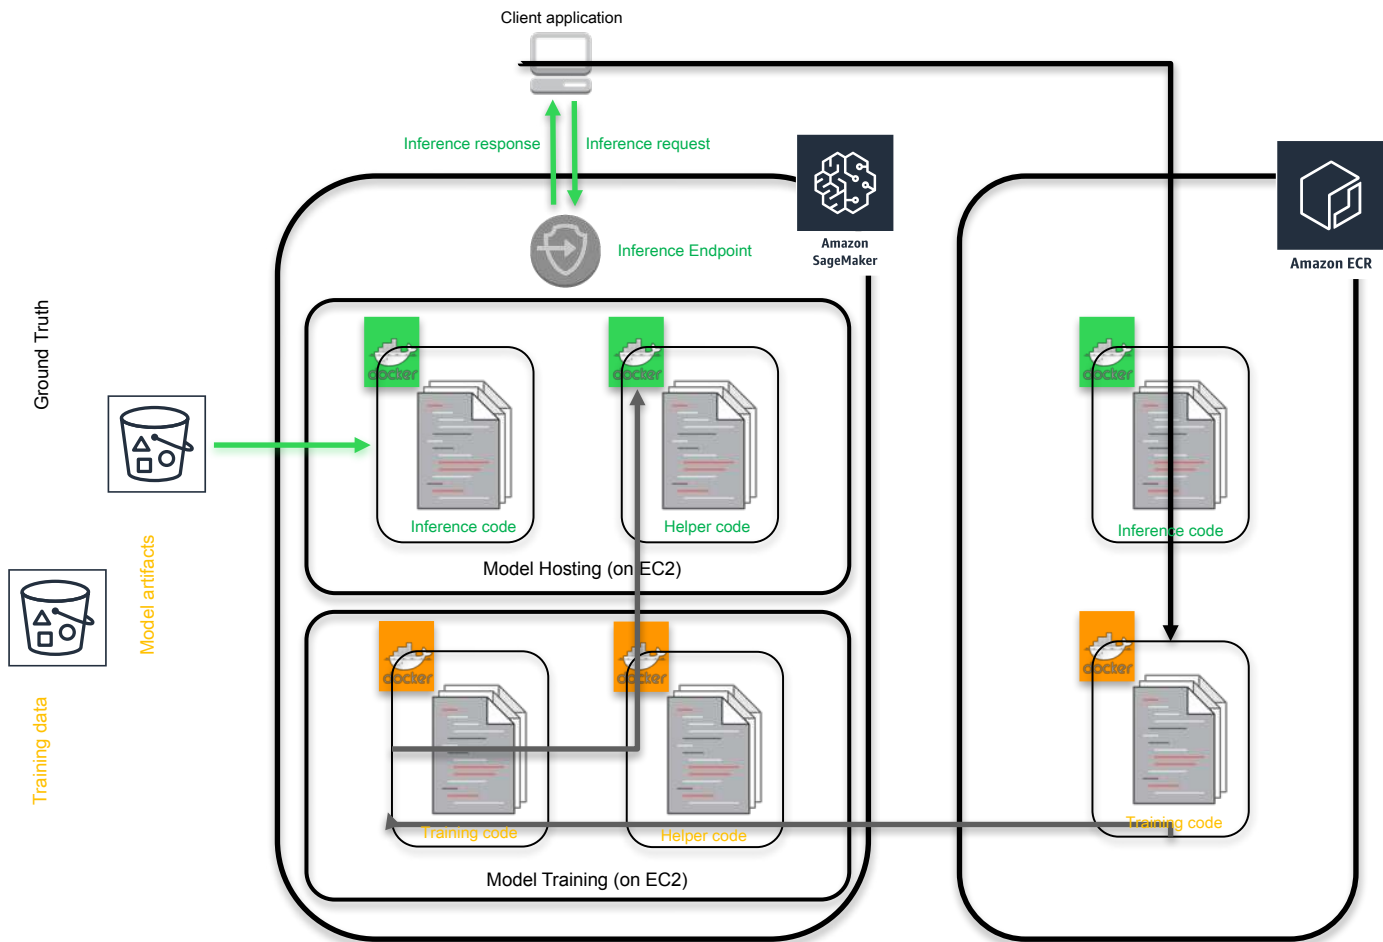
An ensemble demand forecasting model, for intermittent data

Working with Amazon SageMaker



The Amazon SageMaker API

- Python SDK **orchestrating** all Amazon SageMaker activity
 - High-level objects for **algorithm selection**, **training**, **deploying**, **automatic model tuning**, etc.
 - **Spark SDK** (Python & Scala)
- AWS CLI: *'aws sagemaker'*
- AWS SDK: boto3, etc.



Model options



Training code

Factorization Machines
Linear Learner
Principal Component
Analysis
K-Means Clustering
XGBoost
And more

Built-in Algorithms



Bring Your Own Script



Bring Your Own Container

Built-in algorithms

Built-in algorithms

orange: supervised, yellow: unsupervised

Linear Learner: regression, classification	Image Classification: Deep Learning (ResNet)
Factorization Machines: regression, classification, recommendation	Object Detection (SSD): Deep Learning (VGG or ResNet)
K-Nearest Neighbors: non-parametric regression and classification	Neural Topic Model: topic modeling
XGBoost: regression, classification, ranking https://github.com/dmlc/xgboost	Latent Dirichlet Allocation: topic modeling (mostly)
K-Means: clustering	Blazing Text: GPU-based Word2Vec, and text classification
Principal Component Analysis: dimensionality reduction	Sequence to Sequence: machine translation, speech to text and more
Random Cut Forest: anomaly detection	DeepAR: time-series forecasting (RNN)
Object2Vec: general-purpose embedding	IP Insights: usage patterns for IP addresses
Semantic Segmentation: Deep Learning	

Demo:

Image classification with Caltech-256

<https://gitlab.com/juliensimon/dlnotebooks/sagemaker/>

Blazing Text

BlazingText: Scaling and Accelerating Word2Vec using Multiple GPUs

Saurabh Gupta
Amazon Web Services
gsaur@amazon.com

Vineet Khare
Amazon Web Services
vkhare@amazon.com

<https://dl.acm.org/citation.cfm?id=3146354>



Demo:

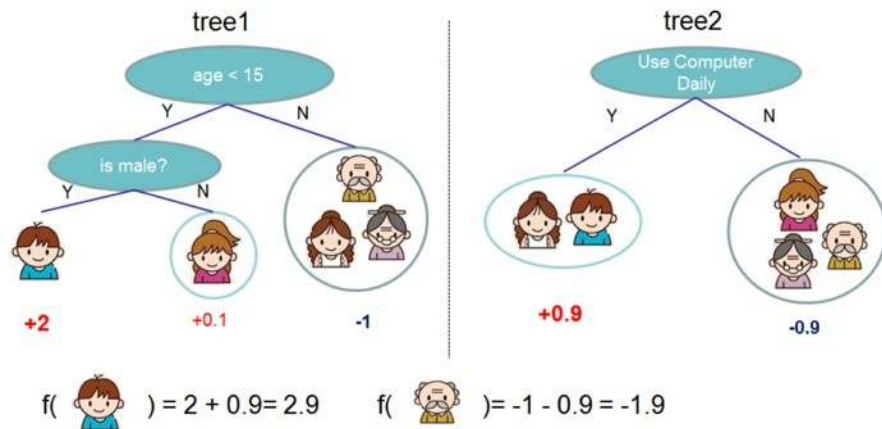
Text Classification with BlazingText

https://github.com/aws-labs/amazon-sagemaker-examples/tree/master/introduction_to_amazon_algorithms/blazingtext_text_classification_dbpedia

XGBoost



- Open Source project
- Popular tree-based algorithm for regression, classification and ranking
- Builds a collection of trees.
- Handles missing values and sparse data
- Supports distributed training
- Can work with data sets larger than RAM



<https://github.com/dmlc/xgboost>

<https://xgboost.readthedocs.io/en/latest/>

<https://arxiv.org/abs/1603.02754>

Demo: XGBoost

AWS re:Invent 2018 workshop

<https://gitlab.com/juliensimon/ent321>

Built-in frameworks



Demo:

Keras/TensorFlow CNN on CIFAR-10

https://github.com/aws-labs/amazon-sagemaker-examples/blob/master/sagemaker-python-sdk/tensorflow_keras_cifar10/tensorflow_keras_CIFAR10.ipynb

Demo:


Sentiment analysis with Apache MXNet

https://github.com/aws-labs/amazon-sagemaker-examples/blob/master/sagemaker-python-sdk/mxnet_sentiment_analysis_with_gluon.ipynb

Amazon SageMaker




Collect and
prepare training
data



Choose and
optimize your
ML algorithm




Set up and
manage
environments
for training



Train and
Tune ML Models



Deploy models
in production



Scale and manage
the production
environment

Build

Train

Deploy

FREE TIER

Getting started

<http://aws.amazon.com/free>

<https://ml.aws>

<https://aws.amazon.com/sagemaker>

<https://github.com/aws/sagemaker-python-sdk>

<https://github.com/aws/sagemaker-spark>

<https://github.com/aws-labs/amazon-sagemaker-examples>

<https://gitlab.com/juliensimon/ent321>

<https://medium.com/@julsimon>

<https://gitlab.com/juliensimon/dlnotebooks>

DEV DAY

Thank you!

Julien Simon

Global Evangelist, AI & Machine Learning

@julsimon

<https://medium.com/@julsimon>

