

An Introduction to Gluon

Julien Simon
Global Technical Evangelist, AI & Machine Learning, AWS
@julsimon

Apache MXNet and Gluon

Apache MXNet



- Open source software library for Deep Learning
- Natively implemented in C++
- Built-in support for many network architectures: FC, CNN, LSTM, etc.
- **Symbolic** API: Python, Scala, Clojure, R, Julia, Perl, Java (inference only)
- **Imperative** API: **Gluon** (Python), with computer vision and natural language processing toolkits


Symbolic execution

‘define then run’

```
A = Variable('A')
B = Variable('B')
C = B * A
D = C + 1
f = compile(D)
d = f(A=np.ones(10),
```

```
B=np.ones(10)*2)
```

C can share memory with D
because C is deleted later



PRO

S

- More chances for optimization
- Language independent
- E.g. TensorFlow, Theano, Caffe, MXNet

CON

S

- Less flexible
- ‘Black box’ training

Apache MXNet : *Symbol* API

```
import mxnet as mx
```

```
train_iter = mx.io.MNISTIter(shuffle=True)
val_iter = mx.io.MNISTIter(image="./t10k-images-idx3-ubyte", label="./t10k-labels-idx1-ubyte")
```

```
data = mx.sym.Variable('data')
data = mx.sym.Flatten(data=data)
fc1 = mx.sym.FullyConnected(data=data, name='fc1', num_hidden=512)
act1 = mx.sym.Activation(data=fc1, name='relu1', act_type="relu")
drop1= mx.sym.Dropout(data=act1,p=0.2)
fc2 = mx.sym.FullyConnected(data=drop1, name='fc2', num_hidden = 256)
act2 = mx.sym.Activation(data=fc2, name='relu2', act_type="relu")
drop2= mx.sym.Dropout(data=act2,p=0.2)
fc3 = mx.sym.FullyConnected(data=drop2, name='fc3', num_hidden=10)
mlp = mx.sym.SoftmaxOutput(data=fc3, name='softmax')
```

```
mod = mx.mod.Module(mlp, context=mx.cpu(0))
mod.bind(data_shapes=train_iter.provide_data, label_shapes=train_iter.provide_label)
mod.init_params(initializer=mx.init.Xavier())
mod.init_optimizer('adam', optimizer_params= (('learning_rate', 0.1),))

mod.fit(train_iter, eval_data=val_iter, num_epoch=50)
```

Define
a network

Train an
optimized
version

Imperative execution

‘define by run’

```
import numpy as np
a = np.ones(10)
b = np.ones(10) * 2
c = b * a
d = c + 1
```

PRO S

- Straightforward and flexible.
- Take advantage of language native features (loop, condition, debugger).
- E.g. Numpy, PyTorch, Gluon API

CON S

- Harder to optimize

Apache MXNet : Gluon API

```
import mxnet as mx
from mxnet import gluon, autograd, ndarray
```

```
train_data = ...
test_data = ...
```

```
net = gluon.nn.Sequential()
with net.name_scope():
    net.add(gluon.nn.Dense(128, activation="relu"))
    net.add(gluon.nn.Dense(64, activation="relu"))
    net.add(gluon.nn.Dense(10))
```

```
net.collect_params().initialize(mx.init.Normal(sigma=0.05))
softmax_cross_entropy = gluon.loss.SoftmaxCrossEntropyLoss()
trainer = gluon.Trainer(net.collect_params(), 'sgd', {'learning_rate': 0.1})
```

```
epochs = 10
for e in range(epochs):
    for i, (data, label) in enumerate(train_data):
        data = data.as_in_context(mx.cpu()).reshape((-1, 784))
        label = label.as_in_context(mx.cpu())
        with autograd.record():
            output = net(data)
            loss = softmax_cross_entropy(output, label)
            loss.backward()
        trainer.step(data.shape[0])
```

Define a
network

Train using
the same
network

Demo

<https://gitlab.com/juliensimon/dlnotebooks/blob/master/mxnet/06%20-%20MNIST%20with%20Gluon%20API.ipynb>


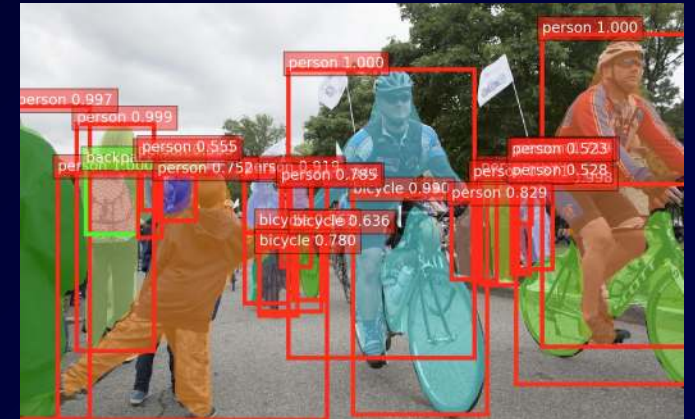
GluonCV and GluonNLP

GluonCV

<https://gluon-cv.mxnet.io>

<https://github.com/dmlc/gluon-cv>

- State-of-the-art **deep learning tools for computer vision**
 - Pre-trained models
 - Training and fine-tuning scripts
 - Prototype products, validate new ideas and learn computer vision
- Image classification: 50+ models
- Object detection: Faster RCNN, SSD, Yolo-v3
- Semantic segmentation: FCN, PSP, DeepLab v3
- Instance segmentation: Mask RCNN
- Pose estimation: Simple Pose
- Person re-identification (Market1501 dataset)
- GANs: Wasserstein GAN, Super Resolution GAN, CycleGAN



Demos

gluon-cv/docs/tutorials

GluonNLP

<https://gluon-nlp.mxnet.io><https://github.com/dmlc/gluon-nlp>

- State-of-the-art **deep learning tools for natural language processing**
 - Pre-trained models and embeddings
 - Training and fine-tuning scripts
 - Prototype products, validate new ideas and learn NLP
- Word embeddings: Word2Vec, FastText, GloVE, BERT
- Machine translation: GNMT, Transformer
- Sentiment analysis: TextCNN
- Text classification: FastText
- Language models
- Text generation
- Natural language inference
- Parsing

Demos

Deep Learning on AWS

Amazon SageMaker



AWS Deep Learning AMI AWS Deep Learning containers*



Amazon EC2



* Tensorflow and MXNet only
PyTorch coming soon

Apache MXNet on Amazon SageMaker: a first-class citizen

- Built-in containers for training and prediction
 - Code available on **Github**: <https://github.com/aws/sagemaker-mxnet-container>
 - Build it, run it on your own machine, customize it, push it to Amazon ECR, etc.
 - Supported versions: 0.12.1, 1.0.0, 1.1.0, 1.2.1, 1.3.0
- Advanced features
 - **Local mode**: train on the notebook instance for faster experimentation
 - **Script mode**: use the same TensorFlow as on your local machine
 - **Distributed training**: zero setup!
 - **Pipe mode**: stream large datasets directly from Amazon S3
 - **Keras support** (tf.keras.* and keras.*)

Demo

Getting started

<http://aws.amazon.com/free>

<https://aws.amazon.com/sagemaker>

<https://github.com/aws/sagemaker-python-sdk>

<https://github.com/aws-labs/amazon-sagemaker-examples>

<https://medium.com/@julsimon>

<https://gitlab.com/juliensimon/dlnotebooks>

Thank you!

Julien Simon
Global Technical Evangelist, AI & Machine Learning, AWS
@julsimon