



INNOVATE
ONLINE CONFERENCE

MACHINE LEARNING
AND AI EDITORIAL



INNOVATE
ONLINE CONFERENCE

MACHINE LEARNING
AND AI EDITORIAL

|

Deep Dive on Amazon SageMaker

Julien Simon
Global Evangelist, AI & Machine Learning
Amazon Web Services

@julsimon

<https://medium.com/@julsimon>

Agenda

1. Storage

- Amazon S3 & Pipe mode
- Amazon EFS **NEW!**
- Amazon FSx for Lustre **NEW!**

2. Training

- Distributed Training
- Managed Spot Training **NEW!**

3. Model tuning

4. Deployment

- Model compilation: Amazon SageMaker Neo
- Cost optimization: Amazon Elastic Inference

Storage



INNOVATE
ONLINE CONFERENCE

MACHINE LEARNING
AND AI EDITORIAL

Passing datasets to algorithms

- Amazon SageMaker algorithms accept input data from channels
 - A channel is a **named input source defining a dataset**
 - At least one, up to twenty: Training, validation, test, etc.
- Channel object
 - Name
 - Data source: `S3DataSource` or `FileSystemDataSource`
 - Data format: CSV, RecordIO, etc.
 - Compression type
 - Input mode: File or Pipe (S3 only)
- The list of channels is passed to `CreateTrainingJob`
- SageMaker Python SDK: `Estimator.fit()` receives a dictionary of
 - `sagemaker.inputs.s3_input` for S3
 - `sagemaker.inputs.FileSystemInput` for EFS/FSx

Storing your dataset in Amazon S3

- Simplest option
 - `sagemaker.session.default_bucket()`, `sagemaker.session.upload_data()`
- S3DataSource
 - Location: URI
 - Type: Prefix, manifest, augmented manifest
 - Distribution: Fully replicated (training instances receive the full dataset), or sharded ($1/n^{\text{th}}$ of the dataset)
- Input mode: **File** or **Pipe**?
 - File: **Copy** the dataset to each training instance (full or $1/n^{\text{th}}$)
 - Pipe: **Stream** directly from S3
 - Training **starts faster** and **runs faster**
 - No need to provision lots of **storage** on training instances
 - Train on **arbitrary large datasets**, as they don't need to be fully stored or loaded in RAM any longer
- Pipe mode is supported by **most built-in algorithms** and can be implemented in TensorFlow, **Apache MXNet**, etc.
 - <https://aws.amazon.com/blogs/machine-learning/using-pipe-input-mode-for-amazon-sagemaker-algorithms/>
 - <https://aws.amazon.com/blogs/machine-learning/accelerate-model-training-using-faster-pipe-mode-on-amazon-sagemaker/>



Demo:

A quick look at Pipe Mode with TensorFlow

https://github.com/aws-labs/amazon-sagemaker-examples/tree/master/sagemaker-python-sdk/tensorflow_script_mode_pipe_mode



Storing your dataset in Amazon EFS

<https://aws.amazon.com/blogs/machine-learning/speed-up-training-on-amazon-sagemaker-using-amazon-efs-or-amazon-fsx-for-lustre-file-systems/>

- If your organization is **sharing data over NFS**, this is a good option
 - Shared datasets, notebooks, etc.
 - Train directly from EFS, **no data movement** required
- Training instances must run in a **VPC**, and open **port 2049** (NFS)
- **FileSystemDataSource**
 - Filesystem id: provided by EFS
 - Type: 'EFS'
 - Directory path
 - Access type: Read-only or read-write

```
estimator = TensorFlow(
    entry_point='tensorflow_mnist/mnist.py',
    role='SageMakerRole',
    train_instance_count=1,
    train_instance_type='ml.c4.xlarge',
    subnets=['subnet-1', 'subnet-2'],
    security_group_ids=['sg-1'])

file_system_input = FileSystemInput(
    file_system_id='fs-1',
    file_system_type='EFS',
    directory_path='/tensorflow',
    file_system_access_mode='ro')

estimator.fit(inputs=file_system_input)
```



Amazon FSx for Lustre

Fully managed Lustre file system for compute-intensive workloads


<https://aws.amazon.com/fsx/lustre/>



Massively scalable
performance



Seamless access to
your data repositories



Simple
and fully managed



Native file
system interface



Cost-optimized for
compute-intensive workloads

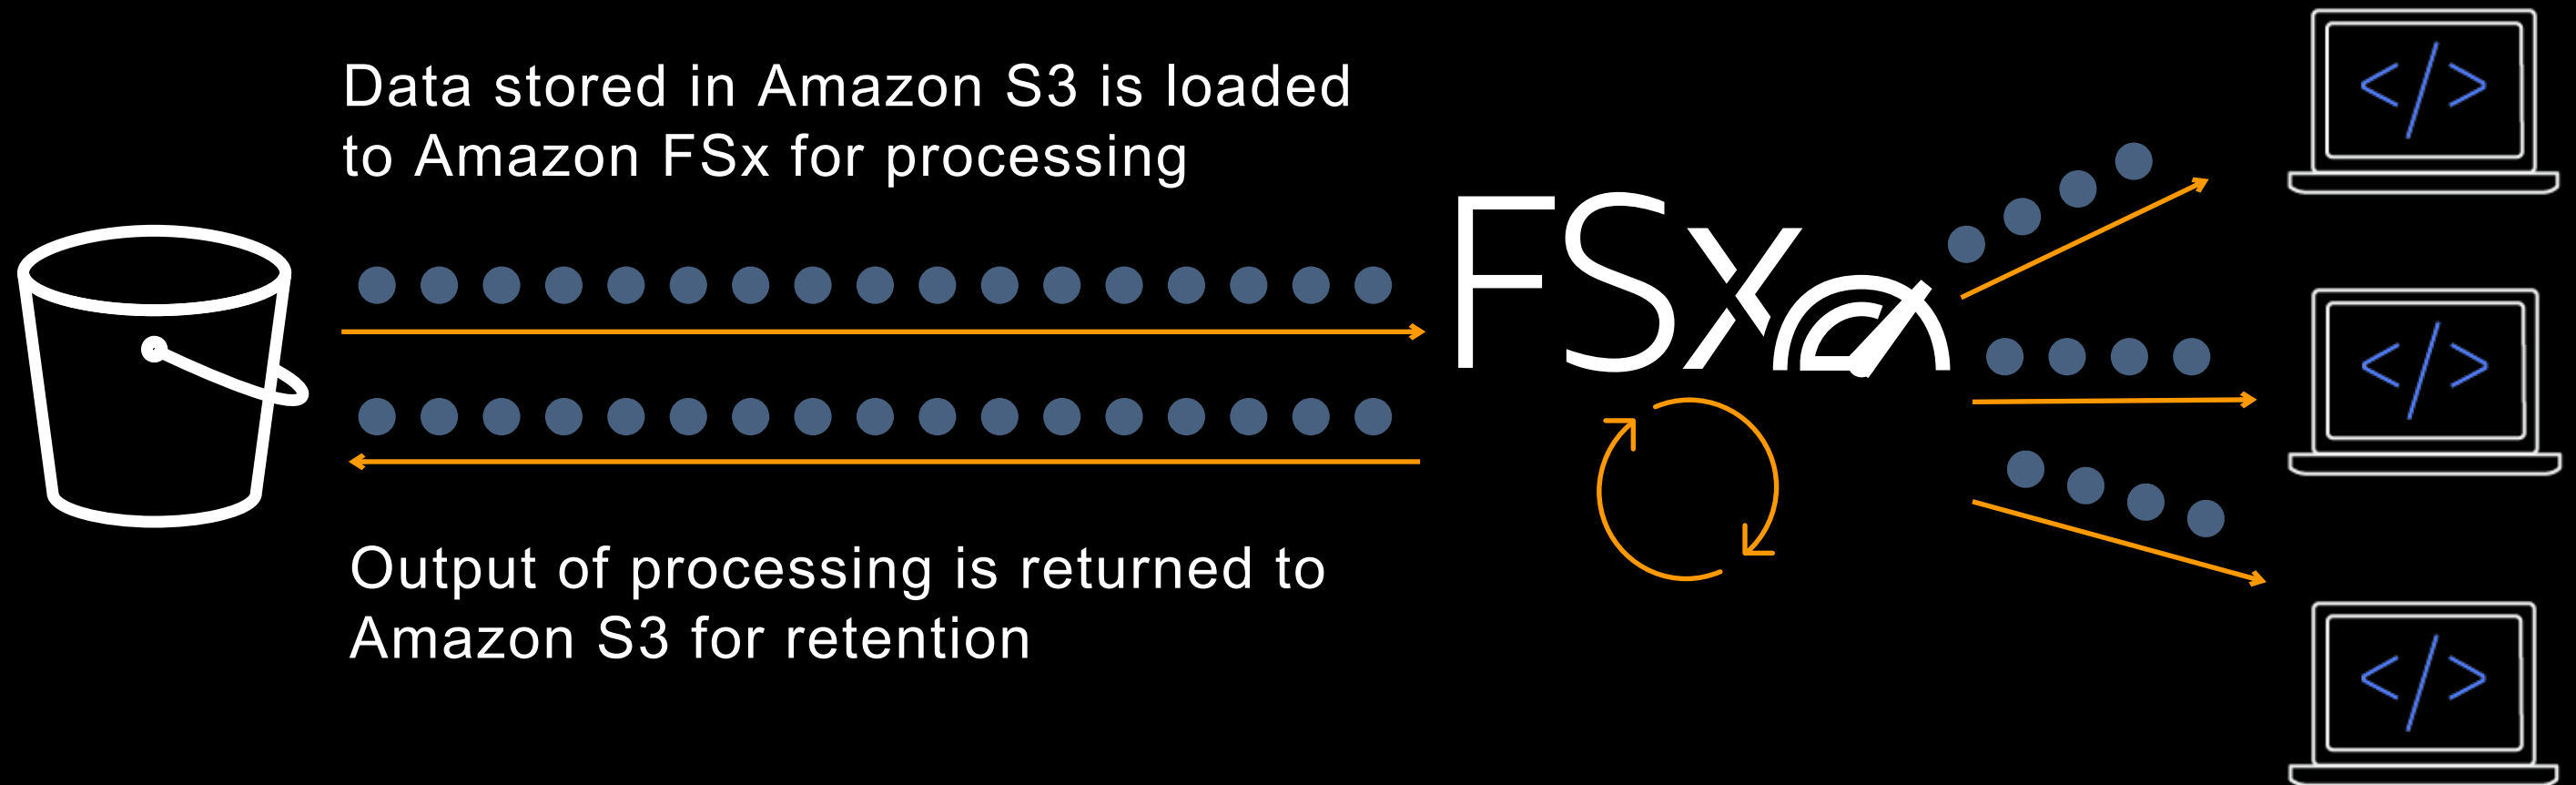


Secure
and compliant



Seamless integration with Amazon S3

Link your Amazon S3 dataset to your Amazon FSx for Lustre file system, then...



When your workload finishes, simply delete your file system

Storing your dataset in Amazon FSx for Lustre

<https://aws.amazon.com/blogs/machine-learning/speed-up-training-on-amazon-sagemaker-using-amazon-efs-or-amazon-fsx-for-lustre-file-systems/>

- Best option for **high-performance**, **low-latency** training
- Create an FSx file system, link it to your S3 bucket, train
- Delete the file system when you're done
- Training instances must run in a **VPC**, and open **port 998** (Lustre)
- **FileSystemDataSource**
 - Filesystem id: provided by FSx
 - Type: 'FSxLustre'
 - Directory path
 - Access type: Read-only or read-write

```
estimator = TensorFlow(
    entry_point='tensorflow_mnist/mnist.py',
    role='SageMakerRole',
    train_instance_count=1,
    train_instance_type='ml.c4.xlarge',
    subnets=['subnet-1', 'subnet-2'],
    security_group_ids=['sg-1'])

file_system_input = FileSystemInput(
    file_system_id='fs-2',
    file_system_type='FSxLustre',
    directory_path='/fsx/tensorflow',
    file_system_access_mode='ro')

estimator.fit(inputs=file_system_input)
```



Training



INNOVATE
ONLINE CONFERENCE

MACHINE LEARNING
AND AI EDITORIAL

Distributed Training

- Natively available for most built-in algorithms
- Natively available for TensorFlow, Apache MXNet, etc.
- You need to implement it yourself if you use a custom container
- Zoom on TensorFlow: Two modes available
 - **Parameter Server**
 - Asynchronous gradient averaging and weight distribution
 - All instances talk to each other: Networking can become a bottleneck and slow down training
 - **Horovod**
 - Based on Ring-AllReduce algorithm
 - More efficient communication helps scale near-linearly to 256 GPUs
 - <https://github.com/aws-samples/sagemaker-horovod-distributed-training>



Demo:


A quick look at TensorFlow with Horovod

https://github.com/aws-labs/amazon-sagemaker-examples/tree/master/sagemaker-python-sdk/tensorflow_script_mode_horovod




Amazon EC2 P3dn


<https://aws.amazon.com/blogs/aws/new-ec2-p3dn-gpu-instances-with-100-gbps-networking-local-nvme-storage-for-faster-machine-learning-p3-price-reduction/>



Reduce machine
learning training time



Better GPU
utilization



Support larger, more
complex models

100Gbps of
networking
bandwidth

KEY
FEATURES

8 NVIDIA Tesla V100 GPUs	32GB of memory per GPU (2x more)
-----------------------------	---

96 Intel Skylake vCPUs (50% more than P3) with AVX-512



Managed Spot Training

<https://aws.amazon.com/blogs/aws/managed-spot-training-save-up-to-90-on-your-amazon-sagemaker-training-jobs/>

- Save up to **90%** on training costs
- Fully managed: Obtain spot instances, start training, handle interruptions
- Implement **checkpointing** to resume interrupted jobs
 - Available in built-in algorithms for computer vision
 - Default behavior in TensorFlow
 - If checkpointing is not implemented, the training job is restarted from **scratch**
 - You get billed for data download only once
- CreateTrainingJob
 - **EnableManagedSpotTraining** = true
 - MaxWaitTimeInSeconds**
= MaxRuntimeInSeconds
+ time waiting for spot instances

```
estimator = TensorFlow(
    entry_point='script.py',
    role='SageMakerRole',
    train_instance_count=1,
    train_instance_type='ml.p3.2xlarge',
    train_use_spot_instances=True,
    train_max_run=3600,
    train_max_wait=7200)

estimator.fit(inputs)
```



Demo:

Fashion-MNIST classification with Keras/TensorFlow

- + Script Mode
- + Managed Spot Training
- + Elastic Inference

<https://aws.amazon.com/blogs/machine-learning/train-and-deploy-keras-models-with-tensorflow-and-apache-mxnet-on-amazon-sagemaker/>

<https://gitlab.com/juliensimon/dlnotebooks/tree/master/keras/05-keras-blog-post>



Tips to speed up training

- Scale out with **distributed training**
- Pick the best format for your dataset
 - Use **protobuf** instead of CSV or JSON
 - <https://github.com/aws/sagemaker-python-sdk/blob/master/src/sagemaker/amazon/common.py>
 - Pack samples into **record-based files**
 - TFRecord (TensorFlow) or RecordIO (MXNet)
 - Splitting in 100MB files looks like the sweet spot
 - Protobuf-encoded + RecordIO ❤️
- Amazon S3: Use **Pipe mode** for large datasets **NEW!**
- Monitor **CPU/GPU usage and network throughput** in Amazon CloudWatch

Model tuning

The never-ending quest for hyperparameters

XGBoost

Tree depth
Max leaf nodes
Gamma
Eta
Lambda
Alpha
...

Neural networks

Number of layers
Hidden layer width
Learning rate
Embedding dimensions
Dropout
...

Finding the optimal set of hyperparameters

<https://aws.amazon.com/blogs/machine-learning/amazon-sagemaker-automatic-model-tuning-now-supports-random-search-and-hyperparameter-scaling/>

1. **Manual search:** “I know what I’m doing”
2. **Grid search:** “X marks the spot”
 - Typically training hundreds of models
 - Slow and expensive
3. **Random search:** “Spray and pray”
 - Works better and faster than Grid Search
 - But... but... but... it’s random!
4. **Hyperparameter optimization (HPO):** Use machine learning
 - Requires fewer training jobs
 - Gaussian Process Regression and Bayesian Optimization



Demo:

HPO with Keras

<https://gitlab.com/juliensimon/dlnotebooks/tree/master/keras/04-fashion-mnist-sagemaker-advanced>




Deployment



INNOVATE
ONLINE CONFERENCE

MACHINE LEARNING
AND AI EDITORIAL

Predictions drive
complexity and
cost in production



Model optimization is extremely complex

mxnet

TensorFlow

PYTORCH

intel

nvidia

arm

Other
architectures

aws

Amazon SageMaker Neo

<https://aws.amazon.com/blogs/aws/amazon-sagemaker-neo-train-your-machine-learning-models-once-run-them-anywhere/>




Get accuracy
and performance



Automatic
optimization



Broad framework
support



Broad hardware
support

Integrated with
Amazon EC2
and Amazon
SageMaker

**KEY
FEATURES**
Open-source runtime and
compiler; 1/10th the size of
original frameworks

<https://github.com/neo-ai>



Compiling ResNet-50 for the Raspberry Pi

Configure the compilation job

```
{
  "RoleArn": $ROLE_ARN,
  "InputConfig": {
    "S3Uri": "s3://jsimon-neo/model.tar.gz",
    "DataInputConfig": "{\"data\": [1, 3, 224, 224]}",
    "Framework": "MXNET"
  },
  "OutputConfig": {
    "S3OutputLocation": "s3://jsimon-neo/",
    "TargetDevice": "rasp3b"
  },
  "StoppingCondition": {
    "MaxRuntimeInSeconds": 300
  }
}
```

Compile the model

```
$ aws sagemaker create-compilation-job
--cli-input-json file://config.json
--compilation-job-name resnet50-mxnet-pi
```

```
$ aws s3 cp s3://jsimon-neo/model-
rasp3b.tar.gz .
```

```
$ gtar tfz model-rasp3b.tar.gz
compiled.params
compiled_model.json
compiled.so
```

Predict with the compiled model


```
from dlr import DLModel
model = DLModel('resnet50', input_shape,
output_shape, device)
out = model.run(input_data)
```




Right-sizing your inference infrastructure

- **Statistical ML models, small DL models, dev/test**
 - CPU instances (C5) deliver the best cost/performance ratio
- **Very large DL models**
 - GPU instances (P2 or P3) should work best, especially if you need high throughput
 - If not, C5n could be a reasonable alternative
- But **what about everything in between?**
 - Mid-sized models
 - NLP models
 - Low-throughput, low-latency workloads
 - « Too slow on CPU, not cost-effective on GPU » ?

Are you making the most of your GPU infrastructure?



Low utilization and high costs




One size does not fit all

Amazon Elastic Inference


<https://aws.amazon.com/blogs/aws/amazon-elastic-inference-gpu-powered-deep-learning-inference-acceleration/>



Reduce GPU inference costs up to 75%



Match capacity to demand



Available between 1 and 32 TFLOPs per accelerator

KEY FEATURES

Integrated with Amazon EC2 and Amazon SageMaker

Support for TensorFlow and Apache MXNet

Single and mixed-precision operations



Demo:

Fashion-MNIST classification with Keras/TensorFlow

- + Script Mode
- + Managed Spot Training
- + Elastic Inference

<https://aws.amazon.com/blogs/machine-learning/train-and-deploy-keras-models-with-tensorflow-and-apache-mxnet-on-amazon-sagemaker/>

<https://gitlab.com/juliensimon/dlnotebooks/tree/master/keras/05-keras-blog-post>



Getting started

<http://aws.amazon.com/free>

<https://ml.aws>

<https://aws.amazon.com/sagemaker>

<https://github.com/aws/sagemaker-python-sdk>

<https://github.com/aws/sagemaker-spark>

<https://github.com/aws-labs/amazon-sagemaker-examples>

<https://gitlab.com/juliensimon/dlnotebooks>

Thank you!

Julien Simon
Global Evangelist, AI & Machine Learning
Amazon Web Services

@julsimon

<https://medium.com/@julsimon>



INNOVATE
ONLINE CONFERENCE

MACHINE LEARNING
AND AI EDITORIAL