DEV DAY

DEV DAY

MLT4

An Introduction to Generative Adversarial Networks

Julien Simon
Global Evangelist, AI & Machine Learning
Amazon Web Services

@julsimon

What are Generative Adversarial Networks?

- Most Machine Learning problems try to make sense of a real-life dataset
  - Tabular data: regression, classification, clustering
  - Text: sentiment analysis, entity extraction
  - Image: object detection, image segmentation

- Generative Adversarial Networks are fundamentally different
  - Goodfellow, 2014 https://arxiv.org/abs/1406.2661

- A GAN uses a real-life dataset to generate new lookalike samples
  - Not based on statistical analysis, e.g. filling in missing values during feature engineering

DEV DAY

A Light Introduction to GAN Theory

Adversarial Networks

- Two deep learning networks compete against one another

- The Generator network creates new samples from random data, not from the data set
  - Start from a random vector (say, 100 bytes) and generate a tensor with the same shape as dataset samples (say, 3x256x256)

- The Discriminator learns how to classify samples
  - Real ones, coming from the data set
  - Fake ones, generated by the Generator
  - Typical Deep Learning process: mini-batch training, backpropagation, etc.

« How do you create something from nothing? »

« How can fake samples look remotely convincing? »

« How does the Generator learn? »

Let’s take about convolution for a minute

- Convolution extracts features from complex data samples
- It reduces the total amount of information
- Samples can shrink, keeping only the « good stuff »
  - Padding and stride
- So we know how to go from (3, 256, 256) to [0,1]
  - From images to probabilities
- Can we do the opposite?

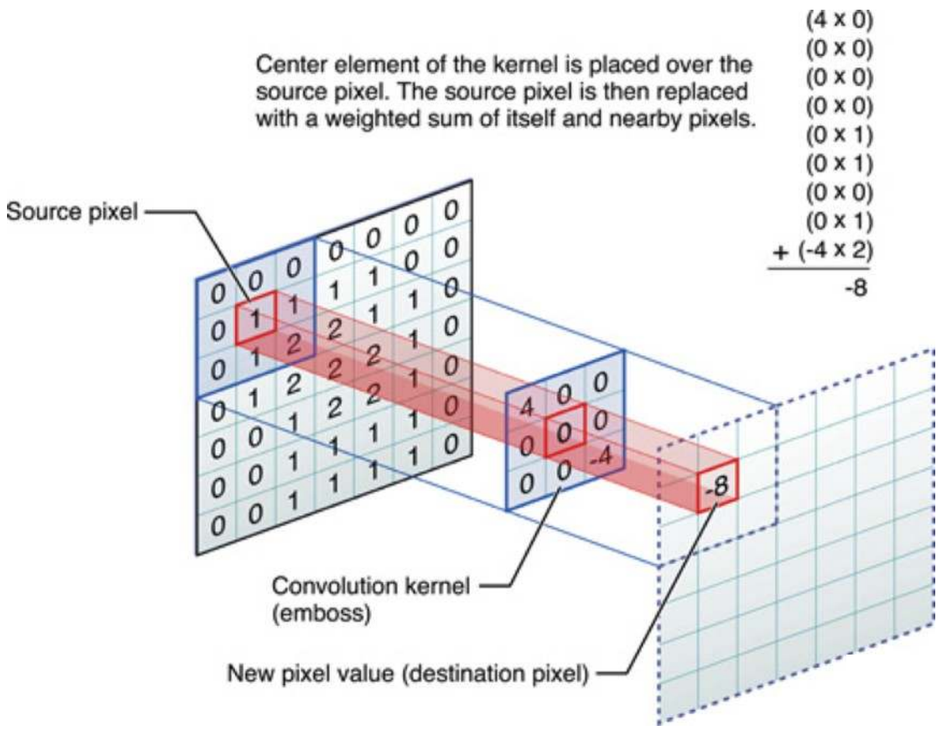

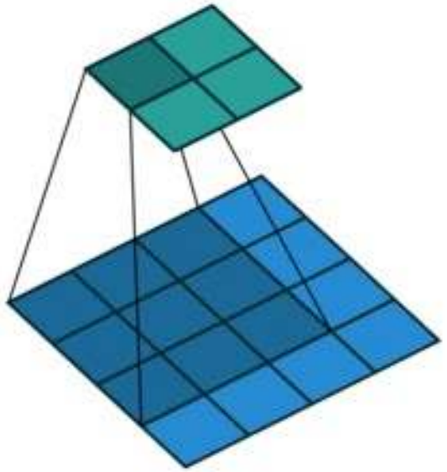

Transposed convolution

- Transposed convolution is the reverse operation
- It increases the total amount of information
- Samples grow
- So we know how to go from (100,) to (3, 256, 256)

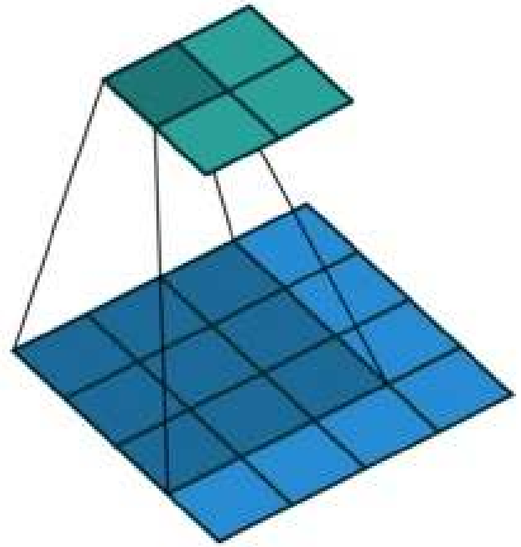

The Deep Convolutional Generator (DCGAN)

https://arxiv.org/abs/1511.06434

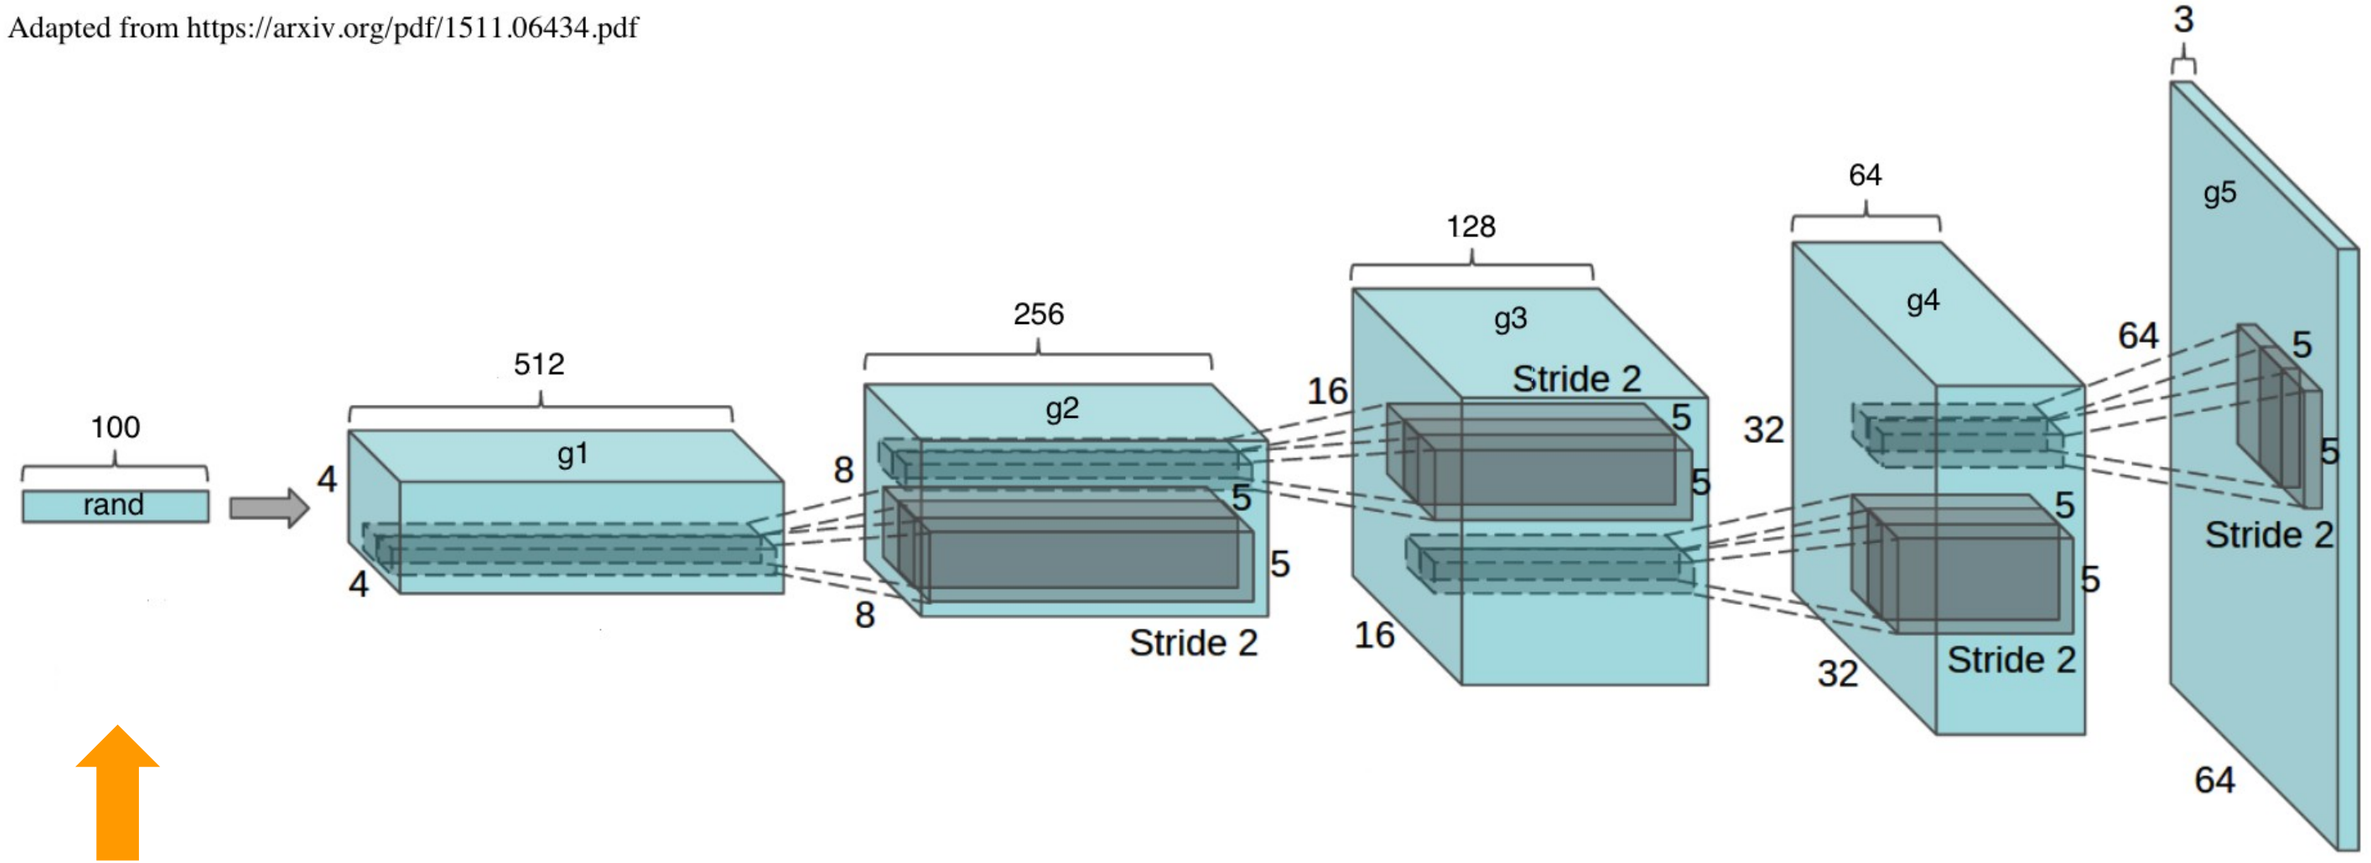

The Deep Convolutional Discriminator (DCGAN)

https://arxiv.org/abs/1511.06434

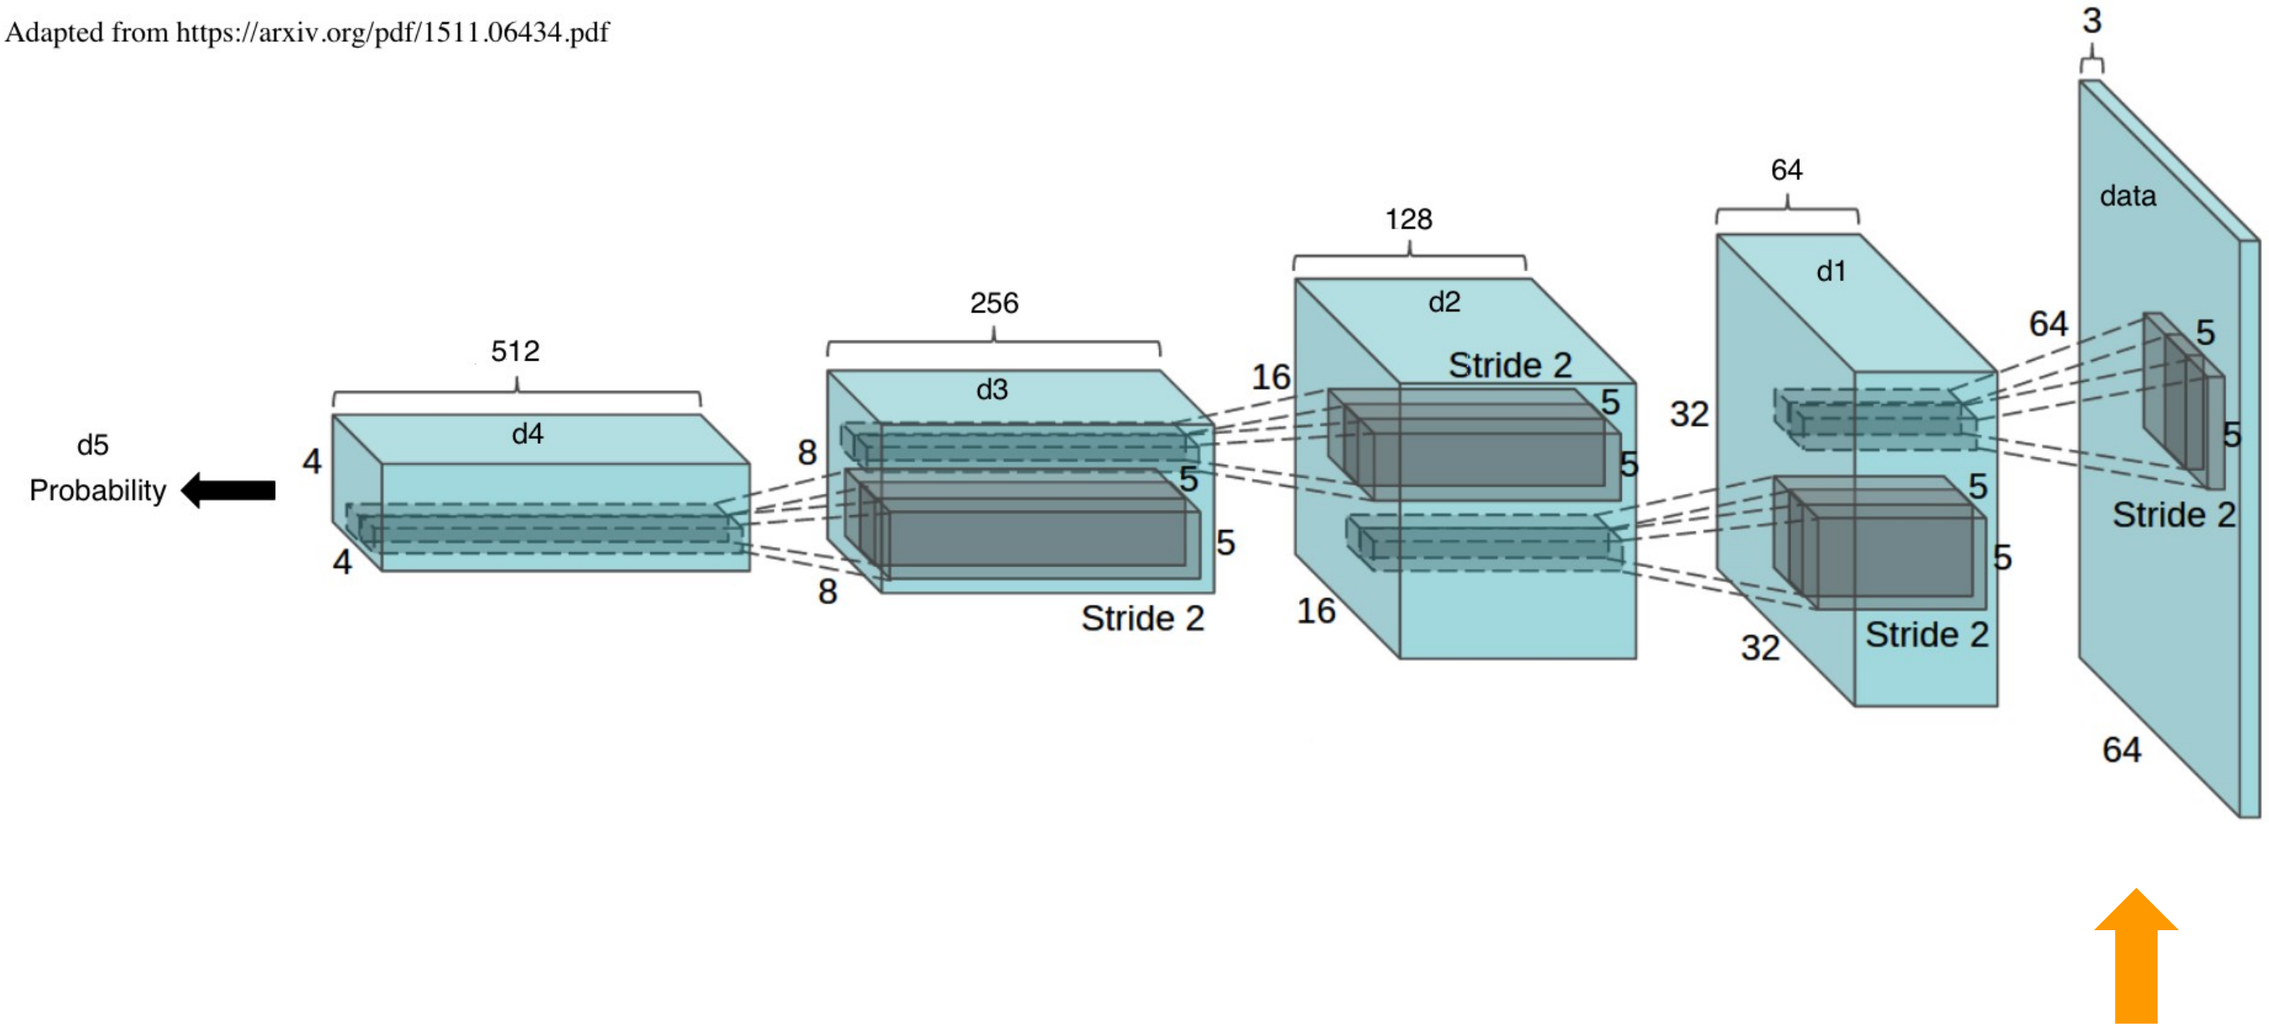

Training DCGAN

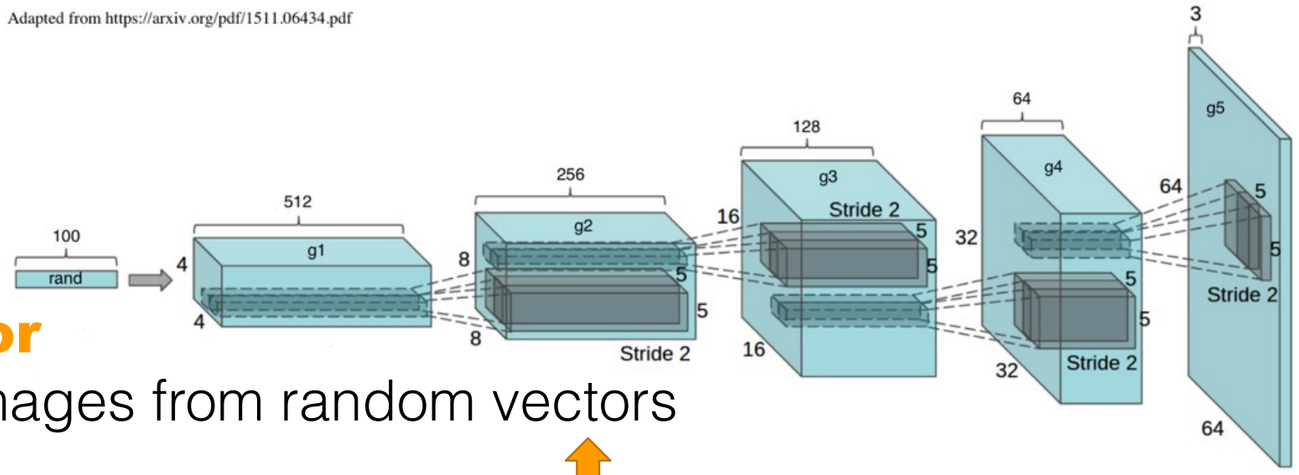

Generator

Building images from random vectors

Gradient updates

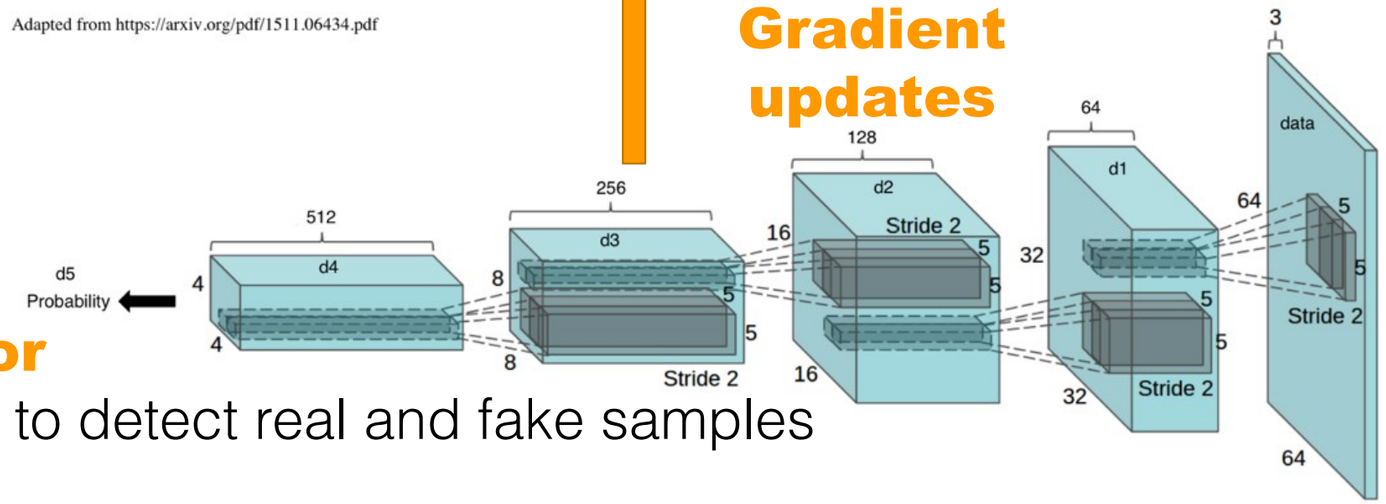

Detector

Learning to detect real and fake samples

Fake images

Real images

https://medium.com/@julsimon/generative-adversarial-networks-on-apache-mxnet-part-1-b6d39e6b5df1

DEV DAY

Demo:
training DCGAN with MXNet Gluon

https://gitlab.com/juliensimon/dlnotebooks/-/blob/master/mxnet/07%20-%20DCGAN%20Gluon.ipynb

DEV DAY

Sample Projects

GAN: Welcome to the (un)real world, Neo

TF

Generating new ”celebrity” faces https://github.com/tkarras/progressive_growing_of_gans

April 2018

PyTorch

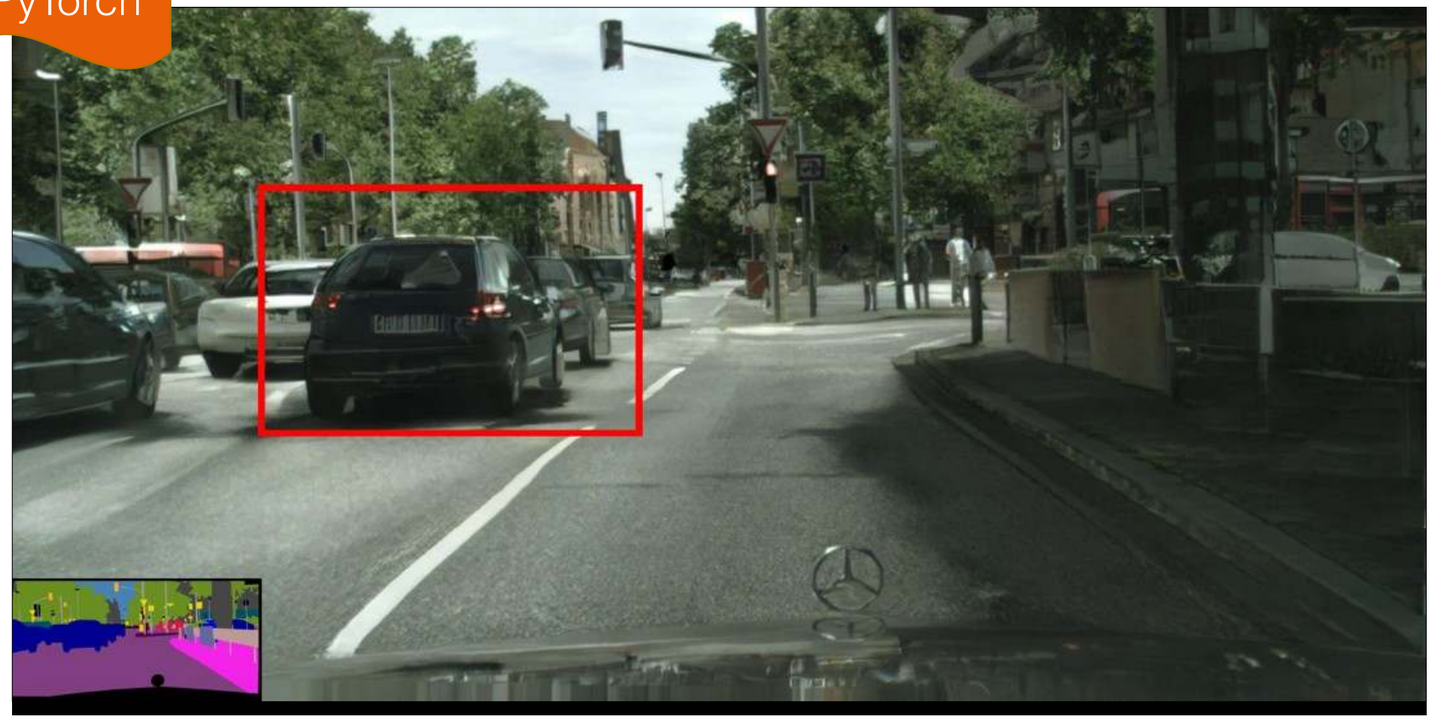

From semantic map to 2048x1024 picture

https://tcwang0509.github.io/pix2pixHD/

November 2017

More face generation with GANs

Controlled Image Generation with TL-GAN
https://github.com/SummitKwan/transparent_latent_gan
October 2018

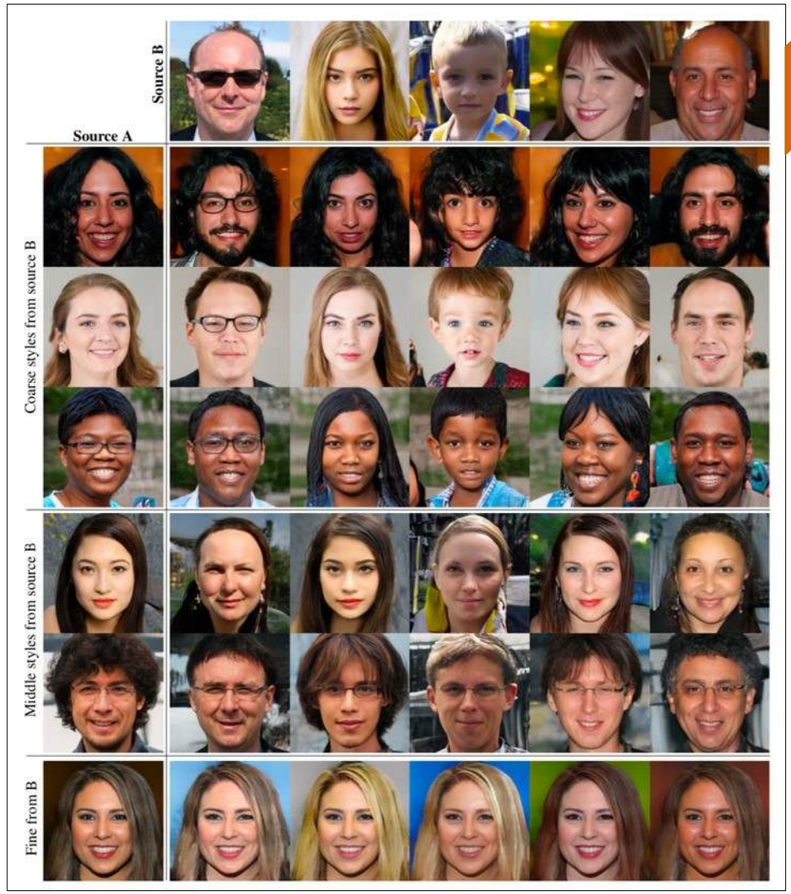

TF

Applying the style of a face to another face https://github.com/NVlabs/stylegan
https://www.youtube.com/watch?v=kSLJriaOumA
March 2019

GAN: Everybody dance now

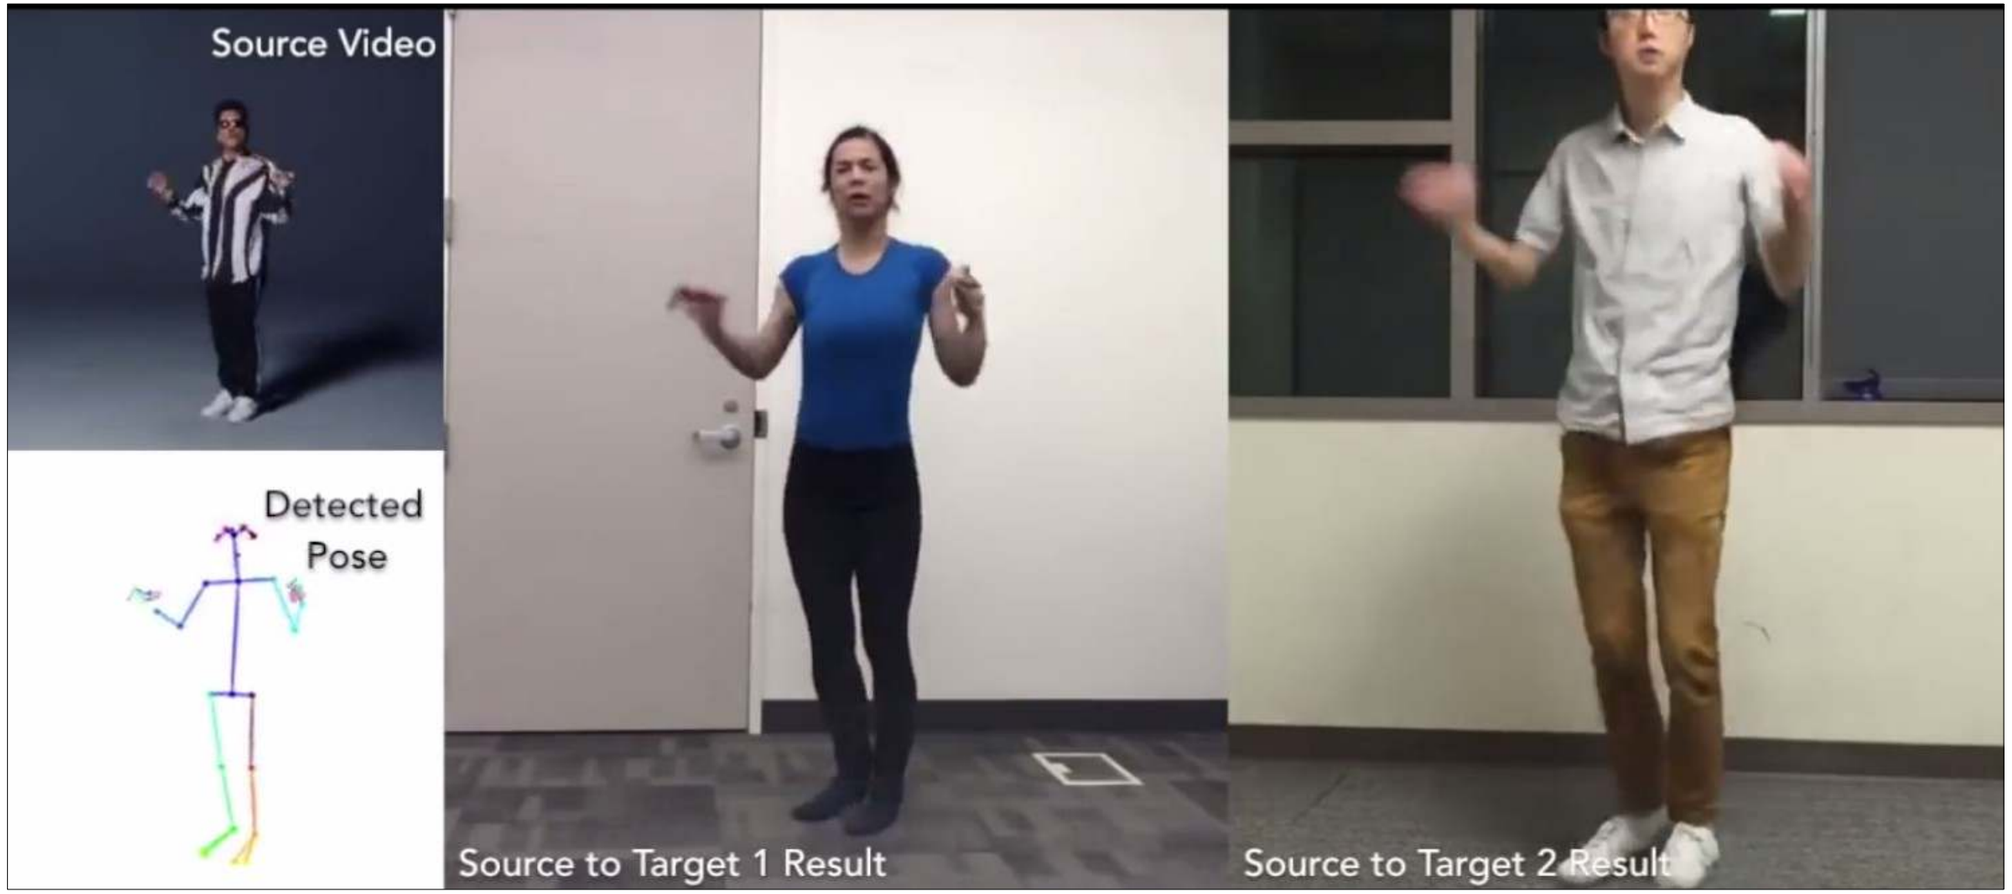

https://arxiv.org/abs/1808.07371
https://www.youtube.com/watch?v=PCBTZh41Ris
August 2018

Autodesk - Generative Design

https://www.youtube.com/watch?v=A31A8KDC9S4

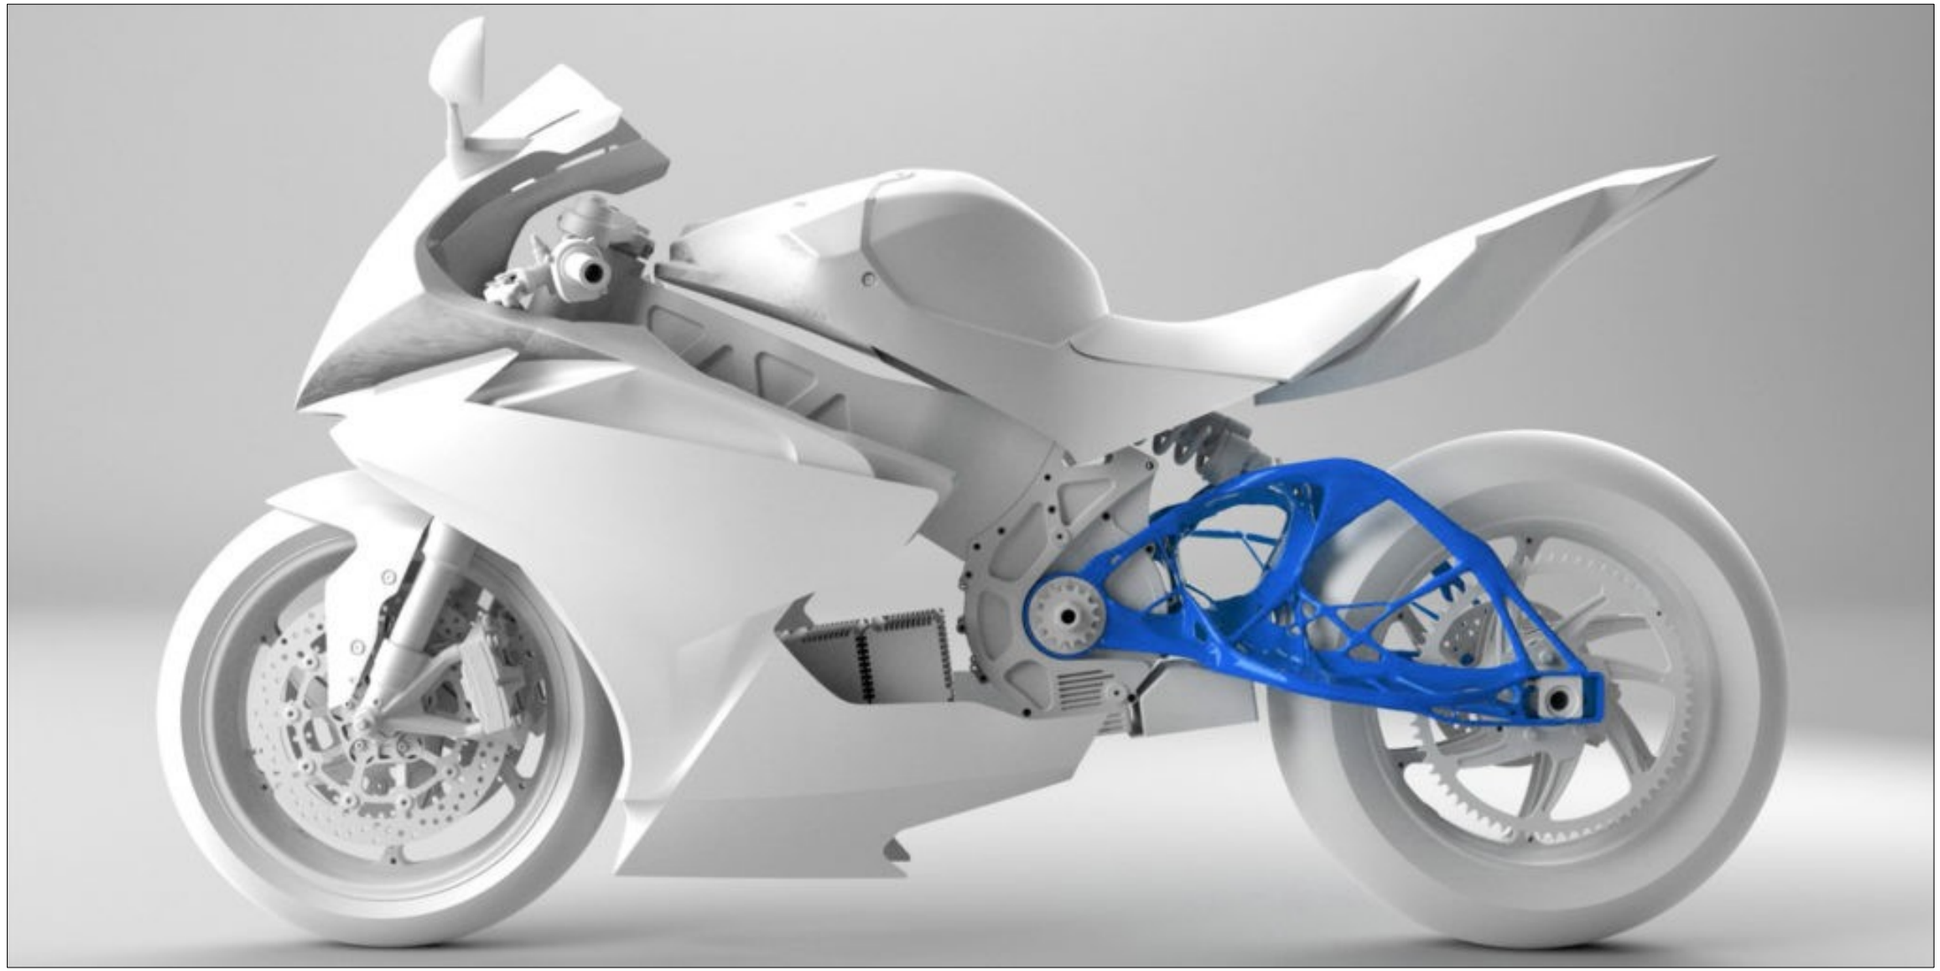

AWS DeepComposer

https://aws.amazon.com/deepcomposer
https://github.com/aws-samples/aws-deepcomposer-samples

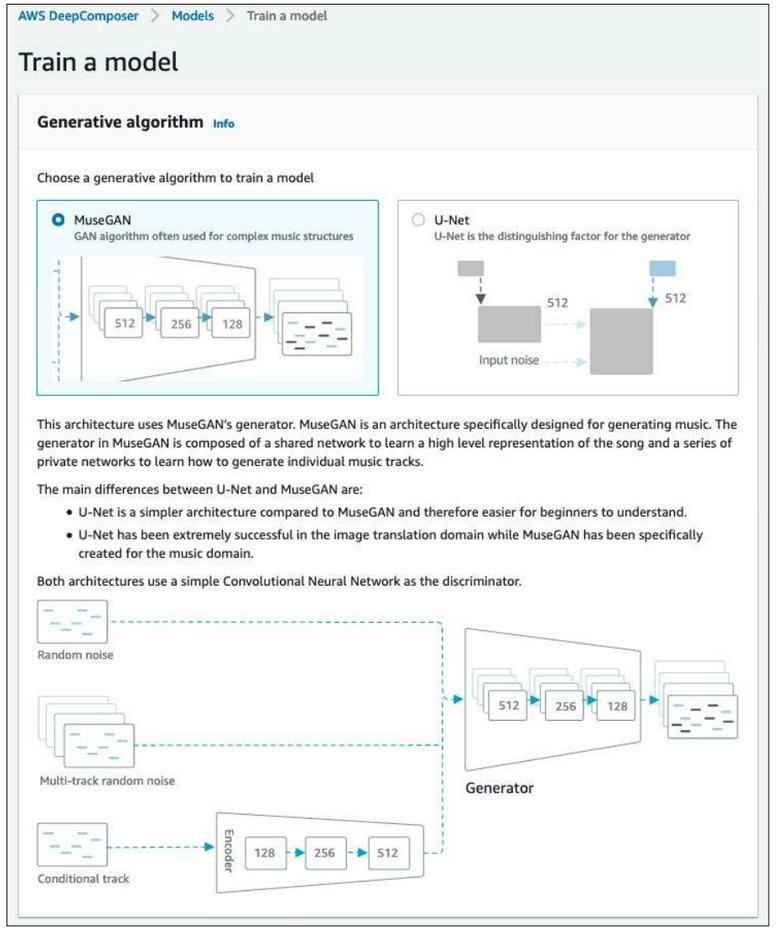

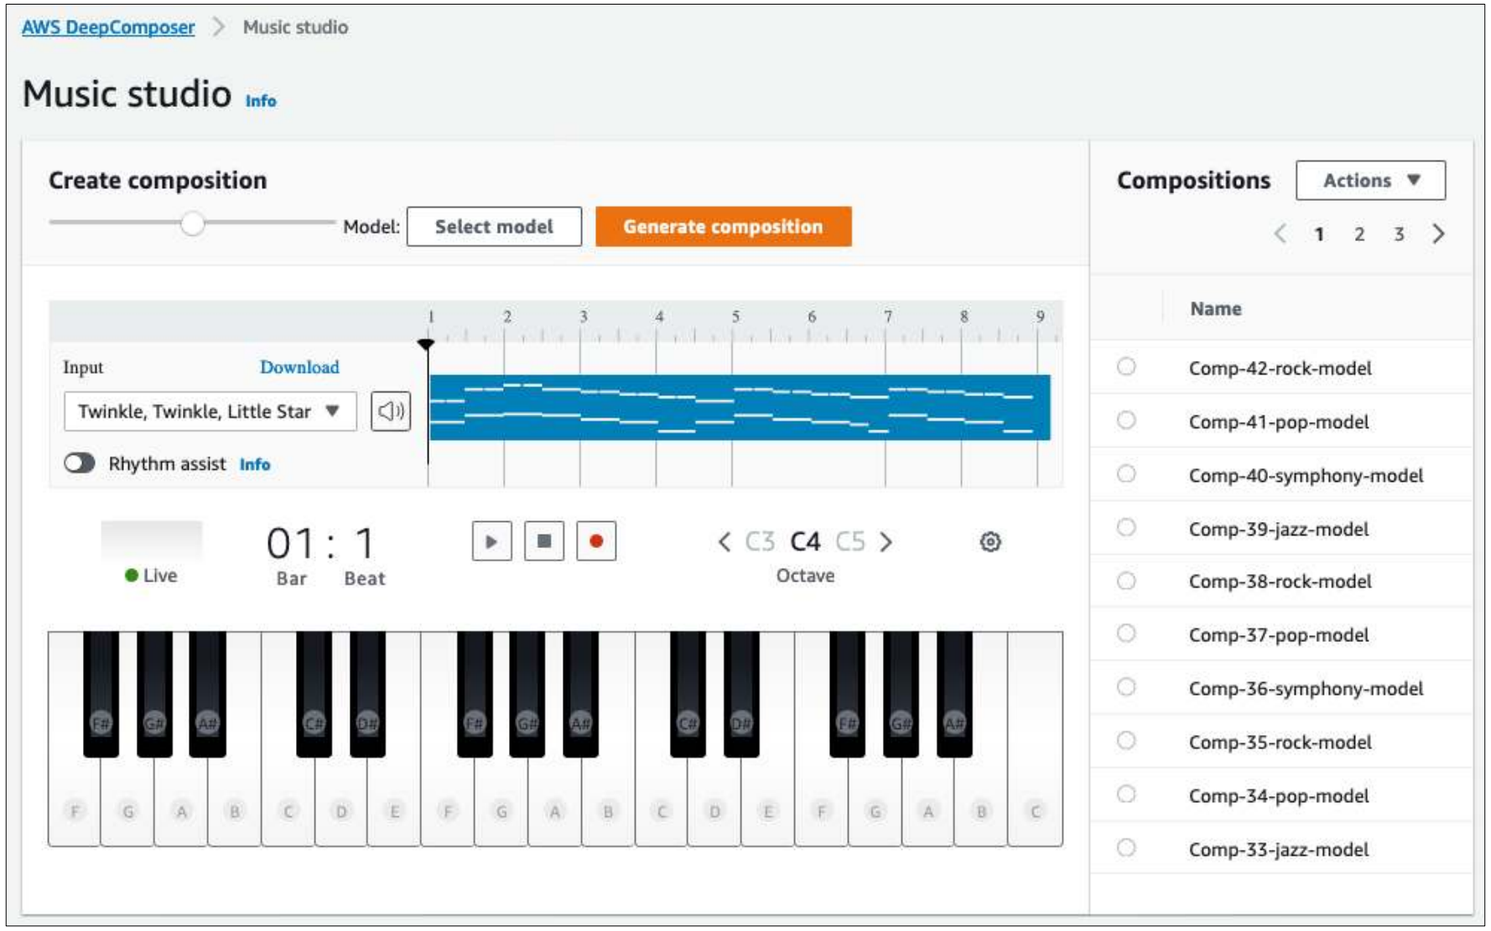

DEV DAY

Demo

DEV DAY

Getting started

Resources

https://aws.training/machinelearning
https://deeplearning.ai
https://fast.ai

http://www.deeplearningbook.org/
https://d2l.ai/

https://gitlab.com/juliensimon/dlnotebooks

DEV DAY

Thank you!