

A Pragmatic Introduction to Natural Language Processing models

Julien Simon
Global Evangelist, AI & Machine Learning

@julsimon

<https://medium.com/@julsimon>






Problem statement

- Since the dawn of AI, NLP has been a **major** research field
 - Text classification, machine translation, text generation, chat bots, voice assistants, etc.
- NLP apps require a **language model** in order to predict the next word
 - Given a sequence of words (w_1, \dots, w_n) , predict w_{n+1} that has the highest probability
- Vocabulary size can be **hundreds of thousands** of words
... in **millions of documents**
- Can we build a **compact mathematical representation** of language, that would help us with a variety of downstream NLP tasks?

« *You shall know a word by the company it keeps* », Firth (1957)

- Initial idea: build **word vectors** from co-occurrence counts
 - Also called **embeddings**
 - High dimensional: at least 50, up to 300
 - Much more compact than an $n*n$ matrix, which would be extremely sparse
- Words with **similar meanings** should have **similar vectors**
 - "car" \approx "automobile" \approx "sedan"
- The **distance** between related vectors should be similar
 - distance("Paris", "France") \approx distance("Berlin", "Germany")
 - distance("hot", "hotter") \approx distance("cold", "colder")



Source: GloVe

High-level view

1. Start from a **large text corpus** (100s of millions of words, even billions)
 2. Preprocess the corpus
 - **Tokenize**: « hello, world! » → « <BOS>hello<SP>world<SP>!<EOS> »
 - **Multi-word entities**: « Rio de Janeiro » → « rio_de_janeiro »
 3. Build the **vocabulary**
 4. Learn **vector representations** for all words
- ... or simply use (or fine-tune) **pre-trained vectors** (more on this later)



Algorithms

Word2Vec (2013)

<https://arxiv.org/abs/1301.3781>

<https://code.google.com/archive/p/word2vec/>

- **Continuous bag of words (CBOW):**
 - the model **predicts the current word** from surrounding context words
 - Word **order** doesn't matter (hence the 'bag of words')
- **Skipgram**
 - the model **uses the current word to predict** the surrounding window of context words
 - This may work better when little data is available
- CBOW trains faster, skipgram is more accurate
- C code, based on shallow neural network







Global Vectors aka GloVe (2014)

<https://nlp.stanford.edu/projects/glove/>

<https://github.com/stanfordnlp/GloVe>

<https://www.quora.com/How-is-GloVe-different-from-word2vec>

- Performance generally similar to Word2Vec
- Several pre-trained embedding collections:
up to 840 billion tokens, 2.2 million vocabulary, 300 dimensions
- C code, based on **matrix factorization**

nearest neighbors of <i>frog</i>	Litoria	Leptodactylidae	Rana	Eleutherodactylus
Pictures				

Source: GloVe

FastText (2016)

<https://arxiv.org/abs/1607.04606> + <https://arxiv.org/abs/1802.06893>

<https://fasttext.cc/>

<https://www.quora.com/What-is-the-main-difference-between-word2vec-and-fastText>

- Extension of Word2Vec: each word is a **set of subwords**, aka **n-grams**
 - « Computer », n=5 : <START>Comp , compu, omput, mpute, puter, uter<END>
 - A word vector is the **average** of its **subword vectors**
- Good for **unknown/mispelled words**, as they share **subwords** with known words
 - « Computerized » and « Cmputer » should be close to « Computer »
- Three modes
 - Unsupervised learning: **compute embeddings** (294 languages available)
 - Supervised learning: **use / fine-tune embeddings** for multi-label, multi-class classification
 - Also language detection for 170 languages
- Multithreaded C++ code, with Python API

BlazingText (2017)

<https://dl.acm.org/citation.cfm?id=3146354>

<https://aws.amazon.com/blogs/machine-learning/enhanced-text-classification-and-word-vectors-using-amazon-sagemaker-blazingtext/>

- Amazon-invented algorithm, available in Amazon SageMaker
- Extends FastText with GPU capabilities
- Unsupervised learning: word vectors
 - 20x faster
 - CBOW and skip-gram with subword support
 - Batch skip-gram for distributed training
- Supervised learning: text classification
 - 100x faster
 - Models are compatible with FastText

	Word2Vec (unsupervised learning)			Text Classification (supervised learning)
Modes	Skip-gram (supports subwords)	CBOW (supports subwords)	batch_skipgram	supervised
Single CPU instance	✓	✓	✓	✓
Single GPU instance (with 1 or more GPUs)	✓	✓		✓*
Multiple CPU instances			✓	

Limitations of Word2Vec (and family)

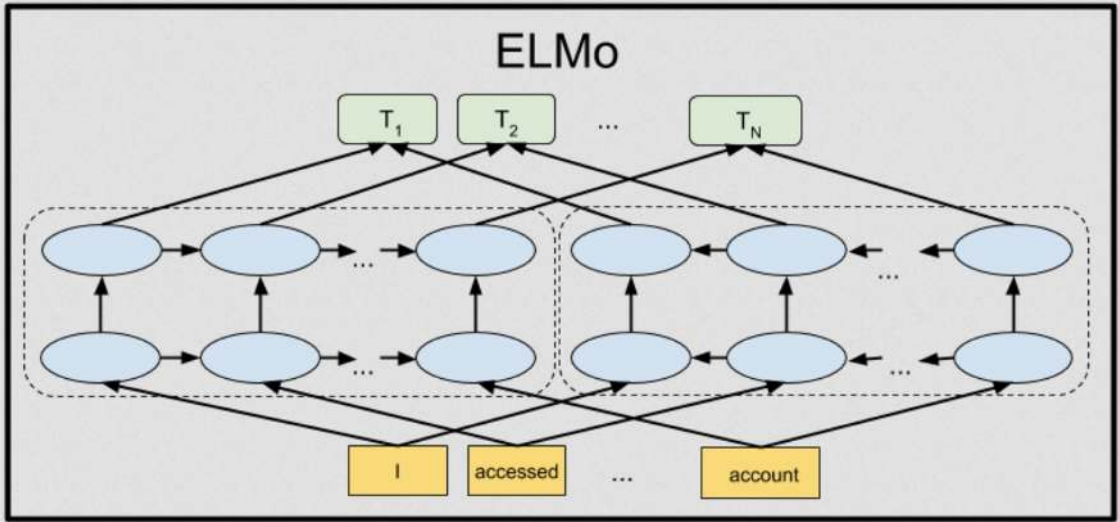
- Some words have different meanings (aka **polysemy**)
 - « *Kevin, stop throwing **rocks!*** » vs. « *Machine Learning **rocks*** »
 - Word2Vec encodes the **different meanings** of a word into the **same vector**
- Bidirectional context is not taken into account
 - Previous words (**left-to-right**) and next words (**right-to-left**)

Embeddings from Language Models aka ELMo (02/2018)



<https://arxiv.org/abs/1802.05365>
<https://allennlp.org/elmo>

- ELMo generates a context-aware vector for each word
 - Character-level CNN + two unidirectional LSTMs
 - Bidirectional context, no cheating possible (can't peek at the next word)
 - “Deep” embeddings, reflecting output from all layers
- No vocabulary, no vector file:
you need to use the model itself
- Reference implementation with TensorFlow



Source: Google

Source		Nearest Neighbors
GloVe	play	playing, game, games, played, players, plays, player, Play, football, multiplayer
	Chico Ruiz made a spectacular <u>play</u> on Alusik 's grounder {...}	Kieffer , the only junior in the group , was commended for his ability to hit in the clutch , as well as his all-round excellent play .
biLM	Olivia De Havilland signed to do a Broadway <u>play</u> for Garson {...}	{...} they were actors who had been handed fat roles in a successful <u>play</u> , and had talent enough to fill the roles competently , with nice understatement .

Source: ELMo paper

Bidirectional Encoder Representations from Transformers aka BERT (10/2018)


<https://arxiv.org/abs/1810.04805>

<https://github.com/google-research/bert>

<https://www.quora.com/What-are-the-main-differences-between-the-word-embeddings-of-ELMo-BERT-Word2vec-and-GloVe>



- BERT improves on ELMo
 - Replace LSTM with Transformers, which deal better with long-term dependencies
 - True bidirectional architecture: left-to-right and right-to-left contexts are learned by the same network
 - 15% of words are randomly masked during training to prevent cheating
- Pre-trained models: BERT Base and BERT Large
 - Masked word prediction
 - Next sentence prediction
- Reference implementation with TensorFlow



Source:
Google

Limitations of BERT

- BERT cannot handle more than 512 input tokens
- BERT masks words during training, but not during fine-tuning (aka training/fine-tuning discrepancy)
- BERT isn't trained to predict the next word, so it's not great at text generation
- BERT doesn't learn dependencies for masked words
 - Train « I am going to <MASK> my <MASK> » on « walk » / « dog », « eat » / « breakfast », and « debug » / « code ».
 - BERT could legitimately predict « I am going to eat my code » or « I am going to debug my dog » :-/

XLNet (06/2019)

<https://arxiv.org/abs/1906.08237>

<https://github.com/zihangdai/xlnet>

- XLNet beats BERT at 20 tasks
- XLNet uses **bidirectional context**, but words are **randomly permuted**
 - No cheating possible
 - No masking required
- **XLNet Base** and **XLNet Large**
- Reference implementation with TensorFlow



07/2019: ERNIE 2.0 (Baidu)

beats BERT & XLNet

<https://github.com/PaddlePaddle/ERNIE/>

Summing things up

- Word2Vec and friends
 - Try **pre-trained embeddings** first
 - Check that the training corpus is similar to your own data
 - Same language, similar vocabulary
 - Remember that subword models will help with unknown / misspelled words
 - If you have exotic requirements AND lots of data, training is not expensive
- EIMo, BERT, XLNet
 - Training is very **expensive**: several days using several GPUs
 - Fine-tuning is **cheap**: just a few GPU hours for SOTA results
 - Fine-tuning scripts and pre-trained models are available: start there!
- The "best model" is the one that works best on **your business problem**

Demos

Demo: finding similarities and analogies with Gluon NLP and pre-trained GloVe embeddings

<https://gitlab.com/juliensimon/dlnotebooks/gluonnlp/>

Demo: embeddings with ELMo on TensorFlow

<https://gitlab.com/juliensimon/dlnotebooks/blob/master/nlp/ELMO%20TensorFlow.ipynb>

Getting started on AWS

<https://ml.aws>

<https://aws.amazon.com/marketplace/solutions/machine-learning/natural-language-processing> → off the shelf algos and models that might just save you from this madness ;)

<https://aws.amazon.com/sagemaker>

<https://github.com/aws-labs/amazon-sagemaker-examples>

Thank you!

Julien Simon
Global Evangelist, AI & Machine Learning

@julsimon

<https://medium.com/@julsimon>

