



AWS re:Invent

AIM410-R

Deep learning applications with TensorFlow, featuring Mobileye

Julien Simon

Global Evangelist, AI/ML
Amazon Web Services
@julsimon

Chaim Rand

ML Algorithm Developer
Mobileye

Agenda

TensorFlow on AWS

Customer case study: Mobileye

Demo: TensorFlow on Amazon SageMaker

Getting started

TensorFlow


<https://www.tensorflow.org>



- Main API in **Python**, with support for Javascript, Java, C++
- TensorFlow 1.x: **symbolic execution**
 - ‘Define then run’: build a graph, optimize it, feed data, and compute
 - Low-level API: variables, placeholders, tensor operations
 - High-level API: *tf.estimator.**
 - Keras library: *Sequential* and *Functional* API, predefined layers
- TensorFlow 2.0: **imperative execution** (aka eager execution)
 - ‘Define by run’: normal Python code, similar to numpy
 - Run it, inspect it, debug it
 - Keras is the preferred API

AWS: The platform of choice for TensorFlow

<https://aws.amazon.com/tensorflow/>



89% of all deep learning workloads in the cloud run on AWS

85% of all TensorFlow workloads in the cloud run on AWS

Source: Nucleus Research, T147, October 2019

TensorFlow: a first-class citizen on Amazon SageMaker


- Built-in TensorFlow containers for **training** and **prediction**
 - Code available on Github: <https://github.com/aws/sagemaker-tensorflow-containers>
 - Build it, run it on your own machine, customize it, etc.
 - Versions : 1.4.1 → 1.15 (2.0 coming soon)
- Not just TensorFlow
 - **Standard tools**: TensorBoard, TensorFlow Serving
 - **SageMaker features**: Local Mode, Script Mode, Model Tuning, Spot Training, Pipe Mode, Amazon EFS & Amazon FSx for Lustre, Amazon Elastic Inference, etc.
 - **Performance optimizations**: GPUs and CPUs (AWS, Intel MKL-DNN library)
 - **Distributed training**: Parameter Server and Horovod

Amazon SageMaker


re:Invent 2019 announcements



First fully integrated
development
environment (IDE) for
machine learning
SageMaker Studio



Enhanced notebook experience
with quick-start &
easy collaboration
SageMaker Notebooks
(preview)



Automatic debugging,
analysis, and alerting
SageMaker
Debugger



Experiment management system to
organize, track, & compare
thousands of experiments
SageMaker Experiments



Model monitoring to detect
deviation in quality & take
corrective actions
SageMaker
Model Monitor



Automatic generation of
ML models with
full visibility & control
SageMaker
Autopilot

Amazon SageMaker at Mobileye

Chaim Rand

ML Algorithm Developer
Mobileye



An Intel
Company

Making Amazon SageMaker work for you

- Story of how we adopted Amazon SageMaker
- Ways in which Amazon SageMaker accelerated our development process
- Challenges we faced and how we overcame them



Spoiler:

Adopting Amazon SageMaker enabled us to reduce our development by up to 10X (from several months to under a week)

Chapter 1 - Introduction

A bit about Mobileye

Founded in 1999 by Prof. Amnon Shashua and Mr. Ziv Aviram

Goal: use computer vision-based technologies to

- revolutionize the transportation industry
- make roads safer
- and **save lives**

Acquired by Intel in 2017 for **\$15.3** billion



A bit about Mobileye

We develop a range of software products, deployed on a proprietary family of computer chips named **EyeQ**

Leading supplier of software that enables advanced driver-assistance systems (ADAS)—deployed in over **40 million** vehicles. (Includes adaptive cruise control, collision avoidance, lane departure warning, ...)

Over **25 automaker partners**, including most of the world's largest



Some of our technologies

Rely on **monocular camera perception**

- Reduces cost
- Other sensors can be used for redundancy when working on high level autonomous driving



Some of our Sensing Technologies

Use deep neural networks (DNNs) to detect:

Road users – including vehicles and pedestrians

Road semantics – including traffic lights, traffic signs, on-road arrows, stop-lines, and crosswalks


Road boundaries – any delimiter of the drivable area, its 3D structure and semantics, including curbs, cones and debris

Road geometry – driving paths and surface profile, including speed bumps and roadside ditches



Day (1/5/21)
egoSpeed: 8.86 yawRate: 0.00 expDTime: 0.03 egoAccel: -0.34 Day gfi: 1

mode=Day
distance



Scene Segmentation



From ADAS to autonomous

The key to full autonomy relies on three technological pillars



Sensing



Mapping



Driving policy

From ADAS to autonomous

The key to full autonomy relies on three technological pillars



Sensing



Mapping



Driving policy


Identify all artifacts in our surrounding environment

From ADAS to autonomous

The key to full autonomy relies on three technological pillars



Sensing



Mapping




Driving policy

Localize the vehicle on a map of the surrounding environment

- Up to $<10\text{cm}$ precision

From ADAS to autonomous

The key to full autonomy relies on three technological pillars



Sensing



Mapping



Driving policy

Decide what actions to take based on the input from Sensing and Mapping

Autonomous
ACTIVE

12
km/h

23° 



Autonomous
ACTIVE

12
km/h

23° 



From ADAS to autonomous

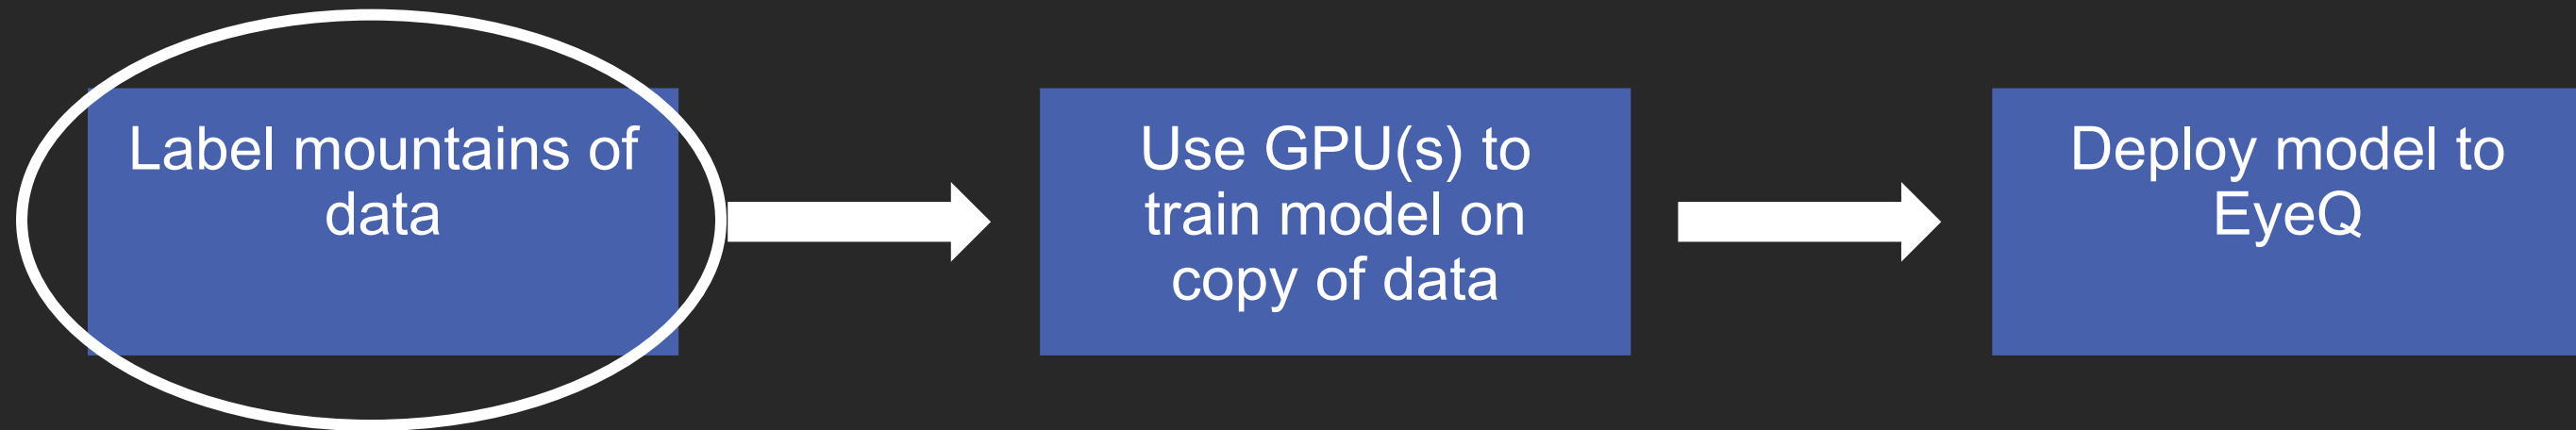
Learn more:

- From Mobileye VP Tal Babaioff
 - **AUT307 - Navigating the winding road toward driverless mobility**
- Online <https://www.mobileye.com>

Chapter 2 – Enter Amazon SageMaker


DNN training challenge

- Variety of different DNN architectures
- Mountains of data: a typical model may train on **up to 200TB** of data
- Simplified development cycle:



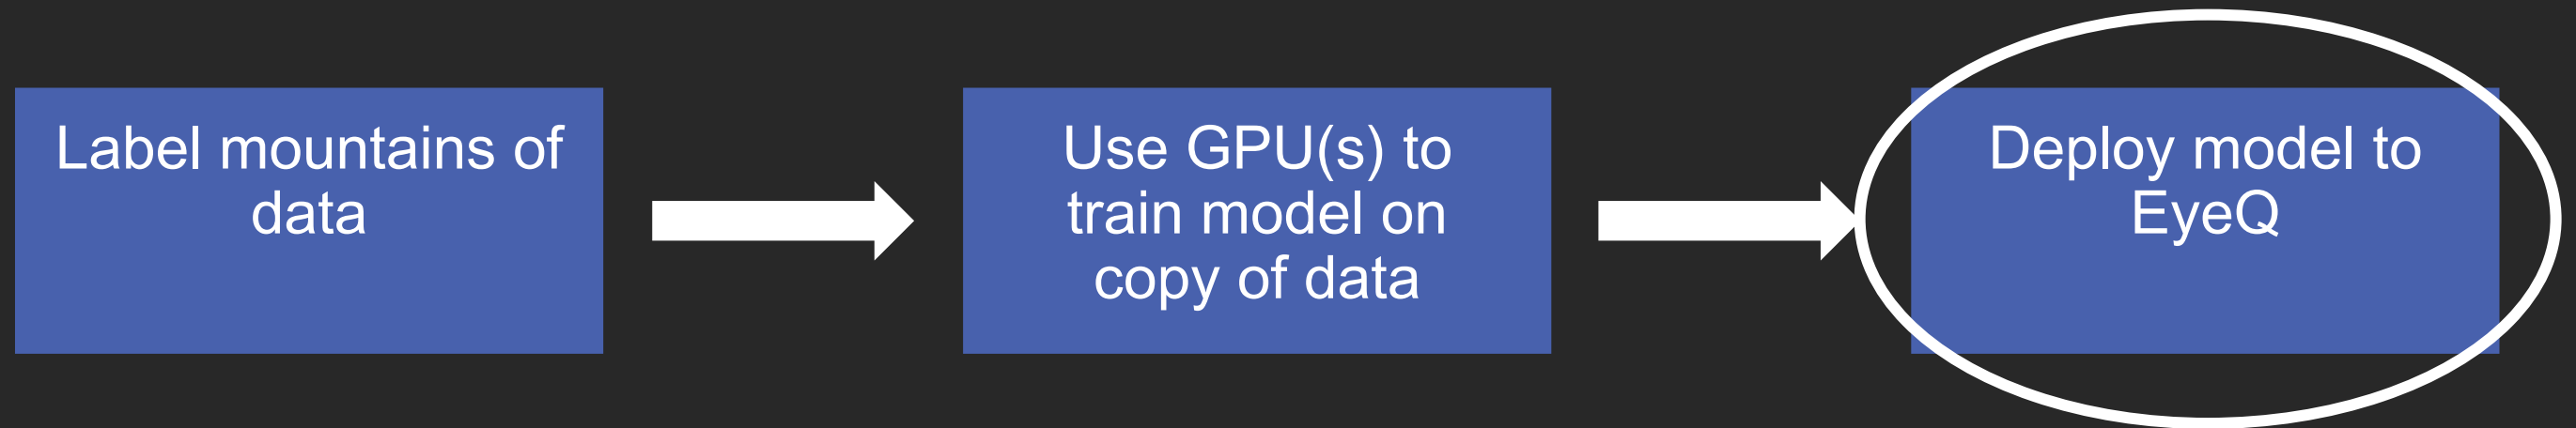
DNN training challenge

- Variety of different DNN architectures
- Mountains of data: a typical model may train on **up to 200TB** of data
- Simplified development cycle:



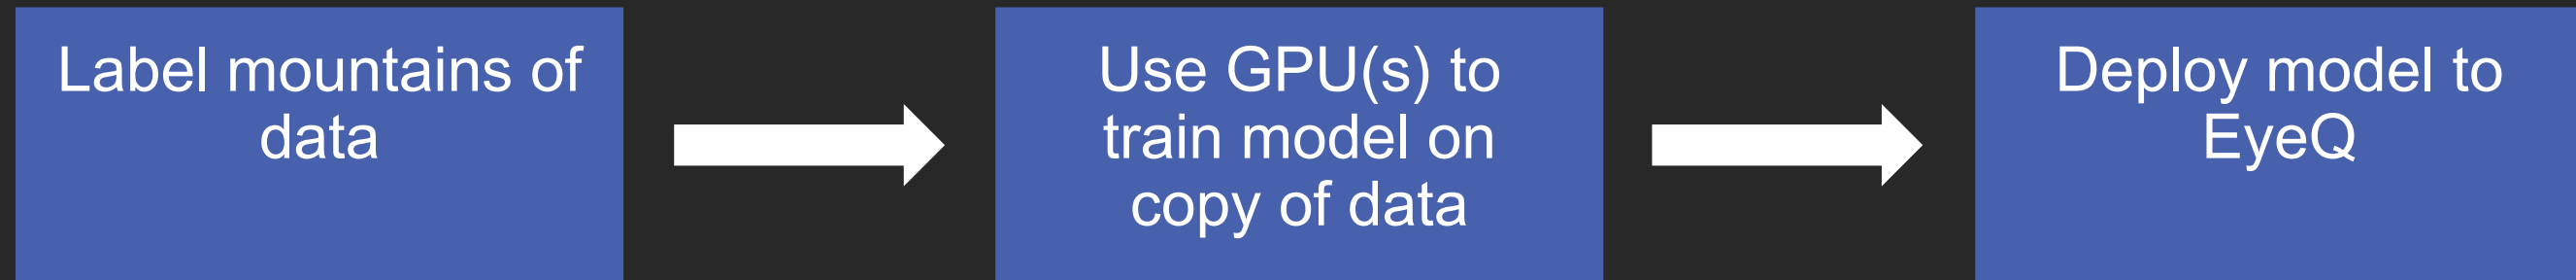
DNN training challenge

- Variety of different DNN architectures
- Mountains of data: a typical model may train on **up to 200TB** of data
- Simplified development cycle:



DNN training challenge

- Variety of different DNN architectures
- Mountains of data: a typical model may train on **up to 200TB** of data
- Simplified development cycle:



- Historically performed on premise

Training on premises

- Obvious **drawbacks** to training on premises
 - Limit to number of GPU instances
 - Limitations to data capacity
 - Challenge of staying up to date with latest HW and SW
- Any alternative must comply with our **development pipeline**, and must keep our IP **safe**
- Enter Amazon SageMaker ...


What I like about Amazon SageMaker

- **Unconstrained capacity** – means I can spin up as many training sessions as I need
- **Instance type variety**
 - Multi-generation CPU and GPU support
 - Up to date with latest HW and SW (tuned to maximize efficiency)
 - Distributed training support
- **Learning curve**
 - Well documented with many samples
- **Secure** – means data security team will let me use it
- **Pipe Mode** support


Chapter 3 – Amazon SageMaker Pipe Mode

What is Amazon SageMaker Pipe Mode?

- Enables **streaming** training data from Amazon S3 to training instances
 - Instead of copying data to each training device




Full distribution



Pipe Mode

Benefits of Pipe Mode

- Enables **streaming** training data from Amazon S3 to training instances
- Removes limitation on dataset size
- No need to have dedicated (and costly) storage for each training instance or to download data to each instance, and no delay to training start
- Enables **decoupling** of data from training instances
- Multiple training instances can all pull from the same dataset in S3



Pipe Mode with TensorFlow

- Lest you should fear the need to have to manage the data stream on your own...
- Amazon SageMaker wraps the pipe with an **implementation** of the *tf.data.Datasets* API
 - Complete with all the standard dataset functions
 - Ready to be fed directly into your model

Setting up Pipe Mode

```
from sagemaker.tensorflow import TensorFlow

tensorflow = TensorFlow(
    entry_point='pipemode.py',
    input_mode='Pipe', ...)

pipes = {'train': 's3://sagemaker-path-to-data'}

tensorflow.fit(pipes)
```

PipeModeDataset initialization

```
def parse(record):  
    feature = {'label': tf.FixedLenSequenceFeature([], tf.int64, allow_missing=True),  
               'image_raw': tf.FixedLenFeature([], tf.string)}  
    features = tf.parse_single_example(record, feature)  
    image = tf.decode_raw(features['image_raw'], tf.uint8)  
    label = features['label']  
    return {"image": image}, label # This is what will be fed into your model  
  
ds = PipeModeDataset("train", record_format='TFRecord')  
  
ds = ds.apply(map_and_batch(parse, batch_size=32, num_parallel_batches=2))  
  
return ds
```


Chapter 4 – Pipe Mode Challenges

Pipe Mode challenges

- Converting data to supported format
- Sequential nature of pipe mode
- Pipe number limitation
 - Currently stands at 20

Pipe Mode challenges

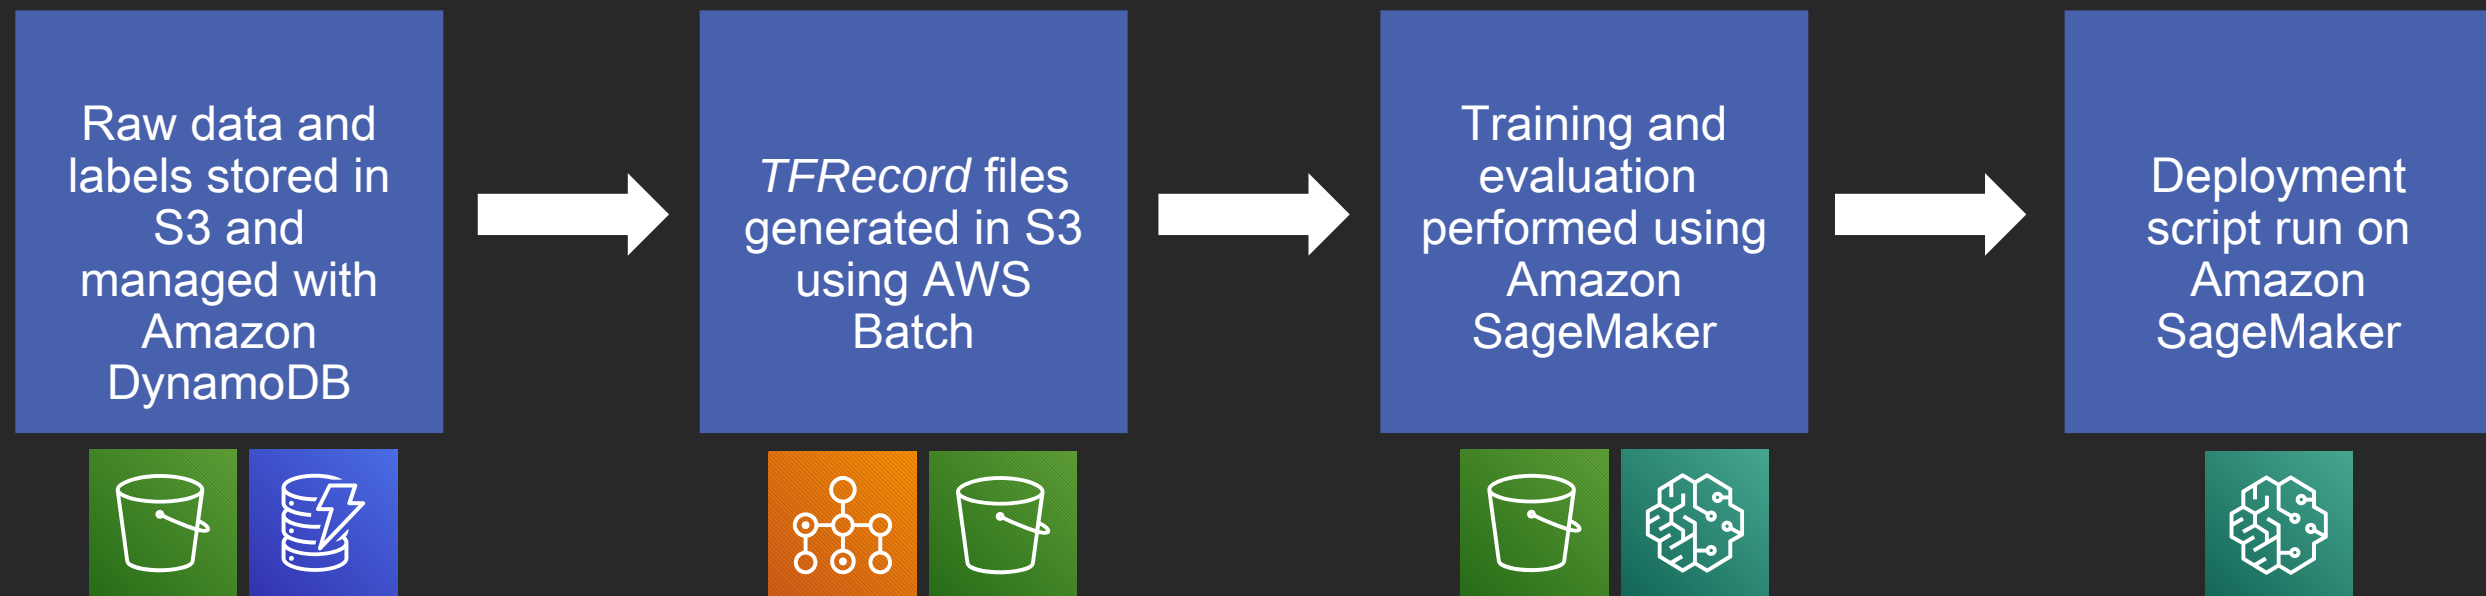
- Converting data to supported format
 - Amazon SageMaker's *PipeModeDataset* supports a limited number of formats
 - Format conversion of “mountains” of data to *TFRecord* format seemed daunting

Generating data in *TFRecord* format

- Turned out to be a **blessing** in disguise
 - Task was performed using AWS Batch
 - Used up to hundreds of thousands of vCPUs in parallel. (Each created a 100MB TFRecord file.)
 - Tip: split generated *TFRecord* files into 100MB chunks
 - Accelerated data preparation time significantly (from days to hours)

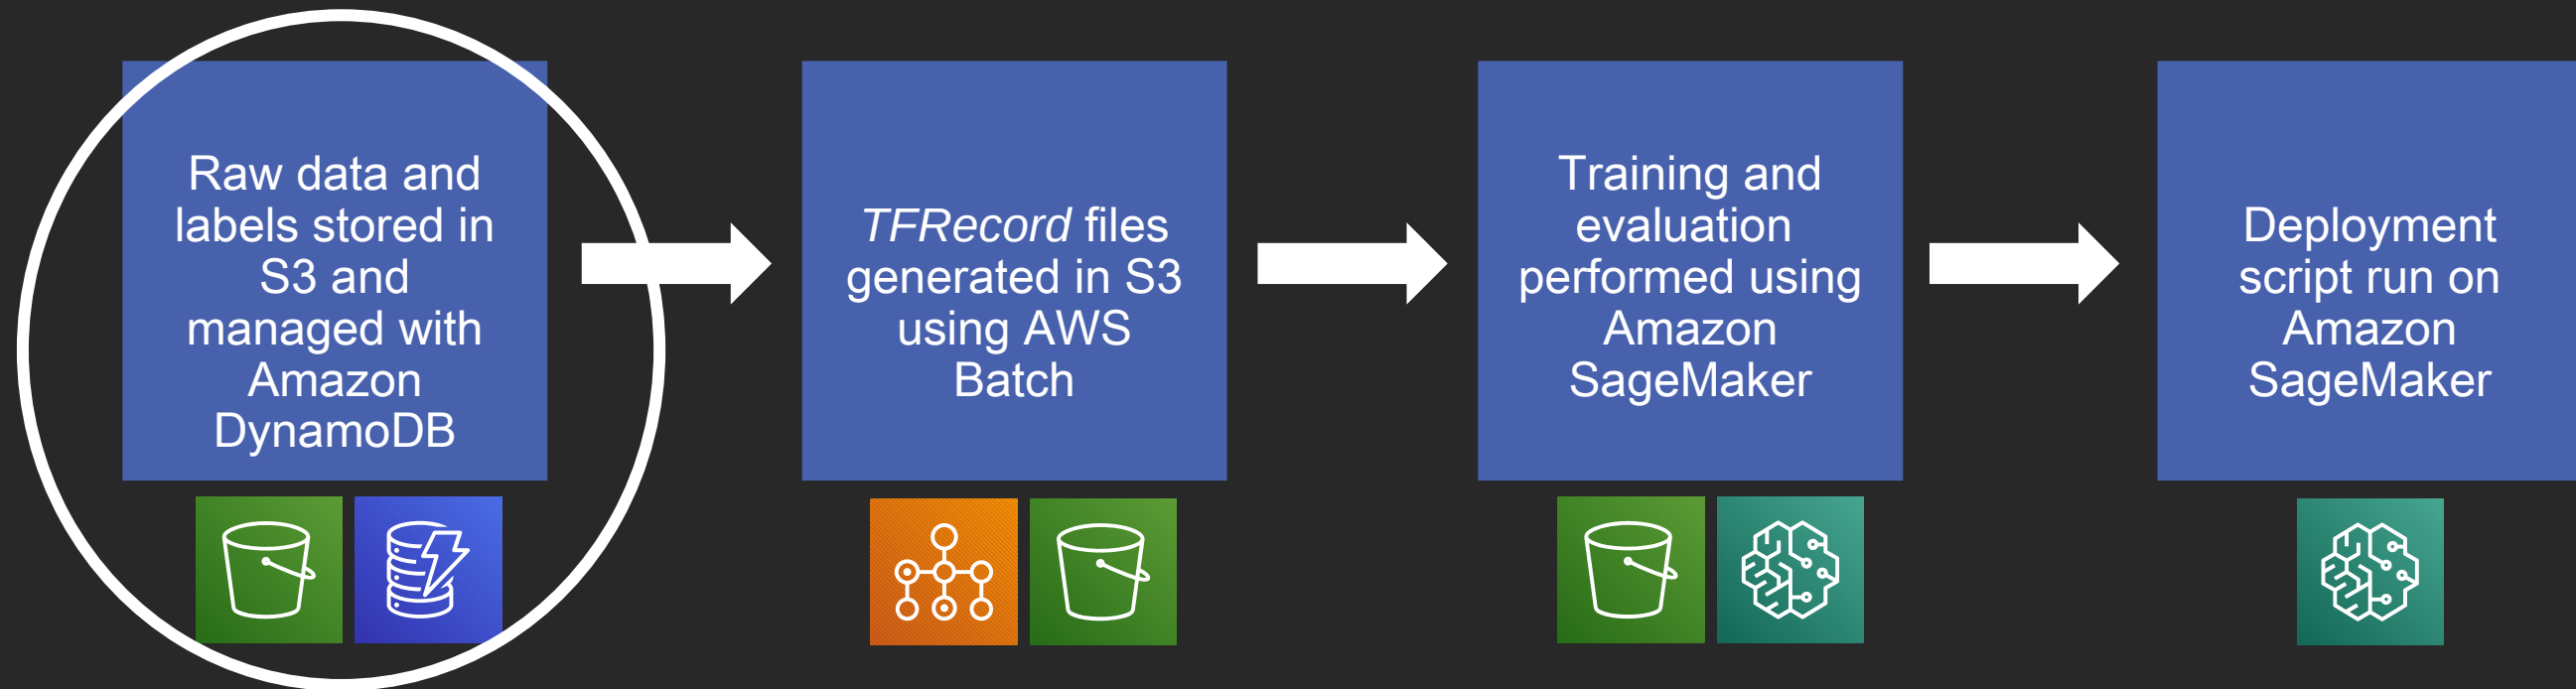
Generating data in *TfRecord* format

- One more step to moving end-to-end development to cloud:



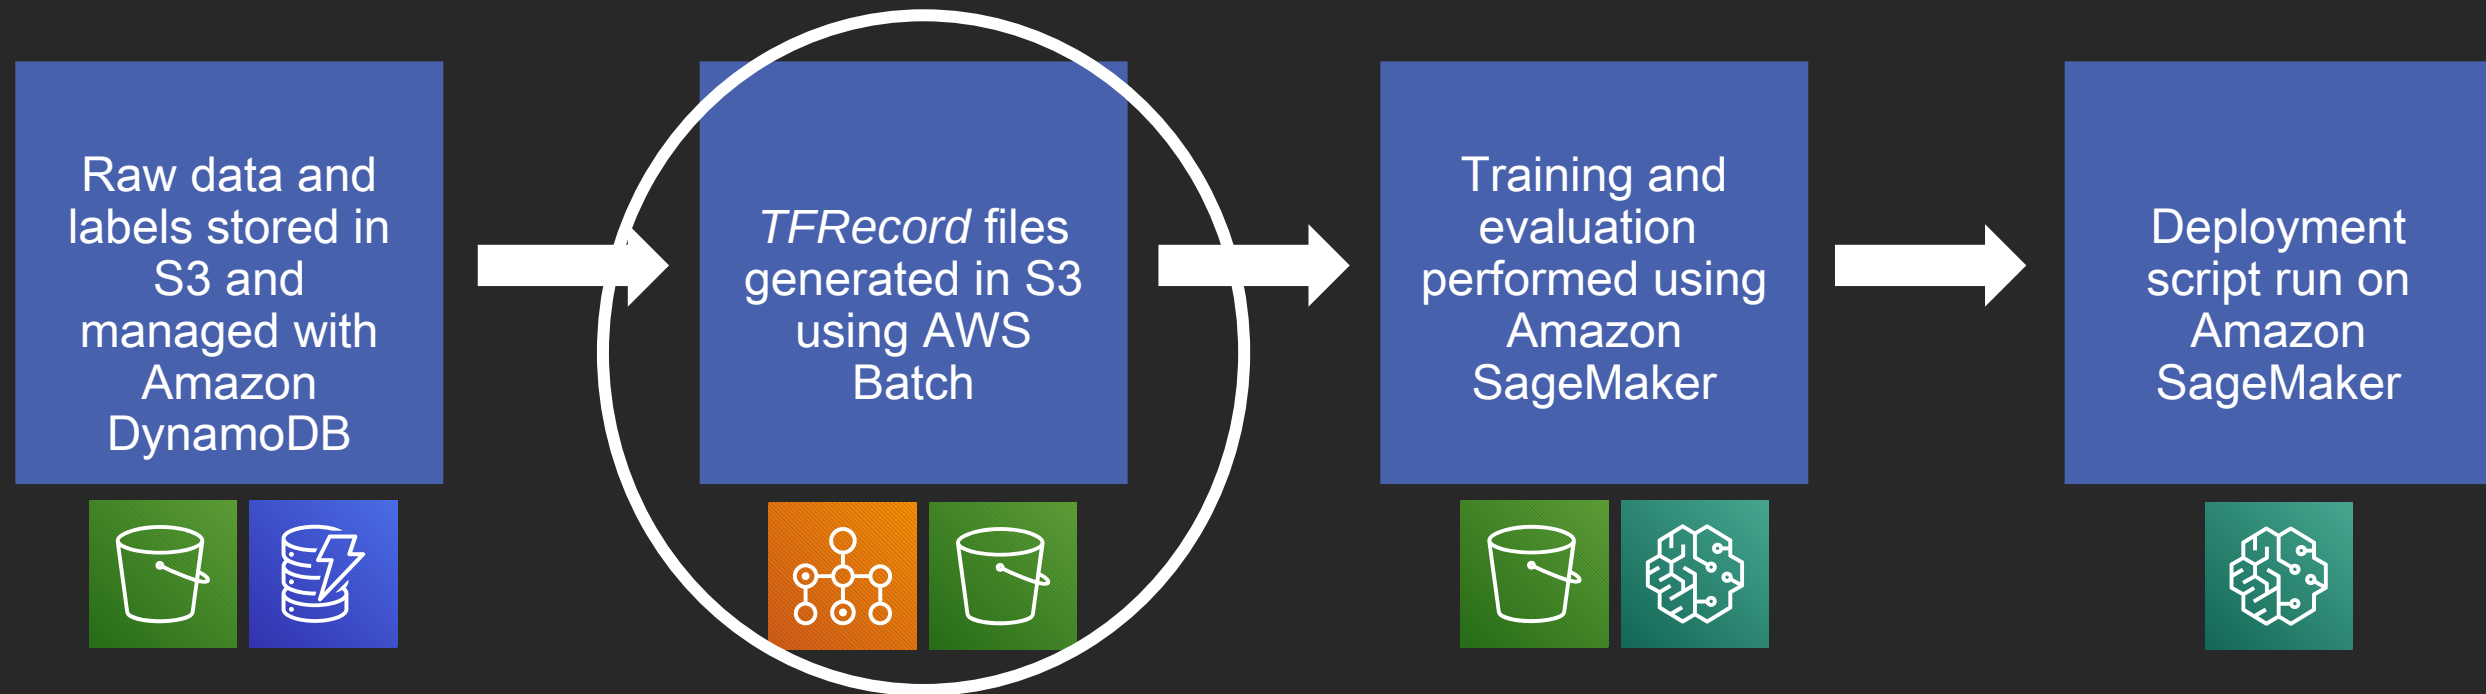
Generating data in *TfRecord* format

- One more step to moving end-to-end development to cloud:



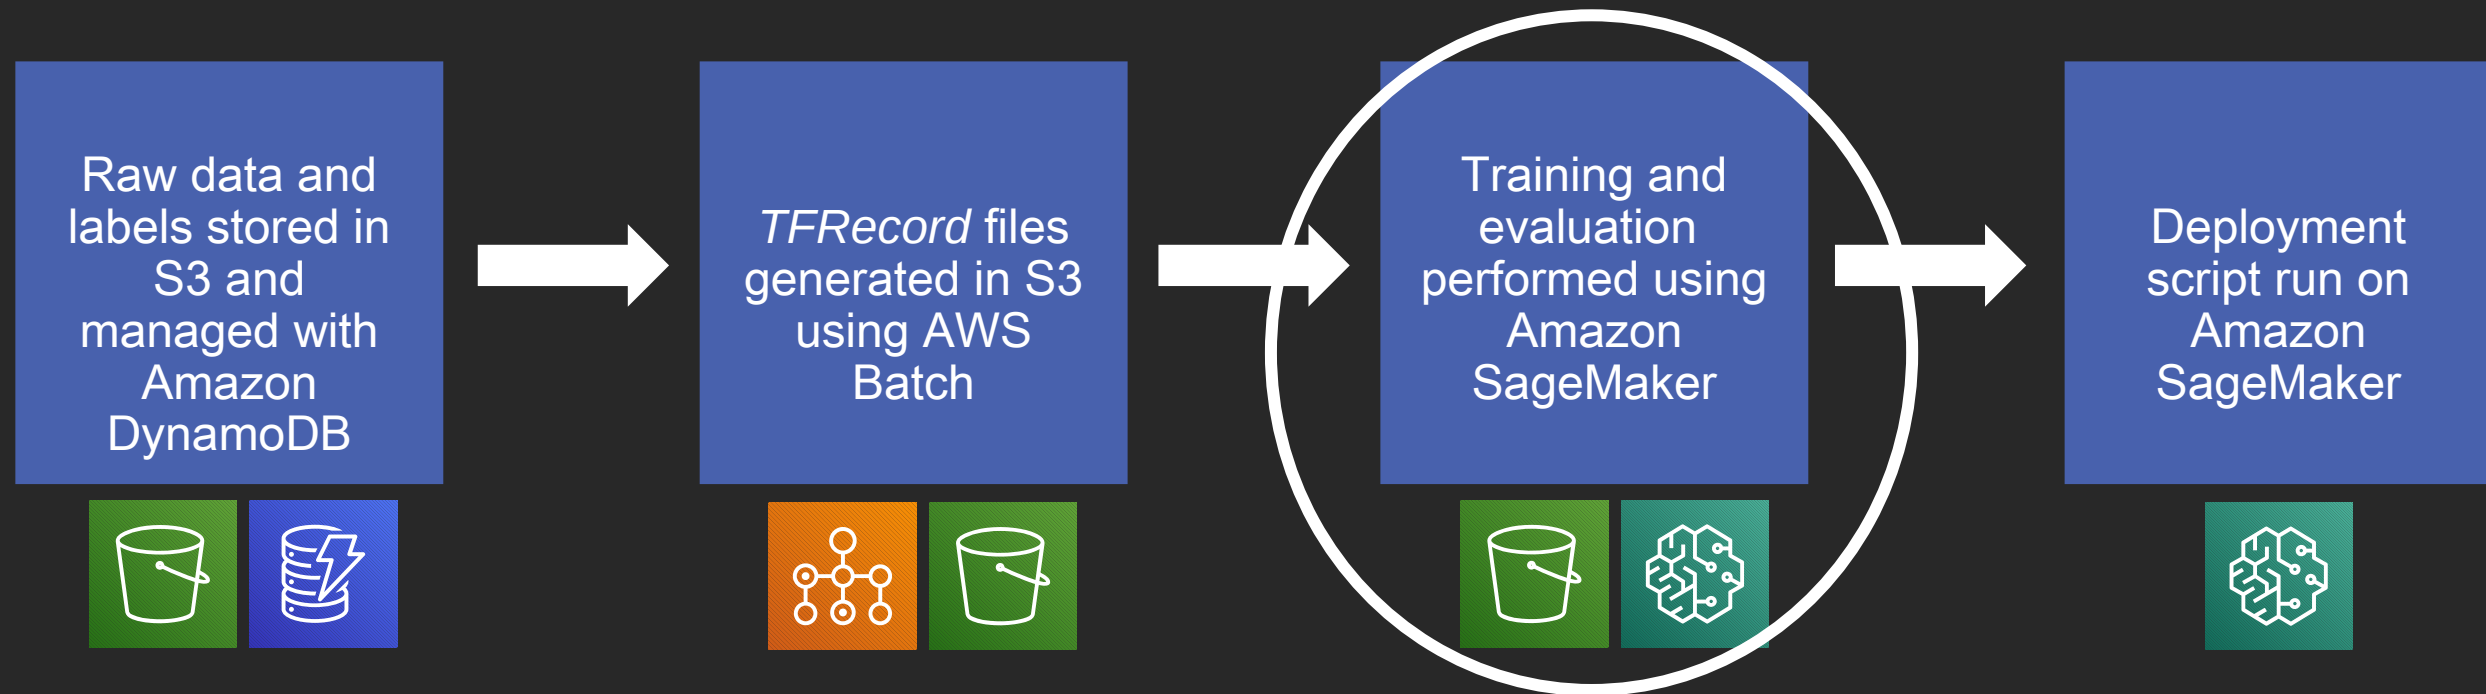
Generating data in *TfRecord* format

- One more step to moving end-to-end development to cloud:



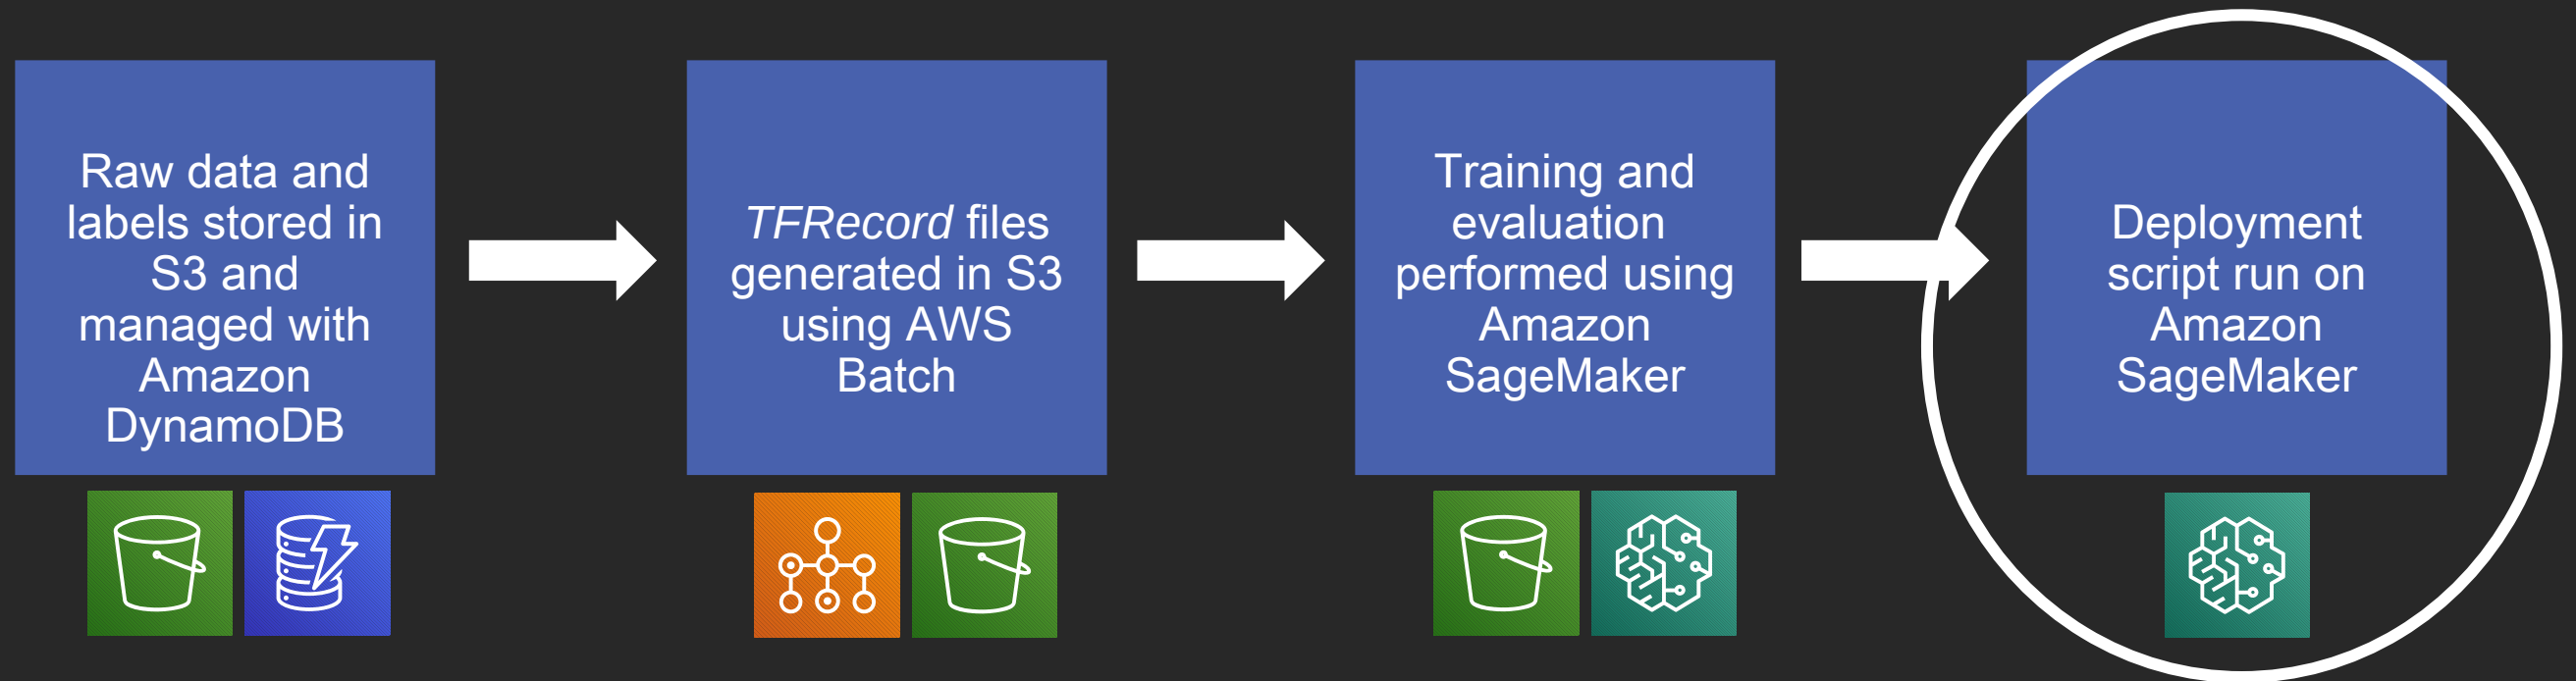
Generating data in *TfRecord* format

- One more step to moving end-to-end development to cloud:




Generating data in *TfRecord* format

- One more step to moving end-to-end development to cloud:



Pipe mode challenges


- Sequential nature of pipe mode
 - Overcoming the lack of random access to full dataset
 - Relied on for shuffling and boosting



Overcoming the lack of random access to data

- Challenge: **shuffling**

- In the pre-Amazon SageMaker era, we relied on random access to the data to ensure shuffling



- Solution: introduce shuffling on a number of levels

- Using Amazon Sagemaker *ShuffleConfig* class to shuffle *TFRecord* files for each epoch

```
train_data = s3_input('s3://sagemaker-path-to-train-data',  
                      shuffle_config=ShuffleConfig(seed))
```


```
pipes = {'train':train_data}
```

- Using TF dataset shuffle

Overcoming the lack of random access to data

- Challenge: **boosting**


- Solution for increasing representation of underrepresented data
 - E.g., pink cars in a vehicle detection DNN
- In the pre-Amazon SageMaker era, we relied on random access to the data to perform boosting



- Solution:


- Underrepresented data can be separated and fed using a dedicated pipe
- *Dataset APIs* can be used to associate a weight with each pipe and interleave the pipes:

```
ds = tf.data.sample_from_datasets(datasets, weights)
```



Pipe Mode challenges

- Overcoming the limitation on number of pipes
- Why might I need more than 20 pipes?
 - Boost **underrepresented data**



Pipe Mode challenges

- Overcoming the limitation on number of pipes
- Why might I need more than 20 pipes?
 - Boost **underrepresented data**
 - Distributed training with **Horovod**, multiply by number of GPUs

Overcoming the limitation on number of pipes

- Solution: use Pipe Mode **manifest files**
 - An alternative way to configure an Amazon SageMaker pipe
 - Replace path prefix with a **list of files**
 - Include the same file **multiple times** to increase its weight
 - Have finer control over the data used for training

```
data = s3_input('s3://path-to-manifest-file',  
               s3_data_type='ManifestFile',  
               shuffle_config=ShuffleConfig(seed))
```

Chapter 5 – Closing Remarks

Other considerations when moving to Amazon SageMaker

- Blog post: <https://bit.ly/2sLRJb5>
- **Distributed training**
 - How to choose the instance type with the optimal number of GPUs
 - Should one use TensorFlow distributed training or Horovod?
 - How does one maximize resource utilization?
- How to take advantage of **Spot Instances** to reduce cost
 - How to ensure that mid-train termination won't have negative impact
- Using **TensorBoard** and other benchmarking tools
- Debugging on a remote environment
 - Tip: make sure your model compiles and runs locally before running on cloud

Summary

- Adopting Amazon SageMaker has **boosted** our DNN development capabilities
- Pipe Mode offers a **compelling** solution for training with large datasets
- Reduce development time **significantly (~10X)**
 - Scalability enables you to test multiple models in parallel
 - Integrating with other AWS services can provide further acceleration
- Challenges were overcome using advanced Amazon SageMaker features and TF dataset APIs
- You can make Amazon SageMaker work for **you** too

Amazon SageMaker at Mobileye

Chaim Rand


ML Algorithm Developer
Mobileye

Demo

<https://gitlab.com/juliensimon/aim410>

Amazon SageMaker

Build, train, deploy machine learning models quickly at scale



**Amazon
SageMaker**

Ground
Truth

ML
Marketplace

Algorithms &
Frameworks

NEW!
Quick-start
notebooks

Training &
Tuning

Reinforcement
Learning

NEW!
SageMaker Studio
IDE

NEW!
Experiments

NEW!
Debugger

NEW!
Autopilot

Neo

Deployment &
Hosting

NEW!
Monitoring

Getting started

<http://aws.amazon.com/free>

<https://aws.amazon.com/tensorflow/>

<https://aws.amazon.com/sagemaker>

<https://github.com/aws/sagemaker-python-sdk>

https://sagemaker.readthedocs.io/en/stable/using_tf.html

<https://github.com/aws-labs/amazon-sagemaker-examples>

<https://gitlab.com/juliensimon/dlnotebooks>

Related breakouts

[[AUT307](#)] [Navigating the winding road toward driverless mobility, with Mobileye]
Dec 4, 4:00 PM – Aria, Level 1 West, Bristlecone 9 Red

[[AIM410R1](#)] [Deep learning applications with TensorFlow, with Fannie Mae]
Dec 5, 1:00 PM – Venetian, Level 3, Lido 3002


[[AIM307](#)] [Amazon SageMaker deep dive: A modular solution for machine learning]
Dec 4, 1:45 PM – Venetian, Level 3, Lido 3005

Learn ML with AWS Training and Certification

The same training that our own developers use, now available on demand



Role-based ML learning paths for developers, data scientists, data platform engineers, and business decision makers



70+ free digital ML courses from AWS experts let you learn from real-world challenges tackled at AWS



Validate expertise with the
AWS Certified Machine Learning - Specialty exam

Visit <https://aws.training/machinelearning>

Thank you!



Please complete the
session survey in the mobile
app.