

Opening keynote

Julien Simon

Global Evangelist, AI & Machine Learning, AWS

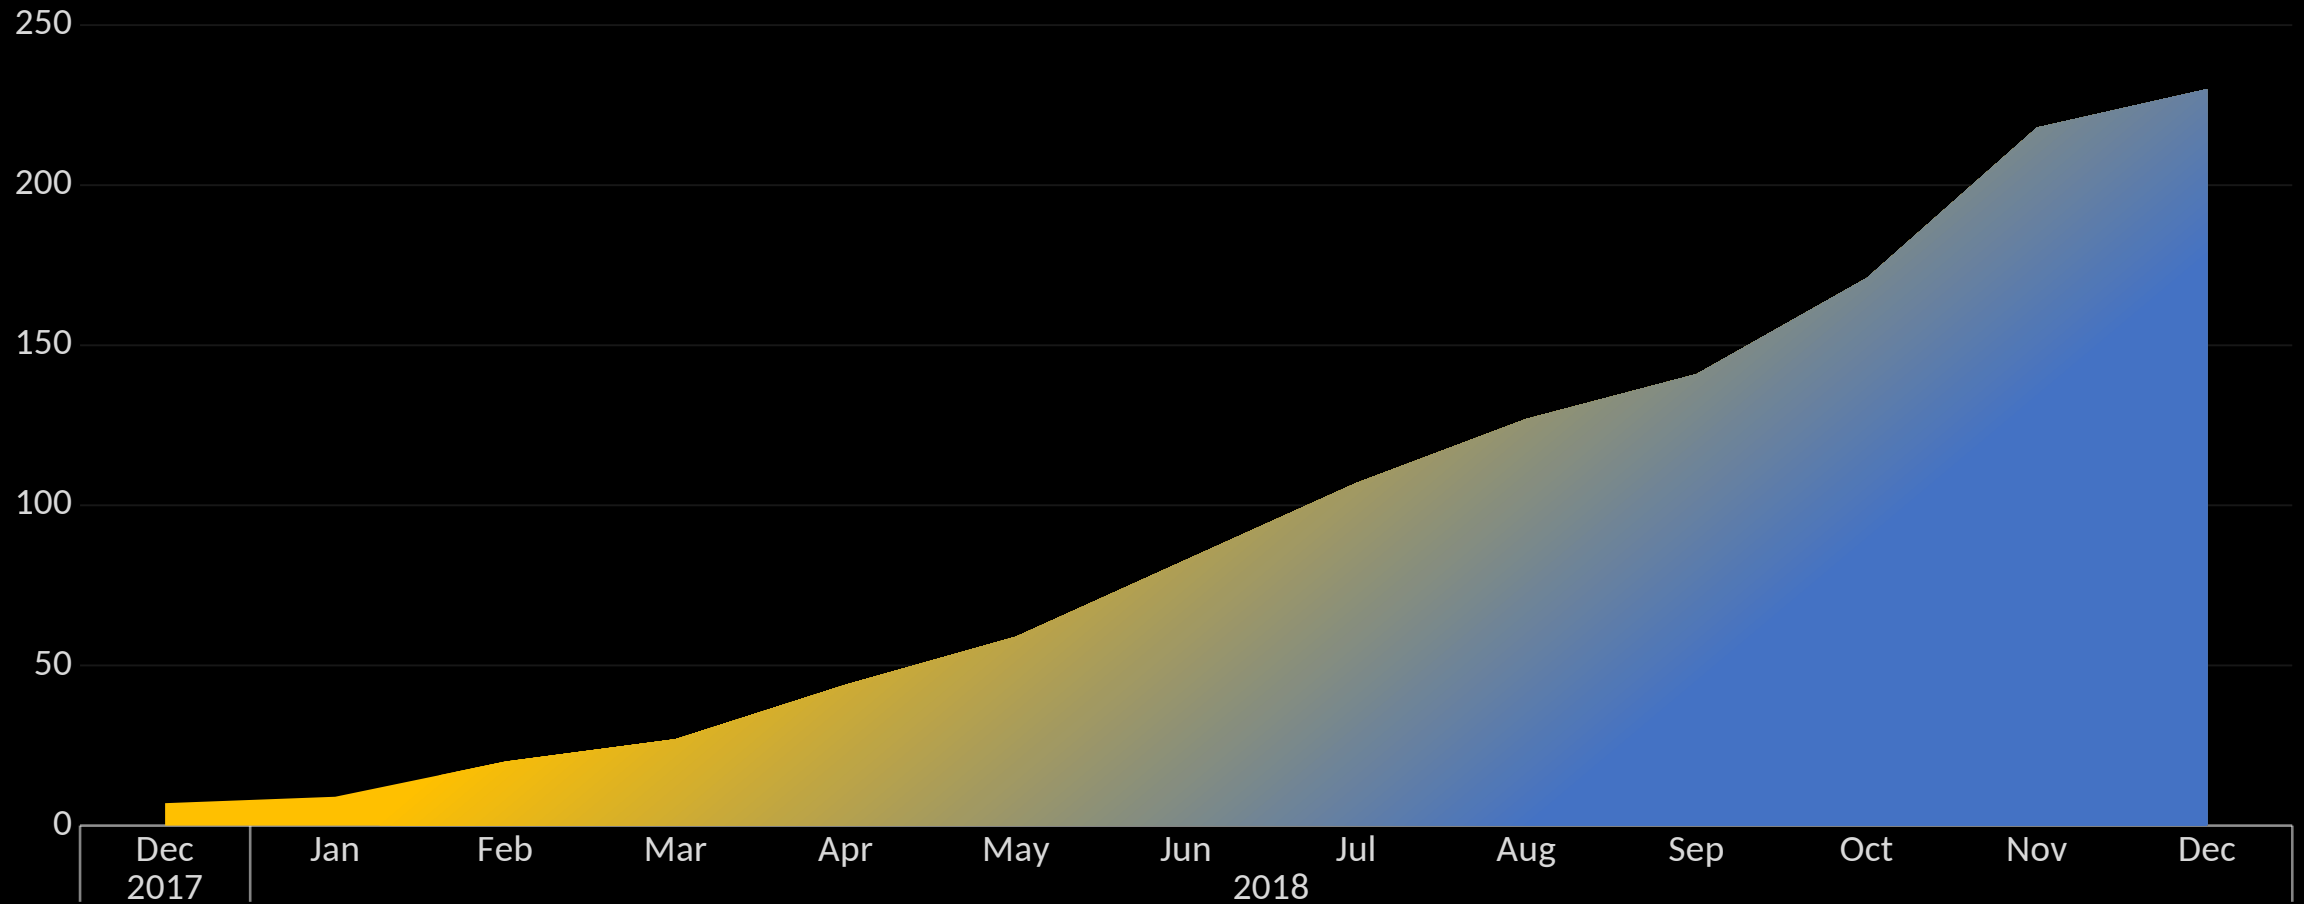
@julsimon

<https://medium.com/@julsimon>

Our mission at AWS

Put machine learning in the
hands of every developer


Over 200 launches since re:Invent 2017



The Amazon ML Stack


AI SERVICES

(App developers with little knowledge of ML)




ML SERVICES

(ML developers and data scientists)



ML FRAMEWORKS & INFRASTRUCTURE

(ML researchers and academics)



Machine Learning University


<https://aws.training/machinelearning>



Uses the same
materials used to train
Amazon developers



Foundational knowledge
with
real-world application



Structured
courses and
specialist certification

We're focused on solving the toughest challenges that hold back success with machine learning

Three challenges facing ML world today

1

Flexibility & Cost

Optimized TensorFlow
MXNet Dynamic Training
Amazon Elastic Inference

2

Data

Amazon SageMaker Ground Truth
Amazon SageMaker RL

3

Ease of Use

AWS Marketplace for Machine Learning
Amazon SageMaker Neo
Amazon Textract
Amazon Forecast
Amazon Personalize
Amazon Comprehend Medical

Let's look under the hood into some of the new features in each layer of the stack.

Three challenges we're focused on today

1

Flexibility & Cost

Optimized TensorFlow
MXNet Dynamic Training
Amazon Elastic Inference

2

Data

Amazon SageMaker Ground Truth
Amazon SageMaker RL

3

Ease of Use

AWS Marketplace for Machine Learning
Amazon SageMaker Neo
Amazon Textract
Amazon Forecast
Amazon Personalize
Amazon Comprehend Medical

The Amazon ML Stack

AI SERVICES

(App developers with little knowledge of ML)



Amazon
Rekognition
Image



Amazon
Rekognition
Video



Amazon
Textract



Amazon
Polly



Amazon
Transcribe



Amazon
Translate



Amazon
Comprehend
& Comprehend
Medical



Amazon
Lex



Amazon
Forecast



Amazon
Personalize

Vision

Speech

Language

Chatbots

Forecasting

Recommendations

ML SERVICES

(ML developers and data scientists)



Amazon
SageMaker

Ground Truth

Notebooks

Algorithms

AWS Marketplace

Supervised
Learning

Unsupervised
Learning

Reinforcement
Learning

Training

Optimization
(Neo)

Deployment

Hosting

ML FRAMEWORKS & INFRASTRUCTURE

(ML researchers and academics)

Frameworks



Interfaces




Infrastructure



Amazon
EC2 P3
& P3DN



Amazon
EC2 C5



FPGAs



AWS
Greengrass



Amazon
Elastic
Inference




Amazon
Inferentia

AWS is framework agnostic


Choose from popular frameworks



Run them fully managed



Or run them yourself



AWS: The platform of choice to run TensorFlow



85% of all
TensorFlow
workloads in the
cloud runs on AWS

Source: Nucleus Research, November 2018

The best place to run TensorFlow

Stock
TensorFlow

65%

scaling efficiency
with 256 GPUs

The best place to run TensorFlow

Stock
TensorFlow

65%

scaling efficiency
with 256 GPUs



AWS-Optimized
TensorFlow

90%

scaling efficiency
with 256 GPUs

Available with
Amazon SageMaker
and the AWS Deep
Learning AMIs



HOROVOD

Apache MXNet: deep learning for enterprise developers



Start with off-the-shelf toolkits

- Gluon CV and Gluon NLP

Fast and scalable training

- Keras-MXNet up to 2x faster than Keras-TensorFlow
- Near-linear scalability up to 256 GPUs
- Dynamic training

Easy deployment

- Java/Scala APIs
- MXNet Model Server

Apache MXNet customer example: Chick-fil-A

TRACK WAFFLE FRY FRESHNESS


Identify fries that have exceeded the hold time

USES DEEP LEARNING

Computer vision model for object detection and tracking

BUILT THE MODEL USING MXNet


A team of enterprise developers with no ML expertise




Amazon Elastic Inference



Lower inference costs
up to 75%



Match capacity
to demand



Available between 1 to 32
TFLOPS

KEY FEATURES

Support for TensorFlow,
Apache MXNet, and ONNX
with PyTorch coming soon

Integrated with
Amazon EC2,
Amazon SageMaker,
and Amazon DL AMIs

Single and
mixed-precision
operations

Helping developers learn computer vision

AWS DeepLens: the world's first deep learning-enabled video camera for developers



- Purpose-built for ML-skills development
- Fully programmable & customizable
- Build custom Amazon SageMaker models
- 10-minutes to your first deep learning project

Three challenges we're focused on today

1

Flexibility & Cost

Optimized TensorFlow
MXNet Dynamic Training
Amazon Elastic Inference

2

Data

Amazon SageMaker Ground Truth
Amazon SageMaker RL

3

Ease of

Use

AWS Marketplace for Machine Learning
Amazon SageMaker Neo
Amazon Textract
Amazon Forecast
Amazon Personalize
Amazon Comprehend Medical

The Amazon ML Stack

AI SERVICES

(App developers with little knowledge of ML)



Amazon
Rekognition
Image



Amazon
Rekognition
Video



Amazon
Textract



Amazon
Polly



Amazon
Transcribe



Amazon
Translate



Amazon
Comprehend
& Comprehend
Medical



Amazon
Lex



Amazon
Forecast



Amazon
Personalize

Vision

Speech

Language

Chatbots

Forecasting

Recommendations

ML SERVICES

(ML developers and data scientists)



Amazon
SageMaker

Ground Truth

Notebooks

Algorithms

AWS Marketplace

Supervised
Learning

Unsupervised
Learning

Reinforcement
Learning

Training

Optimization
(Neo)


Deployment

Hosting

ML FRAMEWORKS & INFRASTRUCTURE

(ML researchers and academics)

Frameworks



Interfaces



Infrastructure



Amazon
EC2 P3
& P3DN



Amazon
EC2 C5



FPGAs



AWS
Greengrass



Amazon
Elastic
Inference




Amazon
Inferentia


Labeling data sets is too time-consuming

Amazon SageMaker Ground Truth

Label machine learning training data easily and accurately



Quickly label
training data



Easily integrate
human labelers



Get accurate
results

KEY FEATURES

Automatic labeling via
machine learning

Ready-made and
custom workflows

Private and public
human workforce

Label
management



CHANGE
HEALTHCARE



gumgum

aws INNOVATE
ONLINE CONFERENCE | AI EDITION

How do you teach machine learning models to make decisions when there is no training data?

Reinforcement Learning




Learn by
interacting with
the real world



Model the real-
world problem
as a simulation
environment



Trial and error
Observe results



Optimize
learning strategy
to maximize
long-term
reward



Model learns
how to make
complex
decisions

Amazon SageMaker RL

RL toolkits that provide RL agent algorithm implementations

RL-Coach

DQN

PPO

A3C

Rainbow

...

RL-Ray RLLib

DQN

PPO

IMPALA

A3C

...

Open AI Baselines

DQN

PPO

...

...

Amazon SageMaker deep learning frameworks

Tensorflow

MxNet

PyTorch

Chainer

 SageMaker supported

 Customer BYO

Customers are using Amazon SageMaker RL



GE Healthcare

HONDA

SIXT

mixi

amazon

SyntheticGestalt

Scientific Research by Artificially Intelligent Agents



Tradelegs

Helping developers learn RL

AWS DeepRacer: a fully autonomous 1/18th-scale race car



- Build machine learning models in Amazon SageMaker
- Train, test, and iterate on the track using the AWS DeepRacer 3D racing simulator
- Compete in the world's first global autonomous racing league, either at AWS Summits or at virtual events
- Race for prizes and a chance to advance to win the coveted AWS DeepRacer Cup

Three challenges we're focused on today

1

Flexibility & Cost

Optimized TensorFlow
MXNet Dynamic Training
Amazon Elastic Inference

2

Data

Amazon SageMaker Ground Truth
Amazon SageMaker RL

3

Ease of Use

AWS Marketplace for Machine Learning
Amazon SageMaker Neo
Amazon Textract
Amazon Forecast
Amazon Personalize
Amazon Comprehend Medical

Deploying models across multiple platforms is too time-consuming

We all want machine learning everywhere



Autonomous
vehicles



Smart
agriculture




Predictive
maintenance



Robotics



Speech and sound
recognition



Video
security

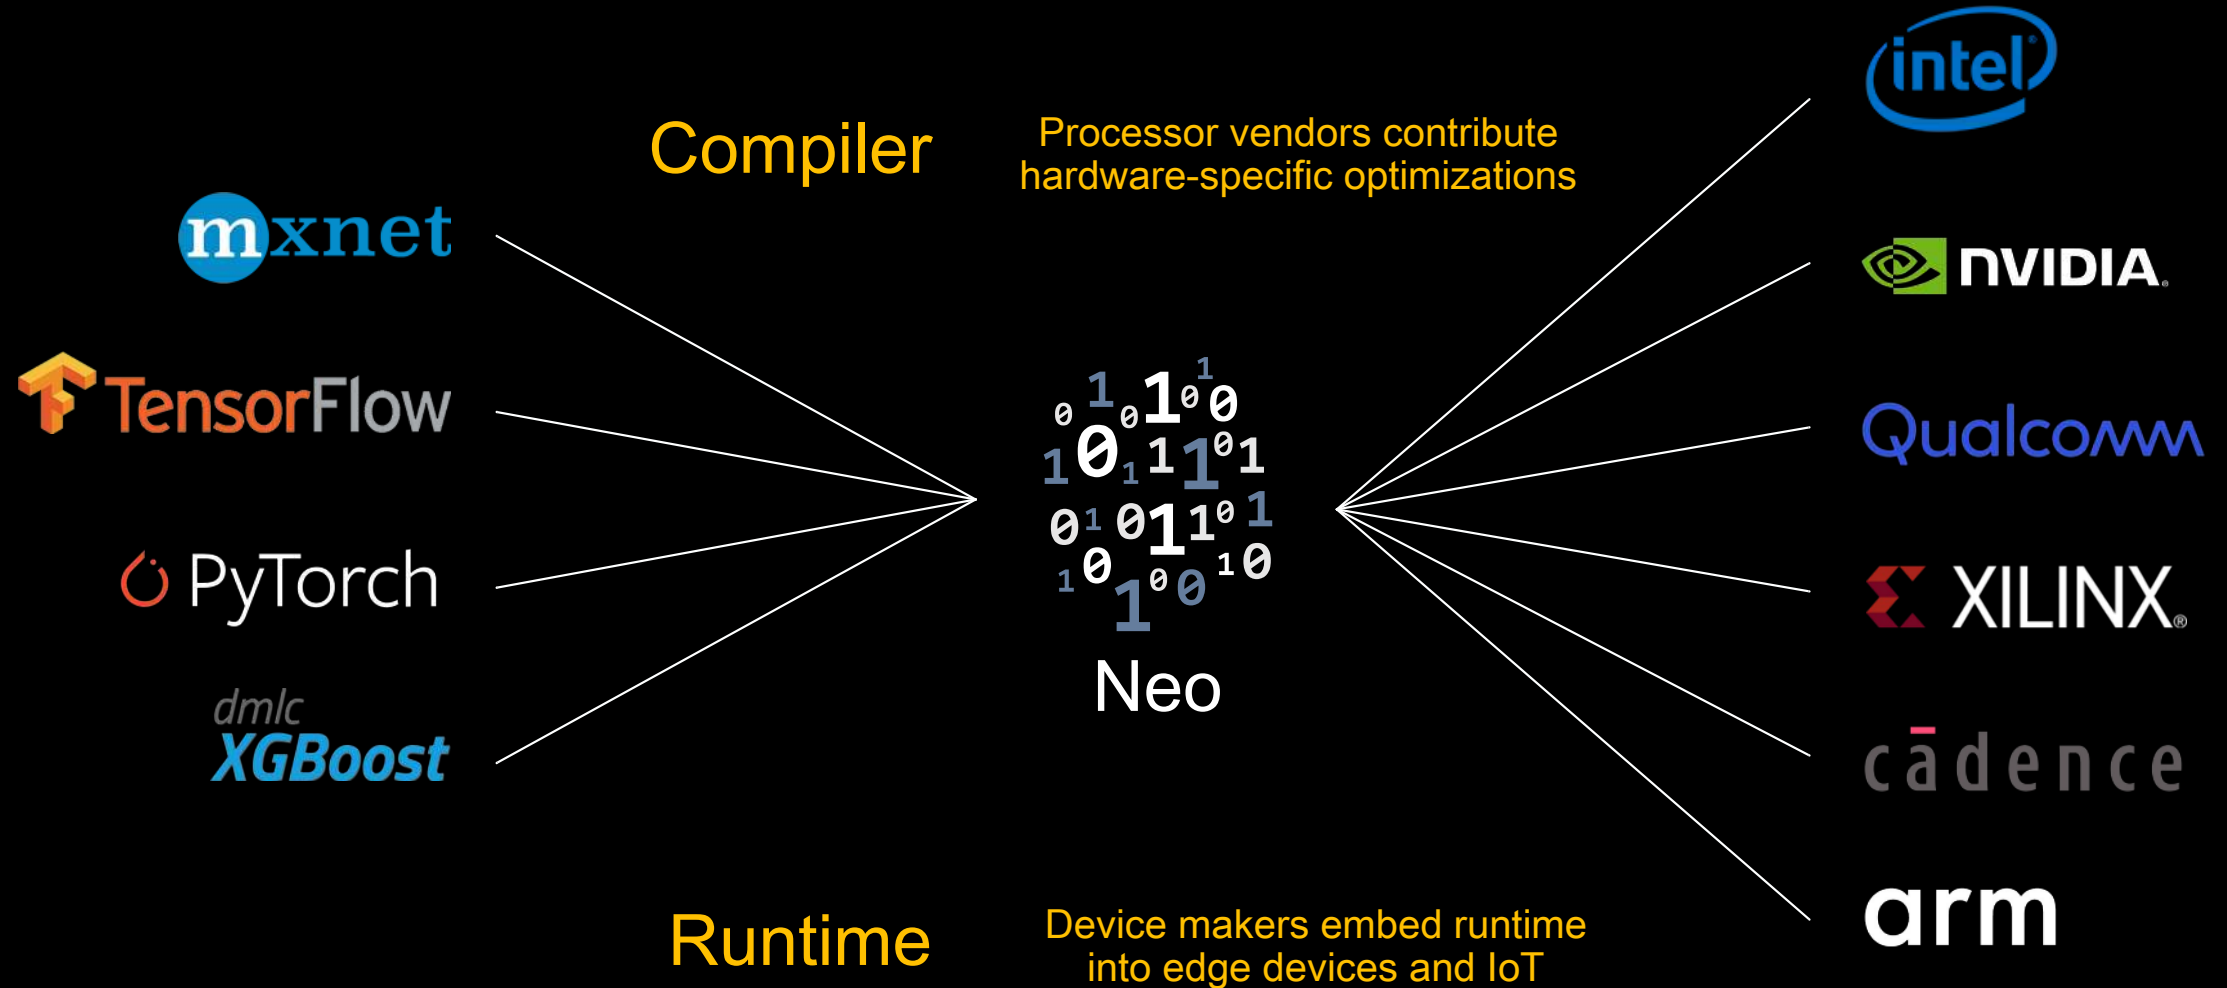


Anomaly
detection



More

Neo: train once, run anywhere



The Amazon ML Stack

AI SERVICES

(App developers with little knowledge of ML)



Amazon
Rekognition
Image



Amazon
Rekognition
Video



Amazon
Textract



Amazon
Polly



Amazon
Transcribe



Amazon
Translate



Amazon
Comprehend
& Comprehend
Medical



Amazon
Lex



Amazon
Forecast



Amazon
Personalize

Vision

Speech

Language

Chatbots

Forecasting

Recommendations

ML SERVICES

(ML developers and data scientists)



Amazon
SageMaker

Ground Truth

Notebooks
Algorithms
AWS Marketplace


Supervised Learning
Unsupervised Learning
Reinforcement Learning

Training
Optimization (Neo)

Deployment
Hosting

ML FRAMEWORKS & INFRASTRUCTURE

(ML researchers and academics)



Amazon
EC2 P3
& P3DN



Amazon
EC2 C5



FPGAs



AWS
Greengrass




Amazon
Elastic
Inference




Amazon
Inferentia

Amazon Textract (preview)

Extract text and data from virtually any document



Extract data quickly and accurately



Eliminate manual effort



Lower document processing costs

KEY FEATURES

Optical Character
Recognition
(OCR)

Key-value pair
detection

Table
detection

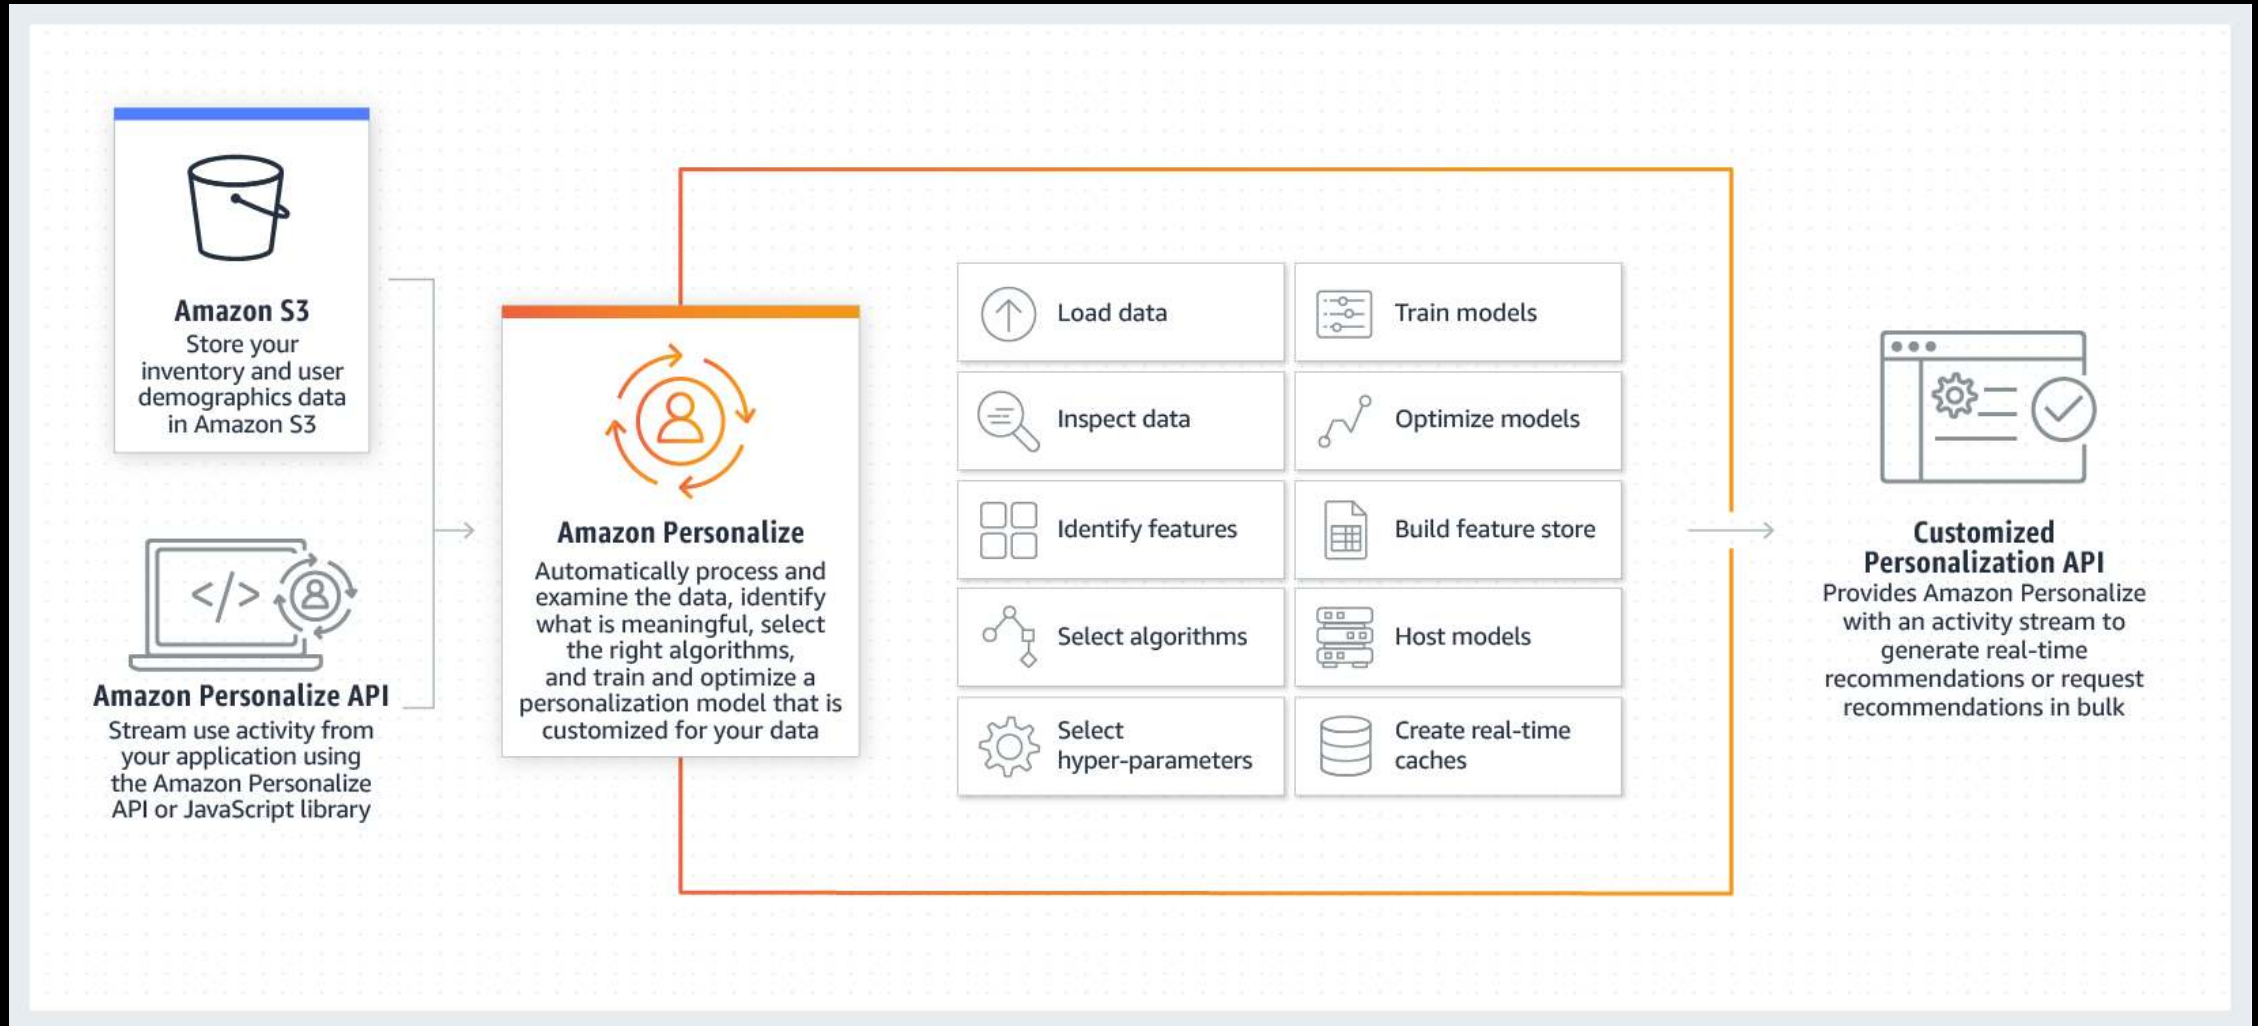
Adjustable confidence
thresholds

Bounding box
coordinates

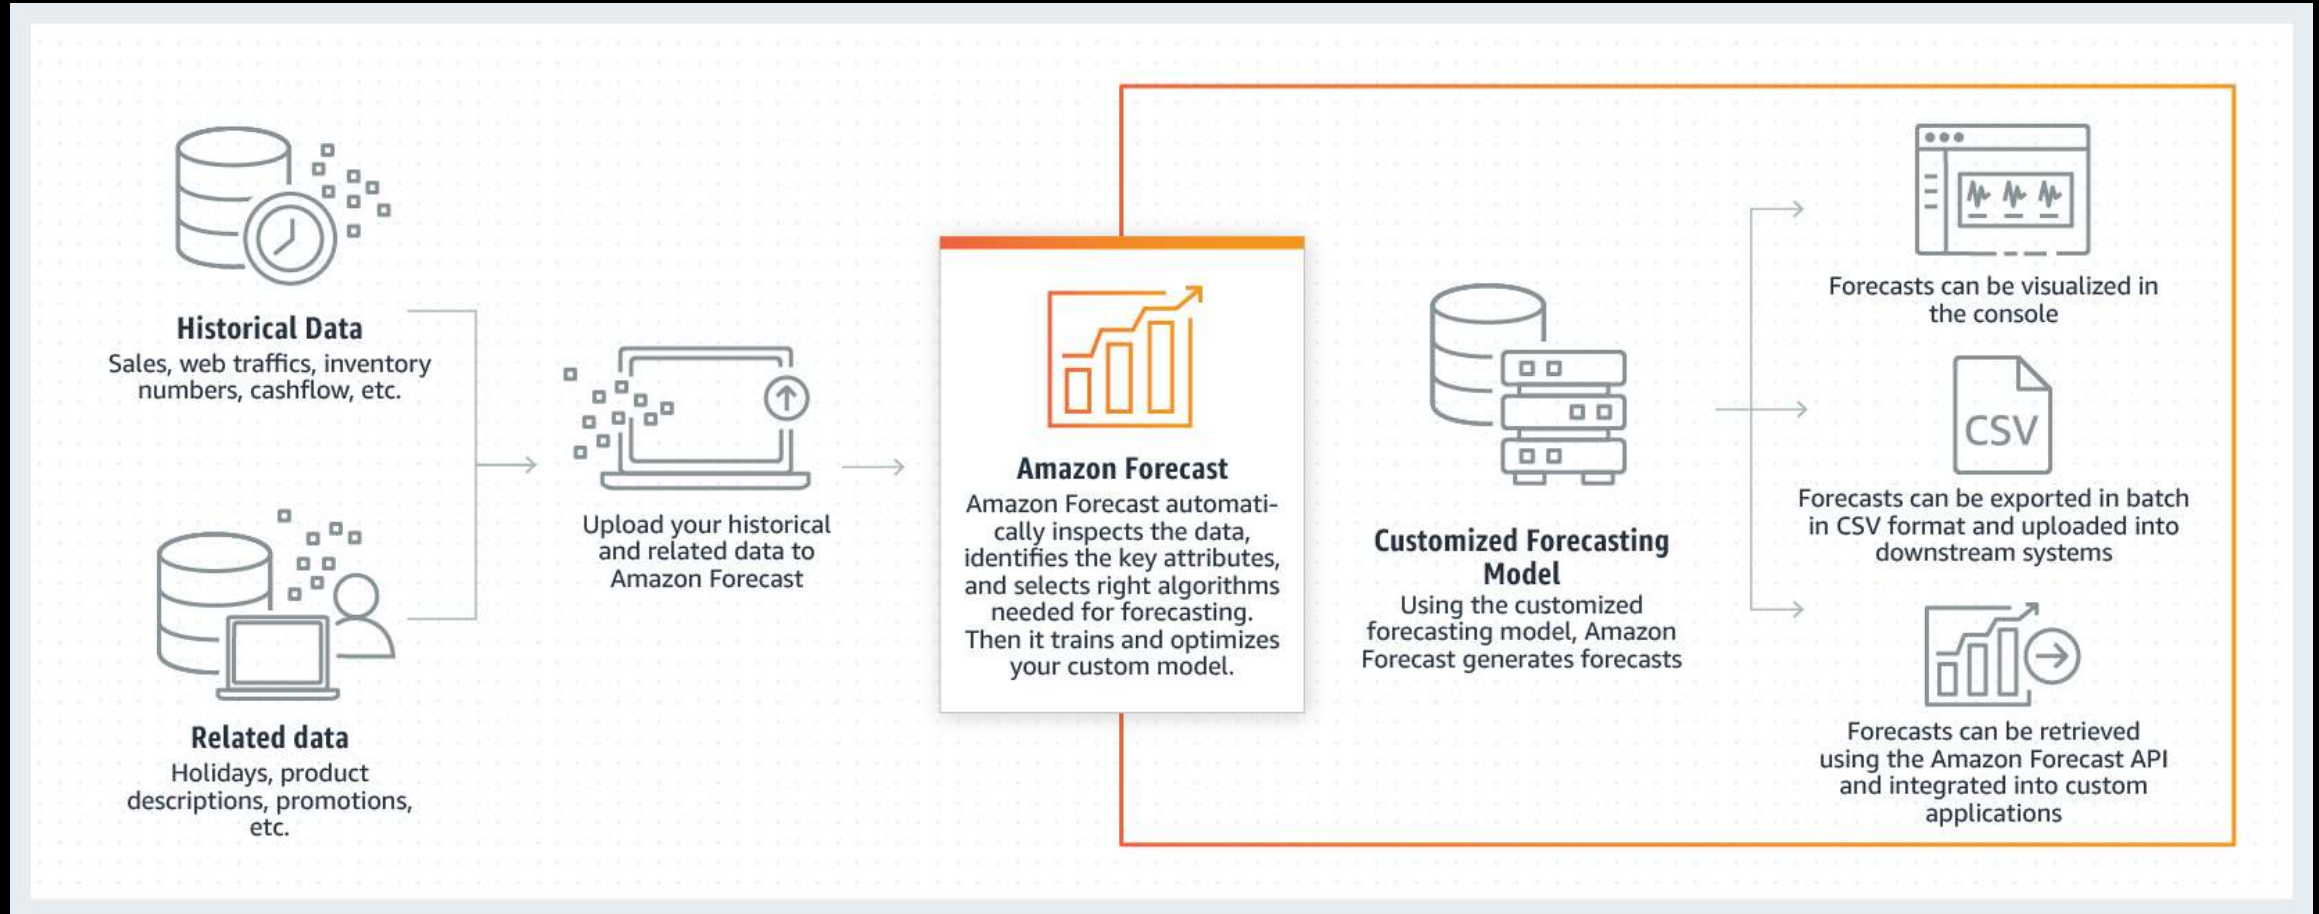
No ML experience
required



Amazon Personalize (preview)



Amazon Forecast (preview)



Amazon Comprehend Medical

Input text

Pt is 40yo mo
HPI : Sleeping
Meds : Vyvanse
HEENT : Bogg
erythematous
Follow-up as :

407 of 10000 char

40yo
0.99+ score

software e
0.98 score

Sleeping t
0.81 score

Clonidine
0.98 score

Rash
0.99+ score

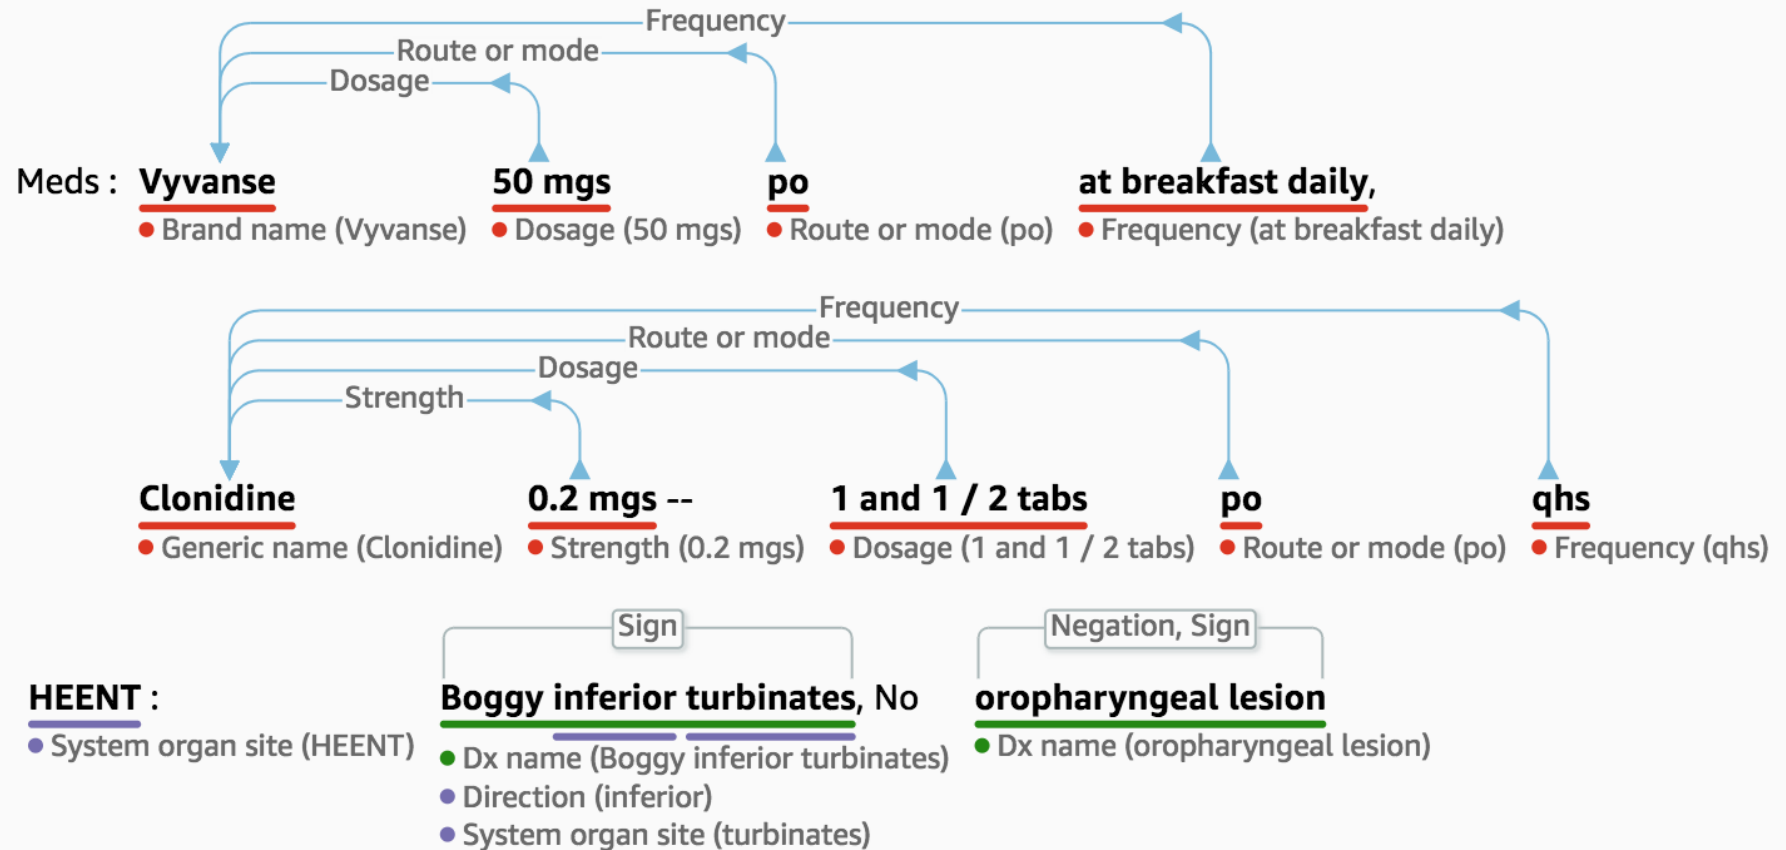
face
0.98 score

leg
0.99+ score

Age

Protected health information

-



Three challenges facing ML world today

1

Flexibility & Cost

Optimized TensorFlow
MXNet Dynamic Training
Amazon Elastic Inference

2

Data

Amazon SageMaker Ground Truth
Amazon SageMaker RL

3


Ease of Use

AWS Marketplace for Machine Learning
Amazon SageMaker Neo
Amazon Textract
Amazon Forecast
Amazon Personalize
Amazon Comprehend Medical

The Amazon ML Stack


AI SERVICES

(App developers with little knowledge of ML)




ML SERVICES

(ML developers and data scientists)



ML FRAMEWORKS & INFRASTRUCTURE

(ML researchers and academics)



Track 1	Track 2	Track 3
An Introduction to Machine Learning (100 Level)	Use Natural Language and Text Analysis to Power More Intelligent Applications (200 Level)	Get Started with Amazon SageMaker: Build, Train, and Deploy Machine Learning Models in the Cloud (200 Level)
Break	Break	Break
Get Started with Amazon SageMaker: Build, Train, and Deploy Machine Learning Models in the Cloud (200 Level)	Power New Application Experiences with Image, Video, and Facial Analysis (200 Level)	Remove Your Biggest Challenge to Getting Started with ML: Using Amazon SageMaker Ground Truth to Label Data at 70% Less Cost (200 Level)
Break	Break	Break
Use AWS DeepRacer to Get Rolling with Machine Learning (200 Level)	Integrate Amazon.com's ML Techniques for Customer Recommendations and Forecasting Workflows into Your Applications (200 Level)	Learn to Build What's Next with Reinforcement Learning on Amazon SageMaker (200 Level)

Thank you!

Julien Simon

Global Evangelist, AI & Machine Learning, AWS

@julsimon

<https://medium.com/@julsimon>