

Getting started with Machine Learning on AWS

Julien Simon
Global Evangelist, AI & Machine Learning, AWS
[@julsimon](#)

A primer on Machine Learning

AI vs. Machine Learning vs. Deep Learning

Artificial Intelligence: design software applications which exhibit human-like behavior, e.g. speech, natural language processing, reasoning or intuition

Machine Learning: using **statistical algorithms**, teach machines to learn from **featurized data** without being explicitly programmed

Deep Learning: using **neural networks algorithms**, teach machines to learn from **complex data** where features **cannot** be easily expressed

Types of Machine Learning

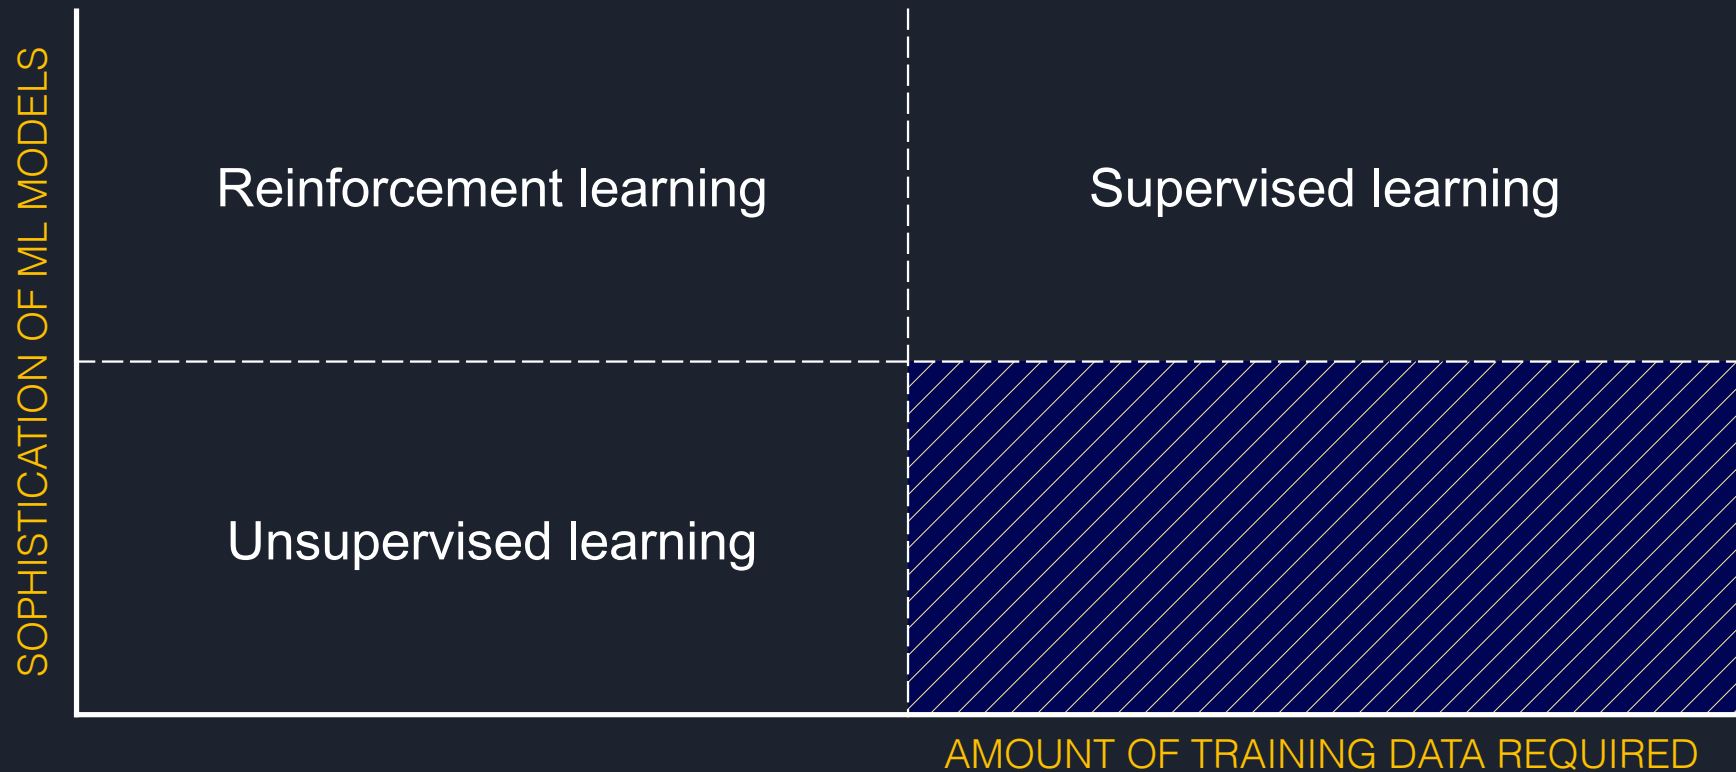
Supervised learning

- Run an algorithm on a **labeled** data set.
- The model learns how to correctly predict the **right answer**.
- Regression and classification are examples of supervised learning.

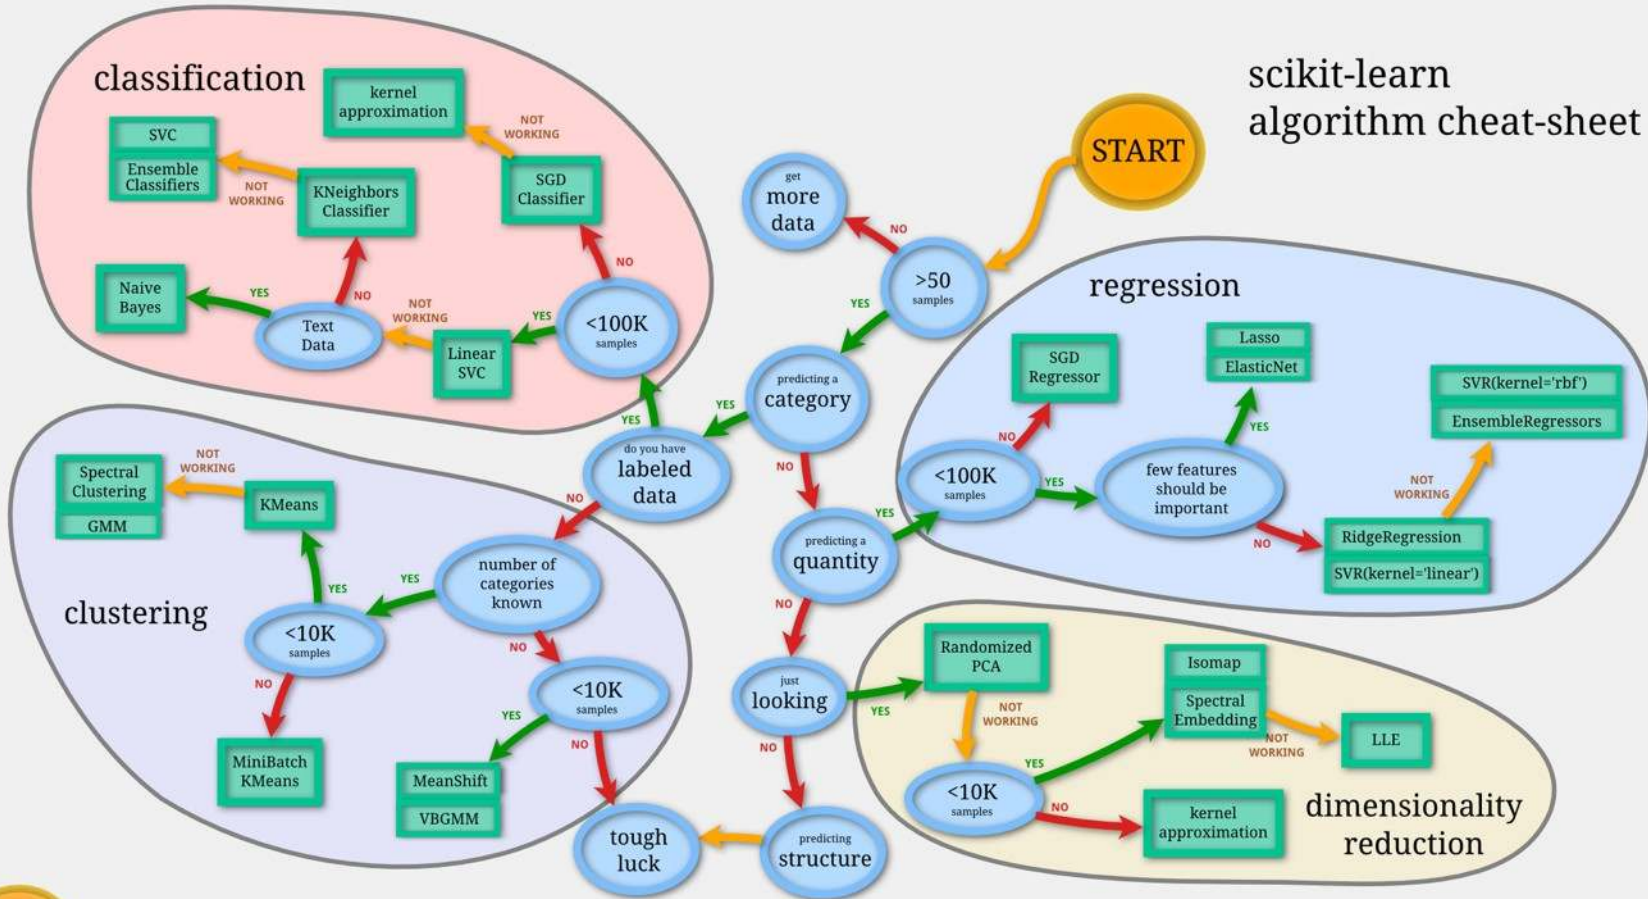
Unsupervised learning

- Run an algorithm on an **unlabeled** data set.
- The model learns **patterns** and organizes samples accordingly.
- Clustering and topic modeling are examples of unsupervised learning.

Types of Machine Learning

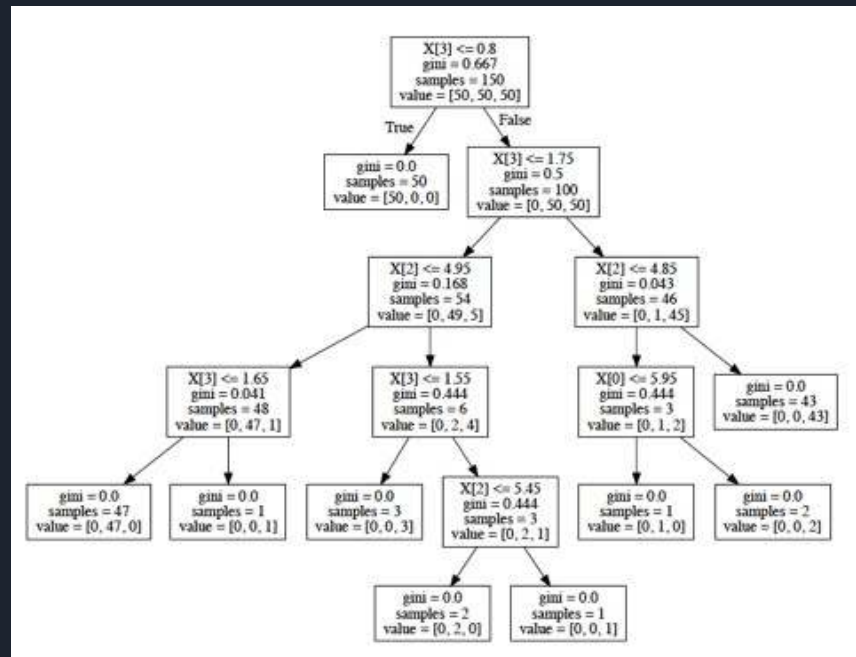


scikit-learn algorithm cheat-sheet



Algorithm example: decision trees

- Supervised learning algorithm
- Goal: build a **decision tree** for regression or classification
- Data set : **features** + target attribute (value or class identifier)
- Intuition: find the “best” **feature thresholds** to go left or right
- “Easy” to interpret
- Advanced variants with multiple trees:
Random Forests, XGBoost, etc.



```

graph TD
    BP((Business Problem)) --> MPF[ML problem framing]
    MPF --> DC[Data collection]
    DC --> DI[Data integration]
    DI --> DPC[Data preparation and cleaning]
    DPC --> DVA[Data visualization and analysis]
    DVA -- Feature augmentation --> DC
    DVA --> FE[Feature engineering]
    FE --> MPT[Model training and parameter tuning]
    MPT --> ME[Model evaluation]
    ME --> D{Are business goals met?}
    D -- NO --> DVA
    D -- YES --> MD[Model deployment]
    MD --> P((Predictions))
    P --> MD[Monitoring and debugging]
    MD -- Re-training --> DC
  
```


Putting your Machine Learning Projects on the right track

1 - Set expectations

- What is the **business** question you're trying to answer?
 - One sentence on the whiteboard
 - Must be **quantifiable**
- Do you have (enough) **data** that could help?
- Involve everyone and come to a **common** understanding
 - Business, IT, Data Engineering, Data Science, Ops, etc.

« We want to see what this technology can do for us »

« We have tons of relational data, surely we can do something with it »

« I read this cool article about FooBar ML, we ought to try it »



2 - Define clear metrics

- What is the **business metric** showing success?
- What's the **baseline** (human and IT)?
- What would be a **significant** and **reasonable** improvement?
- What would be **reasonable** further improvements?

« The confusion matrix for our support ticket classifier has significantly improved ». **Huh?**
« P90 time-to-resolution is now under 24 hours ». **Err....**
« Misclassified emails have gone down 5.3% using the latest model ». **So?**
« The latest survey shows that 'very happy' **customers** are up 9.2% ». **Woohoo!**


3 - Assess needs (not wants) and skills

- Building a data set describing the problem?
- Cleaning, preparing and curating it?
- Writing and tweaking ML algorithms?
- Managing ML infrastructure?



4 - Pick the best tool for the job

- Cost, time to market, accuracy: **pick two**
- The least expensive and fastest option won't probably be the most accurate.
 - Maybe enough to **get started**, and **learn** more about the problem.
- Improving accuracy will take **increasingly** more time and money.
 - Diminishing returns! Know when to stop.
- Keep an eye on **actionable** state of the art advances, ignore the rest
 - Transfer learning
 - AutoML



5 - Use proven best practices

- No, things are **not** different this time.
- AI / ML is **software engineering**
 - Dev, test, QA, documentation, Agile, versioning, etc.
 - Involve all teams
- Sandbox tests are nice, but truth is in **production**
 - Get there fast, as often as needed
 - CI / CD and automation are required
 - Devops for ML



Universal
Pictures


6 - Iterate, iterate, iterate

aka Boyd's Law (1960)

- Start **small**
- Try the **simple** things first
- Go to production **quickly**
- Observe prediction **errors**
- Act: fix data set? Add more data? Tweak the algo? Try another algo?
- Repeat until accuracy gains become **irrelevant**
- Move to the **next** project

Machine Learning at Amazon

Machine Learning innovation at Amazon



Our mission at AWS

Put machine learning in the
hands of every developer

MACHINE LEARNING IS HAPPENING IN COMPANIES OF EVERY SIZE AND INDUSTRY














Tens of thousands customers have chosen AWS for their ML workloads | More than twice as many customers using ML than any other cloud provider




The AWS ML Stack

Broadest and most complete set of Machine Learning capabilities









AI SERVICES

VISION	SPEECH		TEXT			SEARCH	CHATBOTS	PERSONALIZATION	FORECASTING	FRAUD	DEVELOPMENT	CONTACT CENTERS
 Amazon Rekognition	 Amazon Polly	 Amazon Transcribe <i>+Medical</i>	 Amazon Comprehend <i>+Medical</i>	 Amazon Translate	 Amazon Textract	 Amazon Kendra	 Amazon Lex	 Amazon Personalize	 Amazon Forecast	 Amazon Fraud Detector	 Amazon CodeGuru	 Contact Lens <i>For Amazon Connect</i>

ML SERVICES

 Amazon SageMaker	Ground Truth	ML Marketplace	SageMaker Studio IDE								Neo	A2I
			Built-in algorithms	Notebooks	Experiments	Processing & Model Evaluation	Model training & tuning	Debugger	Autopilot	Model hosting	Model Monitor	

ML FRAMEWORKS & INFRASTRUCTURE














 TensorFlow  mxnet  PYTORCH	 GLUON  Keras  fast.ai  PyTorch Lightning  DeepGraphLibrary	Deep Learning AMLs & Containers	GPUs & CPUs	ElasticInference	Inferentia	FPGA
---	---	---------------------------------	-------------	------------------	------------	------

AI Services


The AWS ML Stack

Broadest and most complete set of Machine Learning capabilities





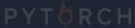


AI SERVICES

VISION	SPEECH		TEXT			SEARCH	CHATBOTS	PERSONALIZATION	FORECASTING	FRAUD	DEVELOPMENT	CONTACT CENTERS
 Amazon Rekognition	 Amazon Polly	 Amazon Transcribe <i>+Medical</i>	 Amazon Comprehend <i>+Medical</i>	 Amazon Translate	 Amazon Textract	 Amazon Kendra	 Amazon Lex	 Amazon Personalize	 Amazon Forecast	 Amazon Fraud Detector	 Amazon CodeGuru	 Contact Lens <i>For Amazon Connect</i>

ML SERVICES

 Amazon SageMaker	Ground Truth	ML Marketplace	SageMaker Studio IDE									Neo	A2I
			Built-in algorithms	Notebooks	Experiments	Processing & Model Evaluation	Model training & tuning	Debugger	Autopilot	Model hosting	Model Monitor		

ML FRAMEWORKS & INFRASTRUCTURE














 TensorFlow  mxnet	 GLUON  Keras					
 PYTORCH	 fast.ai  DeepGraphLibrary	Deep Learning AMLs & Containers	GPUs & CPUs	ElasticInference	Inferentia	FPGA

ML Services


The AWS ML Stack

Broadest and most complete set of Machine Learning capabilities








AI SERVICES

	VISION	SPEECH		TEXT			SEARCH	CHATBOTS	PERSONALIZATION	FORECASTING	FRAUD	DEVELOPMENT	CONTACT CENTERS	
	 Amazon Rekognition	 Amazon Polly	 Amazon Transcribe <i>+Medical</i>	 Amazon Comprehend <i>+Medical</i>	 Amazon Translate	 Amazon Textract	 Amazon Kendra	 Amazon Lex	 Amazon Personalize	 Amazon Forecast	 Amazon Fraud Detector	 Amazon CodeGuru	 Contact Lens <i>For Amazon Connect</i>	

ML SERVICES

 Amazon SageMaker	Ground Truth	ML Marketplace	SageMaker Studio IDE								Neo	A2I
			Built-in algorithms	Notebooks	Experiments	Processing & Model Evaluation	Model training & tuning	Debugger	Autopilot	Model hosting	Model Monitor	

ML FRAMEWORKS & INFRASTRUCTURE

 TensorFlow		 mxnet	 GLUON		 Keras	 PYTORCH		 fast.ai	 DeepGraphLibrary	Deep Learning AMIs & Containers	GPUs & CPUs	ElasticInference	Inferentia	FPGA
---	--	--	--	--	--	--	--	--	---	---------------------------------	-------------	------------------	------------	------

The machine learning workflow is iterative and complex

Prepare

Build

Train & Tune

Deploy & Manage

101011010
010101010
000011110



Collect and
prepare
training data

Choose or build an
ML algorithm



Set up and manage
environments
for training



Train, debug, and
tune models



Manage training runs



Deploy
model in
production



Monitor
models



Validate
predictions



Scale and manage
the production
environment

Amazon SageMaker helps you build, train, and deploy models

Prepare

Build

Train & Tune

Deploy & Manage

Web-based IDE for machine learning

Automatically build and train models

Fully managed data processing jobs and data labeling workflows

101011010
010101010
000011110

One-click collaborative notebooks and built-in, high performance algorithms and models



Choose or build an ML algorithm

One-click training



Set up and manage environments for training

Debugging and optimization



Train, debug, and tune models

Visually track and compare experiments




Manage training runs

One-click deployment and autoscaling



Deploy model in production

Automatically spot data drift



Monitor models

Add human review of predictions



Validate predictions

Fully managed with auto-scaling for 75% less



Scale and manage the production environment

Same service and APIs, from experimentation to production

Model options



AWS Marketplace
for Machine
Learning



Training code



Amazon SageMaker
AutoPilot

Factorization Machines
Linear Learner
Principal Component
Analysis
K-Means Clustering
XGBoost

Built-in Algorithms (17)
No ML coding required



Built-in
Frameworks
Bring your own code
Use open source containers



Bring Your Own
Full control, run your
container
R, C++, etc.














Fully managed training, spot instances included

Frameworks and Infrastructure


The AWS ML Stack

Broadest and most complete set of Machine Learning capabilities









AI SERVICES

	VISION	SPEECH		TEXT			SEARCH	CHATBOTS	PERSONALIZATION	FORECASTING	FRAUD	DEVELOPMENT	CONTACT CENTERS	
	 Amazon Rekognition	 Amazon Polly	 Amazon Transcribe <i>+Medical</i>	 Amazon Comprehend <i>+Medical</i>	 Amazon Translate	 Amazon Textract	 Amazon Kendra	 Amazon Lex	 Amazon Personalize	 Amazon Forecast	 Amazon Fraud Detector	 Amazon CodeGuru	 Contact Lens <i>For Amazon Connect</i>	

ML SERVICES

 Amazon SageMaker	Ground Truth	ML Marketplace	SageMaker Studio IDE									Neo	A2I
			Built-in algorithms	Notebooks	Experiments	Processing & Model Evaluation	Model training & tuning	Debugger	Autopilot	Model hosting	Model Monitor		

ML FRAMEWORKS & INFRASTRUCTURE

 TensorFlow		 mxnet	 GLUON		 Keras	Deep Learning AMLs & Containers	GPUs & CPUs	ElasticInference	Inferentia	FPGA
 PYTORCH			 fast.ai		 DeepGraphLibrary					

AWS: The platform of choice for TensorFlow

<https://aws.amazon.com/tensorflow>



89% of all deep learning workloads in the cloud run on AWS

85% of all TensorFlow workloads in the cloud run on AWS

Source: Nucleus Research, T147, October 2019

Amazon Elastic Inference

<https://aws.amazon.com/blogs/machine-learning/optimizing-costs-in-amazon-elastic-inference-with-amazon-tensorflow/>



Lower
inference
costs
up to 75%

Integrated with
Amazon EC2 and
Amazon SageMaker



Match
capacity
to demand

Support for TensorFlow,
Apache MXNet (Incubating)
—PyTorch coming soon



Available between
1 to 32 TFLOPS
per accelerator

Single and mixed-
precision operations

Amazon EC2 Inferentia

- Fast, low-latency inferencing at a very low cost
 - 64 TeraOPS on 16-bit floating point (FP16 and BF16) and mixed-precision data.
 - 128 TeraOPS on 8-bit integer (INT8) data.
- Neuron SDK: <https://github.com/aws/aws-neuron-sdk>
 - Available in Deep Learning AMIs and Deep Learning Containers
 - TensorFlow and Apache MXNet, PyTorch coming soon

<https://ml.aws>

<https://aws.amazon.com/sagemaker>
<https://github.com/aws/sagemaker-python-sdk>
<https://github.com/aws-labs/amazon-sagemaker-examples>

<https://youtube.com/juliensimonfr>
<https://medium.com/@julsimon>

Published August 2020

Discount link for the paper edition on Amazon (US only)
<https://www.amazon.com/gp/mpc/AOHJSZC7A0AV5>

Discount code for the e-book edition on Packt
20SAGEMAKER
<https://www.packtpub.com/product/learn-amazon-sagemaker/9781800208919>

Valid until November 11th

