

DEV DAY



DEV DAY

MLS 2

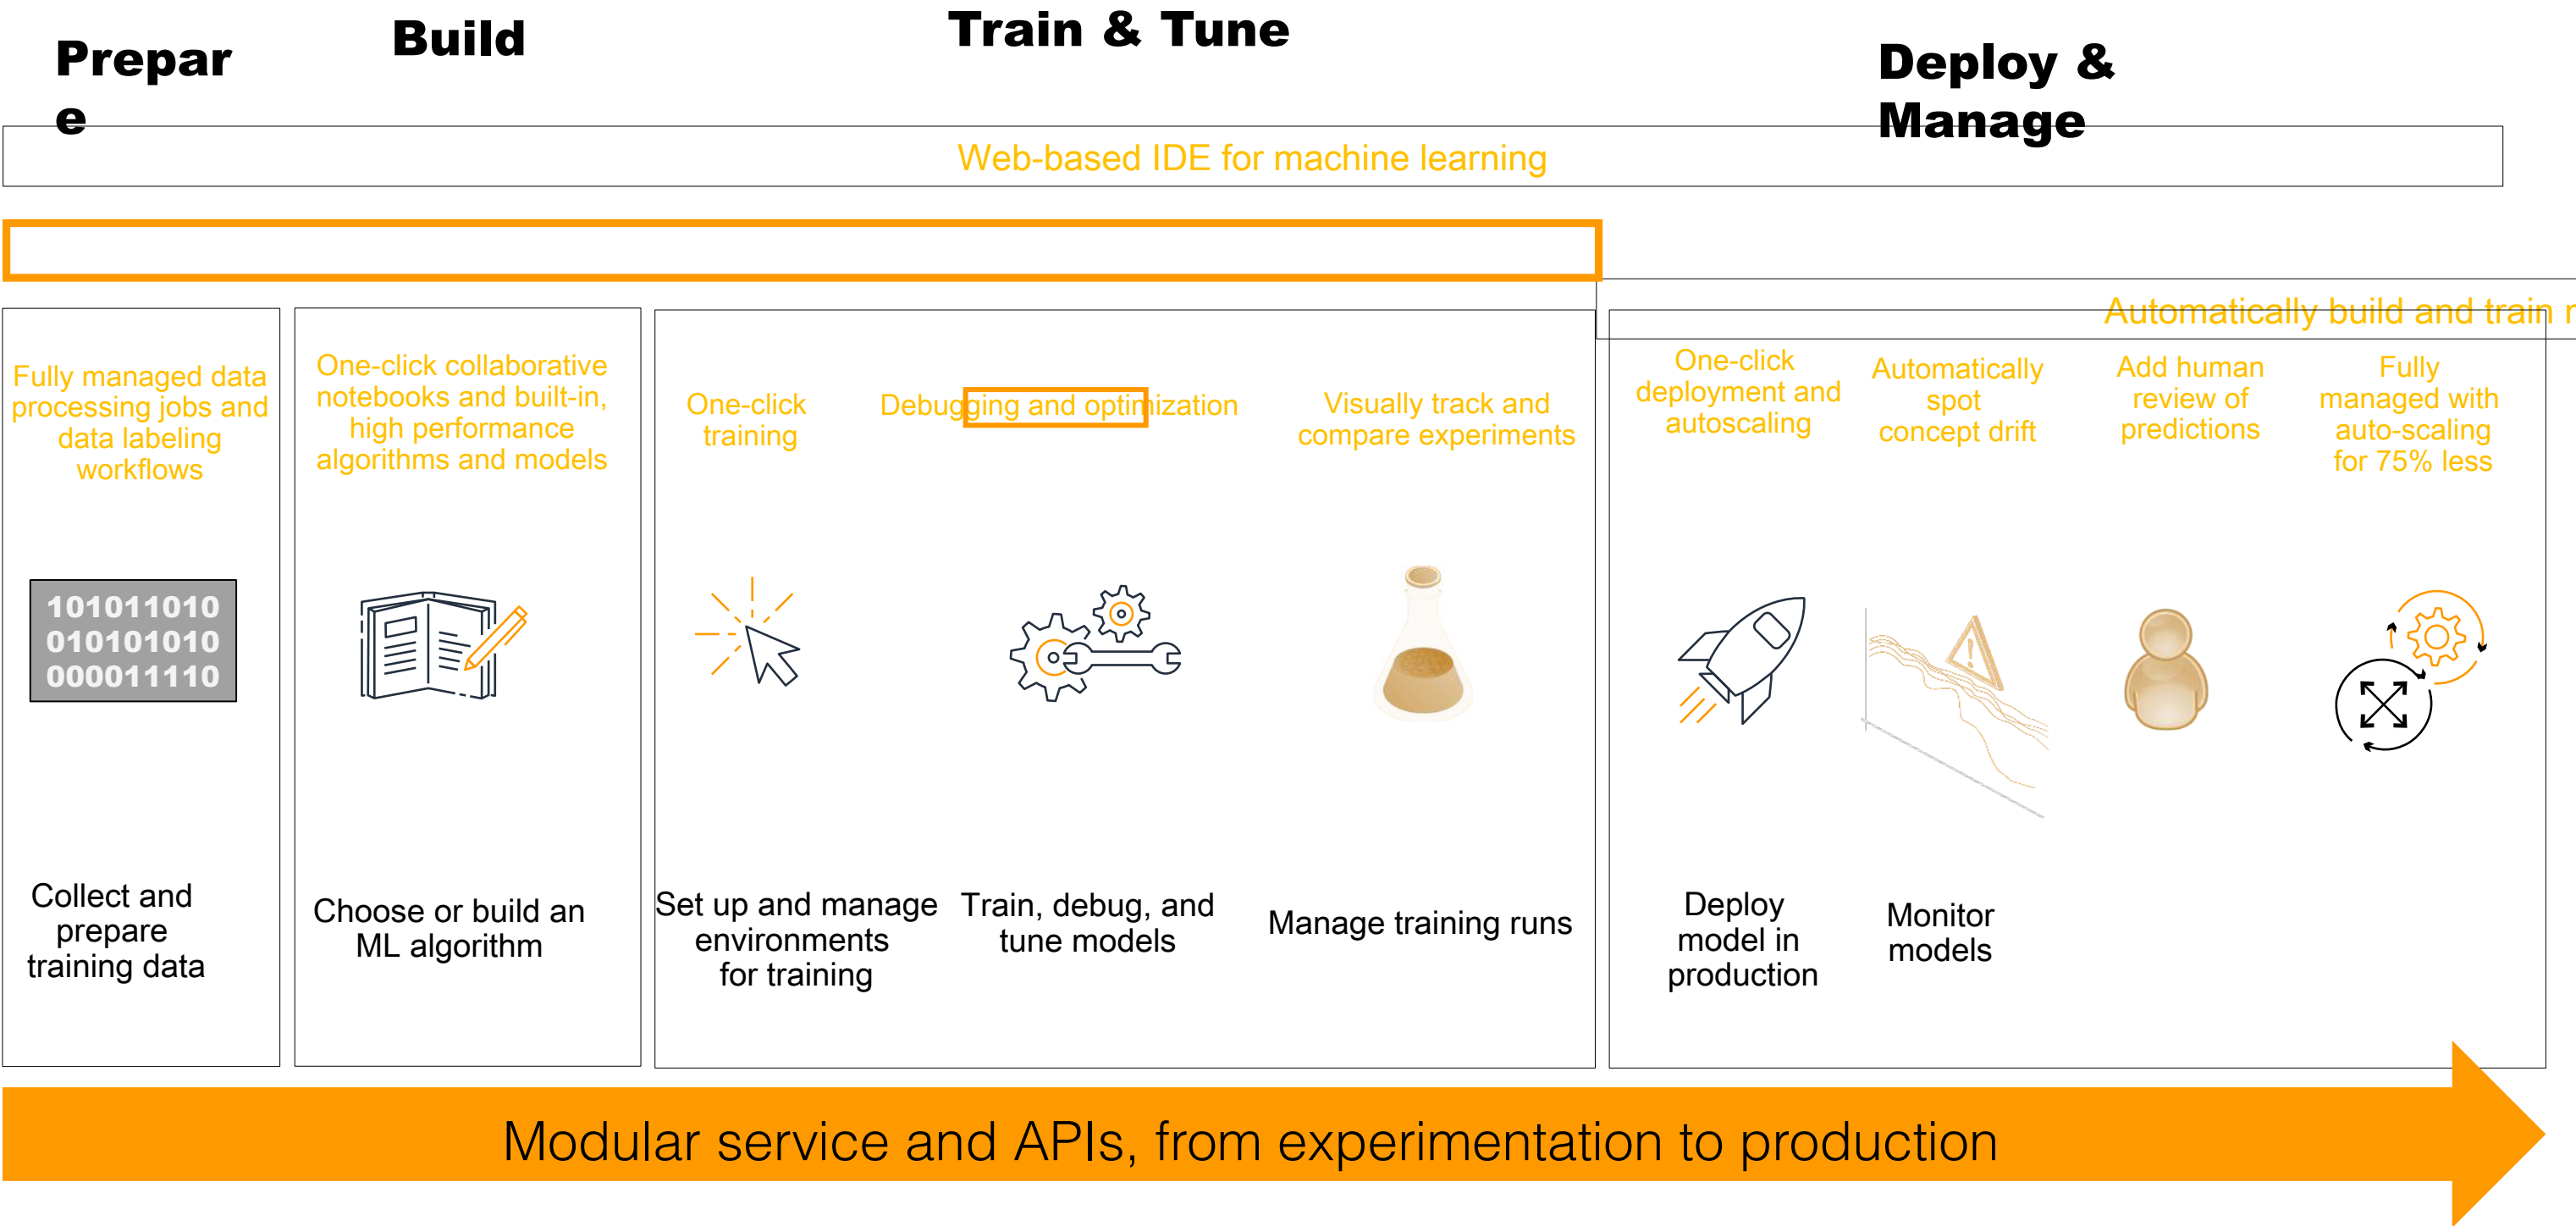
Building Models Automatically with Amazon SageMaker

Julien Simon
Global Evangelist, AI & Machine Learning
Amazon Web Services

@julsimon



Amazon SageMaker helps you build, train, and deploy models



DEV DAY

Automatic Model Tuning with Amazon SageMaker



© 2020, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Algorithms require many hyperparameters

XGBoost

Tree depth
Max leaf nodes
Gamma
Eta
Lambda
Alpha
...

Which ones
are the most
influential?

Which values
should I pick?

How many
combinations
should I try?

Neural Networks

Number of layers
Hidden layer width
Learning rate
Embedding
dimensions
Dropout
...

Setting hyperparameters in Amazon SageMaker

- Built-in algorithms
 - Python **parameters** set on the relevant estimator (*KMeans*, *LinearLearner*, etc.)
`xgb.set_hyperparameters(max_depth=5, eta=0.2, gamma=4)`
- Built-in frameworks
 - *hyperparameters* parameter passed to the relevant estimator (*TensorFlow*, *MXNet*, etc.)
 - This must be a Python **dictionary**
`tf_estimator = TensorFlow(..., hyperparameters={'epochs': 1, 'lr': '0.01'})`
 - Your code must be able to accept them as command-line arguments (**script mode**)
- Bring your own container
 - *hyperparameters* parameter passed to the *Estimator*
 - This must be Python dictionary
 - It's automatically copied inside the container: `/opt/ml/input/config/hyperparameters.json`

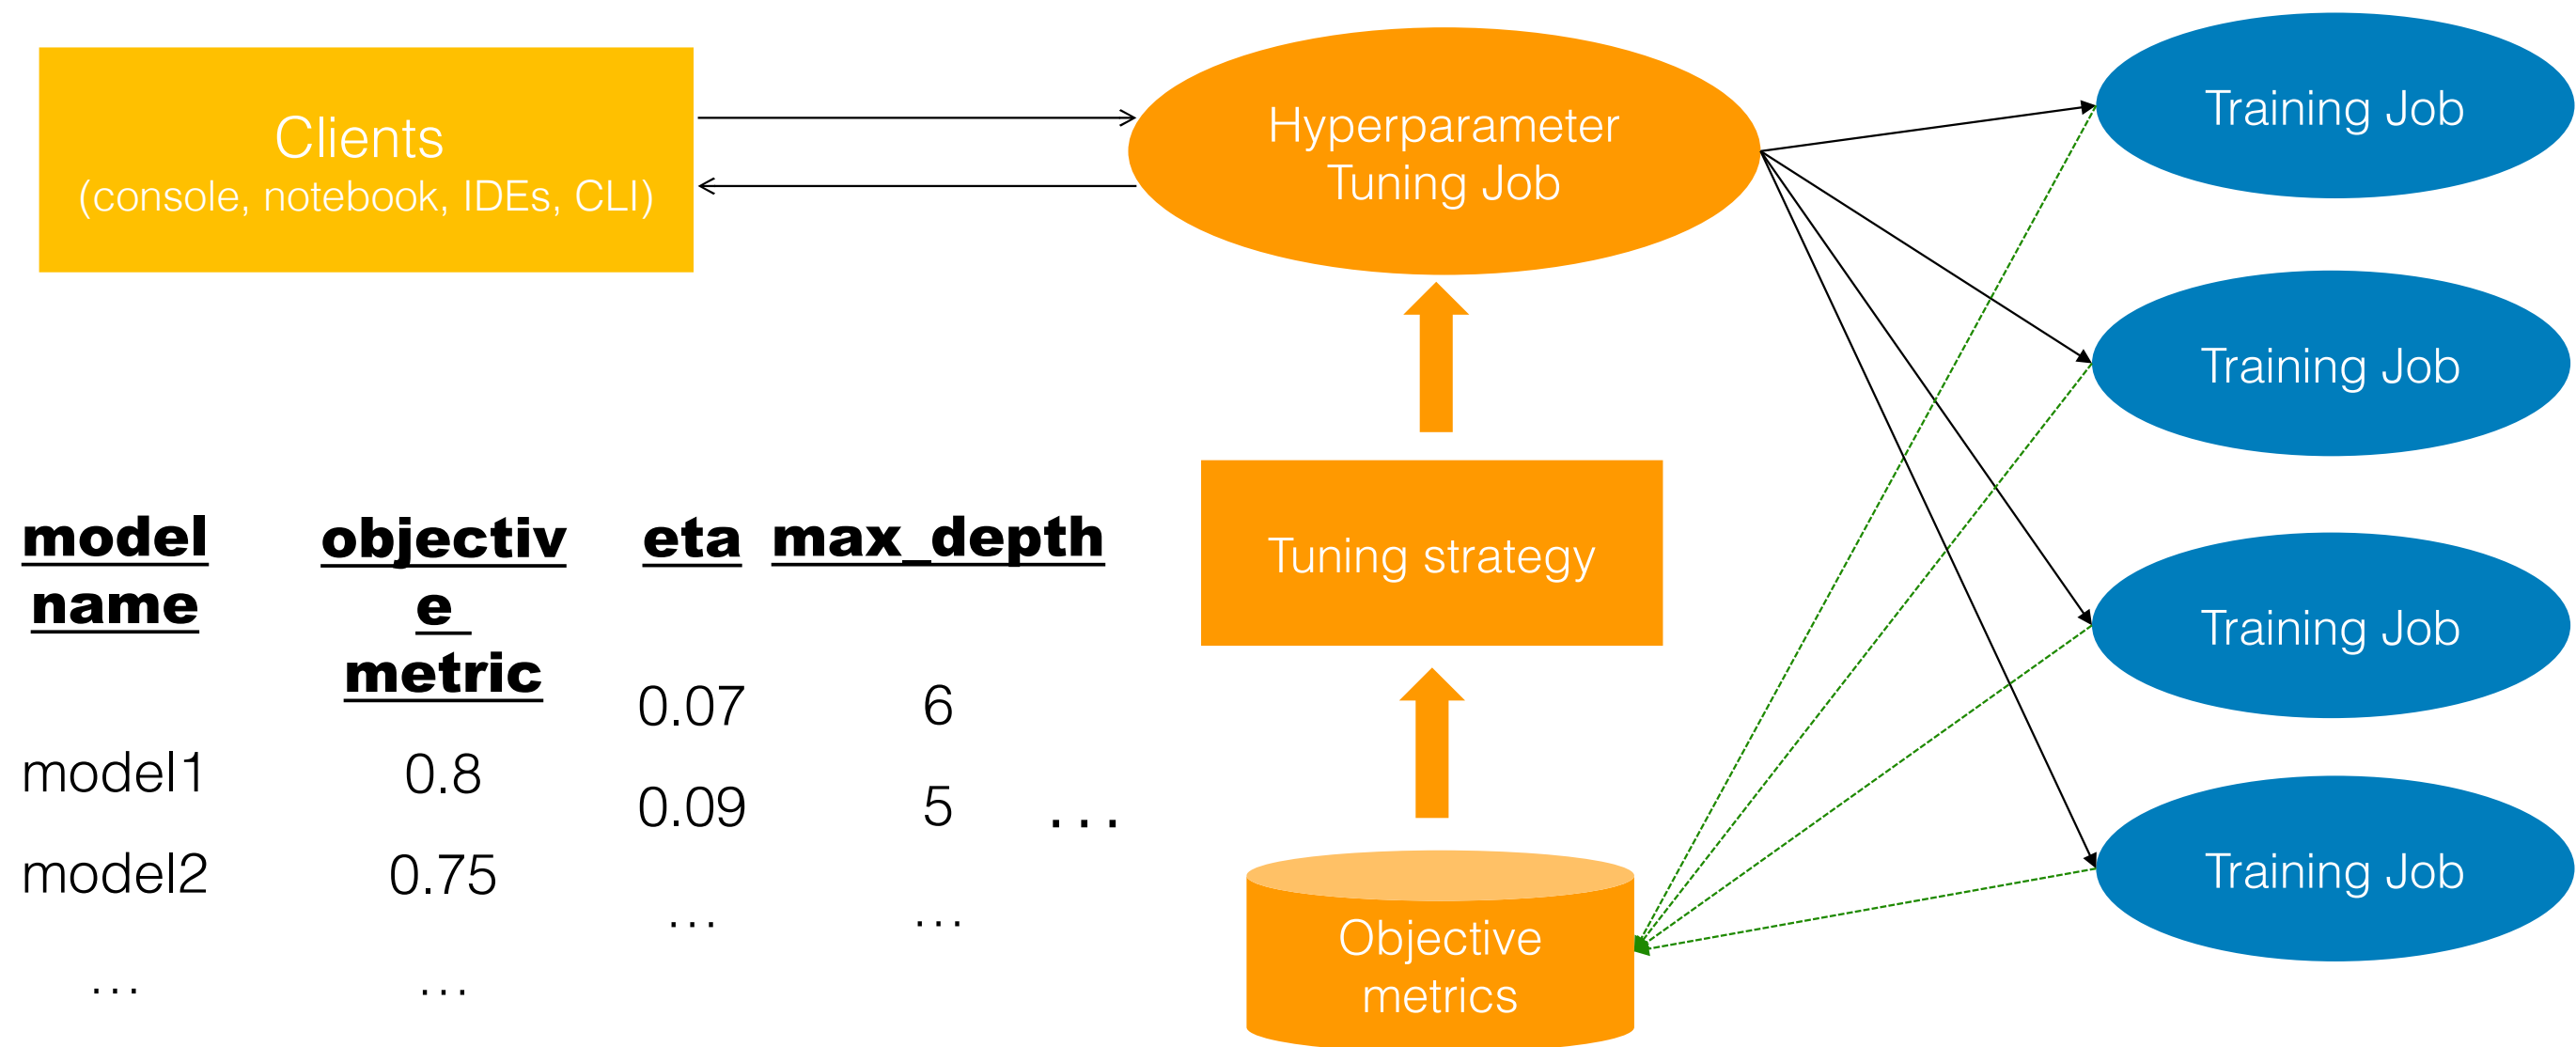
Tactics to find the optimal set of hyperparameters

- **Manual Search**: "I know what I'm doing"
- **Grid Search**: "X marks the spot"
Typically training hundreds of models
Slow and expensive
- **Random Search**: "Spray and pray"
« *Random Search for Hyper-Parameter Optimization* », Bergstra & Bengio, 2012
Works better and faster than Grid Search
But... but... but... it's random!
- **Hyperparameter Optimization**: use ML to predict hyperparameters
Training fewer models
Gaussian Process Regression and Bayesian Optimization
https://docs.aws.amazon.com/en_pv/sagemaker/latest/dg/automatic-model-tuning-how-it-works.html

Automatic Model Tuning in Amazon SageMaker

1. Define an *Estimator* the normal way
2. Define the *metric* to tune on
 - Pre-defined metrics for built-in algorithms and frameworks
 - Or anything present in the training log, provided that you pass a regular expression for it
3. Define *parameter ranges* to explore
 - Type: categorical (avoid if possible), integer, continuous (aka floating point)
 - Range of values
 - Scaling: linear, logarithmic, reverse logarithmic
4. Create an *HyperparameterTuner*
 - *Estimator*, metric, parameters, total number of jobs, number of jobs in parallel
 - Strategy: bayesian (default), or random search
5. Launch the tuning job with *fit()*

Workflow



Automatic Model Tuning in Amazon SageMaker

- You can **view** ongoing tuning jobs in the AWS console
 - List of training jobs
 - Best training job
- You can also **query** their status with the SageMaker SDK
- Calling ***deploy()*** on the *HyperparameterTuner* deploys the best job
 - The best job so far if the tuning job has not yet completed

DEV DAY

Demo



Tips

- Use the **bayesian strategy** for better, faster, cheaper results
 - Most customers use random search as a baseline, to check that bayesian performs better
- Don't run **too many jobs** in parallel
 - This gives the bayesian strategy fewer opportunities to predict
 - Instance limits!
- Don't run **too many jobs**
 - Bayesian typically requires 10x fewer jobs than random
 - Cost vs business benefits (beware of diminishing returns)

Resources on Automatic Model Tuning

Documentation

<https://docs.aws.amazon.com/sagemaker/latest/dg/automatic-model-tuning.html>

<https://sagemaker.readthedocs.io/en/stable/tuner.html>

Notebooks

https://github.com/aws-labs/amazon-sagemaker-examples/tree/master/hyperparameter_tuning

Blog posts

<https://aws.amazon.com/blogs/aws/sagemaker-automatic-model-tuning/>

<https://aws.amazon.com/blogs/machine-learning/amazon-sagemaker-automatic-model-tuning-produces-better-models-faster/>

<https://aws.amazon.com/blogs/machine-learning/amazon-sagemaker-automatic-model-tuning-now-supports-early-stopping-of-training-jobs/>

<https://aws.amazon.com/blogs/machine-learning/amazon-sagemaker-automatic-model-tuning-becomes-more-efficient-with-warm-start-of-hyperparameter-tuning-jobs/>

<https://aws.amazon.com/blogs/machine-learning/amazon-sagemaker-automatic-model-tuning-now-supports-random-search-and-hyperparameter-scaling/>

DEV DAY

AutoML with Amazon SageMaker Autopilot



© 2020, Amazon Web Services, Inc. or its affiliates. All rights reserved.

AutoML

- AutoML aims at automating the process of building a model
 - **Problem identification**: looking at the data set, what class of problem are we trying to solve?
 - **Algorithm selection**: which algorithm is best suited to solve the problem?
 - **Data preprocessing**: how should data be prepared for best results?
 - **Hyperparameter tuning**: what is the optimal set of training parameters?
- Black box vs. white box
 - Black box: the **best model** only
 - Hard to understand the model, impossible to reproduce it manually
 - White box: the **best model**, other **candidates**, full **source code** for preprocessing and training
 - See how the model was built, and keep tweaking for extra performance

AutoML with Amazon SageMaker Autopilot

- SageMaker Autopilot covers all steps
 - **Problem identification**: looking at the data set, what class of problem are we trying to solve?
 - **Algorithm selection**: which algorithm is best suited to solve the problem?
 - **Data preprocessing**: how should data be prepared for best results?
 - **Hyperparameter tuning**: what is the optimal set of training parameters?
- Autopilot is **white box** AutoML
 - You can understand how the model was built, and you can keep tweaking
- Supported algorithms at launch: **regression** and **classification**
 - Linear Learner
 - Factorization Machines
 - KNN
 - XGBoost

AutoML with Amazon SageMaker Autopilot

1. Upload the **unprocessed dataset** to S3
2. Configure the AutoML **job**
 - Location of dataset
 - Completion criteria
3. Launch the job
4. View the list of **candidates** and the **autogenerated notebooks**
5. Deploy the **best candidate** to a real-time endpoint, or use batch transform

DEV DAY

Demo



Resources on Amazon SageMaker AutoPilot

Documentation

<https://docs.aws.amazon.com/sagemaker/latest/dg/autopilot-automate-model-development.html>

<https://sagemaker.readthedocs.io/en/stable/automl.html>

Notebooks

<https://github.com/aws-labs/amazon-sagemaker-examples/tree/master/autopilot>

Blog posts

<https://aws.amazon.com/blogs/aws/amazon-sagemaker-autopilot-fully-managed-automatic-machine-learning/>

For more content:

- AWS blog: <https://aws.amazon.com/blogs/aws/author/julsimon>
- Medium blog: <https://medium.com/@julsimon>
- YouTube: <https://youtube.com/juliensimonfr>
- Podcast: <http://julsimon.buzzsprout.com>
- Twitter: <https://twitter.com/@julsimon>

DEV DAY

Thank you!

Julien Simon
Global Evangelist, AI & Machine Learning
Amazon Web Services

@julsimon

