

Hugging Face on Amazon SageMaker

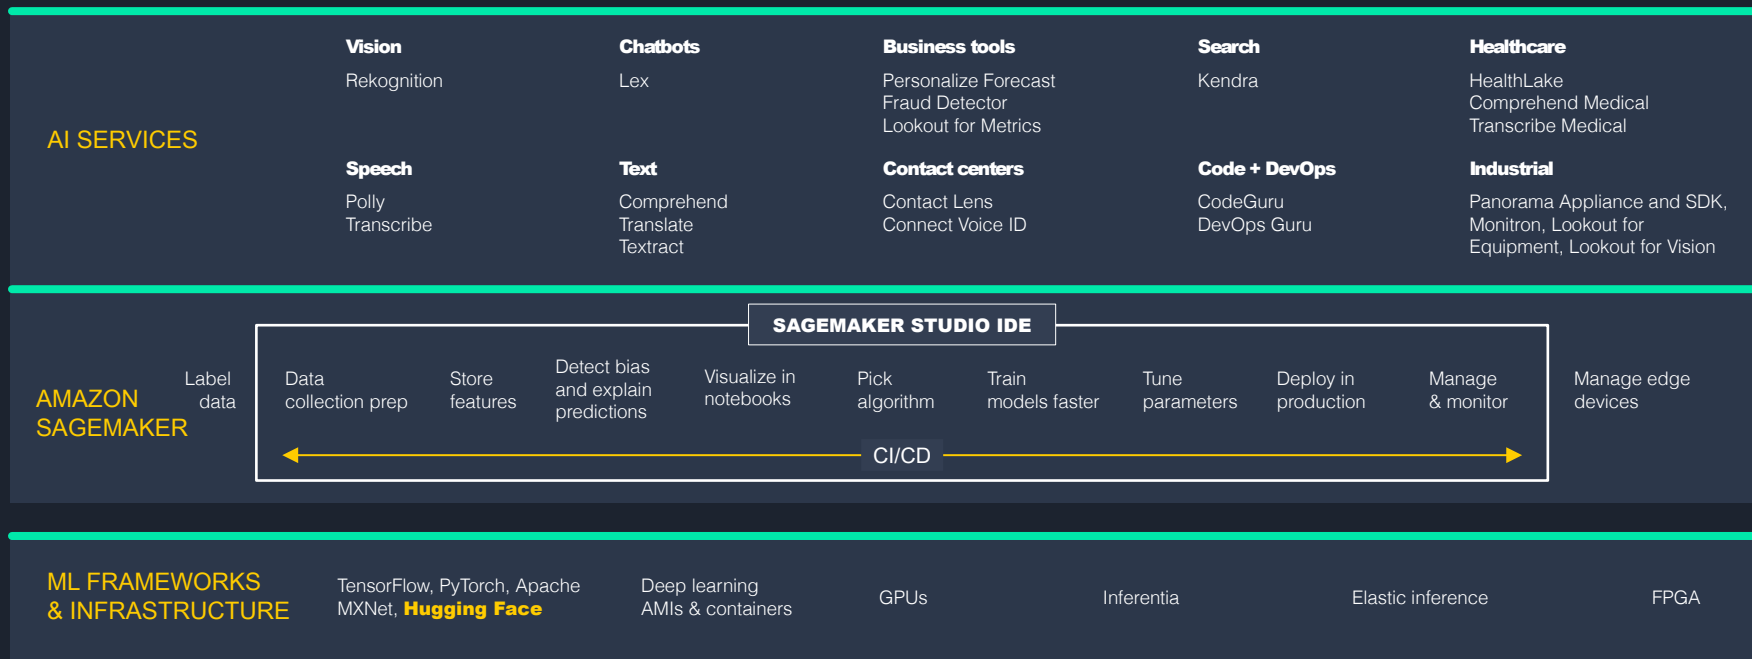
Scale and accelerate your NLP projects with Hugging Face models

AWS ML Vision

Make machine learning accessible to all customers
(developers, data scientists, business users)

The AWS ML stack

Broadest and most complete set of machine learning capabilities



Customers building NLP solutions on AWS



Hugging Face

- Startup founded in 2016
- Based in Paris and New York
- Open source solutions for Natural Language Processing
- \$40 million Series B rounds a few days ago

Hugging Face libraries



Open-source

Datasets, Tokenizers and Transformers



Popular

52k+ GitHub stars (March 2021), 1M+ downloads per month



Intuitive

NLP-specific Python frontends based on PyTorch or TensorFlow



State of the art

Transformer-based models are state-of-the-art, enable transferlearning and scale



Comprehensive

Model zoo with 6000+ model architectures, 160+ languages

What are the Hugging Face libraries?



Open-source

Datasets, Tokenizers and Transformers



Popular

52k+ GitHub stars (March 2021), 1M+ downloads per month



Intuitive

NLP-specific Python frontends based on PyTorch or TensorFlow



State of the art

Transformer-based models are state-of-the-art, enable transferlearning and scale



Comprehensive

Model zoo with 6000+ model architectures, 160+ languages

Introducing a new Hugging Face experience in Amazon SageMaker



A new Deep Learning Container (DLC) for SageMaker, developed with HF (training only for now). Datasets, Transformers and Tokenizers, with SageMaker integration such as SageMaker Debugger and the SageMaker distributed training libraries



A Hugging Face Estimator in the SageMaker SDK to launch NLP scripts on scalable, cost-effective SageMaker training jobs with minimal code changes (script mode), and without any Docker and infrastructure work



An example gallery with samples of Hugging Face scripts on Amazon SageMaker



Supported and maintained by AWS

What are the benefits of running Hugging Face in Amazon SageMaker?



Cost-effective - SageMaker optimizes scale, performance, and costs



MLOps-ready – Monitoring, profiling and debugging (SageMaker Debugger & Profiler), experiment management (SageMaker Experiments), etc.

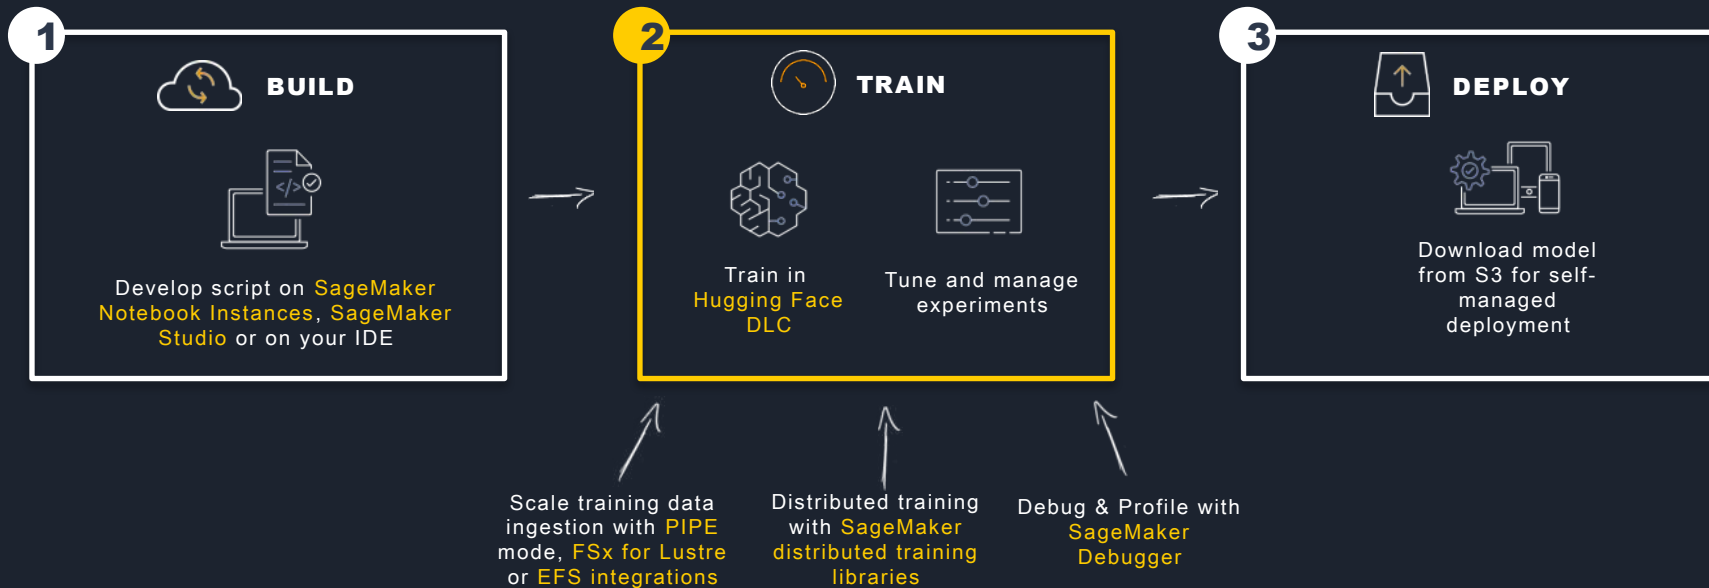


Scalable – Able to run on clusters of GPUs, with efficient data-parallel and model-parallel distribution provided by Amazon SageMaker. Ability to launch several concurrent jobs at the same time with the async mode of the API



Secure – High bar on security, with available mechanisms including encryption at rest and in transit, VPC connectivity and fine-grained IAM permissions

Integrated workflow with Amazon SageMaker



Customers using Hugging Face on SageMaker



Quantum Health

Quantum Health is on a mission to make healthcare navigation smarter, simpler, and more cost-effective for everyone. They use Hugging Face and Amazon SageMaker for use cases like text classification, text summarization, and Q&A to help agents and members.

“We use Hugging Face and Amazon SageMaker a lot for many NLP use cases such as text classification, text summarization, and Q&A with the goal of helping our agents and members. For some use cases, we just use the Hugging Face models directly and for others we fine tune them on SageMaker. We are excited about the integration of Hugging Face Transformers into Amazon SageMaker to make use of the distributed libraries during training to shorten the training time for our larger datasets

Jorge Grisman, NLP Data Scientist, Quantum Health

Customers using Hugging Face on SageMaker



Kustomer

Kustomer is a customer service CRM platform for managing high support volume with ease. They use machine learning models to help customers contextualize conversations, remove time-consuming tasks, and deflect repetitive questions.



In our business, we use machine learning models to help customers contextualize conversations, remove time-consuming tasks, and deflect repetitive questions. We use Hugging Face and Amazon SageMaker extensively, and we are excited about the integration of Hugging Face Transformers into SageMaker since it will simplify the way we fine tune machine learning models for text classification and semantic search

Victor Peinado, ML Software Engineering Manager, Kustomer

Get started with Hugging Face on Amazon SageMaker

- Launch blog post on the AWS Machine Learning blog
- SageMaker docs: <https://docs.aws.amazon.com/sagemaker>
- SageMaker SDK docs: <https://sagemaker.readthedocs.io/>
- Hugging Face examples: <https://github.com/huggingface/notebooks/>